

[\(https://www.whizlabs.com/learn/\)](https://www.whizlabs.com/learn/) > [My Courses \(https://www.whizlabs.com/learn/my-courses\)](https://www.whizlabs.com/learn/my-courses)
 > [AWS Certified Developer Associate \(https://www.whizlabs.com/learn/course/aws-cda-practice-tests#section-1\)](https://www.whizlabs.com/learn/course/aws-cda-practice-tests#section-1)
 > [Practice Test IV \(https://www.whizlabs.com/learn/course/aws-cda-practice-tests/quiz/12748\)](https://www.whizlabs.com/learn/course/aws-cda-practice-tests/quiz/12748) > **Report**

PRACTICE TEST IV

Attempt 2 **Completed on** Tuesday, 29 January 2019, 02:59 PM
Marks Obtained 0 / 60 **Time Taken** 00 H 00 M 05 S
Your score is 0.0% **Result** Fail

Domains / Topics wise Quiz Performance Report

S.No.	Topic	Total Questions	Correct	Incorrect	Unattempted
1	Other	59	0	1	58
2	Refactoring	1	0	0	1

60 Questions	0 Correct	1 Incorrect	59 Unattempted
------------------------	---------------------	-----------------------	--------------------------

Show Answers

All

QUESTION 1 INCORRECT

A startups photo-sharing site is deployed in a VPC. An ELB distributes web traffic across two subnets. ELB session stickiness is configured to use the AWS-generated session cookie, with a session TTL of 5 minutes. The webserver Auto Scaling Group is configured as: min-size=4, max-size=4. The startups preparing for a public launch, by running load-testing software installed on a single EC2 instance running in us-west-2a. After 60 minutes of load-testing, the webserver logs show: Which recommendations can help ensure load-testing HTTP requests are evenly distributed across the four webserver? Choose 2 answers from the options below

WEBSERVER LOGS	# of HTTP requests from load-tester	# of HTTP requests from private beta users
webserver #1 (subnet in us-west-2a):	19,210	434
webserver #2 (subnet in us-west-2a):	21,790	490
webserver #3 (subnet in us-west-2b):	0	410
webserver #4 (subnet in us-west-2b):	0	428

- ☒ **A. Re-configure the load-testing software to re-resolve DNS for each web request.** ✓
☐ **B. Use a 3rd-party load-testing service which offers globally-distributed test clients.** ✓
☒ **C. Configure ELB and Auto Scaling to distribute across us-west-2a and us-west-2c.** ✕
☐ **D. Configure ELB session stickiness to use the app-specific session cookie.**
☐ **E. Launch and run the load-tester EC2 instance from us-east-1 instead.**

Explanation :

Answer - A and B

"If you do not ensure that DNS is re-resolved or use multiple test clients to simulate increased load, the test may continue to hit a single IP address when Elastic Load Balancing has actually allocated many more IP addresses. Because your end users will not all be resolving to that single IP address, your test will not be a realistic sampling of real-world behavior."

Please refer to the below article URL:

- <http://aws.amazon.com/articles/1636185810492479> (<http://aws.amazon.com/articles/1636185810492479>)

Ask our Experts



QUESTION 2 UNATTEMPTED

Company A has recently launched an online commerce site for bicycles on AWS. They have a "Product" DynamoDB table that stores details for each bicycle, such as, manufacturer, color, price, quantity and size to display in the online store. Due to customer demand, they want to include an image for each bicycle along with the existing details.

Which approach below provides the least impact to provisioned throughput on the "Product" table? Choose an answer from the options below

- ☐ A. Create an "Images" DynamoDB table to store the Image with a foreign key constraint to the "Product" table
- ☐ B. Add an image data type to the "Product" table to store the images in binary format
- ☐ C. Serialize the image and store it in multiple DynamoDB tables
- ☒ D. Store the images in Amazon S3 and add an S3 URL pointer to the "Product" table item for each image ✓

Explanation :

Answer – D

This is clearly given in the AWS documentation

Q: When should I use Amazon DynamoDB vs Amazon S3?

Amazon DynamoDB stores structured data, indexed by primary key, and allows low latency read and write access to items ranging from 1 byte up to 400KB. Amazon S3 stores unstructured blobs and suited for storing large objects up to 5 TB. In order to optimize your costs across AWS services, large objects or infrequently accessed data sets should be stored in Amazon S3, while smaller data elements or file pointers (possibly to Amazon S3 objects) are best saved in Amazon DynamoDB.

For more information on DynamoDB, please refer to the link:

- <https://aws.amazon.com/dynamodb/faqs/> (<https://aws.amazon.com/dynamodb/faqs/>)

Ask our Experts



QUESTION 3 UNATTEMPTED

Which DynamoDB limits can be raised by contacting AWS support? Choose 2 answers from the options below

- ☐ A. The number of hash keys per account
- ☐ B. The maximum storage used per account
- ☐ C. The number of tables per account ✓
- ☐ D. The number of local secondary indexes per account
- ☐ E. The number of provisioned throughput units per account ✓

Explanation :

Answer - C and E

Please find the supporting AWS documentation

Provisioned throughputs

An AWS account places some initial maximum limits on the throughput you can provision:

- US East (N. Virginia) Region:
 - Per table – 40,000 read capacity units and 40,000 write capacity units
 - Per account – 80,000 read capacity units and 80,000 write capacity units
- All Other Regions:
 - Per table – 10,000 read capacity units and 10,000 write capacity units
 - Per account – 20,000 read capacity units and 20,000 write capacity units

The provisioned throughput limit includes the sum of the capacity of the table together with the capacity of all of its global secondary indexes.

Tables per account

Tables Per Account

For any AWS account, there is an initial limit of 256 tables per region.

You can request an increase on this limit. For more information, go to <https://aws.amazon.com/support>.

For more information on the service limits for DynamoDB please refer to the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Limits.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Limits.html>)

Ask our Experts



QUESTION 4 UNATTEMPTED

Which of the following items are required to allow an application deployed on an EC2 instance to write data to a DynamoDB table? Assume that no security Keys are allowed to be stored on the EC2 instance. Choose an answer from options below:

- ☐ A. Create an IAM user that allows write access to the DynamoDB table.
- ☐ B. Launch an EC2 instance with the IAM user included in the launch configuration.
- ☒ C. Create an IAM role that allows write access to the DynamoDB table and attach to the instance. ✓
- ☐ D. Create an IAM user and allow programmatic access.

Explanation :

Answer - C

As the question mentions no keys are allowed to be stored on EC2 instance and hence IAM user is incorrect option. Statement from AWS with regard to controlling access to DynamoDB:

Applications running on Amazon EC2 – You can use an IAM role to manage temporary credentials for applications that are running on an EC2 instance and making AWS API requests. This is preferable to storing access keys within the EC2 instance. To assign an AWS role to an EC2 instance and make it available to all of its applications, you create an instance profile that is attached to the instance. An instance profile contains the role and enables programs that are running on the EC2 instance to get temporary credentials. For more information, see Using an IAM Role to Grant Permissions to Applications Running on Amazon EC2 Instances (http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_use_switch-role-ec2.html) in the *IAM User Guide*.

Please refer to the link below for more information:

- <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/authentication-and-access-control.html>
(<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/authentication-and-access-control.html>)

Ask our Experts



QUESTION 5 UNATTEMPTED

Which EC2 API call would you use to retrieve a list of Amazon Machine Images (AMIs)? Choose an answer from

the options below

- ☐ A. DescribeInstances
- ☐ B. You cannot retrieve a list of AMIs as there are over 10,000 AMIs
- ☐ C. GetAMIs
- ☒ D. DescribeImages ✓
- ☐ E. DescribeAMIs

Explanation :

Answer – D

Describes one or more of the images (AMIs, AKIs, and ARIs) available to you. Images available to you include public images, private images that you own, and private images owned by other AWS accounts but for which you have explicit launch permissions.

Additionally, you can use Request Parameter Filters to list images on conditions.

Please read the reference and examples listed at

http://docs.aws.amazon.com/AWSEC2/latest/APIReference/API_DescribeImages.html#API_DescribeImages_Examples

(http://docs.aws.amazon.com/AWSEC2/latest/APIReference/API_DescribeImages.html#API_DescribeImages_Examples)

For more information on DescribeImages, please refer to the below link:

http://docs.aws.amazon.com/AWSEC2/latest/APIReference/API_DescribeImages.html

(http://docs.aws.amazon.com/AWSEC2/latest/APIReference/API_DescribeImages.html)

Ask our Experts



QUESTION 6 UNATTEMPTED

An Amazon S3 bucket, “myawsbucket” is configured with website hosting in Tokyo region, what is the region-specific website endpoint? Choose an answer from the options below

- ☐ A. www.myawsbucket.ap-northeast-1.amazonaws.com
- ☒ B. myawsbucket.s3-website-ap-northeast-1.amazonaws.com ✓
- ☐ C. myawsbucket.amazonaws.com
- ☐ D. myawsbucket.tokyo.amazonaws.com

Explanation :

Answer – B

You have the chance to enable static web site hosting for S3 buckets. This can be done via the properties option for the bucket. The endpoint of the bucket for static hosting will also be configured.

▸ Permissions

▼ Static Website Hosting

You can host your static website entirely on Amazon S3. Once you enable your bucket for static website hosting, all your content is accessible to web browsers via the Amazon S3 website endpoint for your bucket.

Endpoint: devtoolslogging.s3-website-ap-southeast-1.amazonaws.com

Each bucket serves a website namespace (e.g. “www.example.com”). Requests for your host name (e.g. “example.com” or “www.example.com”) can be routed to the contents in your bucket. You can also redirect requests to another host name (e.g. redirect “example.com” to “www.example.com”). See our [walkthrough](#) for how to set up an Amazon S3 static website with your host name.

For more information on S3 buckets, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingBucket.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingBucket.html>)

Ask our Experts



QUESTION 7 UNATTEMPTED

A corporate web application is deployed within an Amazon VPC, and is connected to the corporate data center via IPsec VPN. The application must authenticate against the on-premise LDAP server. Once authenticated, logged-in users can only access an S3 keyspace specific to the user. Which two approaches can satisfy the objectives?

Choose 2 answers from the options below

- ☐ A. The application authenticates against LDAP, and retrieves the name of an IAM role associated with the user. The application then calls the IAM Security Token Service to assume that IAM Role. The application can use the temporary credentials to access the appropriate S3 bucket. ✓
- ☐ B. Develop an identity broker which authenticates against IAM Security Token Service to assume an IAM Role to get temporary AWS security credentials. The application calls the identity broker to get AWS temporary security credentials with access to the appropriate S3 bucket.
- ☐ C. The application authenticates against IAM Security Token Service using the LDAP credentials. The application uses those temporary AWS security credentials to access the appropriate S3 bucket.
- ☐ D. The application authenticates against LDAP. The application then calls the IAM Security Service to login to IAM using the LDAP credentials. The application can use the IAM temporary credentials to access the appropriate S3 bucket.
- ☐ E. Develop an identity broker which authenticates against LDAP, and then calls IAM Security Token Service to get IAM federated user credentials. The application calls the identity broker to get IAM federated user credentials with access to the appropriate S3 bucket. ✓

Explanation :

Answer - A and E

If you already manage user identities outside of AWS, you can use IAM *identity providers* instead of creating IAM users in your AWS account. With an identity provider, you can manage your user identities outside of AWS and give these external user identities permissions to use AWS resources in your account. This is useful if your organization already has its own identity system, such as a corporate user directory. It is also useful if you are creating a mobile app or web application that requires access to AWS resources. Option B is incorrect as it Needs to authenticate against LDAP and not IAM.

Option C is incorrect as Need to authenticate with LDAP

Option D is incorrect as Can't login to IAM using LDAP credentials

For more information on identity providers please visit the below link:

- http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_providers.html
(http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_providers.html)

Ask our Experts



QUESTION 8 UNATTEMPTED

You are providing AWS consulting services for a company developing a new mobile application that will be leveraging Amazon SNS Mobile Push for push notifications. In order to send direct notification messages to individual devices each device registration identifier or token needs to be registered with SNS; however the developers are not sure of the best way to do this. You advise them to

- ☐ A. Call the CreatePlatformEndPoint API function to register multiple device tokens. ✓
- ☐ B. Bulk upload the device tokens contained in a CSV file via the AWS Management Console.
- ☐ C. Let the push notification service (e.g. Amazon Device Messaging) handle the registration.
- ☐ D. Implement a token vending service to handle the registration.

Explanation :

Answer – A

The CreatePlatformEndPoint API creates an endpoint for a device and mobile app on one of the supported push notification services, such as GCM and APNS. CreatePlatformEndpoint requires the PlatformApplicationArn that is returned from CreatePlatformApplication. The EndpointArn that is returned when using CreatePlatformEndpoint can then be used by the Publish action to send a message to a mobile app or by the Subscribe action for subscription to a topic. For more information on the command please visit the below link:

- http://docs.aws.amazon.com/sns/latest/api/API_CreatePlatformEndpoint.html
(http://docs.aws.amazon.com/sns/latest/api/API_CreatePlatformEndpoint.html)

Ask our Experts



QUESTION 9 UNATTEMPTED

What type of block cipher does Amazon S3 offer for server side encryption? Choose an answer from the options below

- ☐ A. RC5
- ☐ B. Blowfish
- ☐ C. Triple DES
- ☒ D. Advanced Encryption Standard ✓

Explanation :

Answer – D

Server-side encryption is about protecting data at rest. Server-side encryption with Amazon S3-managed encryption keys (SSE-S3) employs strong multi-factor encryption. Amazon S3 encrypts each object with a unique key. As an additional safeguard, it encrypts the key itself with a master key that it regularly rotates. Amazon S3 server-side encryption uses one of the strongest block ciphers available, 256-bit Advanced Encryption Standard (AES-256), to encrypt your data.

The object creation REST APIs (see Specifying Server-Side Encryption Using the REST API (<http://docs.aws.amazon.com/AmazonS3/latest/dev/SSEUsingRESTAPI.html>)) provide a request header, x-amz-server-side-encryption that you can use to request server-side encryption.

For more information on S3 encryption, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/serv-side-encryption.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/serv-side-encryption.html>)

Ask our Experts



QUESTION 10 UNATTEMPTED

You have written an application that uses the Elastic Load Balancing service to spread traffic to several web servers. Your users complain that they are sometimes forced to login again in the middle of using your application, after they have already toggled in. This is not behavior you have designed. What is a possible solution to prevent this happening? Choose an answer from the options below

- ☐ A. Use instance memory to save session state.
- ☐ B. Use instance storage to save session state.
- ☐ C. Use EBS to save session state
- ☒ D. Use ElastiCache to save session state. ✓
- ☐ E. Use Glacier to save session slate.

Explanation :

Answer – D

Amazon ElastiCache is a web service that makes it easy to deploy and run Memcached or Redis protocol-compliant server nodes in the cloud (<https://aws.amazon.com/what-is-cloud-computing/>). Amazon ElastiCache improves the performance of web applications by allowing you to retrieve information from a fast, managed, in-memory system, instead of relying entirely on slower disk-based databases. The service simplifies and offloads the management, monitoring and operation of in-memory environments, enabling your engineering resources to focus on developing applications. Using Amazon ElastiCache, you can not only improve load and response times to user actions and queries, but also reduce the cost associated with scaling web applications.

As an example for application session stickiness using Elastic cache, please refer to the below link:

- <https://aws.amazon.com/blogs/developer/elasticache-as-an-asp-net-session-store/>
(<https://aws.amazon.com/blogs/developer/elasticache-as-an-asp-net-session-store/>)

Ask our Experts



QUESTION 11 UNATTEMPTED

Which of the following is an example of a good DynamoDB hash key schema for provisioned throughput efficiency? Choose an answer from the options below

- ☒ A. User ID, where the application has many different users. ✓
- ☐ B. Status Code where most status codes are the same
- ☐ C. Device ID, where one is by far more popular than all the others.
- ☐ D. Game Type, where there are three possible game types

Explanation :

Answer – A

Always have the primary key or hash key for those attributes which will have many values. And only option A fits that requirement.

For more information on DynamoDB tables , please visit the URL:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.html>)

Ask our Experts



QUESTION 12 UNATTEMPTED

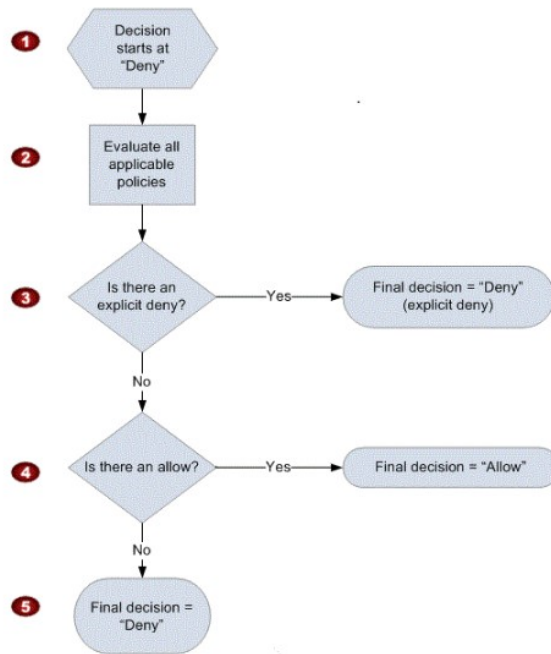
Which of the following are correct statements with policy evaluation logic in AWS Identity and Access Management? Choose 2 answers from the options below

- ☐ A. An explicit deny does not override an explicit allow
- ☐ B. By default, all request are allowed
- ☐ C. An explicit allow overrides default deny. ✓
- ☐ D. An explicit allow overrides an explicit deny
- ☐ E. By default, all requests are denied ✓

Explanation :

Answer – C and E

The below diagram shows the evaluation logic of IAM policies. And as per the evaluation logic, it is clear that the above scenario leads to a default deny.



For more information on the IAM policy evaluation logic, please refer to the link:

- http://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_evaluation-logic.html
(http://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_evaluation-logic.html)

Ask our Experts



QUESTION 13 UNATTEMPTED

Which statements about DynamoDB are true? Choose 2 answers from the options below

- ☐ A. DynamoDB uses optimistic concurrency control ✓
- ☐ B. DynamoDB restricts item access during writes
- ☐ C. DynamoDB uses a pessimistic locking model
- ☐ D. DynamoDB restricts item access during reads
- ☐ E. DynamoDB uses conditional writes for consistency ✓

Explanation :

Answer - A and E

Optimistic locking is a strategy to ensure that the client-side item that you are updating (or deleting) is the same as the item in DynamoDB. If you use this strategy, then your database writes are protected from being overwritten by the writes of others – and vice-versa

To help clients coordinate writes to data items, DynamoDB supports *conditional writes* for PutItem, DeleteItem, and UpdateItem operations. With a conditional write, an operation succeeds only if the item attributes meet one or more expected conditions; otherwise it returns an error

For more information on working with items , please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithItems.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithItems.html>)

Ask our Experts



QUESTION 14 UNATTEMPTED

You are writing to a DynamoDB table and receive the following exception: "ProvisionedThroughputExceededException", though according to your Cloudwatch metrics for the table, you are not exceeding your provisioned throughput. What could be an explanation for this?

- ☐ A. You haven't provisioned enough DynamoDB storage instances
- ☐ B. You're exceeding your capacity on a particular Range Key
- ☒ C. You're exceeding your capacity on a particular Hash Key ✓
- ☐ D. You're exceeding your capacity on a particular Sort Key
- ☐ E. You haven't configured DynamoDB Auto Scaling triggers

Explanation :

Answer – C

"ProvisionedThroughputExceededException" error is caused due to the following reason.

You exceeded your maximum allowed provisioned throughput for a table or for one or more global secondary indexes.

Our question states that according to the cloudwatch metrics we haven't exceeded the throughput of the table. So this error has to do with the throughput exception of global secondary index.

A global secondary index(GSI) consists of Partition key (Hash key- Mandatory) and Range key (Sort key- Optional). So GSI can be a combination of hash key and sort key or hash key itself, Option C seems to be the optimal choice for this question.

Ask our Experts



QUESTION 15 UNATTEMPTED

Company A has an S3 bucket containing premier content that they intend to make available to only paid subscribers of their website. The S3 bucket currently has default permissions of all objects being private to prevent inadvertent exposure of the premier content to non-paying website visitors. How can Company A provide only paid subscribers the ability to download a premier content file in the S3 bucket?

- ☐ A. Apply a bucket policy that grants anonymous users to download the content from the S3 bucket
- ☒ B. Generate a pre-signed object URL for the premier content file when a paid subscriber requests a download ✓
- ☐ C. Add a bucket policy that requires Multi-Factor Authentication for requests to access the S3 bucket objects
- ☐ D. Enable server side encryption on the S3 bucket for data protection against the non-paying website visitors

Explanation :

Answer – B

A pre-signed URL gives you access to the object identified in the URL, provided that the creator of the pre-signed URL has permissions to access that object. That is, if you receive a pre-signed URL to upload an object, you can upload the object only if the creator of the pre-signed URL has the necessary permissions to upload that object.

For more information on Pre-signed urls , please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/PresignedUrlUploadObject.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/PresignedUrlUploadObject.html>)

Ask our Experts



QUESTION 16 UNATTEMPTED

What happens, by default, when one of the resources in a CloudFormation stack cannot be created? Choose an answer from the options below

- ☐ A. The stack creation continues, and the final results indicate which steps failed.
- ☐ B. CloudFormation templates are parsed in advance so stack creation is guaranteed to succeed.
- ☐ C. Previously-created resources are kept but the stack creation terminates.
- ☒ D. Previously-created resources are deleted and the stack creation terminates. ✓

Explanation :

Answer – D

By default, the “automatic rollback on error” feature is enabled. This will cause all AWS resources that AWS CloudFormation created successfully for a stack up to the point where an error occurred to be deleted. This is useful when, for example, you accidentally exceed your default limit of Elastic IP addresses, or you don’t have access to an EC2 AMI you’re trying to run. This feature enables you to rely on the fact that stacks are either fully created, or not at all, which simplifies system administration and layered solutions built on top of AWS CloudFormation.

For more information on cloudformation, please refer to the below link:

- <https://aws.amazon.com/cloudformation/faqs/> (<https://aws.amazon.com/cloudformation/faqs/>)

Ask our Experts



QUESTION 17 UNATTEMPTED

Which of the following are valid arguments for an SNS Publish request? Choose 3 answers from the options below

- ☐ A. Subject ✓
- ☐ B. Language
- ☐ C. Message ✓
- ☐ D. Destination
- ☐ E. TopicArn ✓
- ☐ F. Format

Explanation :

Answer - A,C and E

An example of the notification is shown below , here you can see the different parts of a SNS message.

```
POST / HTTP/1.1
x-amz-sns-message-type: Notification
x-amz-sns-message-id: 22b80b92-fdea-4c2c-8f9d-bdfb0c7bf324
x-amz-sns-topic-arn: arn:aws:sns:us-west-2:123456789012:MyTopic
x-amz-sns-subscription-arn: arn:aws:sns:us-west-2:123456789012:MyTopic:c9135db0-26c4-47e6-8000-000000000000
Content-Length: 773
Content-Type: text/plain; charset=UTF-8
Host: myhost.example.com
Connection: Keep-Alive
User-Agent: Amazon Simple Notification Service Agent

{
  "Type": "Notification",
  "MessageId": "22b80b92-fdea-4c2c-8f9d-bdfb0c7bf324",
  "TopicArn": "arn:aws:sns:us-west-2:123456789012:MyTopic",
  "Subject": "My First Message",
  "Message": "Hello world!",
  "Timestamp": "2012-05-02T00:54:06.655Z",
  "SignatureVersion": "1",
  "Signature": "EXAMPLEw6JRNwmlLFQL4ICB0bnXrdB8C1RMTQFGBqvLpGbM78tJ4etTwC5zU703tS6tGpey...",
  "SigningCertURL": "https://sns.us-west-2.amazonaws.com/SimpleNotificationService-f3ec...",
  "UnsubscribeURL": "https://sns.us-west-2.amazonaws.com/?Action=Unsubscribe&Subscription..."
}
```

For more information on SNS notification format, please refer to the link:

- <http://docs.aws.amazon.com/sns/latest/dg/json-formats.html> (<http://docs.aws.amazon.com/sns/latest/dg/json-formats.html>)

Ask our Experts



QUESTION 18 UNATTEMPTED

A meteorological system monitors 600 temperature gauges, obtaining temperature samples every minute and saving each sample to a DynamoDB table. Each sample involves writing 1 KB of data and the writes are evenly distributed over time. How much write throughput is required for the target table? Choose an answer from the options below.

- ☐ A. 3600 write capacity units
- ☐ B. 1 write capacity unit
- ☒ C. 10 write capacity units ✓
- ☐ D. 60 write capacity units
- ☐ E. 600 write capacity units

Explanation :

Answer – C

Since each item is 1KB and if we near it to the nearest 1KB , we get 1 write capacity. Since we need 600 temperature gauges per minute , it becomes $(600/60) = 10$ per second. This gives is 10 write capacity units.

For more information on DynamoDB provisioned throughput please visit the below link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html>)

Ask our Experts



QUESTION 19 UNATTEMPTED

Which features can be used to restrict access to data in S3? Choose 2 answers from the options below

- ☒ A. Set an S3 Bucket policy. ✓
- ☐ B. Enable IAM Identity Federation.
- ☒ C. Set an S3 ACL on the bucket or the object. ✓
- ☐ D. Use Server-Side encryption to encrypt the data stored in S3 bucket.

Explanation :

Answer – A and C

By default, all Amazon S3 resources—buckets, objects, and related subresources (for example, lifecycleconfiguration and website configuration)—are private: only the resource owner, an AWS account that created it, can access the resource. The resource owner can optionally grant access permissions to others by writing an access policy.

Amazon S3 offers access policy options broadly categorized as resource-based policies and user policies. Access policies you attach to your resources (buckets and objects) are referred to as resource-based policies. For example, bucket policies and access control lists (ACLs) are resource-based policies. You can also attach access policies to users in your account. These are called user policies. You may choose to use resource-based policies, user policies, or some combination of these to manage permissions to your Amazon S3 resources.

For more information on S3 access permissions please visit the below link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/s3-access-control.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/s3-access-control.html>)

Ask our Experts



QUESTION 20 UNATTEMPTED

What AWS products and features can be deployed by Elastic Beanstalk? Choose 3 answers from the options below

- ☒ A. Auto scaling groups ✓
- ☐ B. Route 53 hosted zones
- ☒ C. Elastic Load Balancers ✓
- ☒ D. RDS Instances ✓
- ☐ E. Elastic IP addresses
- ☐ F. SQS Queues

Explanation :

Answer - A , C and D

This is clearly given in the AWS documentation

Q: What are the Cloud resources powering my AWS Elastic Beanstalk application?

AWS Elastic Beanstalk uses proven AWS features and services, such as Amazon EC2, Amazon RDS, Elastic Load Balancing, Auto Scaling, Amazon S3, and Amazon SNS, to create an environment that runs your application. The current version of AWS Elastic Beanstalk uses the Amazon Linux AMI or the Windows Server 2012 R2 AMI.

For more information on elastic beanstalk, please refer to the below URL:

- <https://aws.amazon.com/elasticbeanstalk/faqs/> (<https://aws.amazon.com/elasticbeanstalk/faqs/>)
- <http://docs.aws.amazon.com/elasticbeanstalk/latest/dg/Welcome.html>
(<http://docs.aws.amazon.com/elasticbeanstalk/latest/dg/Welcome.html>)

Ask our Experts



QUESTION 21 UNATTEMPTED

You run an ad-supported photo sharing website using S3 to serve photos to visitors of your site. At some point you find out that other sites have been linking to the photos on your site, causing loss to your business. What is an effective method to mitigate this? Choose an answer from the options below

- ☐ A. Store photos on an EBS volume of the web server
- ☒ B. Remove public read access and use signed URLs with expiry dates. ✓
- ☐ C. Use CloudFront distributions for static content.
- ☐ D. Block the IPs of the offending websites in Security Groups.

Explanation :

Answer - B

A pre-signed URL gives you access to the object identified in the URL, provided that the creator of the pre-signed URL has permissions to access that object. That is, if you receive a pre-signed URL to upload an object, you can upload the object only if the creator of the pre-signed URL has the necessary permissions to upload that object.

For more information on Pre-signed urls , please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/PresignedUrlUploadObject.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/PresignedUrlUploadObject.html>)

Ask our Experts



QUESTION 22 UNATTEMPTED

Games-R-Us is launching a new game app for mobile devices. Users will log into the game using their existing Facebook account and the game will record player data and scoring information directly to a DynamoDB table. What is the most secure approach for signing requests to the DynamoDB API? Choose an answer from the options below

- ☐ A. Create an IAM user with access credentials that are distributed with the mobile app to sign the requests
- ☐ B. Distribute the AWS root account access credentials with the mobile app to sign the requests
- ☒ C. Request temporary security credentials using web identity federation to sign the requests ✓
- ☐ D. Establish cross account access between the mobile app and the DynamoDB table to sign the requests

Explanation :

Answer – C

With web identity federation, you don't need to create custom sign-in code or manage your own user identities. Instead, users of your app can sign in using a well-known identity provider (IdP) –such as Login with Amazon, Facebook, Google, or any other OpenID Connect (OIDC) (<http://openid.net/connect/>)-compatible IdP, receive an authentication token, and then exchange that token for temporary security credentials in AWS that map to an IAM role with permissions to use the resources in your AWS account. Using an IdP helps you keep your AWS account secure, because you don't have to embed and distribute long-term security credentials with your application. For more information on web identity federation , please visit the link:

- http://docs.aws.amazon.com/IAM/latest/UserGuide/fid_roles_providers_oidc.html
(http://docs.aws.amazon.com/IAM/latest/UserGuide/fid_roles_providers_oidc.html)

Ask our Experts



QUESTION 23 UNATTEMPTED

How can you secure data at rest for an existing unencrypted EBS volume? Choose an answer from the options below

- ☐ A. Write the data randomly instead of sequentially.
- ☒ B. Use an encrypted file system on top of the EBS volume. ✓
- ☐ C. Encrypt the volume using the S3 server-side encryption service.
- ☐ D. Create an IAM policy that restricts read and write access to the volume.
- ☐ E. Attach the volume to an instance using EC2's SSL interface.

Explanation :

Answer – B

EBS encryption can be enabled during the creation of an EBS volume. For an existing EBS volume , you need to use the encryption available at the Operating system level URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSEncryption.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSEncryption.html>)

Ask our Experts



QUESTION 24 UNATTEMPTED

You are inserting 1000 new items every second in a DynamoDB table. Once an hour these items are analyzed and then are no longer needed. You need to minimize provisioned throughput, storage, and API calls. Given these requirements, what is the most efficient way to manage these Items after the analysis? Choose an answer from the options below

- ☐ A. Retain the items in a single table
- ☐ B. Delete items individually over a 24 hour period
- ☒ C. Delete the table and create a new table per hour ✓
- ☐ D. Create a new table per hour

Explanation :

Answer – C

The main point of the question is to minimize provisioned throughput.

Option A is wrong because the table will grow and increase provisioned throughput.

Option B is wrong because the table will grow and increase provisioned throughput over the 24 hour period.

Option D is partially correct , you still need to delete the older table.

For more information on DynamoDB provisioned throughput please visit the below link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html>)

Note:

In AWS documentation, it clearly mentioned the following steps as

The following design pattern often handles this kind of scenario effectively:

- Create one table per time period, provisioned with write capacity less than 1,000 write capacity units (WCUs) per partition-key value, and minimum necessary read capacity.
- Before the end of each time period, **rebuild the table for the next period. Just as the current period ends, direct event traffic to the new table. You can assign names to these tables that specify the time periods that they have recorded.**
- As soon as a table is no longer being written to, reduce its provisioned write capacity to 1 WCU and provision whatever read capacity is appropriate. **Reduce the provisioned read capacity of earlier tables as they age, and archive or delete the ones whose contents will rarely or never be needed.**

Please refer the following link to get more information.

- <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/bp-time-series.html#bp-time-series-examples>
(<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/bp-time-series.html#bp-time-series-examples>)

Ask our Experts



QUESTION 25 UNATTEMPTED

Which of the following statements about SWF are true? Choose 3 answers from the options below

- ☒ A. SWF uses deciders and workers to complete tasks ✓
- ☐ B. SWF requires at least 1 EC2 instance per domain
- ☐ C. SWF triggers SNS notifications on task assignment
- ☐ D. SWF requires an S3 bucket for workflow storage
- ☒ E. SWF tasks are assigned once and never duplicated ✓
- ☒ F. SWF workflow executions can last up to a year ✓

Explanation :

Answer - A, E and F

Option A is correct because of the below excerpt from the AWS documentation

In Amazon SWF, tasks represent invocations of logical steps in applications. Tasks are processed by workers which are programs that interact with Amazon SWF to get tasks, process them, and return their results. A worker implements an application processing step. You can build workers in different programming languages and even reuse existing components to quickly create the worker. For example, you can use cloud services, enterprise applications, legacy systems, and even simple scripts to implement workers. By independently controlling the number of workers for processing each type of task, you can control the throughput of your application efficiently.

To coordinate the application execution across workers, you write a program called the decider in your choice of programming language. Option E is correct because of the below excerpt from the AWS documentation

- Communicating and managing the flow of work between your application components. With Amazon SWF, you do not need to design a messaging protocol or worry about lost and duplicated tasks.

Option F is correct because of the below excerpt from the AWS documentation

Q: How long can workflow executions run?

Each workflow execution can run for a maximum of 1 year. Each workflow execution history can grow up to 25,000 events. If your use case requires you to go beyond these limits, you can use features Amazon SWF provides to continue executions and structure your applications using child workflow executions.

For more information on SWF please visit the below link:

- <https://aws.amazon.com/swf/faqs/> (<https://aws.amazon.com/swf/faqs/>)

Ask our Experts



QUESTION 26 UNATTEMPTED

Company D is running their corporate website on Amazon S3 accessed from <http://www.companyd.com>. Their marketing team has published new web fonts to a separate S3 bucket accessed by the S3 endpoint <https://s3-us-west1.amazonaws.com/cdfonts>. While testing the new web fonts, Company D recognized the web fonts are being blocked by the browser. What should Company D do to prevent the web fonts from being blocked by the browser?

- ☐ A. Create a policy on the cdfonts bucket to enable access to everyone
- ☐ B. Add the Content-MD5 header to the request for webfonts in the cdfonts bucket from the website
- ☐ C. Configure the cdfonts bucket to allow cross-origin requests by creating a CORS configuration ✓
- ☐ D. Enable versioning on the cdfonts bucket for each web font

Explanation :

Answer – C

Cross-origin resource sharing (CORS) defines a way for client web applications that are loaded in one domain to interact with resources in a different domain. With CORS support in Amazon S3, you can build rich client-side web applications with Amazon S3 and selectively allow cross-origin access to your Amazon S3 resources.

For more information on S3 CORS configuration, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/cors.html> (<http://docs.aws.amazon.com/AmazonS3/latest/dev/cors.html>)

Ask our Experts



QUESTION 27 UNATTEMPTED

How is provisioned throughput affected by the chosen consistency model when reading data from a DynamoDB table? Choose an answer from the options below

- ☐ A. Strongly consistent reads uses the same amount of provisioned read throughput as eventually consistent reads
- ☐ B. Strongly consistent reads uses variable provisioned read throughput depending on read activity
- ☐ C. Strongly consistent reads uses more provisioned read throughput than eventually consistent reads. ✓
- ☐ D. Strongly consistent reads uses less provisioned read throughput than eventually consistent reads

Explanation :

Answer – C

A unit of Write Capacity enables you to perform one write per second for items of up to 1KB in size. Similarly, a unit of Read Capacity enables you to perform one strongly consistent read per second (or two eventually consistent reads per second) of items of up to 4KB in size.

Hence the strong consistency is more expensive and uses more throughput than the eventual consistency.

For more information on DynamoDB, please visit the below link:

- <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html#HowItWorks.ProvisionedThroughput> (<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html#HowItWorks.ProvisionedThroughput>)

Ask our Experts



QUESTION 28 UNATTEMPTED

Which of the following are valid SNS delivery transports? Choose 2 answers from the options below

- ☒ A. SMS ✓
- ☐ B. DynamoDB
- ☒ C. HTTP ✓
- ☐ D. UDP
- ☐ E. Named Pipes

Explanation :

Answer - A and C

In SNS the following subscription request are possible

- "HTTP", "HTTPS" – Subscribers specify a URL as part of the subscription registration; notifications will be delivered through an HTTP POST to the specified URL.
- "Email", "Email-JSON" – Messages are sent to registered addresses as email. Email-JSON sends notifications as a JSON object, while Email sends text-based email.
- "SQS" – Users can specify an SQS standard queue as the endpoint; Amazon SNS will enqueue a notification message to the specified queue (which subscribers can then process using SQS APIs such as ReceiveMessage, DeleteMessage, etc.). Note that FIFO queues are not currently supported.
- "SMS" – Messages are sent to registered phone numbers as SMS text messages.

For more information on SNS, please visit the below link:

- <https://aws.amazon.com/sns/faqs/> (<https://aws.amazon.com/sns/faqs/>)

Ask our Experts



QUESTION 29 UNATTEMPTED

You have an environment that consists of a public subnet using Amazon VPC and 3 instances that are running in this subnet. These three instances can successfully communicate with other hosts on the Internet. You launch a fourth instance in the same subnet, using the same AMI and security group configuration you used for the others, but find that this instance cannot be accessed from the Internet. What should you do to enable internet access? Choose an answer from the options below

- ☐ A. Deploy a NAT instance into the public subnet.
- ☐ B. Modify the routing table for the public subnet
- ☐ C. Configure a publically routable IP Address In the host OS of the fourth instance.
- ☒ D. Assign an Elastic IP address to the fourth instance. ✓

Explanation :

Answer – D

An *Elastic IP address* is a static IPv4 address designed for dynamic cloud computing. An Elastic IP address is associated with your AWS account. With an Elastic IP address, you can mask the failure of an instance or software by rapidly remapping the address to another instance in your account.

An Elastic IP address is a public IPv4 address, which is reachable from the Internet. If your instance does not have a public IPv4 address, you can associate an Elastic IP address with your instance to enable communication with the Internet; for example, to connect to your instance from your local computer.

For more information on Elastic IP, please visit the below link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/elastic-ip-addresses-eip.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/elastic-ip-addresses-eip.html>)

Ask our Experts



QUESTION 30 UNATTEMPTED

Which of the following services are included at no additional cost with the use of the AWS platform? Choose 2 answers from the options below

- ☐ A. CloudFormation ✓
- ☐ B. Simple Workflow Service
- ☐ C. Elastic Load Balancing
- ☐ D. Elastic Compute Cloud
- ☐ E. Simple Storage Service
- ☐ F. Auto Scaling ✓

Explanation :

Answer - A and F

This is clearly given in the AWS documentation URL:

- <https://aws.amazon.com/cloudformation/pricing/> (<https://aws.amazon.com/cloudformation/pricing/>)

AWS CloudFormation Pricing

There is no additional charge for AWS CloudFormation. You pay for AWS resources (such as Amazon EC2 instances, Elastic Load Balancing load balancers, etc.) created using AWS CloudFormation in the same manner as if you created them manually. You only pay for what you use, as you use it; there are no minimum fees and no required upfront commitments.

URL:

- <https://aws.amazon.com/autoscaling/pricing/> (<https://aws.amazon.com/autoscaling/pricing/>)

Auto Scaling Pricing

Auto Scaling is enabled by [Amazon CloudWatch](#) and carries no additional fees. [Amazon EC2](#) and [Amazon CloudWatch](#) service fees apply and are billed separately. Partial hours are billed as full hours.

Ask our Experts



QUESTION 31 UNATTEMPTED

Which of the following programming languages have an officially supported AWS SDK? Choose 2 answers from

the options below

- ☐ A. PHP ✓
- ☐ B. Pascal
- ☐ C. Java ✓
- ☐ D. SQL
- ☐ E. Perl

Explanation :

Answer - A and C

This is as per the AWS documentation

Java

[Install »](#)
[Documentation »](#)
[Learn more »](#)

.NET

[Install »](#)
[Documentation »](#)
[Learn more »](#)

Node.js

[Install »](#)
[Documentation »](#)
[Learn more »](#)

PHP

[Install »](#)
[Documentation »](#)
[Learn more »](#)

Python

[Install »](#)
[Documentation »](#)
[Learn more »](#)

Ruby

[Install »](#)
[Documentation »](#)
[Learn more »](#)

Go

[Install »](#)
[Documentation »](#)

C++

[Install »](#)
[Documentation »](#)

AWS Mobile SDK

[Install »](#)
[Documentation »](#)

For more information on AWS toolkits available, please refer to the below URL:

- <https://aws.amazon.com/tools/> (<https://aws.amazon.com/tools/>)

Ask our Experts



QUESTION 32 UNATTEMPTED

After launching an instance that you intend to serve as a NAT (Network Address Translation) device in a public subnet you modify your route tables to have the NAT device be the target of internet bound traffic of your private subnet. When you try and make an outbound connection to the Internet from an instance in the private subnet, you are not successful. Which of the following steps could resolve the issue? Choose an answer from the options below

- ☐ A. Attaching a second Elastic Network interface (ENI) to the NAT instance, and placing it in the private subnet
- ☐ B. Attaching an Elastic IP address to the instance in the private subnet
- ☐ C. Attaching a second Elastic Network Interface (ENI) to the instance in the private subnet, and placing it in the public subnet
- ☒ D. Disabling the Source/Destination Check attribute on the NAT instance ✓

Explanation :

Answer – D

You can use a network address translation (NAT) instance in a public subnet in your VPC to enable instances in the private subnet to initiate outbound IPv4 traffic to the Internet or other AWS services, but prevent the instances from receiving inbound traffic initiated by someone on the Internet.

Each EC2 instance performs source/destination checks by default. This means that the instance must be the source or destination of any traffic it sends or receives. However, a NAT instance must be able to send and receive traffic when the source or destination is not itself. Therefore, you must disable source/destination checks on the NAT instance.

For more information on NAT instance, please refer to the below URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_NAT_Instance.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_NAT_Instance.html)

Ask our Experts



QUESTION 33 UNATTEMPTED

Which of the following is chosen as the default region when making an API call with an AWS SDK? Choose an answer from the options below

- ☐ A. ap-northeast-1
- ☐ B. us-west-2
- ☒ C. us-east-1 ✓
- ☐ D. eu-west-1
- ☐ E. us-central-1

Explanation :

Answer – C

The default region for API calls is us-east-1

For more information on regions and AZ's, please refer to the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>)

Ask our Experts



QUESTION 34 UNATTEMPTED

REFACTORING

If an application is storing hourly log files from thousands of instances from a high traffic web site, which naming scheme would give optimal performance on S3? Choose an answer from the options below

- ☐ A. Sequential
- ☒ B. HH-DD-MM-YYYY-log_instanceID ✓
- ☐ C. YYYY-MM-DD-HH-log_instanceID
- ☐ D. instanceID_log-HH-DD-MM-YYYY
- ☐ E. instanceID_log-YYYY-MM-DD-HH

Explanation :

Answer – B

Amazon S3 maintains an index of object key names in each AWS region. Object keys are stored in UTF-8 binary ordering across multiple partitions in the index. The key name dictates which partition the key is stored in. So if it is in an alphabetical sequence it can overwhelm the I/O capacity of the partition.

So introducing some randomness at the beginning of the file name in the naming scheme would give optimal performance.

In this scenario, thousands of instances are generating log files and are storing these log files on to S3 in an hourly manner. All the instance ID-s start with i-O*****. So if we go with this naming convention, thousands of log files are generated in each hour with these characters in the beginning of all the log files which doesn't bring in the randomness required for the file names at the beginning. Hence options D & E are incorrect.

Option A is incorrect as it is mentioning to be Sequential.

Option C is also incorrect since the Year value will take an year to change.

So out of the options given in the question option B seems to be ideal in this scenario.

For more information on request rate considerations please visit the below link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/request-rate-perf-considerations.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/request-rate-perf-considerations.html>)

Note:

Assume that we have some 10,000 instances running and we are taking log files from each instance on an hourly manner. If we go with option B in an hour your 10,000 logs are going to be using the same partition. But in the next hour it will be using another partition to store the 10,000 log files.

But if we go with option C, which used the year, then each hour, these 10,000 files for the entire day, is going to be stored in the same partition. Again for the entire month and for the entire 365 days of the year it will be using the same partition. So based on the options provided option B is to be the most optimal choice.

Ask our Experts



QUESTION 35 UNATTEMPTED

When using a large Scan operation in DynamoDB, what technique can be used to minimize the impact of a scan on a table's provisioned throughput? Choose an answer from the options below

- ☒ A. Set a smaller page size for the scan ✓
- ☐ B. Pre-warm the table by updating all items
- ☐ C. Use parallel scans
- ☐ D. Define a range index on the table

Explanation :

Answer – A

This is clearly given in the ASW documentation

- **Reduce Page Size**

Because a Scan operation reads an entire page (by default, 1 MB), you can reduce the impact of the scan operation by setting a smaller page size. The `scan` operation provides a `Limit` parameter that you can use to set the page size for your request. Each `scan` or `query` request that has a smaller page size uses fewer read operations and creates a "pause" between each request. For example, if each item is 4 KB and you set the page size to 40 items, then a `query` request would consume only 40 strongly consistent read operations or 20 eventually consistent read operations. A larger number of smaller `scan` or `query` operations would allow your other critical requests to succeed without throttling.

- **Isolate Scan Operations**

DynamoDB is designed for easy scalability. As a result, an application can create tables for distinct purposes, possibly even duplicating content across several tables. You want to perform scans on a table that is not taking "mission-critical" traffic. Some applications handle this load by rotating traffic hourly between two tables – one for critical traffic, and one for bookkeeping. Other applications can do this by performing every write on two tables: a "mission-critical" table, and a "shadow" table.

For more information on guidelines on queries and scans in dynamoDB, please refer to the below URL:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/QueryAndScanGuidelines.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/QueryAndScanGuidelines.html>)

Ask our Experts



QUESTION 36 UNATTEMPTED

Which of the following services are key/value stores? Choose 3 answers from the options below

- ☒ A. Amazon ElastiCache ✓
- ☐ B. Simple Notification Service
- ☒ C. DynamoDB ✓

- ☐ D. Simple Workflow Service
- ☐ E. Simple Storage Service ✓

Explanation :

This is provided in the AWS documentation
AWS Elastic cache URL:

- <https://aws.amazon.com/elasticache/faqs/> (<https://aws.amazon.com/elasticache/faqs/>)

Amazon ElastiCache automates common administrative tasks required to operate a distributed in-memory key-value environment. Using Amazon ElastiCache, you can add a caching or in-memory layer to your application architecture in a matter of minutes via a few clicks of the AWS Management Console. Once a cluster is provisioned, Amazon ElastiCache automatically detects and replaces failed nodes, providing a resilient system that mitigates the risk of overloaded databases, which slow website and application load times. Through integration with Amazon CloudWatch monitoring, Amazon ElastiCache provides enhanced visibility into key performance metrics associated with your nodes. Amazon ElastiCache is protocol-compliant with Memcached and Redis, so code, applications, and popular tools that you use today with your existing Memcached or Redis environments will work seamlessly with the service. With the support for clustered configuration in Amazon ElastiCache, you get the benefits of fast, scalable and easy to use managed service that can meet the needs of your most demanding applications. As with all Amazon Web Services, there are no up-front investments required, and you pay only for the resources you use.

AWS DynamoDB

In DynamoDB, an *item* is a collection of attributes. Each attribute has a name and a value. An attribute value can be a scalar, a set, or a document type. For more information, see [Data Types](#).

URL:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithItems.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithItems.html>)

AWS S3

Amazon S3 is a simple key, value store designed to store as many objects as you want. You store these objects in one or more buckets. An object consists of the following:

URL:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingObjects.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingObjects.html>)

Ask our Experts



QUESTION 37 **UNATTEMPTED**

An application stores payroll information nightly in DynamoDB for a large number of employees across hundreds of offices. Item attributes consist of individual name, office identifier, and cumulative daily hours. Managers run reports for ranges of names working in their office. One query is, "Return all Items in this office for names starting with A through E". Which table configuration will result in the lowest impact on provisioned throughput for this query? Choose an answer from the options below

- ☐ A. Configure the table to have a range index on the name attribute, and a hash index on the office identifier ✓
- ☐ B. Configure a hash index on the name attribute and no range index
- ☐ C. Configure the table to have a hash index on the name attribute, and a range index on the office identifier
- ☐ D. Configure a hash index on the office Identifier attribute and no range index

Explanation :

Correct answer is A.

Key points of this question to consider are 'large number of employees across hundreds of offices' and 'reports for ranges of names working in their office.'

Answer B is ruled out because it assumes employee names are unique which might not be the case of 'large number of employees across

hundreds of offices'.

Answer D is obviously incorrect as it doesn't allow an office to have more than one employee. (hash key only model requires unique hashes.)

Answer C makes employees of a single office to be spread over multiple partitions and requires a SCAN operation to generate reports for ranges of names working in a single office. This table configuration requires SCAN operations for reports for ranges of names working in an office. A SCAN operation is not recommended on large tables, because it can consume large amounts of system resources. (see <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SQLtoNoSQL.ReadData.html> (<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SQLtoNoSQL.ReadData.html>))

Answer A puts all employees of a single office in to a single hash. With configuring the table to have a range index on the name attribute, and a hash index on the office identifier, a QUERY operation is sufficient to generate reports like 'Return all Items in this office for names starting with A through E'. This table configuration results in the lowest impact on provisioned throughput. This is the optimal arrangement to accommodate regular 'reports for ranges of names working in an office.'

Reference:

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.html>

(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.html>)

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html>

(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html>)

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SQLtoNoSQL.Rea>

(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SQLtoNoSQL.ReadData.html>)

Ask our Experts



QUESTION 38 UNATTEMPTED

What item operation allows the retrieval of multiple items from a DynamoDB table in a single API call? Choose an answer from the options below

- ☐ A. GetItem
- ☒ B. BatchGetItem ✓
- ☐ C. GetMultipleItems
- ☐ D. GetItemRange

Explanation :

Answer – B

The BatchGetItem operation returns the attributes of one or more items from one or more tables. You identify requested items by primary key.

A single operation can retrieve up to 16 MB of data, which can contain as many as 100 items. BatchGetItem will return a partial result if the response size limit is exceeded, the table's provisioned throughput is exceeded, or an internal processing failure occurs. If a partial result is returned, the operation returns a value for UnprocessedKeys. You can use this value to retry the operation starting with the next item to get.

For More information:

- https://docs.aws.amazon.com/amazondynamodb/latest/APIReference/API_BatchGetItem.html (https://docs.aws.amazon.com/amazondynamodb/latest/APIReference/API_BatchGetItem.html)

Ask our Experts



QUESTION 39 UNATTEMPTED

In DynamoDB, what type of HTTP response codes indicate that a problem was found with the client request sent to the service? Choose an answer from the options below

- ☐ A. 5xx HTTP response code
- ☐ B. 200 HTTP response code
- ☐ C. 306 HTTP response code
- ☐ D. 4xx HTTP response code ✓

Explanation :

Answer – D

This is clearly given in the AWS documentation

HTTP Status Code 400

An HTTP 400 status code indicates a problem with your request, such as authentication failure, missing required parameters, or exceeding a table's provisioned throughput. You will have to fix the issue in your application before submitting the request again.

For more information on DynamoDB errors, please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Programming.Errors.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Programming.Errors.html>)

Ask our Experts



QUESTION 40 UNATTEMPTED

If a message is retrieved from a queue in Amazon SQS, how long is the message inaccessible to other users by default? Choose an answer from the options below

- ☐ A. 30 seconds ✓
- ☐ B. 0 seconds
- ☐ C. 1 hour
- ☐ D. 1 day
- ☐ E. forever

Explanation :

Answer – A

Each queue starts with a default setting of 30 seconds for the visibility timeout. You can change that setting for the entire queue.

Typically, you'll set the visibility timeout to the average time it takes to process and delete a message from the queue. When receiving messages, you can also set a special visibility timeout for the returned messages without changing the overall queue timeout.

For more information on SQS please visit the below link:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 41 UNATTEMPTED

Is DynamoDB suited only for high scale based applications? Choose an answer from the options below

- ☐ A. True
- ☐ B. False ✓

Explanation :

Answer – B

This is clearly given in the AWS documentation

Q: Is DynamoDB only for high-scale applications?

No. DynamoDB offers seamless scaling so you can start small and scale up and down in line with your requirements. If you need fast, predictable performance at any scale then DynamoDB may be the right choice for you.

For more information on DynamoDB, please visit the following URL:

- <https://aws.amazon.com/dynamodb/faqs/> (<https://aws.amazon.com/dynamodb/faqs/>)

Ask our Experts



QUESTION 42 UNATTEMPTED

Can conditional operations be used in a DynamoDB query? Choose an answer from the options below

- ☒ **A. True** ✓
- ☐ **B. False**

Explanation :

Answer – A

This is clearly given in the AWS documentation

Yes, you can specify a condition that must be satisfied for a put, update, or delete operation to be completed on an item. To perform a conditional operation, you can define a ConditionExpression that is constructed from the following:

- Boolean functions: ATTRIBUTE_EXISTS, CONTAINS, and BEGINS_WITH
- Comparison operators: =, <>, <, >, <=, >=, BETWEEN, and IN
- Logical operators: NOT, AND, and OR.

For more information on DynamoDB, please visit the following URL:

- <https://aws.amazon.com/dynamodb/faqs/> (<https://aws.amazon.com/dynamodb/faqs/>)

Ask our Experts



QUESTION 43 UNATTEMPTED

Which of the below logical operators are allowed in a DynamoDB query? Choose 3 answers from the options below

- ☒ **A. NOT** ✓
- ☐ **B. XOR**
- ☒ **C. OR** ✓
- ☒ **D. AND** ✓

Explanation :

Answer – A,C and D

This is clearly given in the AWS documentation

Yes, you can specify a condition that must be satisfied for a put, update, or delete operation to be completed on an item. To perform a conditional operation, you can define a ConditionExpression that is constructed from the following:

- Boolean functions: ATTRIBUTE_EXISTS, CONTAINS, and BEGINS_WITH
- Comparison operators: =, <>, <, >, <=, >=, BETWEEN, and IN
- Logical operators: NOT, AND, and OR.

For more information on DynamoDB, please visit the following URL:

- <https://aws.amazon.com/dynamodb/faqs/> (<https://aws.amazon.com/dynamodb/faqs/>)

Ask our Experts



QUESTION 44 UNATTEMPTED

Which of the below APIs are available for tables in DynamoDB?

Choose 3 answers from the options below

- ☐ A. CreateTable ✓
- ☐ B. UpdateAllTables
- ☐ C. UpdateTable ✓
- ☐ D. ListTables ✓

Explanation :

Answer – A,C and D

This is clearly given in the aws documentation

- CreateTable – Creates a table and specifies the primary index used for data access.
- UpdateTable – Updates the provisioned throughput values for the given table.
- DeleteTable – Deletes a table.
- DescribeTable – Returns table size, status, and index information.
- ListTables – Returns a list of all tables associated with the current account and endpoint.

For more information on DynamoDB, please visit the following URL:

- <https://aws.amazon.com/dynamodb/faqs/> (<https://aws.amazon.com/dynamodb/faqs/>)

Ask our Experts



QUESTION 45 UNATTEMPTED

Your support teams calls you and tells you that the application you developed is giving the error “Your proposed upload exceeds the maximum allowed object size.” error message when uploading an object to S3. What is a possible solution for this? Choose an answer from the options below

- ☐ A. None, Simple Storage Service objects are limited to 5 GB
- ☐ B. Use the multi-part upload API for this object ✓
- ☐ C. Use the large object upload API for this object
- ☐ D. Contact support to increase your object size limit
- ☐ E. Upload to a different region

Explanation :

Answer – B

The Multipart upload API enables you to upload large objects in parts. You can use this API to upload new large objects or make a copy of an existing object (see Operations on Objects (<http://docs.aws.amazon.com/AmazonS3/latest/dev/ObjectOperations.html>)). Multipart uploading is a three-step process: You initiate the upload, you upload the object parts, and after you have uploaded all the parts, you complete the multipart upload. Upon receiving the complete multipart upload request, Amazon S3 constructs the object from the uploaded parts, and you can then access the object just as you would any other object in your bucket.

For more information on S3 Multi Part file upload, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/mpuoverview.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/mpuoverview.html>)

Ask our Experts

**QUESTION 46 UNATTEMPTED**

What are the new type of SQS queues available in some regions? Choose an answer from the options below

- ☐ A. LIFO
- ☐ B. Dynamic
- ☒ C. FIFO ✓
- ☐ D. Static

Explanation :

Answer – C

For more information on SQS , please visit the following URL:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts

**QUESTION 47 UNATTEMPTED**

What configuration does AWS provide to handle unsuccessfully-processed messages in SQS? Choose an answer from the options below

- ☐ A. Retention queues
- ☐ B. Delivery queues
- ☒ C. Dead letter queues ✓
- ☐ D. There is no such option available

Explanation :

Answer – C

This is clearly given in the AWS documentation

Q: How does Amazon SQS handle unsuccessfully-processed messages?

In Amazon SQS, you can use the API or the console to configure *dead letter queues*, which are queues that receive messages from other source queues.

If you make a queue into a dead letter queue, it receives messages after a maximum number of processing attempts cannot be completed. You can use dead letter queues to isolate messages that can't be processed for later analysis.

For more information on SQS , please visit the following URL:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 48 UNATTEMPTED

In SQS does long polling cost more than short polling? Choose an answer from the options below

- ☐ A. True
- ☐ B. False ✓

Explanation :

Answer – B

This is clearly given in the AWS documentation

Q: Is there an additional charge for using Amazon SQS long polling?

No. Long-polling ReceiveMessage calls are billed exactly the same as short-polling ReceiveMessage calls.

For more information on SQS , please visit the following URL:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 49 UNATTEMPTED

What are the 2 actions required to fanout identical messages to multiple SQS queues. Choose 2 answers from the options below

- ☐ A. Use Amazon SNS to create a topic. ✓
- ☐ B. Create and subscribe multiple Amazon SQS standard queues to the Amazon SNS topic. ✓
- ☐ C. Create and subscribe single Amazon SQS standard queues to the Amazon SNS topic.
- ☐ D. Use Cloudformation templates

Explanation :

Answer – A and B

This is clearly given in the AWS documentation

Q: How can I fan out identical messages to multiple Amazon SQS queues?

1. Use Amazon SNS to create a topic.
2. Create and subscribe multiple Amazon SQS standard queues to the Amazon SNS topic.
3. Whenever a message is sent to the Amazon SNS topic, it is fanned out to the Amazon SQS message queues.

Amazon SNS delivers the message to all Amazon SQS message queues subscribed to the topic.

For more information on SQS , please visit the following URL:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 50 UNATTEMPTED

What are the protocol versions supported by SQS. Choose 3 answers from the options below

- ☐ A. TLS1.0 ✓
- ☐ B. TLS1.1 ✓
- ☐ C. TLS1.2 ✓
- ☐ D. TLS1.3

Explanation :

Answer – A,B and C

This is clearly given in the AWS documentation

Amazon SQS supports the HTTP over SSL (HTTPS) and Transport Layer Security (TLS) protocols. Most clients can automatically negotiate to use newer versions of TLS without any code or configuration change. Amazon SQS supports versions 1.0, 1.1, and 1.2 of the Transport Layer Security (TLS) protocol in all regions.

For more information on SQS , please visit the following URL:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 51 UNATTEMPTED

How long can messages live in an SQS queue. Choose the correct answer from the options below

- ☐ A. 12 hours
- ☐ B. 10 days
- ☐ C. 14 days ✓
- ☐ D. 1 year

Explanation :

Answer – C

This is clearly given in the AWS documentation

Q: How long can I keep my messages in Amazon SQS message queues?

Longer message retention provides greater flexibility to allow for longer intervals between message production and consumption.

You can configure the Amazon SQS message retention period to a value from 1 minute to 14 days. The default is 4 days. Once the message retention limit is reached, your messages are automatically deleted.

For more information on SQS , please visit the following URL:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 52 UNATTEMPTED

What is the maximum length of a topic name in SNS? Choose the correct answer from the options below

- ☐ A. 128 characters
- ☐ B. 256 characters ✓
- ☐ C. 512 characters

☐ D. There is no limit

Explanation :

Answer – B

This is clearly given in the AWS documentation

Topic names are limited to 256 characters. Alphanumeric characters plus hyphens (-) and underscores (_) are allowed. Topic names must be unique within an AWS account. After you delete a topic, you can reuse the topic name. When a topic is created, Amazon SNS will assign a unique ARN (Amazon Resource Name) to the topic, which will include the service name (SNS), region, AWS ID of the user and the topic name. The ARN will be returned as part of the API call to create the topic. Whenever a publisher or subscriber needs to perform any action on the topic, they should reference the unique topic ARN.

For more information on SNS , please visit the following URL:

- <https://aws.amazon.com/sns/faqs/> (<https://aws.amazon.com/sns/faqs/>)

Ask our Experts



QUESTION 53 UNATTEMPTED

What are some of the owner operations possible in SNS?

Choose 3 correct answers from the options below

- ☐ A. CreateTopic ✓
- ☐ B. DeleteTopic ✓
- ☐ C. ListSubscriptions
- ☐ D. AddPermission ✓

Explanation :

Answer – A, B and D

This is clearly given in the AWS documentation

Amazon SNS provides a set of simple APIs to enable event notifications for topic owners, subscribers and publishers.

Owner operations:

CreateTopic – Create a new topic.

DeleteTopic – Delete a previously created topic.

ListTopics – List of topics owned by a particular user (AWS ID).

ListSubscriptionsByTopic – List of subscriptions for a particular topic

SetTopicAttributes – Set/modify topic attributes, including setting and modifying publisher/subscriber permissions, transports supported, etc.

GetTopicAttributes – Get/view existing attributes of a topic

AddPermission – Grant access to selected users for the specified actions

RemovePermission – Remove permissions for selected users for the specified actions

Topic names are limited to 256 characters. Alphanumeric characters plus hyphens (-) and underscores (_) are allowed. Topic names must be unique within an AWS account. After you delete a topic, you can reuse the topic name. When a topic is created, Amazon SNS will assign a unique ARN (Amazon Resource Name) to the topic, which will include the service name (SNS), region, AWS ID of the user and the topic name. The ARN will be returned as part of the API call to create the topic. Whenever a publisher or subscriber needs to perform any action on the topic, they should reference the unique topic ARN.

For more information on SNS , please visit the following URL:

- <https://aws.amazon.com/sns/faqs/> (<https://aws.amazon.com/sns/faqs/>)

Ask our Experts



QUESTION 54 UNATTEMPTED

Which of the following SNS information item has the following information - "A Universally Unique Identifier, unique for each notification published". Choose the correct answer from the options below

- ☐ A. TopicARN
- ☒ B. Messageid ✓
- ☐ C. Type
- ☐ D. Message

Explanation :

Answer – B

This is clearly given in the AWS documentation

The notification message sent by Amazon SNS for deliveries over HTTP, HTTPS, Email-JSON and SQS transport protocols will consist of a simple JSON object, which will include the following information:

- MessageId: A Universally Unique Identifier, unique for each notification published.
- Timestamp: The time (in GMT) at which the notification was published.
- TopicArn: The topic to which this message was published
- Type: The type of the delivery message, set to "Notification" for notification deliveries.
- UnsubscribeURL: A link to unsubscribe the end-point from this topic, and prevent receiving any further notifications.
- Message: The payload (body) of the message, as received from the publisher.
- Subject: The Subject field – if one was included as an optional parameter to the publish API call along with the message.
- Signature: Base64-encoded "SHA1withRSA" signature of the Message, MessageId, Subject (if present), Type, Timestamp, and Topic values.
- SignatureVersion: Version of the Amazon SNS signature used.

For more information on SNS , please visit the following URL:

- <https://aws.amazon.com/sns/faqs/> (<https://aws.amazon.com/sns/faqs/>)

Ask our Experts



QUESTION 55 UNATTEMPTED

What is the maximum number of topics allowed per account in SNS? Choose the correct answer from the options below

- ☐ A. 1,000
- ☐ B. 10,000
- ☒ C. 100,000 ✓
- ☐ D. There is no limit

Explanation :

Answer – C

This is clearly given in the AWS documentation

Q: Are there limits to the number of topics or number of subscribers per topic?

By default, SNS offers 10 million subscriptions per topic, and 100,000 topics per account. To request a higher limit, please contact us at <http://aws.amazon.com/support>

For more information on SNS , please visit the following URL:

- <https://aws.amazon.com/sns/faqs/> (<https://aws.amazon.com/sns/faqs/>)

Ask our Experts



QUESTION 56 UNATTEMPTED

Company C is currently hosting their corporate site in an Amazon S3 bucket with Static Website Hosting enabled. Currently, when visitors go to <http://www.companyc.com> the index.html page is returned. Company C now would like a new page welcome.html to be returned when a visitor enters <http://www.companyc.com> in the browser. Which of the following steps will allow Company C to meet this requirement? Choose 2 answers from the options below

- ☒ A. Upload an html page named welcome.html to their S3 bucket ✓
- ☐ B. Create a welcome subfolder in their S3 bucket
- ☒ C. Set the Index Document property to welcome.html ✓
- ☐ D. Move the index.html page to a welcome subfolder
- ☐ E. Set the Error Document property to welcome.html

Explanation :

Answer - A and C

When you go to S3, make sure that static website hosting is enabled. Also upload the index.html document , make sure it is Public and also ensure that it is put as the Index document.

▼ Static Website Hosting

You can host your static website entirely on Amazon S3. Once you enable your bucket for static website hosting, all your content is accessible to web browsers via the Amazon S3 website endpoint for your bucket.

Endpoint: `unique-bucket-1021.s3-website-ap-southeast-1.amazonaws.com`

Each bucket serves a website namespace (e.g. "www.example.com"). Requests for your host name (e.g. "example.com" or "www.example.com") can be routed to the contents in your bucket. You can also redirect request to another host name (e.g. redirect "example.com" to "www.example.com"). See our [walkthrough](#) for how to set up an Amazon S3 static website with your host name.

☐ Do not enable website hosting

☒ Enable website hosting

Index Document:

Error Document:

► **Edit Redirection Rules:** You can set custom rules to automatically redirect web page requests for specific content.

For more information on all S3 web site hosting please refer to the below URL:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/WebsiteHosting.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/WebsiteHosting.html>)

Ask our Experts



QUESTION 57 UNATTEMPTED

In SNS, push notifications are allowed for multiple platforms.

- ☒ A. True ✓
- ☐ B. False

Explanation :

Answer – A

This is clearly given in the AWS documentation

Q: Which push notifications platforms are supported?

Currently, the following push notifications platforms are supported:

- Amazon Device Messaging (ADM)
- Apple Push Notification Service (APNS)
- Google Cloud Messaging (GCM)
- Windows Push Notification Service (WNS) for Windows 8+ and Windows Phone 8.1+
- Microsoft Push Notification Service (MPNS) for Windows Phone 7+
- Baidu Cloud Push for Android devices in China

For more information on SNS , please visit the following URL:

- <https://aws.amazon.com/sns/faqs/> (<https://aws.amazon.com/sns/faqs/>)

Ask our Experts



QUESTION 58 UNATTEMPTED

Your Manager wants to have a facility to have templates from EC2 instances launched in any region without much effort. When it comes to Amazon Machine images (AMIs) which of the below statements are true when it comes to a given public AMI. Choose an answer from the options below

- ☐ A. be used to launch EC2 Instances in any AWS region.
- ☐ B. only be used to launch EC2 instances in the same country as the AMI is stored.
- ☒ C. only be used to launch EC2 instances in the same AWS region as the AMI is stored. ✓
- ☐ D. only be used to launch EC2 instances in the same AWS availability zone as the AMI is stored

Explanation :

Answer – C

AMI's can only be shared within a region. To make them available across regions , you need to copy them across regions.

You can copy an Amazon Machine Image (AMI) within or across an AWS region using the AWS Management Console, the AWS command line tools or SDKs, or the Amazon EC2 API, all of which support the CopyImageaction. You can copy both Amazon EBS-backed AMIs and instance store-backed AMIs. You can copy AMIs with encrypted snapshots and encrypted AMIs.

For more information on how to copy AMI's, please refer to the below link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>)

Ask our Experts



QUESTION 59 UNATTEMPTED

When a Simple Queue Service message triggers a task that takes 5 minutes to complete, which process below will result in successful processing of the message and remove it from the queue while minimizing the chances of duplicate processing? Choose an answer from the options below

- ☐ A. Retrieve the message with an increased visibility timeout, delete the message from the queue, process the message
- ☐ B. Retrieve the message with increased DelaySeconds, process the message, delete the message from the queue

- ☐ C. Retrieve the message with an increased visibility timeout, process the message, delete the message from the queue ✓
- ☐ D. Retrieve the message with increased DelaySeconds, delete the message from the queue, process the message

Explanation :

Answer - C

Since the default visibility timeout is less than 5 minutes , call the API to increase the timeout and then process and delete the message accordingly.

For more information on visibility timeout , please refer to the below link:

- <http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-visibility-timeout.html> (<http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-visibility-timeout.html>)

Ask our Experts



QUESTION 60 UNATTEMPTED

What is the format of structured notification messages sent by Amazon SNS? Choose an answer from the options below

- ☐ A. An XML object containing MessageId, UnsubscribeURL, Subject, Message and other values
- ☐ B. A JSON object containing MessageId, DuplicateFlag, Message and other values
- ☐ C. An XML object containing MessageId, DuplicateFlag, Message and other values
- ☐ D. A JSON object containing MessageId, unsubscribeURL, Subject, Message and other values ✓

Explanation :

Answer - D

An example of the notification is shown below , here you can see the different parts of a SNS message.

```
POST / HTTP/1.1
x-amz-sns-message-type: Notification
x-amz-sns-message-id: 22b80b92-fdea-4c2c-8f9d-bdfb0c7bf324
x-amz-sns-topic-arn: arn:aws:sns:us-west-2:123456789012:MyTopic
x-amz-sns-subscription-arn: arn:aws:sns:us-west-2:123456789012:MyTopic:c9135db0-26c4-47e
Content-Length: 773
Content-Type: text/plain; charset=UTF-8
Host: myhost.example.com
Connection: Keep-Alive
User-Agent: Amazon Simple Notification Service Agent

{
  "Type" : "Notification",
  "MessageId" : "22b80b92-fdea-4c2c-8f9d-bdfb0c7bf324",
  "TopicArn" : "arn:aws:sns:us-west-2:123456789012:MyTopic",
  "Subject" : "My First Message",
  "Message" : "Hello world!",
  "Timestamp" : "2012-05-02T00:54:06.655Z",
  "SignatureVersion" : "1",
  "Signature" : "EXAMPLEEv6JRNvmlLFQL4ICB0bnXrdB8C1RMTQFGBqVlpGbM78tJ4etTwC5zU703tS6tGpey",
  "SigningCertURL" : "https://sns.us-west-2.amazonaws.com/SimpleNotificationService-f3ec",
  "UnsubscribeURL" : "https://sns.us-west-2.amazonaws.com/?Action=Unsubscribe&Subscripti"
}
```

For more information on SNS notification format, please refer to the link:

- <http://docs.aws.amazon.com/sns/latest/dg/json-formats.html> (<http://docs.aws.amazon.com/sns/latest/dg/json-formats.html>)

Ask our Experts



Certification

☛ Cloud Certification

(<https://www.whizlabs.com/cloud-certification-training-courses/>) us)

☛ Java Certification

(<https://www.whizlabs.com/oracle/java-certifications/>)

☛ PM Certification

(<https://www.whizlabs.com/project-management-certifications/>)

☛ Big Data Certification

(<https://www.whizlabs.com/big-data-certifications/>)

Company

☛ Support

(<https://help.whizlabs.com/hc/en-us>)

☛ Discussions

(<http://ask.whizlabs.com/>)

☛ Blog

(<https://www.whizlabs.com/blog/>)

Mobile App

🤖 Android Coming Soon

🍏 iOS Coming Soon

Follow us



(<https://www.facebook.com/whizlabs.software/>)



(<https://in.linkedin.com/company/whizlabs-software>)



(<https://twitter.com/whizlabs?lang=en>)



(<https://plus.google.com/+WhizlabsSoftware>)