



[Home](https://www.whizlabs.com/learn) (<https://www.whizlabs.com/learn>) > [My Courses](https://www.whizlabs.com/learn/my-courses) (<https://www.whizlabs.com/learn/my-courses>)  
> [AWS Certified Developer Associate](https://www.whizlabs.com/learn/course/aws-cda-practice-tests#section-1) (<https://www.whizlabs.com/learn/course/aws-cda-practice-tests#section-1>)  
> [Practice Test V](https://www.whizlabs.com/learn/course/aws-cda-practice-tests/quiz/13231) (<https://www.whizlabs.com/learn/course/aws-cda-practice-tests/quiz/13231>) > **Report**

## PRACTICE TEST V

<b>Attempt</b>	1	<b>Completed on</b>	Tuesday , 29 January 2019 , 04:24 PM
<b>Marks Obtained</b>	2 / 60	<b>Time Taken</b>	00 H 00 M 41 S
<b>Your score is</b>	3.33%	<b>Result</b>	Fail

### Domains / Topics wise Quiz Performance Report

S.No.	Topic	Total Questions	Correct	Incorrect	Unattempted
1	Other	60	2	0	58

<b>60</b> Questions	<b>2</b> Correct	<b>0</b> Incorrect	<b>58</b> Unattempted	Show Answers	All	▼
------------------------	---------------------	-----------------------	--------------------------	--------------	-----	---

### QUESTION 1 CORRECT

Is a best practice to allow code to run on an EC2 instance via a role rather than access keys? Choose the correct answer from the options below

- ☒ A. True ✓
- ☐ B. False

#### Explanation :

Answer – A

The best practise for IAM is to create roles which has specific access to an AWS service and then give the user permission to the AWS service via the role.

To get the role in place , follow the below steps

Step 1) Create a role which has the required ELB access

## Set Role Name

Enter a role name. You cannot edit the role name after the role is created.

Role Name

ELBAccess

Maximum 64 characters. Use alphanumeric and '+,=, @, \_' characters

Step 2) You need to provide permissions to the underlying EC2 instances in the Elastic Load Balancer

### Select Role Type

**AWS Service Roles**

Amazon EC2

Allows EC2 instances to call AWS services on your behalf.

Select

<input type="checkbox"/>		AmazonEC2ContainerService...	0	2015-04-09 20:14 UTC+0400	2016-08-11 17:08 UTC+0400
<input checked="" type="checkbox"/>		AmazonEC2FullAccess	0	2015-02-06 22:40 UTC+0400	2015-02-06 22:40 UTC+0400

For the best practises on IAM policies, please visit the link:

- <http://docs.aws.amazon.com/IAM/latest/UserGuide/best-practices.html>  
(<http://docs.aws.amazon.com/IAM/latest/UserGuide/best-practices.html>)

Ask our Experts



QUESTION 2 CORRECT

A user named john has an AWS account. He has created an AMI in the us-east1 region. Is it possible for him to share the AMI with another user account in AWS?  
Choose the correct answer from the options below

- ☒ A. True ✓
- ☐ B. False

#### Explanation :

Answer – A

This is clearly mentioned in the AWS documentation.

## Sharing an AMI with Specific AWS Accounts

You can share an AMI with specific AWS accounts without making the AMI public. All you need are the AWS account IDs.

AMIs are a regional resource. Therefore, sharing an AMI makes it available in that region. To make an AMI available in a different region, copy the AMI to the region and then share it. For more information, see [Copying an AMI](#).

For more information on how to share AMI's, please refer to the below link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/sharingamis-explicit.html>  
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/sharingamis-explicit.html>)

Ask our Experts



### QUESTION 3 UNATTEMPTED

What is the Size limit for a root device for the instances that are based on instance Store backed AMI's. Choose an answer from the options below?

- ☐ A. 16 TiB
- ☒ B. 10 GiB ✓
- ☐ C. 100 GiB
- ☐ D. 50 GiB

#### Explanation :

Answer - B

As per the AWS documentation, the Size limit for Amazon Instance Store backed instances is 10 GiB.

Characteristic	Amazon EBS-Backed	Amazon Instance Store-Backed
Boot time	Usually less than 1 minute	Usually less than 5 minutes
Size limit	16 TiB	10 GiB
Root device volume	Amazon EBS volume	Instance store volume
Data persistence	By default, the root volume is deleted when the instance terminates.* Data on any other Amazon EBS volumes persists after instance termination by default. Data on any instance store volumes persists only during the life of the instance.	Data on any instance store volumes persists only during the life of the instance. Data on any Amazon EBS volumes persists after instance termination by default.

For more information on Instance type differences, please refer to the below link:



- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>  
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>)

Ask our Experts



QUESTION 4 UNATTEMPTED

Which API call is used to Bundle an Amazon instance store-backed Windows instance? Choose the correct answer from the options below

- ☐ A. AllocateInstance
- ☐ B. CreateImage
- ☒ C. BundleInstance ✓
- ☐ D. ami-register-image

**Explanation :**

Answer – C

This is given in the AWS documentation

**BundleInstance**

Bundles an Amazon instance store-backed Windows instance.

During bundling, only the root device volume (C:\) is bundled. Data on other instance store volumes is not preserved.

**Note**

This action is not applicable for Linux/Unix instances or Windows instances that are backed by Amazon EBS.

For more information, see [Creating an Instance Store-Backed Windows AMI](#).

For more information on BundleInstance, please refer to the below link:

- [http://docs.aws.amazon.com/AWSEC2/latest/APIReference/API\\_BundleInstance.html](http://docs.aws.amazon.com/AWSEC2/latest/APIReference/API_BundleInstance.html)  
([http://docs.aws.amazon.com/AWSEC2/latest/APIReference/API\\_BundleInstance.html](http://docs.aws.amazon.com/AWSEC2/latest/APIReference/API_BundleInstance.html))

Ask our Experts



QUESTION 5 UNATTEMPTED



What is one difference between an Amazon EBS-backed and an instance-store backed instance? Choose the correct answer from the options below

- ☐ A. Instance-store backed instances can be stopped and restarted
- ☐ B. The size limit for Instance store volumes is 16 TiB
- ☐ C. Virtual Private Cloud requires EBS backed instances
- ☐ D. The instance type and kernel for an EBS backed AMI can easily be changed as compared to an Instance store-backed AMI. ✓

**Explanation :**

Answer – D

This is clearly mentioned in the AWS documentation

Characteristic	Amazon EBS-Backed	Amazon Instance Store-Backed
Boot time	Usually less than 1 minute	Usually less than 5 minutes
Size limit	16 TiB	10 GiB
Root device volume	Amazon EBS volume	Instance store volume
Data persistence	By default, the root volume is deleted when the instance terminates.* Data on any other Amazon EBS volumes persists after instance termination by default. Data on any instance store volumes persists only during the life of the instance.	Data on any instance store volumes persists only during the life of the instance. Data on any Amazon EBS volumes persists after instance termination by default.
Upgrading	The instance type, kernel, RAM disk, and user data can be changed while the instance is stopped.	Instance attributes are fixed for the life of an instance.
Charges	You're charged for instance usage, Amazon EBS volume usage, and storing your AMI as an Amazon EBS snapshot.	You're charged for instance usage and storing your AMI in Amazon S3
AMI creation/bundling	Uses a single command/call	Requires installation and use of AMI tools
Stopped state	Can be placed in stopped state where instance is not running, but the root volume is persisted in Amazon EBS	Cannot be in stopped state; instances are running or terminated

For more information on Instance type differences, please refer to the below link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>  
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>)

Ask our Experts



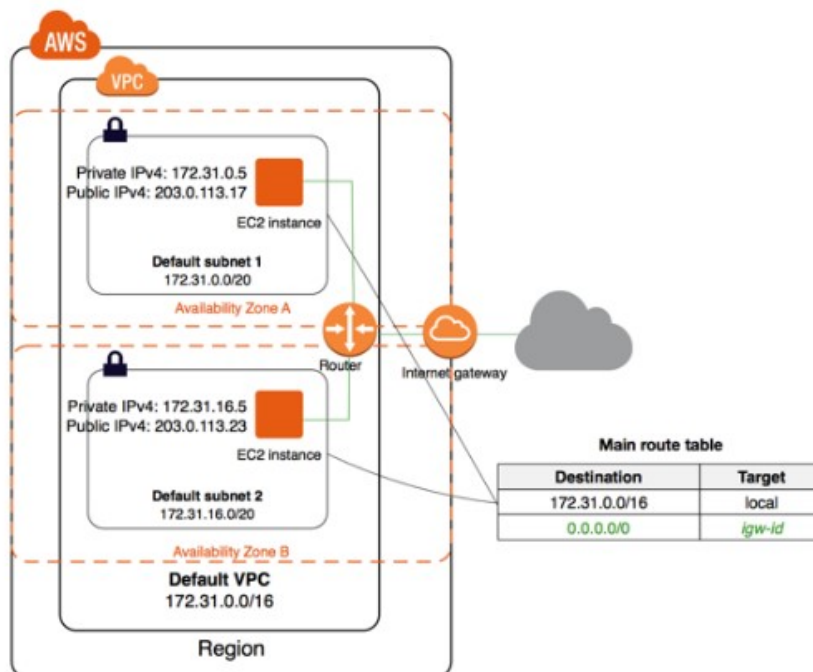
Is a public IP address required for an instance to be accessed from the internet in a VPC? Choose the correct answer from the options below

- ☐ A. True ✓
- ☐ B. False

#### Explanation :

Answer – A

The below example shows a VPC which has an EC2 instance in a subnet which has an internet gateway. You can see that in order to get to the internet, it needs to have a public IP address.



For more information on the default VPC, please refer to the below link:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/default-vpc.html>  
(<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/default-vpc.html>)

Ask our Experts



QUESTION 7 UNATTEMPTED

What is the hourly rate to run an Autoscaling service? Choose the correct answer from the options below



- ☐ A. .002/hour
- ☐ B. .01/hour
- ☒ C. Free ✓
- ☐ D. .05/hour

**Explanation :**

Answer – C

This is clearly mentioned in the AWS documentation that there are no charges for using the Autoscaling service as it is.

## Pricing

### Auto Scaling Pricing

Auto Scaling is enabled by Amazon CloudWatch and carries no additional fees. Amazon EC2 and Amazon CloudWatch service fees apply and are billed separately. Partial hours are billed as full hours.

For more information on Autoscaling, please refer to the below link:

- <https://aws.amazon.com/autoscaling/> (<https://aws.amazon.com/autoscaling/>)

Ask our Experts



**QUESTION 8      UNATTEMPTED**

What is required for a subnet to control the flow of traffic in a subnet? Choose one answer from the options below?

- ☒ A. Route table ✓
- ☐ B. Subnet table
- ☐ C. VPC table
- ☐ D. Route53

**Explanation :**

Answer - A

This is clearly mentioned in the AWS documentation that Route tables are used to control the flow of traffic in a subnet.



A *route table* contains a set of rules, called *routes*, that are used to determine where network traffic is directed.

Each subnet in your VPC must be associated with a route table; the table controls the routing for the subnet. A subnet can only be associated with one route table at a time, but you can associate multiple subnets with the same route table.

For more information on Route tables, please refer to the below link:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Route\\_Tables.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Route_Tables.html)  
([http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Route\\_Tables.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Route_Tables.html))

Ask our Experts



#### QUESTION 9 UNATTEMPTED

Which service provided by AWS allows one to have a dedicated private connection from a remote network to your VPC? Choose the correct answer from the options below

- ☒ A. AWS Direct Connect ✓
- ☐ B. Server
- ☐ C. Elastic IP Address
- ☐ D. None of the above

#### Explanation :

Answer – A

- AWS Direct Connect - AWS Direct Connect provides a dedicated private connection from a remote network to your VPC. You can combine this connection with an AWS hardware VPN connection to create an IPsec-encrypted connection.

For more information on AWS direct connect, please refer to the below link:

- <https://aws.amazon.com/directconnect/> (<https://aws.amazon.com/directconnect/>)

Ask our Experts



#### QUESTION 10 UNATTEMPTED

As per the IAM decision logic what is the first step of access permissions for any resource in AWS. Choose the correct answer from the options below

- ☒ A. A default deny ✓



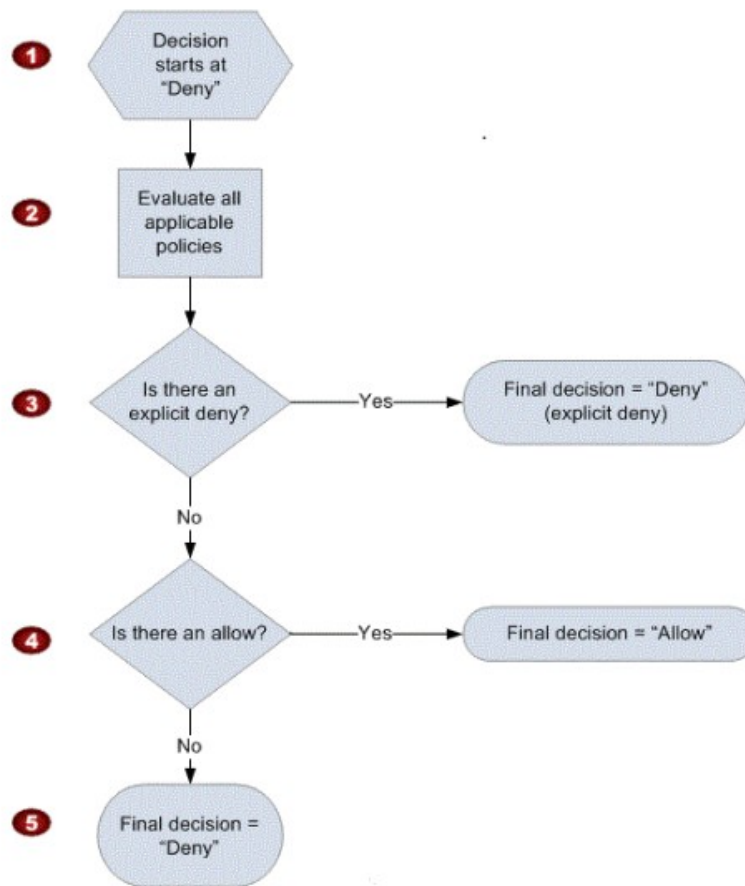


- ☐ B. An explicit deny
- ☐ C. An allow
- ☐ D. An explicit allow

**Explanation :**

Answer – A

The below diagram shows the evaluation logic of IAM policies. And as per the evaluation logic, it is clear that the above scenario leads to a default deny.



For more information on the IAM policy evaluation logic, please refer to the link:

- [http://docs.aws.amazon.com/IAM/latest/UserGuide/reference\\_policies\\_evaluation-logic.html](http://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_evaluation-logic.html)  
([http://docs.aws.amazon.com/IAM/latest/UserGuide/reference\\_policies\\_evaluation-logic.html](http://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_evaluation-logic.html))

Ask our Experts



QUESTION 11 UNATTEMPTED

What is one of the common pattern scenario's when it comes to the combination of

SNS and SQS? Choose the correct answer from the options below

- ☐ A. fanin
- ☒ B. fanout ✓
- ☐ C. Neither option A and B
- ☐ D. Both option A and B

**Explanation :**

Answer - B

One common design pattern is called “fanout.” In this pattern, a message published to an SNS topic is distributed to a number of SQS queues in parallel. By using this pattern, you can build applications that take advantage parallel, asynchronous processing

For more information on SNS and SQS integration, please refer to the link:

- <https://aws.amazon.com/blogs/aws/queues-and-notifications-now-best-friends/>  
(<https://aws.amazon.com/blogs/aws/queues-and-notifications-now-best-friends/>)

Ask our Experts



QUESTION 12 UNATTEMPTED

What are the different S3 Encryption techniques?

- ☐ A. Use Server-Side Encryption with Amazon S3-Managed Keys
- ☐ B. Server-Side Encryption with AWS KMS-Managed Keys
- ☐ C. Server-Side Encryption with Customer-Provided Keys
- ☒ D. All of the above ✓

**Explanation :**

Answer – D

There are various ways to encrypt objects in S3.

**Use Server-Side Encryption with Amazon S3-Managed Keys (SSE-S3)** – Each object is encrypted with a unique key employing strong multi-factor encryption. As an additional safeguard, it encrypts the key itself with a master key that it regularly rotates. Amazon S3 server-side encryption uses one of the strongest block ciphers available, 256-bit Advanced Encryption Standard (AES-256), to encrypt your data.

**Use Server-Side Encryption with AWS KMS-Managed Keys (SSE-KMS)** – Similar to SSE-S3, but with some additional benefits along with some additional charges for using this service. There are separate permissions for the use of an envelope key (that is, a key that protects your data's encryption key) that provides added protection against unauthorized access of your objects in S3. SSE-KMS also provides

you with an audit trail of when your key was used and by whom. Additionally, you have the option to create and manage encryption keys yourself, or use a default key that is unique to you, the service you're using, and the region you're working in.

**Use Server-Side Encryption with Customer-Provided Keys (SSE-C)** – You manage the encryption keys and Amazon S3 manages the encryption, as it writes to disks, and decryption, when you access your objects.

For more information on S3 encryption, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/serv-side-encryption.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/serv-side-encryption.html>)

Ask our Experts



QUESTION 13 UNATTEMPTED

Which of the following would you not expect to see in an SNS message body?  
Choose the correct answer from the options below

- ☐ A. Signature
- ☐ B. MessageId
- ☐ C. SigningCertURL
- ☒ D. SNSID ✓

**Explanation :**

Answer – D

An example of the notification is shown below, There is no SNSID

```
POST / HTTP/1.1
x-amz-sns-message-type: Notification
x-amz-sns-message-id: 22b80b92-fdea-4c2c-8f9d-bdfb0c7bf324
x-amz-sns-topic-arn: arn:aws:sns:us-west-2:123456789012:MyTopic
x-amz-sns-subscription-arn: arn:aws:sns:us-west-2:123456789012:MyTopic:c9135db0-26c4-47ec
Content-Length: 773
Content-Type: text/plain; charset=UTF-8
Host: myhost.example.com
Connection: Keep-Alive
User-Agent: Amazon Simple Notification Service Agent

{
  "Type" : "Notification",
  "MessageId" : "22b80b92-fdea-4c2c-8f9d-bdfb0c7bf324",
  "TopicArn" : "arn:aws:sns:us-west-2:123456789012:MyTopic",
  "Subject" : "My First Message",
  "Message" : "Hello world!",
  "Timestamp" : "2012-05-02T00:54:06.655Z",
  "SignatureVersion" : "1",
  "Signature" : "EXAMPLEw6JRNwm1LFQL4ICB0bnXrdB8ClRMTQFGBqwLpGbM78tJ4etTwC5zU703tS6tGpey",
  "SigningCertURL" : "https://sns.us-west-2.amazonaws.com/SimpleNotificationService-f3ec",
  "UnsubscribeURL" : "https://sns.us-west-2.amazonaws.com/?Action=Unsubscribe&Subscription",
}
```

For more information on SNS notification format, please refer to the link:

- <http://docs.aws.amazon.com/sns/latest/dg/json-formats.html>  
(<http://docs.aws.amazon.com/sns/latest/dg/json-formats.html>)

Ask our Experts



QUESTION 14 UNATTEMPTED

You are explaining SNS message attributes to an IT admin. What specifically do you need to tell the IT admin about the message attributes of SNS?

Choose the correct answer from the options below

- ☐ A. Name, type, and value can be empty or null but the message body cannot be empty or null.
- ☐ B. Name, type, and value must not be empty or null but the message body can be empty or null.
- ☐ C. Name, type, and value can be empty or null and the message body can be empty or null.
- ☐ D. Name, type, and value must not be empty or null and the message body shouldn't be empty or null either. ✓

**Explanation :**

Answer – D

Amazon Simple Notification Service (Amazon SNS) provides support for delivery of message attributes to Amazon SQS endpoints. Message attributes allow you to provide structured metadata items (such as timestamps, geospatial data, signatures, and identifiers) about the message. Also the requirement for each attribute to be not NULL in addition to the message body is given in the AWS documentation.

For more information on SNS message attributes, please refer to the link:

- <http://docs.aws.amazon.com/sns/latest/dg/SNSMessageAttributes.html>  
(<http://docs.aws.amazon.com/sns/latest/dg/SNSMessageAttributes.html>)

Ask our Experts



QUESTION 15 UNATTEMPTED

An IT admin has enabled long polling in their SQS queue. What must be done for long polling to be enabled in SQS? Choose the correct answer from the options below ^

- ☐ A. Create a dead letter queue.
- ☐ B. Set the message size to 256KB
- ☐ C. Set the ReceiveMessageWaitTimeSeconds property of the queue to 0 seconds
- ☒ D. Set the ReceiveMessageWaitTimeSeconds property of the queue to 20 seconds



#### Explanation :

Answer – D

Amazon SQS long polling is a way to retrieve messages from your Amazon SQS queues. While the regular short polling returns immediately, even if the message queue being polled is empty, long polling doesn't return a response until a message arrives in the message queue, or the long poll times out. Long polling makes it inexpensive to retrieve messages from your Amazon SQS queue as soon as the messages are available. Using long polling might reduce the cost of using SQS, because you can reduce the number of empty receives

For more information on Long polling, please refer to the link:

- <http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-long-polling.html>  
(<http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-long-polling.html>)

Ask our Experts



#### QUESTION 16 UNATTEMPTED

What is the minimum value that can be configured for Amazon SQS MaximumMessageSize attribute?

- ☒ A. 1KB ✓
- ☐ B. 128KB
- ☐ C. 16 KB
- ☐ D. 256KB

#### Explanation :

Answer – A

This is clearly specified in the AWS documentation.



To configure the maximum message size, use the console or the `SetQueueAttributes` method to set the `MaximumMessageSize` attribute. This attribute specifies the limit on bytes that an Amazon SQS message can contain. Set this limit to a value between 1,024 bytes (1 KB), and 262,144 bytes (256 KB). For more information, see [Using Amazon SQS Message Attributes](#) in the *Amazon SQS Developer Guide*.

To send messages larger than 256 KB, use the [Amazon SQS Extended Client Library for Java](#). This library lets you send an Amazon SQS message that contains a reference to a message payload in Amazon S3 that can be as large as 2 GB. For FIFO queues, see "Are Amazon SQS FIFO queues compatible with the Amazon SQS Buffered Asynchronous Client, the Amazon SQS Extended Client Library for Java, or the Amazon SQS Java Message Service (JMS) Client?"

For more information on AWS SQS, please refer to the link:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



## QUESTION 17 UNATTEMPTED

In SQS when the `ReceiveMessageWaitTimeSeconds` is set to a value of 20 seconds, what is the functionality being served in SQS? Choose the correct answer from the options below

- ☐ A. Short polling
- ☐ B. SQS retention
- ☒ C. Long polling ✓
- ☐ D. None of the above

### Explanation :

Answer – C

Amazon SQS long polling is a way to retrieve messages from your Amazon SQS queues. While the regular short polling returns immediately, even if the message queue being polled is empty, long polling doesn't return a response until a message arrives in the message queue, or the long poll times out. Long polling makes it inexpensive to retrieve messages from your Amazon SQS queue as soon as the messages are available. Using long polling might reduce the cost of using SQS, because you can reduce the number of empty receives

To enable long polling you need to set the value of `ReceiveMessageWaitTimeSeconds` to greater than 0 and less than or equal to 20 seconds.

For more information on Long polling, please refer to the link:

- <http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-long-polling.html>  
(<http://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-long-polling.html>)

Ask our Experts



## QUESTION 18

UNATTEMPTED

An IT admin is worried about the number of message queues being created in the SQS in their AWS account? Is there a limit, if yes, what is the correct limit? Choose the correct answer from the options below

- ☐ A. 50
- ☐ B. 100
- ☐ C. 200
- ☒ D. There is no limit ✓

**Explanation :**

Answer – D

As per the AWS documentation, there is no limit on the number of queues.

**Q: How many message queues can I create?**

You can create any number of message queues.

For more information on AWS SQS, please refer to the link:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



## QUESTION 19

UNATTEMPTED

Is it possible for duplicate messages to be received within SQS? Choose the correct answer from the options below.

- ☒ A. True ✓
- ☐ B. False

**Explanation :**

Answer – A

Amazon SQS has a *minimum* deduplication interval of 5 minutes. Retrying `SendMessage` requests after the deduplication interval expires can introduce duplicate messages into the queue. For example, a mobile device in a car sends messages whose order is important. If the car loses cellular connectivity for a period of time before receiving an acknowledgement, retrying the request after regaining cellular connectivity can create a duplicate.

Amazon SWF keeps track of all tasks and events in an application. Amazon SQS requires you to implement your own application-level tracking, especially if your application uses multiple queues. For more information on SQS, please visit the link:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 20 UNATTEMPTED

What can be used to deploy workers and deciders in Amazon SWF? Choose the correct answer from the options below.

- ☐ A. Amazon EC2 instances
- ☐ B. Amazon Lambda
- ☐ C. On-premise machines
- ☒ D. All of the above ✓

Explanation :

Answer – D

This is clearly mentioned in the AWS documentation

To coordinate the application execution across workers, you write a program called the decider in your choice of programming language. The separation of processing steps and their coordination makes it possible to manage your application in a controlled manner and give you the flexibility to deploy, run, scale and update them independently. You can choose to deploy workers and deciders either in the [cloud](#) (e.g. Amazon EC2 or Lambda) or on machines behind corporate firewalls. Because of the decoupling of workers and deciders, your business logic can be dynamic and your application can be quickly updated to accommodate new requirements. For example, you can remove, skip, or retry tasks and create new application flows simply by changing the decider.

For more information on SWF , please visit the link:

- <https://aws.amazon.com/swf/faqs/> (<https://aws.amazon.com/swf/faqs/>)

Ask our Experts



QUESTION 21 UNATTEMPTED

What are some of the benefits of using AWS SWF? Choose 2 correct answer from the options below.

- ☐ A. Integrating a range of programs and components. ✓





- ☐ B. Automate workflows that include human tasks ✓
- ☐ C. Load balancing between applications
- ☐ D. Routing between applications

#### Explanation :

Answer – A and B

Amazon SWF can be used to address many challenges that arise while building applications with distributed components. For example, you can use Amazon SWF and the accompanying AWS Flow Framework for:

- Writing your applications as asynchronous programs using simple programming constructs that abstract details such as initiating tasks to run remotely and tracking the program's runtime state.
- Maintaining your application's execution state (e.g. which steps have completed, which ones are running, etc.). You do not have to use databases, custom systems, or ad hoc solutions to keep execution state.
- Communicating and managing the flow of work between your application components. With Amazon SWF, you do not need to design a messaging protocol or worry about lost and duplicated tasks.
- Centralizing the coordination of steps in your application. Your coordination logic does not have to be scattered across different components, but can be encapsulated in a single program.
- Integrating a range of programs and components, including legacy systems and 3rd party cloud services, into your applications. By allowing your application flexibility in where and in what combination the application components are deployed, Amazon SWF helps you gradually migrate application components from private data centers to public cloud infrastructure without disrupting the application availability or performance.
- Automating workflows that include long-running human tasks (e.g. approvals, reviews, investigations, etc.) Amazon SWF reliably tracks the status of processing steps that run up to several days or months.

For more information on SWF , please visit the link:

- <https://aws.amazon.com/swf/faqs/> (<https://aws.amazon.com/swf/faqs/>)

Ask our Experts



QUESTION 22

UNATTEMPTED

What is one of the major use cases for using SWF?

- ☐ A. Storing files In S3.
- ☐ B. Migrating components from the datacenter to the cloud ✓
- ☐ C. Storing files In database.
- ☐ D. Provisioning servers.



### Explanation :

Answer – B

This is clearly mentioned in the AWS documentation

Use case #3: Migrating components from the datacenter to the cloud. Business critical operations are hosted in a private datacenter but need to be moved entirely to the cloud without causing disruptions.

With Amazon SWF: Amazon SWF-based applications can combine workers that wrap components running in the datacenter with workers that run in the cloud. To transition a datacenter worker seamlessly, new workers of the same type are first deployed in the cloud. The workers in the datacenter continue to run as usual, along with the new cloud-based workers. The cloud-based workers are tested and validated by routing a portion of the load through them. During this testing, the application is not disrupted because the workers in the datacenter continue to run. After successful testing, the workers in the datacenter are gradually stopped and those in the cloud are scaled up, so that the workers are eventually run entirely in the cloud. This process can be repeated for all other workers in the datacenter so that the application moves entirely to the cloud. If for some business reason, certain processing steps must continue to be performed in the private data center, those workers can continue to run in the private data center and still participate in the application.

### Q: What are the different ways to access SWF?

You can access SWF in any of the following ways:

- AWS SDK for Java, Ruby, .NET, and PHP
- AWS Flow Framework for Java (Included in the AWS SDK for Java)
- Amazon SWF web service APIs
- AWS Management Console

For more information on SWF , please visit the link:

- <https://aws.amazon.com/swf/faqs/> (<https://aws.amazon.com/swf/faqs/>)

Ask our Experts



QUESTION 23

UNATTEMPTED

For S3 server side encryption what needs to be included in the API header call?  
Choose the correct answer from the options below

- ☐ A. AES256
- ☐ B. amz-server-side-encryption
- ☒ C. x-amz-server-side-encryption ✓
- ☐ D. server-side-encryption

### Explanation :

Answer – C

When using server-side encryption with customer-provided encryption keys (SSE-C), you must



provide encryption key information using the following request headers.

The object creation REST APIs provide a request header, x-amz-server-side-encryption that you can use to request server-side encryption.

For more information on server side encryption , please refer to the below URL:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingServerSideEncryption.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingServerSideEncryption.html>)

Ask our Experts



## QUESTION 24 UNATTEMPTED

What are the different ways to access SWF?

- ☐ A. AWS SDK for Java
- ☐ B. AWS Management console
- ☐ C. AWS SWF API's
- ☒ D. All of the above ✓
- ☐ E. Both A and B
- ☐ F. Both B and C
- ☐ G. Both A and C

### Explanation :

Answer – D

This is clearly mentioned in the AWS documentation

#### **Q: What are the different ways to access SWF?**

You can access SWF in any of the following ways:

- AWS SDK for Java, Ruby, .NET, and PHP
- AWS Flow Framework for Java (Included in the AWS SDK for Java)
- Amazon SWF web service APIs
- AWS Management Console

For more information on SWF , please visit the link:

- <https://aws.amazon.com/swf/faqs/> (<https://aws.amazon.com/swf/faqs/>)

Ask our Experts



What in AWS can be used to restrict access to SWF?

- ☐ A. ACL
- ☐ B. SWF Roles
- ☒ C. IAM ✓
- ☐ D. None of the above

**Explanation :**

Answer – C

This is clearly mentioned in the AWS documentation

Name, type, and value must not be empty or null. In addition, the message body should not be empty or null. All parts of the message attribute, including name, type, and value, are included in the message size restriction, which is currently 256 KB (262,144 bytes).

For more information on SWF, please visit the link:

- <https://aws.amazon.com/swf/faqs/> (<https://aws.amazon.com/swf/faqs/>)

Ask our Experts



What is the maximum number of SWF workflows/activity types allowed in each SWF domain? Choose the correct answer from the options below.

- ☐ A. 1000
- ☐ B. 100
- ☒ C. 10,000 ✓
- ☐ D. 100,000

**Explanation :**

Answer – C

You can have a maximum of 10,000 workflow and activity types (in total) that are either registered or deprecated in each domain. You can have a maximum of 100 Amazon SWF domains (including registered and deprecated domains) in your AWS account.

For more information on SWF , please visit the link:



- <https://aws.amazon.com/swf/faqs/> (<https://aws.amazon.com/swf/faqs/>)

Ask our Experts



QUESTION 27 UNATTEMPTED

Which of the following is the best option for creating a primary key in DynamoDB?  
Select the BEST possible answer.

- ☐ A. student\_id where there are many students. ✓
- ☐ B. course\_name where there are few courses.
- ☐ C. department\_id where there are only several departments.
- ☐ D. None of the above

Explanation :

Answer – A

When defining primary keys , you should always use a many to few principle and only Option A follows that principle.

For more information on dynamoDB , please visit the link:

- <https://aws.amazon.com/dynamodb/faqs/> (<https://aws.amazon.com/dynamodb/faqs/>)

Ask our Experts



QUESTION 28 UNATTEMPTED

Which of the below features in DynamoDB can be used to increment or decrement the value of an existing attribute without interfering with other write requests?  
Choose the correct answer from the options below.

- ☐ A. Increment write
- ☐ B. Increment read
- ☐ C. Atomic counters ✓
- ☐ D. Number counters

Explanation :



Answer – C

DynamoDB supports *atomic counters*, where you use the UpdateItem operation to increment or decrement the value of an existing attribute without interfering with other write requests. (All write requests are applied in the order in which they were received.) For example, a web application might want to maintain a counter per visitor to their site. In this case, the application would need to increment this counter regardless of its current value.

For more information on working with items in dynamoDB , please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithItems.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithItems.html>)

Ask our Experts



QUESTION 29

UNATTEMPTED

Company A is writing 5 items to the products table every second. Each item is 20Kb in size. What would be the required provisioned write throughput for best performance? Choose the correct answer from the options below.

- ☒ A. 100 ✓
- ☐ B. 160
- ☐ C. 155
- ☐ D. 16

**Explanation :**

Answer – A

For write capacity, the rule is to divide the item size by 1KB. Hence we need to divide 20 by 1 which gives us 20 to the nearest 1KB. Since we are writing 5 items per second , we need to multiply  $20 \times 5 = 100$ .

For more information on working with tables in dynamoDB , please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.html>)

Ask our Experts



QUESTION 30

UNATTEMPTED

What is the maximum size of an item in DynamoDB? Choose the correct answer from the options below.



- ☐ A. 1 MB
- ☒ B. 400 KB ✓
- ☐ C. 1 GB
- ☐ D. 10 GB

**Explanation :**

Answer – B

This is clearly given in the AWS documentation

**BatchGetItem**

A single `BatchGetItem` operation can retrieve a maximum of 100 items. The total size of all the items retrieved cannot exceed 16 MB.

For more information on DynamoDB limits, please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Limits.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Limits.html>)

Ask our Experts



**QUESTION 31      UNATTEMPTED**

What is the maximum number of items that the `BatchGetItem` API retrieve from DynamoDB? Choose the correct answer from the options below.

- ☐ A. 10
- ☒ B. 100 ✓
- ☐ C. 1000
- ☐ D. There is no limit

**Explanation :**

Answer – B

The `BatchGetItem` operation returns the attributes of one or more items from one or more tables. You identify requested items by primary key.

A single operation can retrieve up to 16 MB of data, which can contain as many as 100 items. `BatchGetItem` will return a partial result if the response size limit is exceeded, the table's provisioned throughput is exceeded, or an internal processing failure occurs. If a partial result is returned, the operation returns a value for `UnprocessedKeys`. You can use this value to retry the operation starting with the next item to get.

#### Important

If you request more than 100 items `BatchGetItem` will return a `ValidationException` with the message "Too many items requested for the BatchGetItem call".

- [https://docs.aws.amazon.com/amazondynamodb/latest/APIReference/API\\_BatchGetItem.html](https://docs.aws.amazon.com/amazondynamodb/latest/APIReference/API_BatchGetItem.html)  
([https://docs.aws.amazon.com/amazondynamodb/latest/APIReference/API\\_BatchGetItem.html](https://docs.aws.amazon.com/amazondynamodb/latest/APIReference/API_BatchGetItem.html))

Ask our Experts



#### QUESTION 32 UNATTEMPTED

What is the maximum limit of data that can be retrieved by a scan operation in DynamoDB? Choose the correct answer from the options below.

- ☒ A. 1 MB ✓
- ☐ B. 10 MB
- ☐ C. 400 KB
- ☐ D. There is no limit

#### Explanation :

Answer – A

This is clearly given in the AWS documentation

**Scan**

The result set from a `scan` is limited to 1 MB per call. You can use the `LastEvaluatedKey` from the scan response to retrieve more results.

For more information on DynamoDB limits, please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Limits.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Limits.html>)

Ask our Experts



#### QUESTION 33 UNATTEMPTED

True or False. A local secondary index has the same partition key as the primary key?

- ☒ A. True ✓
- ☐ B. False





### Explanation :

Answer – A

- Global secondary index

(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/GSI.html>) – an index with a partition key and a sort key that can be different from those on the base table. A global secondary index is considered "global" because queries on the index can span all of the data in the base table, across all partitions.

- Local secondary index

(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/LSI.html>) – an index that has the same partition key as the base table, but a different sort key. A local secondary index is "local" in the sense that every partition of a local secondary index is scoped to a base table partition that has the same partition key value.

For more information on DynamoDB Indexes, please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SecondaryIndexes.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SecondaryIndexes.html>)

Ask our Experts



QUESTION 34

UNATTEMPTED

Which of the following is a valid S3 bucket name? Choose the correct answer from the options below

- ☐ A. .example.com
- ☒ B. example.com ✓
- ☐ C. -example.com
- ☐ D. example.-com

### Explanation :

Answer – B

Some of the naming restrictions for buckets are given below

- Bucket names must be at least 3 and no more than 63 characters long.
- Bucket names must be a series of one or more labels. Adjacent labels are separated by a single period (.). Bucket names can contain lowercase letters, numbers, and hyphens. Each label must start and end with a lowercase letter or a number.
- Bucket names must not be formatted as an IP address (e.g., 192.168.5.4).

- When using virtual hosted-style buckets with SSL, the SSL wildcard certificate only matches buckets that do not contain periods. To work around this, use HTTP or write your own certificate verification logic. We recommend that you do not use periods (".") in bucket names.

For more information on S3 bucket naming conventions, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/BucketRestrictions.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/BucketRestrictions.html>)

Ask our Experts



QUESTION 35 UNATTEMPTED

A Global Secondary Index can have different partition key and sort key from those of its base table. True or False.

- ☒ A. True ✓
- ☐ B. False

**Explanation :**

Answer – A

GSI can also use the same partition key as the base table. Even in the AWS documentation they say that GSI can use a different partition key and sort key. However anywhere in the document they are not saying that it has to be different.



AWS Console Test:

Create DynamoDB table

Tutorial ?

DynamoDB is a schema-less database that only requires a table name and primary key. The table's primary key is made up of one or two attributes that uniquely identify items, partition the data, and sort data within each partition.

Table name\*Employee ?

Primary key\* Partition key

EmpIDNumber ?

✓ Add sort key

DeptIDNumber ?

Table settings

Default settings provide the fastest way to get started with your table. You can modify these default settings now or after your table has been created.

Use default settings

Secondary indexes

Name	Type	Partition key	Sort key	Projected Attributes
<a href="#">+ Add index</a>				

Provisioned capacity

Read capacity units

Table5

Write capacity units

5

Estimated cost \$6.61 / month (Capacity calculator)

Add index

Primary key\* Partition key

EmpIDNumber ?

✓ Add sort key

SectionIDString ?

Index name\*EmpID-SectionID-index ?

Projected attributesAll ?

Create as Local Secondary Index ?

CancelAdd index

"EmpID" is my base table partition key, which I'm using in both GSI and LSI.

Primary key\* Partition key

EmpIDNumber ?

✓ Add sort key

DeptIDNumber ?

Table settings

Default settings provide the fastest way to get started with your table. You can modify these default settings now or after your table has been created.

Use default settings

Secondary indexes

Name	Type	Partition key	Sort key	Projected Attributes	
EmpID-S	GSI	EmpID (Num)	SectionID	ALL	✕
EmpID-G	LSI	EmpID (Num)	Gender	ALL	✕

[+ Add index](#)

Table is created successfully.

Filter by table name X								Viewing 1 of 1 Tables	
Name	Status	Partition key	Sort key	Indexes	Total read capacity	Total write capacity	Auto Scaling		
Employee	Active	EmplID (Number)	DeptID (Number)	2	10	10	READ_AND_WRITE		

Employee Close

Overview Items Metrics Alarms Capacity **Indexes** Backups Triggers Access control Tags

Create index Delete index

Name	Status	Type	Partition key	Sort key	Attributes	Read capacity	Write capacity	Size	Item
EmplID-SectionID-index	Active	GSI	EmplID (Number)	SectionID (String)	ALL	5	5	0	0
EmplID-Gender-index	N/A	LSI	EmplID (Number)	Gender (String)	ALL	-	-	0	0

So now we have seen that we can create GSI with the same partition key as the Base table.

- Global secondary index

(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/GSI.html>) – an index with a partition key and a sort key that can be different from those on the base table. A global secondary index is considered "global" because queries on the index can span all of the data in the base table, across all partitions.

- Local secondary index

(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/LSI.html>) – an index that has the same partition key as the base table, but a different sort key. A local secondary index is "local" in the sense that every partition of a local secondary index is scoped to a base table partition that has the same partition key value.

For more information on DynamoDB Indexes, please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SecondaryIndexes.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SecondaryIndexes.html>)

Ask our Experts



## QUESTION 36 UNATTEMPTED

What is used in DynamoDB to help a write operation only succeed if the item attributes meet one or more expected conditions? Choose answer from the options below.



- ☒ A. Conditional Writes ✓
- ☐ B. Eventual Consistency
- ☐ C. Extra API read calls to determine if the data was updated before the update call is made
- ☐ D. Atomic Counters

#### Explanation :

Answer – A

To help clients coordinate writes to data items, DynamoDB supports *conditional writes* for PutItem, DeleteItem, and UpdateItem operations. With a conditional write, an operation succeeds only if the item attributes meet one or more expected conditions; otherwise it returns an error

For more information on working with items , please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithItems.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithItems.html>)

Ask our Experts



#### QUESTION 37 UNATTEMPTED

True or False. Can a global secondary index created at the same time as the table creation.

- ☒ A. True ✓
- ☐ B. False

#### Explanation :

Answer - A

This is clearly given in the AWS documentation

<b>Online Index Operations</b>	Global secondary indexes can be created at the same time that you create a table. You can also add a new global secondary index to an existing table, or delete an existing global secondary index. For more information, see <a href="#">Managing Global Secondary Indexes</a> .	Local secondary indexes are created at the same time that you create a table. You cannot add a local secondary index to an existing table, nor can you delete any local secondary indexes that currently exist.
--------------------------------	---	---

For more information on DynamoDB Indexes, please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SecondaryIndexes.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SecondaryIndexes.html>)

Ask our Experts



QUESTION 38

UNATTEMPTED

Company B has a DynamoDB table where the average item size is 20KB. Company B anticipates the application will read 10 items from the table per second using eventually consistent reads. How much read capacity throughput should they provision? Choose the correct answer from the options below.

- ☐ A. 50
- ☐ B. 100
- ☒ C. 25 ✓
- ☐ D. 20

**Explanation :**

Answer – C

For read capacity, the rule is to divide the item size by 4KB. Hence we need to divide 20 by 4 which gives us 5 to the nearest 4KB. Since we are writing 10 items per second, we need to multiply  $10 \times 5 = 50$ . Since it is eventual consistency, we need to divide by 2 which gives us 25.

For more information on working with tables in dynamoDB , please visit the link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.html>)

Ask our Experts



QUESTION 39

UNATTEMPTED

What can be used in DynamoDB as part of the Query API call to filter results based on the values of primary keys? Choose an answer from the options below

- ☒ A. Expressions ✓
- ☐ B. Conditions
- ☐ C. Query API
- ☐ D. Scan API



### Explanation :

Answer - A

This is clearly provided in the AWS documentation

#### Q: Are expressions supported for key conditions?

Yes, you can specify an expression as part of the Query API call to filter results based on values of primary keys on a table using the `KeyConditionExpression` parameter.

For more information on DynamoDB please refer to the below link:

- <https://aws.amazon.com/dynamodb/faqs/> (<https://aws.amazon.com/dynamodb/faqs/>)

Ask our Experts



## QUESTION 40 UNATTEMPTED

What is the max size of an item that corresponds to a single write capacity unit.  
(When you create a table or index in Amazon DynamoDB, you must specify your capacity requirements for read and write activity)?.

Choose an answer from the options below.

- ☒ A. 1 KB ✓
- ☐ B. 4 KB
- ☐ C. 2 KB
- ☐ D. 8 KB

### Explanation :

Answer – A

This is clearly given in the AWS documentation

In DynamoDB, you specify provisioned throughput requirements in terms of *capacity units*. Use the following guidelines to determine your provisioned throughput:

- One *read capacity unit* represents one strongly consistent read per second, or two eventually consistent reads per second, for items up to 4 KB in size. If you need to read an item that is larger than 4 KB, DynamoDB will need to consume additional read capacity units. The total number of read capacity units required depends on the item size, and whether you want an eventually consistent or strongly consistent read.
- One *write capacity unit* represents one write per second for items up to 1 KB in size. If you need to write an item that is larger than 1 KB, DynamoDB will need to consume additional write capacity units. The total number of write capacity units required depends on the item size.

For more information on DynamoDB provisioned throughput please refer to the below link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput>.  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput>.ht

Ask our Experts



QUESTION 41 UNATTEMPTED

Which of the below functions is used in Cloudformation to return the value corresponding to keys in a two-level map that is declared in the Mappings section? Choose an answer from the options below

- ☐ A. Fn::GetAtt
- ☐ B. Fn::GetMap
- ☒ C. Fn::FindInMap ✓
- ☐ D. Fn::ReadMap

Explanation :

Answer - C

This is clearly given in the AWS documentation

**Fn::FindInMap**

Filter View: All ▼

The intrinsic function `Fn::FindInMap` returns the value corresponding to keys in a two-level map that is declared in the `Mappings` section.

For more information on the function please refer to the below link:

- <http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-findinmap.html>  
(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-findinmap.html>)

Ask our Experts



QUESTION 42 UNATTEMPTED

Which of the below functions is used in Cloudformation to append a set of values into a single value? Choose an answer from the options below





- ☐ A. Fn::GetAtt
- ☐ B. Fn::Combine
- ☐ C. Fn::Join ✓
- ☐ D. Fn::Select

**Explanation :**

Answer - C

This is clearly given in the AWS documentation

**Fn::Join**

Filter View: All ▼

The intrinsic function `Fn::Join` appends a set of values into a single value, separated by the specified delimiter. If a delimiter is the empty string, the set of values are concatenated with no delimiter.

For more information on the function please refer to the below link:

- <http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-join.html>  
(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-join.html>)

Ask our Experts



**QUESTION 43 UNATTEMPTED**

Which of the below functions is used in Cloudformation to retrieve an object from a set of objects? Choose an answer from the options below

- ☐ A. Fn::GetAtt
- ☐ B. Fn::Combine
- ☐ C. Fn::Join
- ☐ D. Fn::Select ✓

**Explanation :**

Answer - D

This is clearly given in the AWS documentation



## Fn::Select

Filter View: All ▼

The intrinsic function `Fn::Select` returns a single object from a list of objects by index.

### Important

`Fn::Select` does not check for null values or if the index is out of bounds of the array. Both conditions will result in a stack error, so you should be certain that the index you choose is valid, and that the list contains non-null values.

For more information on the function please refer to the below link:

- <http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-select.html>  
(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-select.html>)

Ask our Experts



## QUESTION 44 UNATTEMPTED

What result would you expect from the `Fn::Join` function in the following line in a CloudFormation template? Choose an answer from the options below

```
"Fn::Join": [ ".", [ "1", "2", "3" ] ]
```

- ☒ A. 1.2.3 ✓
- ☐ B. .1.2.3
- ☐ C. .1.2.3.
- ☐ D. 1.2.3.

### Explanation :

Answer – A

The intrinsic function `Fn::Join` appends a set of values into a single value, separated by the specified delimiter. If a delimiter is the empty string, the set of values are concatenated with no delimiter.

For more information on `Fn::Join` function please refer to the below link:

- <http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-join.html>  
(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-join.html>)

Ask our Experts



An administrator is getting an error while trying to create a new bucket in S3? You feel that bucket limit has been crossed. What is the bucket limit per account in AWS?

Choose the correct answer from the options below

- ☒ A. 100 ✓
- ☐ B. 50
- ☐ C. 1000
- ☐ D. 150

**Explanation :**

Answer – A

This is clearly mentioned in the AWS documentation.

**Amazon Simple Storage Service (Amazon S3) Limits**

Resource	Default Limit
Buckets	100 per account

For more information on AWS service limitations, please visit the link:

- [http://docs.aws.amazon.com/general/latest/gr/aws\\_service\\_limits.html](http://docs.aws.amazon.com/general/latest/gr/aws_service_limits.html)  
([http://docs.aws.amazon.com/general/latest/gr/aws\\_service\\_limits.html](http://docs.aws.amazon.com/general/latest/gr/aws_service_limits.html))

Ask our Experts



What result would you expect from the Fn::Select function in the following line in a CloudFormation template? Choose an answer from the options below

```
{ "Fn::Select": [ "1", [ "1", "2", "3", "4" ] ] }
```

- ☐ A. 1
- ☒ B. 2 ✓
- ☐ C. 3
- ☐ D. 4



### Explanation :

Answer – B

The intrinsic function Fn::Select returns a single object from a list of objects by index. In this case, Fn::Select returns 2 which is Answer B.

For more information on intrinsic functions, Please check

<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference.html> (<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference.html>)

Intrinsic function Fn::Select reference:

<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-select.html> (<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-select.html>)

Ask our Experts



QUESTION 47

UNATTEMPTED

What result would you expect from the Fn::split function in the following line in a CloudFormation template? Choose an answer from the options below

```
{ "Fn::Split" : [ "|", "a|b|c" ] }
```

- ☒ A. ["a", "b", "c"] ✓
- ☐ B. ["a|", "b|", "c|"]
- ☐ C. "abc"
- ☐ D. ["a"]

### Explanation :

Answer – A

To split a string into a list of string values so that you can select an element from the resulting string list, use the Fn::Split intrinsic function. Specify the location of splits with a delimiter, such as , (a comma).

After you split a string, use the Fn::Select

(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-select.html>) function to pick a specific element.

For more information on intrinsic functions , please visit the below URL:

- <https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-split.html>
- (<https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/intrinsic-function-reference-split.html>)

Ask our Experts



QUESTION 48 UNATTEMPTED

Which of the below API calls can be used to get information about stacks based on a specific filter. Choose an answer from the options below

- ☒ A. ListStacks ✓
- ☐ B. GetStacks
- ☐ C. GetAllStacks
- ☐ D. ListFilterStacks

Explanation :

Answer – A

This is clearly given in the AWS documentation

**ListStacks**

Returns the summary information for stacks whose status matches the specified StackStatusFilter. Summary information for stacks that have been deleted is kept for 90 days after the stack is deleted. If no StackStatusFilter is specified, summary information for all stacks is returned (including existing stacks and stacks that have been deleted).

For more information on the API, please visit the below URL:

- [http://docs.aws.amazon.com/AWSCloudFormation/latest/APIReference/API\\_ListStacks.html](http://docs.aws.amazon.com/AWSCloudFormation/latest/APIReference/API_ListStacks.html)  
([http://docs.aws.amazon.com/AWSCloudFormation/latest/APIReference/API\\_ListStacks.html](http://docs.aws.amazon.com/AWSCloudFormation/latest/APIReference/API_ListStacks.html))

Ask our Experts



QUESTION 49 UNATTEMPTED

When designing Cloudformation templates, is the Resources section optional?

- ☐ A. True
- ☒ B. False ✓



### Explanation :

Answer – B

If you refer to the documentation, you will see that Resources is the only mandatory field Specifies the stack resources and their properties, such as an Amazon Elastic Compute Cloud instance or an Amazon Simple Storage Service bucket.

For more information on cloudformation templates, please refer to the below link:

- <http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/template-anatomy.html>  
(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/template-anatomy.html>)

Ask our Experts



## QUESTION 50 UNATTEMPTED

When a failure occurs in a stack creation in Cloudformation, does a rollback occur?

- ☒ A. True ✓
- ☐ B. False

### Explanation :

Answer – A

By default, the “automatic rollback on error” feature is enabled. This will cause all AWS resources that AWS CloudFormation created successfully for a stack up to the point where an error occurred to be deleted. This is useful when, for example, you accidentally exceed your default limit of Elastic IP addresses, or you don’t have access to an EC2 AMI you’re trying to run. This feature enables you to rely on the fact that stacks are either fully created, or not at all, which simplifies system administration and layered solutions built on top of AWS CloudFormation.

For more information on cloudformation, please refer to the below link:

- <https://aws.amazon.com/cloudformation/faqs/>  
(<https://aws.amazon.com/cloudformation/faqs/>)

Ask our Experts



## QUESTION 51 UNATTEMPTED

Which of the below configuration softwares can Cloudformation be used with?

Choose 2 answers from the options below

- ☐ A. Chef ✓



- ☐ B. Ansible
- ☐ C. Powershell
- ☐ D. Puppet ✓

**Explanation :**

Answer – A and D

This is clearly mentioned in the AWS documentation

**Q: Can I use AWS CloudFormation with Chef?**

Yes. AWS CloudFormation can be used to bootstrap both the Chef Server and Chef Client software on your EC2 instances. For more details and a how-to see [Integrating AWS CloudFormation with Chef](#).

**Q: Can I use AWS CloudFormation with Puppet?**

Yes. AWS CloudFormation can be used to bootstrap both the Puppet Master and Puppet Client software on your EC2 instances. For more details and a how-to see [Integrating AWS CloudFormation with Puppet](#).

For more information on cloudformation, please refer to the below link:

- <https://aws.amazon.com/cloudformation/faqs/>  
(<https://aws.amazon.com/cloudformation/faqs/>)

Ask our Experts



QUESTION 52

UNATTEMPTED

In DynamoDB you get the following error

“ItemCollectionSizeLimitExceededException”. What could be the possible cause for this error. Choose an answer from the options below

- ☐ A. Number of tables have exceeded in DynamoDB
- ☐ B. Number of Indexes have increased
- ☐ C. Size of a group of items with the same partition key value has exceeded 10GB. ✓
- ☐ D. None of the above

**Explanation :**

Answer - C

This is clearly given in the AWS documentation



### ItemCollectionSizeLimitExceededException

Message: *Collection size exceeded.*

For a table with a local secondary index, a group of items with the same partition key value has exceeded the maximum size limit of 10 GB. For more information on item collections, see [Item Collections](#).

For more information on DynamoDB errors, please refer to the below link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Programming.Errors.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Programming.Errors.html>)

Ask our Experts



QUESTION 53

UNATTEMPTED

One DynamoDB strongly consistent read capacity unit is equal to two eventual consistent read per second. Choose the correct answer from the options below

- ☐ A. True ✓
- ☐ B. False

#### Explanation :

Answer – A

This is clearly mentioned in the AWS documentation.

- One *read capacity unit* represents one strongly consistent read per second, or two eventually consistent reads per second, for items up to 4 KB in size. If you need to read an item that is larger than 4 KB, DynamoDB will need to consume additional read capacity units. The total number of read capacity units required depends on the item size, and whether you want an eventually consistent or strongly consistent read.

For more information on DynamoDB read consistency please visit the below link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html>  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.html>)

Ask our Experts



QUESTION 54

UNATTEMPTED





Your items are 15KB in size and you want to have 100 strongly consistent reads per second. How many read capacity units do you need to provision? Choose the correct answer from the options below

- ☐ A. 100
- ☐ B. 200
- ☐ C. 300
- ☒ D. 400 ✓

**Explanation :**

Answer – D

Since each item is 15KB and if we near it to the nearest 4KB , we get 4. Since we need 100 strongly consistent , it becomes  $(100 * 4) = 400$ .

For more information on DynamoDB provisioned throughput please visit the below link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput>.  
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ProvisionedThroughput.ht>

Ask our Experts



QUESTION 55

UNATTEMPTED

Is the default visibility timeout for an SQS queue 1 minute?

- ☐ A. True
- ☒ B. False ✓

**Explanation :**

Answer – B

Each queue starts with a default setting of 30 seconds for the visibility timeout. You can change that setting for the entire queue. Typically, you'll set the visibility timeout to the average time it takes to process and delete a message from the queue. When receiving messages, you can also set a special visibility timeout for the returned messages without changing the overall queue timeout.

For more information on SQS please visit the below link:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



You decide to create a bucket on AWS S3 called 'newbucket' and then perform the following actions in the order that they are listed here.

- You upload a file to the bucket called 'file1'
- You upload a file called 'file2'
- You enable versioning on the bucket

Which of the following is true for your bucket 'newbucket'? Choose the correct answer from the options below

- ☐ A. There will be 1 version ID for file1 and file2.
- ☐ B. There will be 1 version ID for file1 and null version ID for file2.
- ☐ C. There will be 1 version ID for file2 and null version ID for file1.
- ☒ D. The version ID will be null for both file1 and file2 ✓

#### Explanation :

Answer – D

Any objects uploaded prior to versioning will have the version ID as NULL.

This is clearly mentioned in the AWS documentation.

- Objects stored in your bucket before you set the versioning state have a version ID of `null`. When you enable versioning, existing objects in your bucket do not change. What changes is how Amazon S3 handles the objects in future requests. For more information, see [Managing Objects in a Versioning-Enabled Bucket](#).

For more information on S3 versioning, please visit the below link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/Versioning.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/Versioning.html>)

Ask our Experts



Is Example.com a valid S3 bucket name?

- ☐ A. True
- ☒ B. False ✓



### Explanation :

Answer – B

Below are the naming conventions for S3 buckets

- The bucket name can be between 3 and 63 characters long, and can contain only lower-case characters, numbers, periods, and dashes.
- Each label in the bucket name must start with a lowercase letter or number.
- The bucket name cannot contain underscores, end with a dash, have consecutive periods, or use dashes adjacent to periods.
- The bucket name cannot be formatted as an IP address (198.51.100.24).

For more information on the S3 naming conventions please visit the below link:

- <http://docs.aws.amazon.com/awsccloudtrail/latest/userguide/cloudtrail-s3-bucket-naming-requirements.html> (<http://docs.aws.amazon.com/awsccloudtrail/latest/userguide/cloudtrail-s3-bucket-naming-requirements.html>)

Ask our Experts



QUESTION 58

UNATTEMPTED

Which of the descriptions below best describes what the following bucket policy does?

```
{
  "Version": "2012-10-17",
  "Id": "Statement1",
  "Statement": [
    {
      "Sid": "Statement2",
      "Effect": "Allow",
      "Principal": "*",
      "Action": [
        "s3:GetObject",
        "s3:PutObject"
      ],
      "Resource": "arn:aws:s3:::mybucket/*",
      "Condition": {
        "StringLike": {
          "aws:Referer": [
            "http://www.example.com/*",
```



```
"http://www.demo.com/*"  
  ]  
}  
}  
}  
]  
}
```

Choose the correct answer from the options below

- ☐ A. It allows read or write actions on the bucket 'mybucket'
- ☐ B. It allows read access to the bucket 'mybucket' but only if it is accessed from example.com or demo.com
- ☐ C. It allows unlimited access to the bucket 'mybucket'
- ☐ D. It allows read or write access to the bucket 'mybucket' but only if it is accessed from example.com or demo.com ✓

#### Explanation :

Answer – D

The PutObject allows one to put objects in an S3 bucket.

s3:ListMultipartUploadParts	List Parts
s3:PutObject	PUT Object, POST Object, Initiate Multipart Upload, Upload Part, Complete Multipart Upload, PUT Object - Copy

For more information on S3 bucket policy examples, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/example-bucket-policies.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/example-bucket-policies.html>)

Ask our Experts



QUESTION 59 UNATTEMPTED

While working with the AWS API you receive the following error message: 409 Conflict. What might be the cause of this error?

- ☐ A. BadDigest
- ☐ B. User does not have proper permissions to make the API call



- ☒ C. Bucket is not empty ✓
- ☐ D. Bucket name does not exist

**Explanation :**

Answer – C

This is clearly provided as part of the S3 error codes in the S3 documentation.

BucketNotEmpty	The bucket you tried to delete is not empty.	409 Conflict	Client
----------------	--	--------------	--------

For more information on S3 Error codes, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/API/ErrorResponse.html#ErrorCodeList>  
(<http://docs.aws.amazon.com/AmazonS3/latest/API/ErrorResponse.html#ErrorCodeList>)

Ask our Experts



**QUESTION 60 UNATTEMPTED**

What is used in S3 to enable client web applications that are loaded in one domain to interact with resources in a different domain? Choose the correct answer from the options below

- ☒ A. CORS Configuration ✓
- ☐ B. Public Object Permissions
- ☐ C. Public ACL Permissions
- ☐ D. None of the above

**Explanation :**

Answer – A

Cross-origin resource sharing (CORS) defines a way for client web applications that are loaded in one domain to interact with resources in a different domain. With CORS support in Amazon S3, you can build rich client-side web applications with Amazon S3 and selectively allow cross-origin access to your Amazon S3 resources.

For more information on S3 CORS configuration, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/cors.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/cors.html>)

Ask our Experts



---

Finish Review (<https://www.whizlabs.com/learn/course/aws-cda-practice-tests/quiz/13231>)

## Certification

- 🔗 Cloud Certification  
(<https://www.whizlabs.com/cloud-certification-training-courses/>)
- 🔗 Java Certification  
(<https://www.whizlabs.com/oracle-java-certifications/>)
- 🔗 PM Certification  
(<https://www.whizlabs.com/project-management-certifications/>)
- 🔗 Big Data Certification  
(<https://www.whizlabs.com/big-data-certifications/>)

## Mobile App

 Android Coming Soon

 iOS Coming Soon

## Company

- 🔗 Support  
(<https://help.whizlabs.com/hc/en-us>)
- 🔗 Discussions (<http://ask.whizlabs.com/>)
- 🔗 Blog (<https://www.whizlabs.com/blog/>)



## Follow us



(<https://www.facebook.com/whizlabs.software/>)



(<https://in.linkedin.com/company/whizlabs-software>)



(<https://twitter.com/whizlabs?lang=en>)



(<https://plus.google.com/+WhizlabsSoftware>)

---

© Copyright 2018. Whizlabs Software Pvt. Ltd. All Rights Reserved.

