



- [🏠 \(https://www.whizlabs.com/learn\)](https://www.whizlabs.com/learn) > [My Courses \(https://www.whizlabs.com/learn/my-courses\)](https://www.whizlabs.com/learn/my-courses)
> [AWS Certified Solutions Architect Associate \(https://www.whizlabs.com/learn/course/aws-csaa-practice-tests#section-1\)](https://www.whizlabs.com/learn/course/aws-csaa-practice-tests#section-1)
> [New Diagnostic Test \(https://www.whizlabs.com/learn/course/aws-csaa-practice-tests/quiz/13954\)](https://www.whizlabs.com/learn/course/aws-csaa-practice-tests/quiz/13954) > **Report**

NEW DIAGNOSTIC TEST

Attempt 2
Marks Obtained 45 / 60
Your score is 75%

Completed on Friday , 25 January 2019 , 03:18 PM
Time Taken 00 H 43 M 57 S
Result Pass

Domains / Topics wise Quiz Performance Report

S.No.	Topic	Total Questions	Correct	Incorrect	Unattempted
1	Other	60	45	15	0

60 Questions	45 Correct	15 Incorrect	0 Unattempted
------------------------	----------------------	------------------------	-------------------------

Show Answers

All	▼
-----	---

QUESTION 1 CORRECT

In AWS what is used for encrypting and decrypting login information to EC2 instances.

☐ A. Templates

- ☐ B. AMI's
- ☒ C. Key pairs ✓
- ☐ D. None of the above

Explanation :

Answer - C

Amazon EC2 uses public-key cryptography to encrypt and decrypt login information. Public-key cryptography uses a public key to encrypt a piece of data, such as a password, then the recipient uses the private key to decrypt the data. The public and private keys are known as a key pair.

To log in to your instance, you must create a key pair, specify the name of the key pair when you launch the instance, and provide the private key when you connect to the instance. Linux instances have no password, and you use a key pair to log in using SSH. With Windows instances, you use a key pair to obtain the administrator password and then log in using RDP.

For more information on key pairs, please visit the below url

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-key-pairs.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-key-pairs.html>)
(<https://aws.amazon.com/cloudfront/>)

Ask our Experts



QUESTION 2 CORRECT

For DynamoDB, what are the scenario's in which you would want to enable Cross-region replication?

- ☐ A. Live data migration
- ☐ B. Easier Traffic management
- ☐ C. Disaster Recovery
- ☒ D. All of the above ✓

Explanation :

Answer – D

From the AWS Documentation, it clearly states the reason for why you would want to enable Cross-

Region Replication

Q: When should I use cross-region replication?

You can use cross-region replication for the following scenarios.

- **Efficient disaster recovery:** By replicating tables in multiple data centers, you can switch over to using DynamoDB tables from another region in case a data center failure occurs.
- **Faster reads:** If you have customers in multiple regions, you can deliver data faster by reading a DynamoDB table from the closest AWS data center.
- **Easier traffic management:** You can use replicas to distribute the read workload across tables and thereby consume less read capacity in the master table.
- **Easy regional migration:** By creating a read replica in a new region and then promoting the replica to be a master, you migrate your application to that region more easily.
- **Live data migration:** To move a DynamoDB table from one region to another, you can create a replica of the table from the source region in the destination region. When the tables are in sync, you can switch your application to write to the destination region.

For more information on DynamoDB, please visit the url

- <https://aws.amazon.com/dynamodb/faqs/> (<https://aws.amazon.com/dynamodb/faqs/>)

Ask our Experts



QUESTION 3 CORRECT

You have launched two web servers in private subnet and one ELB (internet facing) in public subnet in your VPC. Yet, you are still unable to access your web application through the internet, which of the following would likely the cause of this? Choose two correct options

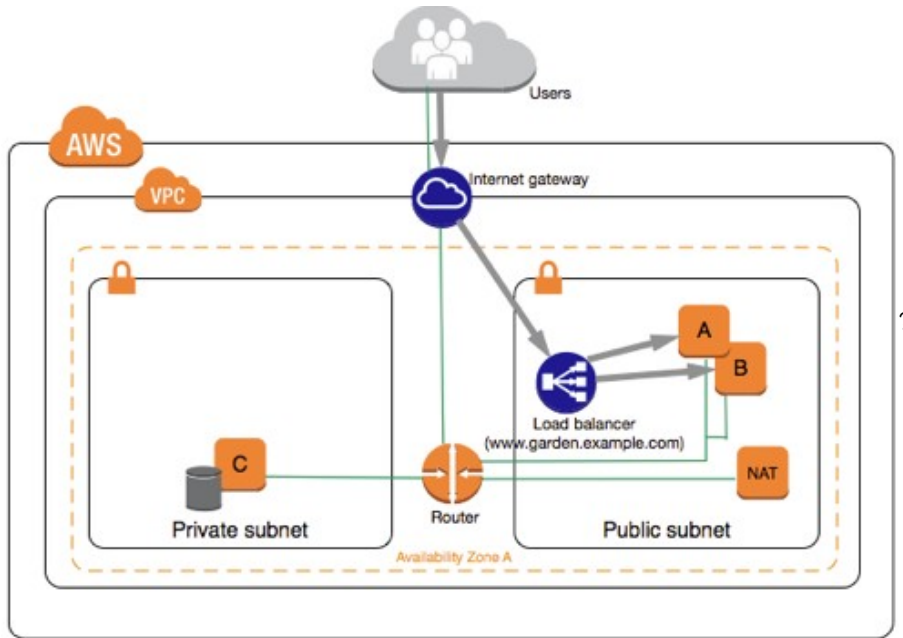
- ☐ A. Web server must be launched inside public subnet and not private subnet.
- ☒ B. Route table for public subnet is not configured to route to VPC internet gateway.
✓
- ☐ C. No elastic IP is assigned to web servers.

- ✓ **D. Add a rule on the instance's security group to allow traffic from the ELB's Security Group. ✓**

Explanation :

Answer - B and D

In order for the EC2 or ELBs to be accessible from internet, we would need to configure the route table for public subnet to route traffic to VPC internet gateway. For example:



For information on VPC Route Tables and VPC Internet Gateway, please visit the link:

- https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Route_Tables.html
(https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Route_Tables.html)
- https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Internet_Gateway.html
(https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Internet_Gateway.html)

Scenario is explained in following link with detail:

- <https://aws.amazon.com/premiumsupport/knowledge-center/public-load-balancer-private-ec2/> (<https://aws.amazon.com/premiumsupport/knowledge-center/public-load-balancer-private-ec2/>)

Ask our Experts



QUESTION 4**CORRECT**

You are developing a highly available web application using stateless web servers. Which services are suitable for storing session state data?

Choose 3 answers

- ☐ A. Amazon CloudWatch
- ☒ B. Amazon Relational Database Service (RDS) ✓
- ☐ C. Elastic Load Balancing
- ☒ D. Amazon ElastiCache ✓
- ☐ E. AWS Storage Gateway
- ☒ F. Amazon DynamoDB ✓

Explanation :

Answer – B, D and F

Please find the AWS Documentation references for Elastic Cache and DynamoDB. Relational databases have always been a source for storing session data.

Amazon ElastiCache is a web service that makes it easy to deploy, operate, and scale an in-memory data store or cache in the cloud (<https://aws.amazon.com/what-is-cloud-computing/>). The service improves the performance of web applications by allowing you to retrieve information from fast, managed, in-memory data stores, instead of relying entirely on slower disk-based databases.

For more information on Elastic Cache , please refer to the below link

- <https://aws.amazon.com/elasticache/> (<https://aws.amazon.com/elasticache/>)

An example of managing session state via DynamoDB is given below

- <http://docs.aws.amazon.com/sdk-for-net/v2/developer-guide/dynamodb-session-net-sdk.html> (<http://docs.aws.amazon.com/sdk-for-net/v2/developer-guide/dynamodb-session-net-sdk.html>)

Ask our Experts



QUESTION 5

CORRECT

Which of the following is mandatory when defining a cloudformation template?

- ☒ A. Resources ✓
- ☐ B. Parameters
- ☐ C. Outputs
- ☐ D. Mappings

Explanation :

Answer – A

This is clearly given in the aws documentation

Resources (required)

Specifies the stack resources and their properties, such as an Amazon Elastic Compute Cloud instance or an Amazon Simple Storage Service bucket. You can refer to resources in the `Resources` and `Outputs` sections of the template.

For more information on Cloudformation templates, please visit the url

<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/template-anatomy.html>

(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/template-anatomy.html>)

Ask our Experts



QUESTION 6

MARKED AS REVIEW

CORRECT

In IAM, what is the representation of a person or service ?

- ☒ A. User ✓
- ☐ B. Group
- ☐ C. Team
- ☐ D. Role

Explanation :

Answer – A

An IAM *user* is an entity that you create in AWS. The IAM user represents the person or service who uses the IAM user to interact with AWS

An IAM *group* is a collection of IAM users. You can use groups to specify permissions for a collection of users, which can make those permissions easier to manage for those users

An IAM *role* is very similar to a user, in that it is an identity with permission policies that determine what the identity can and cannot do in AWS

For more information on IAM entities, please visit the url

<http://docs.aws.amazon.com/IAM/latest/UserGuide/id.html>

(<http://docs.aws.amazon.com/IAM/latest/UserGuide/id.html>)

Note: AWS says that "An IAM *user* is an entity that you create in AWS to represent the person or service that uses it to interact with AWS. A user in AWS consists of a name and credentials."

Option A is meant to be IAM user, not just user. I can understand that its bit unclear. However AWS frames questions like this. So we are keeping it like that for consistency.

Please check the below link to know more about it.

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_users.html

(https://docs.aws.amazon.com/IAM/latest/UserGuide/id_users.html)

Note: In AWS terminology, User relates to person/service(Single) But a role relates to MANY people/services.

Ask our Experts



QUESTION 7 INCORRECT

Which of the following is not supported by AWS Import/Export? Choose 2 options

- ☐ A. Import to Amazon S3
- ☐ B. Export from Amazon S3
- ☒ C. Import to Amazon EBS ✕
- ☐ D. Import to Amazon Glacier ✓

✓ E. Export from Amazon Glacier ✓

Explanation :

Answer – D and E

NOTE: Import-export is no longer used. This question is for old version of exam.

This will be retired once AWS removes its old version of exam.

The AWS documentation mentions the following

AWS Import/Export accelerates transferring data between the AWS cloud and portable storage devices that you mail to us. AWS Import/Export is a good choice if you have 16 terabytes (TB) or less of data to import into Amazon Simple Storage Service or Amazon Elastic Block Store (Amazon EBS).

You can also export data from Amazon S3 with AWS Import/Export.

For more information on AWS Import/Export , please refer to the below link:

- <https://docs.aws.amazon.com/amazonglacier/latest/dev/uploading-an-archive.html>
(<https://docs.aws.amazon.com/amazonglacier/latest/dev/uploading-an-archive.html>)
- <http://docs.aws.amazon.com/AWSImportExport/latest/DG/whatisdisk.html>
(<http://docs.aws.amazon.com/AWSImportExport/latest/DG/whatisdisk.html>)

Since October 2017, Import/Export disk service, **no longer support the Direct import into Amazon Glacier.**

Before Amazon Glacier data can be exported it needs to be restored to Amazon S3 using the S3 Lifecycle Restore feature.

More information is available at:

<https://docs.aws.amazon.com/AWSImportExport/latest/DG/AWSImportExport-dg.pdf>
(<https://docs.aws.amazon.com/AWSImportExport/latest/DG/AWSImportExport-dg.pdf>)

Ask our Experts



QUESTION 8 CORRECT

Your company is hosting their application in AWS and there is a need in analyzing the data collected from daily operation. Currently there are 4 reserved instances that are running the analytics task. As a solution architect of the company, you are asked to propose a cost effective solution that would help to expedite the analytics process. Which of the following instance type is the best for this situation?

○ A. Spot Instances ✓

- ☐ B. On-Demand Instances
- ☐ C. Dedicated Instances
- ☐ D. This is not possible in AWS

Explanation :

Answer – A

This is clearly given in the aws documentation

Improve Application Throughput

You can run and scale applications such as stateless web services, image rendering, big data analytics and massively parallel computations on Spot instances. Since Spot instances typically cost 50-90% less, you can increase your compute capacity by 2-10x within the same budget.

For more information on Spot Instances, please visit the link –

<https://aws.amazon.com/ec2/spot/> (<https://aws.amazon.com/ec2/spot/>)

Ask our Experts



QUESTION 9 CORRECT

Which of the below are incremental backups of your EBS volumes? Choose one answer from the options given below.

- ☐ A. Volumes
- ☐ B. State Manager
- ☐ C. Placement Groups
- ☒ D. Snapshots ✓

Explanation :

Answer - D

You can easily create a snapshot from a volume while the instance is running and the volume is in use. You can do this from the EC2 dashboard.



For more information on EBS snapshots, please visit the link -

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>

(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>)

Ask our Experts



QUESTION 10

MARKED AS REVIEW

INCORRECT

There is a requirement to host a NoSQL database with a need for low latency.

Which class of instances from the below list should they choose from ?

- ☒ A. T2 ✕
- ☐ B. I2 ✓
- ☐ C. T1
- ☐ D. G2

Explanation :

Answer - B

I2 instances are optimized to deliver tens of thousands of low-latency, random I/O operations per second (IOPS) to applications. They are well suited for the following scenarios:

- NoSQL databases (for example, Cassandra and MongoDB)
- Clustered databases
- Online transaction processing (OLTP) systems

For more information on I2 instances, please visit the link:

- <https://aws.amazon.com/blogs/aws/amazon-ec2-new-i2-instance-type-available-now/>
(<https://aws.amazon.com/blogs/aws/amazon-ec2-new-i2-instance-type-available-now/>)

Note:

It is now changed to I3... not I2 anymore.

Please note that in the actual exam, some of the questions might not be up-to-date. Especially after the 2017 AWS events, lots of features have been revised and improved. If you do not see a "is correct" answer, please choose the "was correct" answer.

You may refer to the latest EC2 instance type from the link below:

<https://aws.amazon.com/ec2/instance-types/> (<https://aws.amazon.com/ec2/instance-types/>)

Ask our Experts



QUESTION 11

MARKED AS REVIEW

CORRECT

You are designing a site for a new start up which generates cartoon images for people automatically.

Customers will log on to the site, upload an image which is stored in S3. The application then passes a job to AWS SQS and a fleet of EC2 instances poll the queue to receive new processing jobs. These EC2 instances will then turn the picture in to a cartoon and will then need to store the processed job somewhere. Users will typically download the image once (immediately), and then never download the image again.

What is the most commercially feasible method to store the processed images?

- ☐ A. Rather than use S3, store the images inside a BLOB on RDS with Multi-AZ configured for redundancy.
- ☒ B. Store the images on S3 RRS, and create a lifecycle policy to delete the image after 24 hours. ✓
- ☐ C. Store the images on glacier instead of S3.
- ☐ D. Use elastic block storage volumes to store the images.

Explanation :

Answer B

Explanation:

Use the AWS Reduced Redundancy storage to save on costs. The use lifecycle policies to delete the data since it is not required.

For more information on AWS Reduced Redundancy storage , please refer to the below link

- <https://aws.amazon.com/s3/reduced-redundancy/> (<https://aws.amazon.com/s3/reduced-redundancy/>)

The AWS Documentation mentions the following on Lifecycle policies

Lifecycle configuration enables you to specify the lifecycle management of objects in a bucket. The configuration is a set of one or more rules, where each rule defines an action for Amazon S3 to apply to a group of objects

For more information on S3 Lifecycle policies , please refer to the below link

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html> (<http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>)

Ask our Experts



QUESTION 12

INCORRECT

You have a high performance compute application and you need to minimize network latency between EC2 instances as much as possible. What can you do to achieve this?

- ☒ A. Use Elastic Load Balancing to load balance traffic between availability zones ✕
- ☐ B. Create a CloudFront distribution and to cache objects from an S3 bucket at Edge Locations.
- ☐ C. Create a placement group within an Availability Zone and place the EC2 instances within that placement group. ✓
- ☐ D. Deploy your EC2 instances within the same region, but in different subnets and different availability zones so as to maximize redundancy.

Explanation :

Answer - C

The AWS Documentation mentions the following on placement Groups

A *placement group* is a logical grouping of instances within a single Availability Zone. Placement groups are recommended for applications that benefit from low network latency, high network throughput, or both. To provide the lowest latency, and the highest packet-per-second network performance for your placement group, choose an instance type that supports enhanced networking.

For more information on placement groups , please refer to the below link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>)

Ask our Experts



QUESTION 13

CORRECT

Which of the below elements can you manage in the Billing dashboard ?

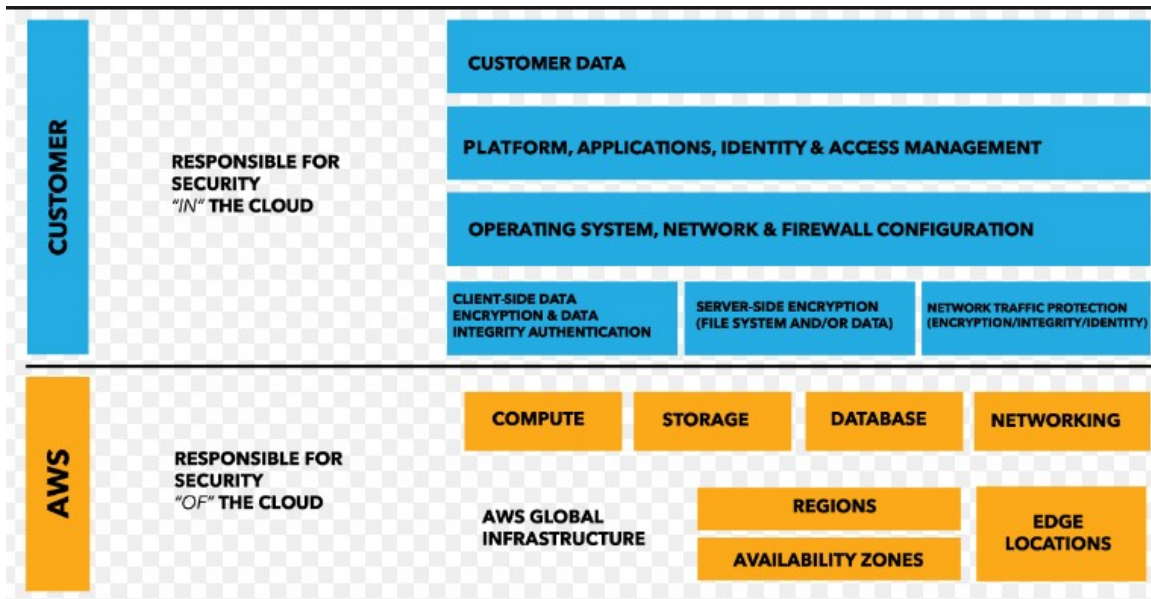
Select 2 options.

- ☒ A. Budgets ✓
- ☐ B. Policies
- ☐ C. Credential Report
- ☒ D. Cost Explorer ✓

Explanation :

Answer – A and D

When you go to your Billing dashboard, below are the set of elements which can be configured.



For more information on AWS cloud billing and pricing, please visit the link
<http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/billing-getting-started.html>
(<http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/billing-getting-started.html>)

Ask our Experts



QUESTION 14 CORRECT

What is the name of the VPC that is automatically created for your AWS account for the first time ?

- ☐ A. Primary VPC
- ☐ B. First VPC
- ☒ C. Default VPC ✓
- ☐ D. Initial VPC

Explanation :

Answer – C

A default VPC is a logically isolated virtual network in the AWS cloud that is automatically created for your AWS account the first time you provision Amazon EC2 resources. When you launch an instance without specifying a subnet-ID, your instance will be launched in your default VPC.

Default VPCs

Q. What is a default VPC?

A default VPC is a logically isolated virtual network in the AWS cloud that is automatically created for your AWS account the first time you provision Amazon EC2 resources. When you launch an instance without specifying a subnet-ID, your instance will be launched in your default VPC.

Q. What are the benefits of a default VPC?

When you launch resources in a default VPC, you can benefit from the advanced networking functionalities of Amazon VPC (EC2-VPC) with the ease of use of Amazon EC2 (EC2-Classic). You can enjoy features such as changing security group membership on the fly, security group egress filtering, multiple IP addresses, and multiple network interfaces without having to explicitly create a VPC and launch instances in the VPC.

For more information on VPC, please refer to the link

- <https://aws.amazon.com/vpc/faqs/> (<https://aws.amazon.com/vpc/faqs/>)

Ask our Experts



QUESTION 15

CORRECT

Which of the following databases support the read replica feature?

Select 3 options.

- ☒ A. MySQL ✓
- ☒ B. MariaDB ✓
- ☒ C. PostgreSQL ✓
- ☐ D. Oracle

Explanation :

Answer – A, B and C

Read replicas are available in Amazon RDS for MySQL, MariaDB, and PostgreSQL. When you create a read replica, you specify an existing DB Instance as the source. Amazon RDS takes a snapshot of the source instance and creates a read-only instance from the snapshot. For MySQL, MariaDB and PostgreSQL, Amazon RDS uses those engines' native asynchronous replication to update the read replica whenever there is a change to the source DB instance.

The read replica operates as a DB instance that allows only read-only connections; applications can connect to a read replica just as they would to any DB instance. Amazon RDS replicates all databases in the source DB instance.

Note: Microsoft SQL and Oracle not supports the read replica.

For more information on rds read replicas, please refer to the link

- <https://aws.amazon.com/rds/details/read-replicas/>
(<https://aws.amazon.com/rds/details/read-replicas/>)

Ask our Experts



QUESTION 16

CORRECT

What can be used from AWS to import existing Virtual Machines Images into AWS?

- ☒ A. VM Import/Export ✓
- ☐ B. AWS S3
- ☐ C. AWS Storage Gateway
- ☐ D. This is not possible in AWS

Explanation :

Answer – A

VM Import/Export enables customers to import Virtual Machine (VM) images in order to create Amazon EC2 instances. Customers can also export previously imported EC2 instances to create VMs. Customers can use VM Import/Export to leverage their previous investments in building VMs by migrating their VMs to Amazon EC2.

Using the Import/Export Tools

VM Import/Export offers several ways to import your virtual machine into Amazon EC2.

The first method is to import your VM image using the AWS CLI tools. To get started, simply:

- Download and [install the AWS Command Line Interface](#).
- Verify that your VM satisfies the [prerequisites for VM Import](#), prepare it for import, and export it from its current environment as an OVA file (or VMDK, VHD, or RAW).
- Upload the VM image to S3 using the AWS CLI. Multi-part uploads will provide improved performance. As an alternative, you can also send the VM image to AWS using the [AWS Import](#) service.
- Once the VM image is uploaded, import your VM using the `ec2 import-image` command. As part of this command, you can specify the licensing model and other parameters for your imported image.
- Use the `ec2 describe-import-image-tasks` command to monitor the import progress.
- Once your import task is completed, you can use the `ec2 run-instances` command to create an Amazon EC2 instance from the AMI generated during the import process.

Alternatively, if you use the VMware vSphere virtualization platform, you can use the AWS Management Portal for vCenter, which provides you a simple graphical user interface to import your virtual machines. You can [learn more about the AWS Management Portal for vCenter here](#).

For more information on AWS EC2, please visit

- <https://aws.amazon.com/ec2/faqs/> (<https://aws.amazon.com/ec2/faqs/>)
- <https://aws.amazon.com/ec2/vm-import/> (<https://aws.amazon.com/ec2/vm-import/>)

Ask our Experts



QUESTION 17

CORRECT

What is the service used by AWS to segregate control over the various AWS services ?

- ☐ A. AWS RDS
- ☐ B. AWS Integrity Management

☒ C. AWS Identity and Access Management ✓

☐ D. Amazon EMR

Explanation :

Answer – C

AWS Identity and Access Management (IAM) is a web service that helps you securely control access to AWS resources for your users. You use IAM to control who can use your AWS resources (authentication) and what resources they can use and in what ways (authorization).

For more information on IAM, please visit:

- <http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>
(<http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>)

Ask our Experts



QUESTION 18

CORRECT

Which of the following programming languages have an officially supported AWS SDK?

Select 2 options.

☒ A. PHP ✓

☐ B. Pascal

☒ C. Java ✓

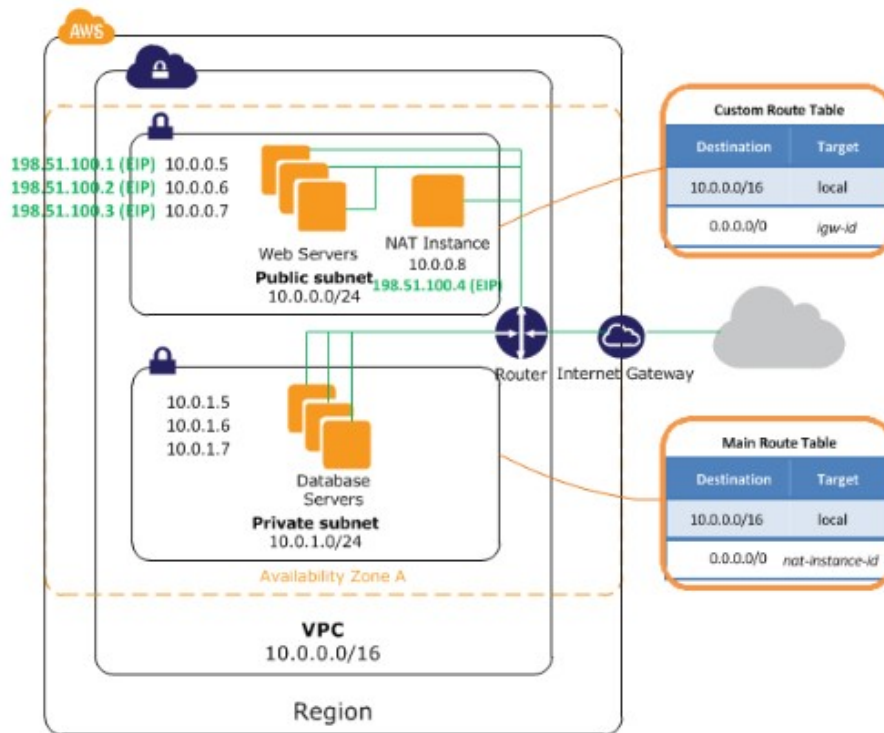
☐ D. SQL

☐ E. Perl

Explanation :

Answer - A and C

This is as per the AWS documentation



For more information on aws toolkits available, please refer to the below url:

- <https://aws.amazon.com/tools/> (<https://aws.amazon.com/tools/>)

Ask our Experts



QUESTION 19 CORRECT

How long can messages live in a SQS queue ?

- ☐ A. 12 hours
- ☐ B. 10 days
- ☒ C. 14 days ✓
- ☐ D. 1 year

Explanation :

Answer – C

This is clearly given in the AWS documentation

#####

Q: How long can I keep my messages in Amazon SQS message queues?

Longer message retention provides greater flexibility to allow for longer intervals between message production and consumption.

You can configure the Amazon SQS message retention period to a value from 1 minute to 14 days.

The default is 4 days. Once the message retention limit is reached, your messages are automatically deleted.

#####

For more information on SQS , please visit the following url

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 20

CORRECT

You working in the media industry and you have created a web application where users will be able to upload photos they create to your website. This web application must be able to call the S3 API in order to be able to function. Where should you store your API credentials whilst maintaining the maximum level of security?

- ☐ A. Save the API credentials to your php files.
- ☒ B. Don't save your API credentials. Instead create a role in IAM and assign this role to an EC2 instance when you first create it. ✓
- ☐ C. Save your API credentials in a public Github repository.
- ☐ D. Pass API credentials to the instance using instance userdata.

Explanation :

Answer – B

Always use IAM Roles for accessing AWS resources from EC2 Instances

The AWS Documentation mentions the following

IAM roles are designed so that your applications can securely make API requests from your instances, without requiring you to manage the security credentials that the applications use. Instead of creating and distributing your AWS credentials, you can delegate permission to make API requests using IAM roles

For more information on IAM Roles for EC2 Instances, please refer to the below link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/iam-roles-for-amazon-ec2.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/iam-roles-for-amazon-ec2.html>)

Ask our Experts



QUESTION 21

CORRECT

What are some of the common causes why you cannot connect to a DB instance on AWS ?

Select 3 options.

- ☐ A. There is a read replica being created, hence you cannot connect
- ☒ B. The DB is still being created ✓
- ☒ C. The local firewall is stopping the communication traffic ✓
- ☒ D. The security groups for the DB are not properly configured. ✓

Explanation :

Answer – B, C and D

There are some steps clearly given in the AWS documentation:

When you cannot connect to a DB instance, the following are common causes:

- The access rules enforced by your local firewall and the ingress IP addresses that you authorized to access your DB instance in the instance's security group are not in sync. The problem is most likely the ingress rules in your security group. By default, DB instances do not allow access; access is granted through a security group. To grant access, you must create your own security group with specific ingress and egress rules for your situation. For more information about setting up a security group, see [Provide Access to the DB Instance in the VPC by Creating a Security Group](#).
- The port you specified when you created the DB instance cannot be used to send or receive communications due to your local firewall restrictions. In this case, check with your network administrator to determine if your network allows the specified port to be used for inbound and outbound communication.
- Your DB instance is still being created and is not yet available. Depending on the size of your DB instance, it can take up to 20 minutes before an instance is available.

For more information on rds troubleshooting please visit the below link

- http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Troubleshooting.html
(http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Troubleshooting.html)

Ask our Experts



QUESTION 22

MARKED AS REVIEW

CORRECT

Which of the following RDS databases safeguard your data by performing a synchronous physical replication to keep data on the standby up-to-date with the primary in a Multi-AZ deployment ?

Select 3 options.

- ☐ A. SQL Server
- ☒ B. MySQL ✓
- ☒ C. Oracle ✓
- ☒ D. MariaDB ✓

Explanation :

Answer – B, C and D

Multi-AZ deployments for the MySQL, MariaDB, Oracle, and PostgreSQL engines utilize synchronous physical replication to keep data on the standby up-to-date with the primary.

Multi-AZ deployments for the SQL Server engine use synchronous logical replication to achieve the same result, employing SQL Server-native Mirroring technology.
(<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZ.html>)

Ask our Experts



QUESTION 23 CORRECT

Which is the service provided by AWS for providing a petabyte-scale data warehouse ?

- ☐ A. Amazon DynamoDB
- ☒ B. Amazon Redshift ✓
- ☐ C. Amazon Kinesis
- ☐ D. Amazon Simple Queue Service

Explanation :

Answer – B

Amazon Redshift is a fast, fully managed, petabyte-scale data warehouse that makes it simple and cost-effective to analyze all your data using your existing business intelligence tools. Start small for \$0.25 per hour with no commitments and scale to petabytes for \$1,000 per terabyte per year, less than a tenth the cost of traditional solutions.

Option A is wrong because it is used as a NOSQL solution.

Option C is wrong because it is used for processing streams and not for storage.

Option D is wrong because it is a de-coupling solution.

For more information on Redshift, please visit the below url

- https://aws.amazon.com/redshift/?nc2=h_m1 (https://aws.amazon.com/redshift/?nc2=h_m1)

Ask our Experts



QUESTION 24 INCORRECT

An image named photo.jpg has been uploaded to a bucket named examplebucket in the us-east-1 region. Which of the below is the right URL to access the image, if it were made public ? Consider that S3 is used as a static website.

- ☐ A. <http://examplebucket.s3-website-us-east-1.amazonaws.com/photo.jpg> ✓
- ☐ B. <http://examplebucket.website-us-east-1.amazonaws.com/photo.jpg>
- ☐ C. <http://examplebucket.s3-us-east-1.amazonaws.com/photo.jpg> ✗
- ☐ D. <http://examplebucket.amazonaws.s3-website-us-east-1./photo.jpg>

Explanation :

Answer - A

The URL for a S3 web site is as shown below

.s3-website-.amazonaws.com

Hence the right option is option A

When you configure a bucket for website hosting, the website is available via the region-specific website endpoint. Website endpoints are different from the endpoints where you send REST API requests. For more information about the differences between the endpoints, see [Key Differences Between the Amazon Website and the REST API Endpoint](#).

The two general forms of an Amazon S3 website endpoint are as follows:

--> `bucket-name.s3-website-region.amazonaws.com`

--> `bucket-name.s3-website.region.amazonaws.com`

Website Endpoints

When you configure a bucket for website hosting, the website is available via the region-specific website endpoint. Website endpoints are different from the endpoints where you send REST API requests. For more information about the differences between the endpoints, see [Key Differences Between the Amazon Website and the REST API Endpoint](#).

The two general forms of an Amazon S3 website endpoint are as follows:

```
bucket-name.s3-website-region.amazonaws.com
```



```
bucket-name.s3-website-region.amazonaws.com
```



Which form is used for the endpoint depends on what Region the bucket is in. For example, if your bucket is named `example-bucket` and it resides in the US East (N. Virginia) region, the website is available at the following Amazon S3 website endpoint:

```
http://examplebucket.s3-website-us-east-1.amazonaws.com/
```



Or, if your bucket is named `example-bucket` and it resides in the EU (Frankfurt) region, the website is available at the following Amazon S3 website endpoint:

```
http://examplebucket.s3-website-eu-central-1.amazonaws.com/
```



For a list of the Amazon S3 website endpoints by Region, see [Amazon Simple Storage Service Website Endpoints in the AWS General Reference](#).

In order for your customers to access content at the website endpoint, you must make all your content publicly readable. To do so, you can use a bucket policy or an ACL on an object to grant the necessary permissions.

Note

Requester Pays buckets or DevPay buckets do not allow access through the website endpoint. Any request to such a bucket receives a `403 Access Denied` response. For more information, see [Requester Pays Buckets](#).

If you have a registered domain, you can add a DNS CNAME entry to point to the Amazon S3 website endpoint. For example, if you have registered domain, `www.example-bucket.com`, you could create a bucket `www.example-bucket.com` and add a DNS CNAME record that points to `www.example-bucket.com.s3-website-region.amazonaws.com`. All requests to `http://www.example-bucket.com` are routed to `www.example-bucket.com.s3-website-region.amazonaws.com`. For more information, see [Virtual Hosting of Buckets](#).

Which form is used for the endpoint depends on what Region the bucket is in. For example, if your bucket is named `example-bucket` and it resides in the US East (N. Virginia) region, the website is available at the following Amazon S3 website endpoint:

For more information on the bucket and the URL format for S3 buckets, please visit the below url:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/WebsiteEndpoints.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/WebsiteEndpoints.html>)
- <http://docs.aws.amazon.com/AmazonS3/latest/dev/HostingWebsiteOnS3Setup.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/HostingWebsiteOnS3Setup.html>)

Ask our Experts



QUESTION 25

CORRECT

A company has an EC2 instance that is hosting a web solution which is mostly used for read-only purposes. The CPU utilization is constantly 100% on the EC2 instance. Which of the below solutions can help alleviate and provide a quick resolution to the problem.

- ☒ A. Use Cloudfront and place the EC2 instance as the origin ✓
- ☐ B. Let the EC2 instance continue to run at 100%, since the AWS environment can handle the load.
- ☐ C. Use SNS to notify the IT admin when it reaches 100% so that they can disconnect some sessions to help alleviate the load
- ☐ D. Use SES to notify the IT admin when it reaches 100% so that they can disconnect some sessions to help alleviate the load

Explanation :

Answer – A

Cloudfront can be used alleviate the load on web based solutions by caching the recent reads in its edge locations and reduce the burden on the EC2 instance.

Amazon CloudFront is a global content delivery network (CDN) service that accelerates delivery of your websites, APIs, video content or other web assets

For more information on AWS Cloudfront please visit the below url

<https://aws.amazon.com/cloudfront/> (<https://aws.amazon.com/cloudfront/>)

Ask our Experts



QUESTION 26

INCORRECT

Which of the mentioned AWS services uses the concept of shards and is uniquely identified group of data records in a stream?

- ☐ A. Cloudfront
- ☒ B. SQS ✗
- ☐ C. Kinesis ✓
- ☐ D. SES

Explanation :

Answer – C

In Amazon Kinesis, a shards is a uniquely identified group of data records in a stream. A stream is composed of one or more shards, each of which provides a fixed unit of capacity. Each shard can support up to 5 transactions per second for reads, up to a maximum total data read rate of 2 MB per second and up to 1,000 records per second for writes, up to a maximum total data write rate of 1 MB per second (including partition keys).

For more information on AWS Kinesis please visit the below url

- <http://docs.aws.amazon.com/streams/latest/dev/key-concepts.html>
(<http://docs.aws.amazon.com/streams/latest/dev/key-concepts.html>)

Ask our Experts



QUESTION 27

CORRECT

Which of the below mentioned services are the building blocks for creating a basic high availability architecture in AWS.


Select 2 options.

- ☒ A. EC2 ✓
- ☐ B. SQS
- ☒ C. Elastic Load Balancer ✓
- ☐ D. Cloudwatch

Explanation :

Answer – A and C

Having EC2 instances hosting your applications in multiple subnets, hence multiple AZ's and placing them behind an ELB is the basic building block of a high availability architecture in AWS.



For more information on High availability and Fault tolerance please visit the below url
https://media.amazonwebservices.com/architecturecenter/AWS_ac_ra_ftha_04.pdf
(https://media.amazonwebservices.com/architecturecenter/AWS_ac_ra_ftha_04.pdf)

Ask our Experts



QUESTION 28

CORRECT

You have a set of EC2 Instances launched via Autoscaling. You now want to change the Instance type for the instances that would be launched in the future via Autoscaling. What would you do in such a case

- ☐ A. Change the Launch configuration to reflect the new instance type
- ☐ B. Change the Autoscaling Group and add the new instance type.
- ☒ C. Create a new Launch Configuration with the new instance type and replace the existing Launch configuration attached to the Autoscaling Group. ✓
- ☐ D. Create a new Launch Configuration with the new instance type and add it along with the existing Launch configuration attached to the Autoscaling Group.

Explanation :

Answer – C

The AWS Documentation mentions the following

When you create an Auto Scaling group, you must specify a launch configuration. You can specify your launch configuration with multiple Auto Scaling groups. However, you can only specify one launch configuration for an Auto Scaling group at a time, and you can't modify a launch configuration after you've created it. Therefore, if you want to change the launch configuration for your Auto Scaling group, you must create a launch configuration and then update your Auto Scaling group with the new launch configuration.

For more information on Launch Configuration for Autoscaling, please refer to the below link

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/LaunchConfiguration.html>
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/LaunchConfiguration.html>)

Ask our Experts



QUESTION 29

CORRECT

Which of the following services provides an object store which can also be used to store files ?

- ☒ A. S3 ✓
- ☐ B. SQS
- ☐ C. SNS
- ☐ D. EC2

Explanation :

Answer – A

Amazon Simple Storage Service is storage for the Internet. It is designed to make web-scale computing easier for developers.

Amazon S3 has a simple web services interface that you can use to store and retrieve any amount of data, at any time, from anywhere on the web. It gives any developer access to the same highly scalable, reliable, fast, inexpensive data storage infrastructure that Amazon uses to run its own global network of web sites.

For more information on S3 please visit the below url

<http://docs.aws.amazon.com/AmazonS3/latest/dev/Welcome.html>

(<http://docs.aws.amazon.com/AmazonS3/latest/dev/Welcome.html>)

Ask our Experts



QUESTION 30

INCORRECT

What are the different types of scale out options available in the AutoScaling service provided by AWS?

Select 3 options.

- ☒ A. Scheduled Scaling ✓

- ☒ B. Dynamic Scaling ✓
- ☐ C. Manual Scaling ✓
- ☒ D. Static Scaling ✗

Explanation :

Answer – A, B and C

This is clearly given in the AWS documentation:

Scale Out

The following scale out events direct the Auto Scaling group to launch EC2 instances and attach them to the group:

- You manually increase the size of the group. For more information, see [Manual Scaling](#).
- You create a scaling policy to automatically increase the size of the group based on a specified increase in demand. For more information, see [Dynamic Scaling](#).
- You set up scaling by schedule to increase the size of the group at a specific time. For more information, see [Scheduled Scaling](#).

For more information on Autoscaling please visit the below URL:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/AutoScalingGroupLifecycle.html>
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/AutoScalingGroupLifecycle.html>)

Ask our Experts



QUESTION 31 CORRECT

Which of the following services provides edge locations that can be used to cache frequently accessed pages of a web application ?

- ☐ A. SQS
- ☒ B. Cloudfront ✓
- ☐ C. Subnets
- ☐ D. EC2

Explanation :

Answer - B

Amazon CloudFront is a global content delivery network (CDN) service that accelerates delivery of your websites, APIs, video content or other web assets

For more information on AWS Cloudfront please visit the below url

<https://aws.amazon.com/cloudfront/> (<https://aws.amazon.com/cloudfront/>)

Ask our Experts



QUESTION 32

CORRECT

You are a security architect working for a large antivirus company. The production environment has recently been moved to AWS and is in a public subnet. You are able to view the production environment over HTTP however when your customers try to update their virus definition files over a custom port, that port is blocked. You log in to the console and you allow traffic in over the custom port. How long will this take to take effect?

- ☐ A. Straight away but to the new instances only.
- ☒ B. Immediately. ✓
- ☐ C. After a few minutes this should take effect.
- ☐ D. Straight away to the new instances, but old instances must be stopped and restarted before the new rules apply.

Explanation :

Answer – B

When you make a change to the Security Groups or Network ACLs, they are applied immediately.

For more information on Security Groups, please refer to the below link:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html)

Ask our Experts



You work for a company who are deploying a hybrid cloud approach. Their legacy servers will remain on premise within their own datacenter however they will need to be able to communicate to the AWS environment over a site to site VPN connection. What do you need to do to establish the VPN connection?

- ☐ A. Connect to the environment using AWS Direct Connect.
- ☐ B. Assign a static routable address to the customer gateway ✓
- ☒ C. Create a dedicated NAT and deploy this to the public subnet. ✕
- ☐ D. Update your route table to add a route for the NAT to 0.0.0.0/0.

Explanation :

Answer – B

This requirement is given in the AWS documentation for the customer gateway. The traffic from the VPC gateway must be able to leave the VPC and traverse through the internet onto the customer gateway. Hence the customer gateway needs to be assigned a static IP that can be routable via the internet.

Internet-routable IP address (static) of the customer gateway's external interface.	Used to create and configure your customer gateway (it's referred to as YOUR_UPLINK_ADDRESS)	The public IP address value must be static. If your customer gateway is behind a network address translation (NAT) device that's enabled for NAT traversal (NAT-T), use the public IP address of your NAT device, and adjust your firewall rules to unblock UDP port 4500.
---	--	--

For more information on VPC Virtual Private connections, please refer to the below link:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_VPN.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_VPN.html)

Ask our Experts



Which is the service provided by AWS for collecting and processing large streams of data in real time?

- ☒ A. Amazon Kinesis ✓
- ☐ B. AWS Data Pipeline
- ☐ C. Amazon AppStream
- ☐ D. Amazon Simple Queue Service

Explanation :

Answer – A

Use Amazon Kinesis Streams to collect and process large streams of data records in real time. You'll create data-processing applications, known as Amazon Kinesis Streams applications. A typical Amazon Kinesis Streams application reads data from an Amazon Kinesis stream as data records. These applications can use the Amazon Kinesis Client Library, and they can run on Amazon EC2 instances. The processed records can be sent to dashboards, used to generate alerts, dynamically change pricing and advertising strategies, or send data to a variety of other AWS services.

For more information on Amazon Kinesis, please visit this URL:

- <http://docs.aws.amazon.com/streams/latest/dev/introduction.html>
(<http://docs.aws.amazon.com/streams/latest/dev/introduction.html>)

Ask our Experts



QUESTION 35

CORRECT

Which of the following criteria are some of which must be met when attaching an EC2 instance to an existing AutoScaling Group ?

Select 3 options.

- ☒ A. The instance is in the running state. ✓
- ☒ B. The AMI used to launch the instance must still exist. ✓
- ☒ C. The instance is not a member of another Auto Scaling group. ✓

☐ D. They should have the same private key

Explanation :

Answer – A, B and C

This is clearly mentioned in the AWS documentation:

Attach EC2 Instances to Your Auto Scaling Group

Auto Scaling provides you with an option to enable Auto Scaling for one or more EC2 instances by attaching them to your existing Auto Scaling group. After the instances are attached, they become a part of the Auto Scaling group.

The instance that you want to attach must meet the following criteria:

- The instance is in the `running` state.
- The AMI used to launch the instance must still exist.
- The instance is not a member of another Auto Scaling group.
- The instance is in the same Availability Zone as the Auto Scaling group.
- If the Auto Scaling group has an attached load balancer, the instance and the load balancer must both be in EC2-Classic or the same VPC. If the Auto Scaling group has an attached target group, the instance and the Application Load Balancer must both be in the same VPC.

For more information on the criteria for attaching an EC2 instance to an existing AutoScaling group please visit the below URL:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/attach-instance-asg.html>
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/attach-instance-asg.html>)

Ask our Experts



QUESTION 36

CORRECT

A company wants to make use of serverless code. Which service in AWS provides such a facility?

- ☐ A. SQS
- ☐ B. Cloudfront
- ☐ C. EC2
- ☒ D. Lambda ✓

Explanation :

Answer – D

AWS Lambda is a compute service that lets you run code without provisioning or managing servers. AWS Lambda executes your code only when needed and scales automatically, from a few requests per day to thousands per second. You pay only for the compute time you consume - there is no charge when your code is not running.

For more information on AWS Lambda please visit the below URL:

- <http://docs.aws.amazon.com/lambda/latest/dg/welcome.html>
(<http://docs.aws.amazon.com/lambda/latest/dg/welcome.html>)

Ask our Experts



QUESTION 37 INCORRECT

A t2.medium EC2 instance type must be launched with what type of Amazon Machine Image (AMI)?

- ☒ A. An Instance store Hardware Virtual Machine AMI ✕
- ☐ B. An Instance store Paravirtual AMI
- ☐ C. An Amazon EBS-backed Hardware Virtual Machine AMI ✓
- ☐ D. An Amazon EBS-backed Paravirtual AMI

Explanation :

Answer – C

The AWS Documentation mentions the below

Linux Amazon Machine Images use one of two types of virtualization: paravirtual (PV) or hardware virtual machine (HVM). The main difference between PV and HVM AMIs is the way in which they boot and whether they can take advantage of special hardware extensions (CPU, network, and storage) for better performance.

The below snapshot also shows the support for the T2 Instance family.

Instance Family	HVM EBS-Backed 64-bit	HVM Instance Store 64-bit	PV EBS-Backed 64-bit	PV Instance Store 64-bit
T2	✓			
M4	✓			
M3	✓	✓	✓	✓
C4	✓			
C3	✓	✓	✓	✓

For more information on the Instance types for Linux AMI's, please refer to the below link

- <https://aws.amazon.com/amazon-linux-ami/instance-type-matrix/>
(<https://aws.amazon.com/amazon-linux-ami/instance-type-matrix/>)

Ask our Experts



QUESTION 38 CORRECT

You are building an automated transcription service in which Amazon EC2 worker instances process an uploaded audio file and generate a text file. You must store both of these files in the same durable storage until the text file is retrieved. You do not know what the storage capacity requirements are. Which storage option is both cost-efficient and scalable?

- ☐ A. Multiple Amazon EBS volume with snapshots
- ☐ B. A single Amazon Glacier vault
- ☒ C. A single Amazon S3 bucket ✓
- ☐ D. Multiple instance stores

Explanation :

Answer – C

The AWS Simple Storage service is the best option for this scenario. The AWS documentation provides the following information on the Simple Storage service

Amazon S3 is object storage built to store and retrieve any amount of data from anywhere – web sites and mobile apps, corporate applications, and data from IoT sensors or devices. It is designed to deliver 99.999999999% durability, and stores data for millions of applications used by market leaders in every industry

For more information on the Simple Storage Service, please refer to the below link

- <https://aws.amazon.com/s3/> (<https://aws.amazon.com/s3/>)

Ask our Experts



QUESTION 39 INCORRECT

You have some EC2 instances hosted in your AWS environment. You have a concern that not all of the EC2 instances are being utilized.

Which of the below mentioned services can help you find underutilized resources in AWS ?

Select 2 options.

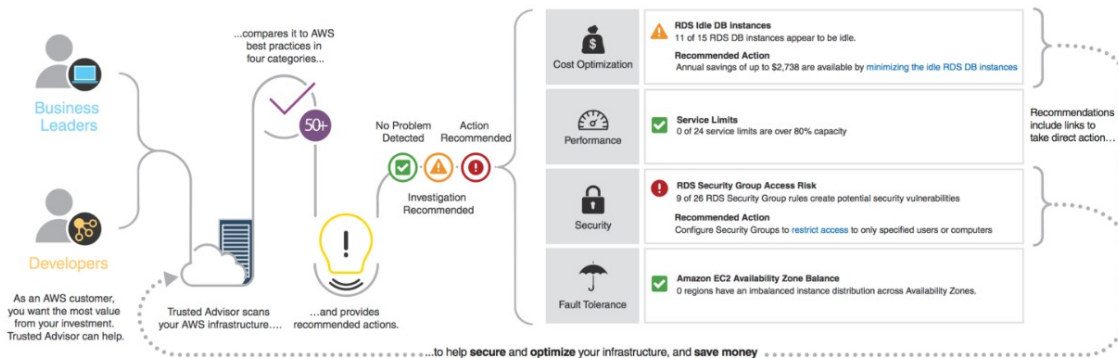
- ☐ A. AWS Cloudwatch ✓
- ☐ B. SNS
- ☒ C. AWS Trusted Advisor ✓
- ☒ D. Cloudtrail ✗

Explanation :

Answer – A and C

The AWS Trusted Advisor can help you identify underutilized resources in AWS.

An Introduction to AWS Trusted Advisor



For more information on AWS trusted advisor please visit the below URL:

- <https://aws.amazon.com/premiumsupport/trustedadvisor/>
(<https://aws.amazon.com/premiumsupport/trustedadvisor/>)

If You look at the Cloudwatch graphs, the CPU utilization of your resources and you can see the trend over time in the graphs.

For more information on AWS Cloudwatch please visit the below URL:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)

Ask our Experts



QUESTION 40

CORRECT

Which of the following features can be used to capture information for outgoing and incoming IP traffic from network interfaces in a VPC.

- ☐ A. AWS Cloudwatch
- ☐ B. AWSEC2
- ☐ C. AWS SQS
- ☒ D. AWS VPC Flow Logs ✓

Explanation :

Answer - D

VPC Flow Logs is a feature that enables you to capture information about the IP traffic going to and from network interfaces in your VPC. Flow log data is stored using Amazon CloudWatch Logs. After you've created a flow log, you can view and retrieve its data in Amazon CloudWatch Logs. For more information on VPC Flowlogs please visit the below URL:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/flow-logs.html>
(<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/flow-logs.html>)

Ask our Experts



QUESTION 41 CORRECT

What are the main benefits of AWS regions?

Select 2 options.

- ☒ **A. Regions allow you to design applications to conform to specific laws and regulations for specific parts of the world. ✓**
- ☐ **B. All regions offer the same service at the same prices.**
- ☐ **C. Regions allow you to choose a location in any country in the world.**
- ☒ **D. Regions allow you to place AWS resources in the area of the world closest to your customers who access those resources. ✓**

Explanation :

Answer – A and D

AWS developer data centers across the world to help develop solutions that are close to the customer as possible. They also have center in core countries to help tie up with the specific laws and regulations for specific parts of the world.

AWS does not have centers in every location of the world, hence option C is invalid.

Services and prices are specific to every region, hence option B is invalid.

For more information on Regions please visit the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html> (<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>)

Ask our Experts



QUESTION 42

CORRECT

Which of the following are ways that users can interface with AWS?

Select 2 options.

- ☐ A. AWS Cloudfront
- ☒ B. AWS CLI ✓
- ☒ C. AWS Console ✓
- ☐ D. AWS Cloudwatch

Explanation :

Answer – B and C

AWS Management Console - The console is a browser-based interface to manage IAM and AWS resources.

AWS Command Line Tools - You can use the AWS command line tools to issue commands at your system's command line to perform IAM and AWS tasks; this can be faster and more convenient than using the console.

For more information on interfacing with AWS please visit the below URL:

- <http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>
(<http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>)

Ask our Experts



QUESTION 43

INCORRECT

What are the two layers of security provided by AWS in the VPC?

- ☐ A. Security Groups and NACLs ✓

- ☐ B. NACLs and DHCP Options
- ☐ C. Route Tables and Internet gateway
- ☒ D. None of the above ✕

Explanation :

Answer – A

This is clearly given in the AWS documentation

Security

Amazon VPC provides features that you can use to increase and monitor the security for your VPC:

- Security groups — Act as a firewall for associated Amazon EC2 instances, controlling both inbound and outbound traffic at the instance level
- Network access control lists (ACLs) — Act as a firewall for associated subnets, controlling both inbound and outbound traffic at the subnet level
- Flow logs — Capture information about the IP traffic going to and from network interfaces in your VPC

For more information on VPC Security please visit the below URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Security.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Security.html)

Ask our Experts



QUESTION 44

CORRECT

A company wants to have a 50 Mbps dedicated connection to its AWS resources. Which of the below services can help fulfil this requirement ?

- ☐ A. Virtual Private Gateway
- ☐ B. Virtual Private Connection (VPN)
- ☒ C. Direct Connect ✓
- ☐ D. Internet Gateway

Explanation :

Answer - C

AWS Direct Connect makes it easy to establish a dedicated network connection from your premises to AWS. Using AWS Direct Connect, you can establish private connectivity between AWS and your datacenter, office, or colocation environment, which in many cases can reduce your network costs, increase bandwidth throughput, and provide a more consistent network experience than Internet-based connections.

For more information on AWS Direct Connect, please visit the below URL:

<https://aws.amazon.com/directconnect/> (<https://aws.amazon.com/directconnect/>)

If you require a port speed less than 1 Gbps, you cannot request a connection using the console.

Instead, contact an APN partner, who will create a hosted connection for you. The hosted connection appears in your AWS Direct Connect console, and must be accepted before use.

Please find an exact explanation in AWS documentation below:

http://docs.aws.amazon.com/directconnect/latest/UserGuide/getting_started.html#ConnectionRequest
(http://docs.aws.amazon.com/directconnect/latest/UserGuide/getting_started.html#ConnectionRequest)

Step 2: Submit AWS Direct Connect Connection Request

You can submit a connection request using the AWS Direct Connect console. Before you begin, ensure that you have the following information:

- The port speed that you require: 1 Gbps or 10 Gbps. You cannot change the port speed after you've created the connection request.
- The AWS Direct Connect location to which to connect.

If you require a port speed less than 1 Gbps, you cannot request a connection using the console. Instead, contact an APN partner, who will create a hosted connection for you. The hosted connection appears in your AWS Direct Connect console, and must be accepted before use. Skip the following procedure and go to [Accept Your Hosted Connection](#).

Ask our Experts



QUESTION 45 INCORRECT

What is the service name in AWS that can display costs in a chart format?

- ☐ A. Cost Explorer ✓
- ☐ B. Cost Allocation Tags
- ☒ C. AWS Budgets ✗
- ☐ D. Payment History

Explanation :

Answer – A

Cost Explorer is a free tool that you can use to view charts of your costs (also known as spend data) for up to the last 13 months, and forecast how much you are likely to spend for the next three months. You can use Cost Explorer to see patterns in how much you spend on AWS resources over time, identify areas that need further inquiry, and see trends that you can use to understand your costs. You can also specify time ranges for the data you want to see, and you can view time data by day or by month.

For more information on Cost Explorer, please visit the URL:

- <http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/cost-explorer-what-is.html>
(<http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/cost-explorer-what-is.html>)

Ask our Experts



QUESTION 46 INCORRECT

In the shared responsibility model, what is the customer not responsible for?

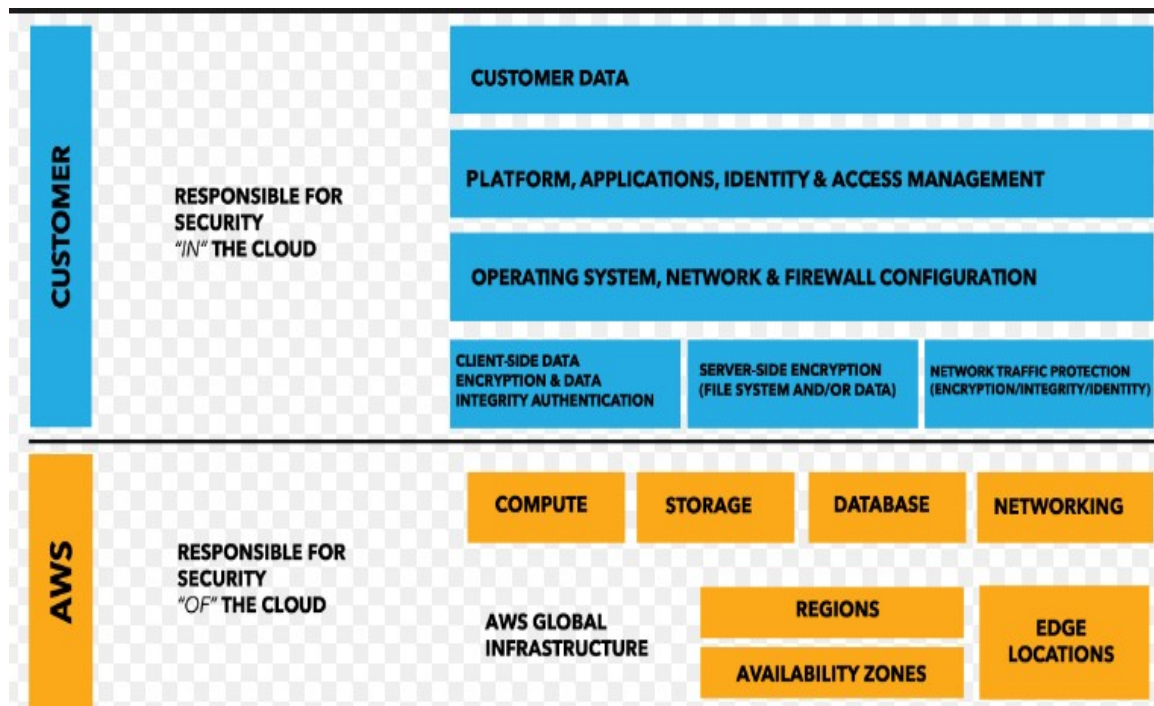
- ☐ A. Edge locations ✓

- ☒ B. Installation of custom firewall software ✕
- ☐ C. Security Groups
- ☐ D. Applying an SSL Certificate to an ELB

Explanation :

Answer - A

AWS has published the Shared Responsibility Model. And the Physical networking comes as part of the responsibility of AWS.



For more information on the Shared Responsibility Model, please refer to the below URL:

- <https://aws.amazon.com/compliance/shared-responsibility-model/>
(<https://aws.amazon.com/compliance/shared-responsibility-model/>)

Ask our Experts



QUESTION 47 INCORRECT

Which of the below options best describes how EBS snapshots work?

- ☒ A. Each snapshot stores the entire volume in S3 ✕
- ☐ B. Snapshots are not possible for EBS volumes
- ☐ C. Snapshots are incremental in nature and are stored in S3 ✓
- ☐ D. Snapshots are stored in DynamoDB

Explanation :

Answer – C

You can back up the data on your EBS volumes to Amazon S3 by taking point-in-time snapshots. Snapshots are incremental backups, which means that only the blocks on the device that have changed after your most recent snapshot are saved. This minimizes the time required to create the snapshot and saves on storage costs. When you delete a snapshot, only the data unique to that snapshot is removed.

For more information on EBS Snapshots, please refer to the below link

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>

(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>)

Ask our Experts



QUESTION 48 CORRECT

You work for a market analysis firm who are designing a new environment. They will ingest large amounts of market data via Kinesis and then analyse this data using Elastic Map Reduce. The data is then imported in to a high performance NoSQL Cassandra database which will run on EC2 and then be accessed by traders from around the world. The database volume itself will sit on EBS volume. They are expecting very high demand during peak times, with an IOPS performance level of approximately 15,000. Which EBS volume should you recommend?

- ☐ A. Magnetic
- ☐ B. General Purpose SSD
- ☒ C. Provisioned IOPS ✓
- ☐ D. Turbo IOPS (TIOPS)

Explanation :

Answer – C

When you are looking at hosting intensive I/O applications such as databases, always look to using IOPS as the preferred storage option.

For more information on the various EBS Volume types, please refer to the below link

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html>)

Ask our Experts



QUESTION 49

CORRECT

Which of the following best describes the main feature of an Elastic Load Balancer (ELB) in AWS?

- ☒ A. To evenly distribute traffic among multiple EC2 instances in separate Availability Zones. ✓
- ☐ B. To evenly distribute traffic among multiple EC2 instances in a single Availability Zones.
- ☐ C. To evenly distribute traffic among multiple EC2 instances in a multiple regions.
- ☐ D. To evenly distribute traffic among multiple EC2 instances in a multiple counties.

Explanation :

Answer – A

Elastic Load Balancing automatically distributes incoming application traffic across multiple Amazon EC2 instances. It enables you to achieve fault tolerance in your applications, seamlessly providing the required amount of load balancing capacity needed to route application traffic. The ELB is best used for EC2 instances multiple AZ's.

You cannot use ELB for distributing traffic across regions.

For more information on AWS ELB, please refer to the below link:

- <https://aws.amazon.com/elasticloadbalancing/>
(<https://aws.amazon.com/elasticloadbalancing/>)

Ask our Experts



QUESTION 50

CORRECT

What are the 2 main components of AutoScaling?

Select 2 options.

- ☒ A. Launch Configuration ✓
- ☐ B. Cloudtrail
- ☐ C. Cloudwatch
- ☒ D. AutoScaling Groups ✓

Explanation :

Answer – A and D

Groups - Your EC2 instances are organized into groups so that they can be treated as a logical unit for the purposes of scaling and management. When you create a group, you can specify its minimum, maximum, and, desired number of EC2 instances.

Launch configurations - Your group uses a launch configuration as a template for its EC2 instances. When you create a launch configuration, you can specify information such as the AMI ID, instance type, key pair, security groups, and block device mapping for your instances.

For more information on AWS Autoscaling, please refer to the below link:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/WhatIsAutoScaling.html>
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/WhatIsAutoScaling.html>)

Ask our Experts



QUESTION 51

CORRECT

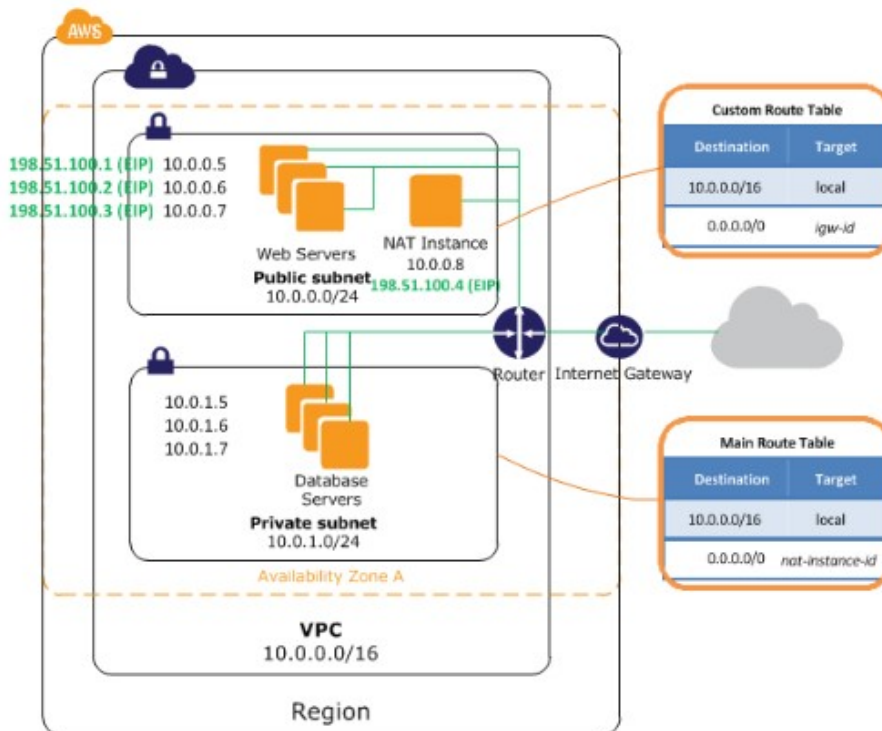
What are two primary requirements of a NAT Instance? Choose the correct answer from the options below:

- ☐ A. A NAT instance must be provisioned into a private subnet, and it must part of the private subnet's route table.
- ☒ B. A NAT instance must be provisioned into a public subnet, and it must part of the private subnet's route table. ✓
- ☐ C. A NAT instance must be provisioned into a private subnet, and does not require a public IP address.
- ☐ D. A NAT instance must be provisioned into a public subnet, and must be combined with a bastion host.

Explanation :

Answer – B

The below snapshot from the AWS documentation shows how the NAT instance is setup.



It needs to be placed in the public subnet and the private subnet should have a route to it.

For more information on AWS NAT Instance, please refer to the below link:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_NAT_Instance.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_NAT_Instance.html)

Ask our Experts



You have just provisioned a fleet of EC2 instances and realized that none of them have a public IP address. What settings would need to be changed for the next fleet of instances to be created with public IP addresses?

- ☒ A. Modify the auto-assign public IP setting on the subnet. ✓
- ☐ B. Modify the auto-assign public IP setting on the instance type.
- ☐ C. Modify the auto-assign public IP setting on the route table.
- ☐ D. Modify the auto-assign public IP setting on the VPC.

Explanation :

Answer – A

This setting is done at the subnet level and if marked as true, all instances launched in that subnet will get a public IP address by default.

Modify Auto-Assign Public IP



Enable auto-assign public IP to automatically request a public IP address for instances launched into this subnet.

☒ Enable auto-assign Public IP

Note: You can override the auto-assign public IP setting for each individual instance at launch time. Regardless of how you've configured the auto-assign public IP feature, you can assign a public IP address to an instance that has a single, new network interface with a device index of eth0.

Cancel

Save

For more information on AWS IP Addressing, please refer to the below link:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-ip-addressing.html>
(<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-ip-addressing.html>)

Ask our Experts



You have created a VPC named test1 using the VPC wizard. You want the instances in the VPC to connect to the internet. You have created an Internet Gateway named IGW1.

You are trying to associate IGW1 to the VPC test1, but the VPC is not listed on the drop-down list. What is the most likely reason for this?

- ☐ A. You need to have a customer gateway defined first before attaching an internet gateway
- ☒ B. You need to have a public subnet defined first before attaching an internet gateway ✕
- ☐ C. You need to have a private subnet defined first before attaching an internet gateway
- ☐ D. An Internet gateway is already attached to the VPC ✓

Explanation :

Answer – D

When you create a VPC using VPC wizard, an IGW is automatically created and associated with the VPC. So test1 VPC will not be listed in the "VPC selection drop down list" since it is already associated.

Attach to VPC

Attach an internet gateway to a VPC to enable communication with the internet. Specify the VPC you would like to attach below.

VPC*

Filter by attributes

VPC ID	Name
No results found.	

* Required

For more information on Internet gateways, please refer to the below link:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Internet_Gateway.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Internet_Gateway.html)

Ask our Experts



QUESTION 54

CORRECT

Which of the following feature can be used to move Objects from S3 standard storage to Amazon Glacier

- ☐ A. S3 Events
- ☐ B. Object Versioning
- ☐ C. Storage Class
- ☒ D. Lifecycle policies ✓

Explanation :

Answer – D

Lifecycle configuration enables you to specify the lifecycle management of objects in a bucket. The configuration is a set of one or more rules, where each rule defines an action for Amazon S3 to apply to a group of objects. These actions can be classified as follows:

- Transition actions – In which you define when objects transition to another storage class. For example, you may choose to transition objects to the STANDARD_IA (IA, for infrequent access) storage class 30 days after creation, or archive objects to the GLACIER storage class one year after creation.
- Expiration actions – In which you specify when the objects expire. Then Amazon S3 deletes the expired objects on your behalf.

For more information on Lifecycle policies, please refer to the below link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>)

Ask our Experts



QUESTION 55

CORRECT

You work for a company that stores records for a minimum of 10 years. Most of these records will never be accessed but must be made available upon request (within a few hours). What is the most cost-effective storage option?

- ☐ A. S3-IA
- ☐ B. Reduced Redundancy Storage (RRS)
- ☒ C. Glacier ✓
- ☐ D. AWS Import/Export

Explanation :

Answer – C

Amazon Glacier is a secure, durable, and extremely low-cost cloud storage service for data archiving and long-term backup. Customers can reliably store large or small amounts of data for as little as \$0.004 per gigabyte per month, a significant savings compared to on-premises solutions. To keep costs low yet suitable for varying retrieval needs, Amazon Glacier provides three options for access to archives, from a few minutes to several hours.

For more information on Amazon Glacier, please refer to the below link:

- <https://aws.amazon.com/glacier/> (<https://aws.amazon.com/glacier/>)

Ask our Experts



QUESTION 56

CORRECT

What are the services from AWS helps to migrate databases to AWS easily ?

- ☐ A. AWS Snowball
- ☐ B. AWS Direct Connect
- ☒ C. AWS Database Migration Service (DMS) ✓
- ☐ D. None of the above

Explanation :

Answer – C

AWS Database Migration Service helps you migrate databases to AWS easily and securely. The source database remains fully operational during the migration, minimizing downtime to applications that rely on the database. The AWS Database Migration Service can migrate your data to and from most widely used commercial and open-source databases.

For more information on aws database migration service, please visit the URL:

- <https://aws.amazon.com/dms> (<https://aws.amazon.com/dms>)

Ask our Experts



QUESTION 57

CORRECT

Your company has petabytes of data that it wants to move from their on-premise location to AWS.

Which of the following can be used to fulfil this requirement?

- ☐ A. AWS VPN
- ☐ B. AWS Migration
- ☐ C. AWS VPC
- ☒ D. AWS Snowball ✓

Explanation :

Answer – D

Snowball is a petabyte-scale data transport solution that uses secure appliances to transfer large amounts of data into and out of the AWS cloud. Using Snowball addresses common challenges with large-scale data transfers including high network costs, long transfer times, and security concerns. Transferring data with Snowball is simple, fast, secure, and can be as little as one-fifth the cost of high-speed Internet.

For more information, please refer the below link:

- <http://docs.aws.amazon.com/snowball/latest/ug/whatisSnowball.html>
(<http://docs.aws.amazon.com/snowball/latest/ug/whatisSnowball.html>)

Ask our Experts



QUESTION 58

INCORRECT

There are multiple issues reported from an EC2 instance hence it is required to analyze the logs files. What can be used in AWS to store and analyze the log files?

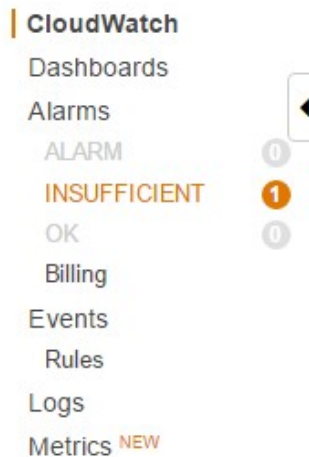
- ☐ A. SQS
- ☒ B. S3 ✕
- ☐ C. Cloudtrail
- ☐ D. Cloudwatch Logs ✓

Explanation :

Answer – D

You can use Amazon CloudWatch Logs to monitor, store, and access your log files from Amazon Elastic Compute Cloud (Amazon EC2) instances, AWS CloudTrail, and other sources. You can then retrieve the associated log data from CloudWatch Log. To enable Cloudwatch logs , follow the below steps

Step 1) Go to the Cloudwatch section and click on Logs



Step 2) Once you create a log group, you can then configure your EC2 instance to send logs to this Log Group.

Create log group

Log Group

Name:

Cancel

Create log group

For more information on Cloudwatch logs please visit the link:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/WhatIsCloudWatchLogs.html>
(<http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/WhatIsCloudWatchLogs.html>)

Ask our Experts



QUESTION 59

CORRECT

You are not able to connect to an EC2 instance via SSH, and you have already verified that the instance has a public IP and the Internet gateway and route tables are in place, what should you check next?

- ☒ A. Adjust the security group to allow traffic to port 22 ✓
- ☐ B. Adjust the security group to allow traffic to port 3389
- ☐ C. Restart the instance since there might be some issue with the instance
- ☐ D. Create a new instance since there might be some issue with the instance

Explanation :

Answer – A

The reason why you cannot connect to the instance is because maybe the SSH protocol has not been enabled in the security group.

Go to your EC2 Security groups, click on the required security groups to make the changes. Go to the Inbound Tab. Ensure that the Inbound rules has a rule for SSH protocol.

Name	Group ID	Group Name	VPC ID
	sg-042ebf79	launch-wizard-5	vpc-6dcc550a

Security Group: sg-042ebf79

Description

Inbound

Outbound

Tags

Edit

Type	Protocol	Port Range
SSH	TCP	22

Ask our Experts



QUESTION 60 CORRECT

Your company VPC has a need to communicate with another company VPC within the same AWS region.

What can be used from AWS to interface between the two VPC?

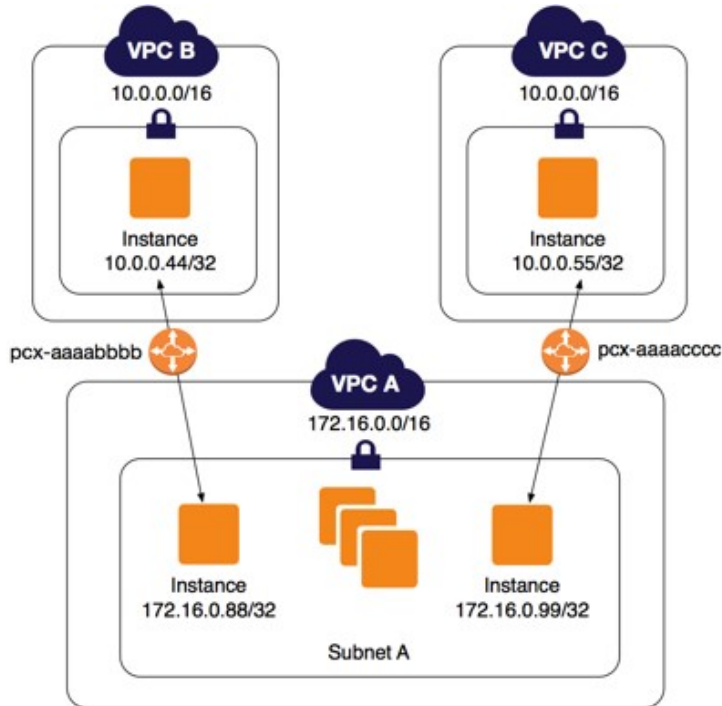
- ☐ A. VPC Connection
- ☐ B. VPN Connection
- ☐ C. Direct Connect
- ☒ D. VPC Peering ✓

Explanation :

Answer - D

A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them using private IP addresses. Instances in either VPC can communicate with each other as if they are within the same network. You can create a VPC peering connection between your own VPCs, or with a VPC in another AWS account.

Amazon EC2 now allows peering relationships to be established between VPCs across different AWS regions. Inter-Region VPC Peering allows VPC resources like EC2 instances, RDS databases and Lambda functions running in different AWS regions to communicate with each other using private IP addresses, without requiring gateways, VPN connections or separate network appliances. The below diagram shows an example of VPC peering. Now please note that VPC B cannot communicate to VPC C because there is no peering between them.



For more information on VPC peering, please visit the URL:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-peering.html>
(<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-peering.html>)

Note: With latest AWS Update, VPC Peering for different region is available now.

Ask our Experts



Finish Review (<https://www.whizlabs.com/learn/course/aws-csaa-practice-tests/quiz/13954>)

Company

- ➔ Support (<https://help.whizlabs.com/hc/en-us>)
- ➔ Discussions (<http://ask.whizlabs.com/>)

Certification

- ➔ Cloud Certification
(<https://www.whizlabs.com/cloud-certification-training-courses/>)
- ➔ Java Certification
(<https://www.whizlabs.com/oracle-java-certifications/>)
- ➔ PM Certification
(<https://www.whizlabs.com/project-management-certifications/>)
- ➔ Big Data Certification
(<https://www.whizlabs.com/big-data-certifications/>)

Mobile App

 Android Coming Soon

 iOS Coming Soon

➔ Blog (<https://www.whizlabs.com/blog/>)

Follow us



(<https://www.facebook.com/whizlabs.software/>)



(<https://in.linkedin.com/company/whizlabs-software>)



(<https://twitter.com/whizlabs?lang=en>)



(<https://plus.google.com/+WhizlabsSoftware>)