



- [Home](https://www.whizlabs.com/learn) (<https://www.whizlabs.com/learn>) > [My Courses](https://www.whizlabs.com/learn/my-courses) (<https://www.whizlabs.com/learn/my-courses>)
> [AWS Certified Solutions Architect Associate](https://www.whizlabs.com/learn/course/aws-csaa-practice-tests#section-1) (<https://www.whizlabs.com/learn/course/aws-csaa-practice-tests#section-1>)
> [New Practice Test VI](https://www.whizlabs.com/learn/course/aws-csaa-practice-tests/quiz/13960) (<https://www.whizlabs.com/learn/course/aws-csaa-practice-tests/quiz/13960>) > **Report**

NEW PRACTICE TEST VI

Attempt 1
Marks Obtained 3 / 60
Your score is 5%

Completed on Monday , 28 January 2019 , 06:53 PM
Time Taken 00 H 02 M 54 S
Result Fail

Domains / Topics wise Quiz Performance Report

S.No.	Topic	Total Questions	Correct	Incorrect	Unattempted
1	Other	60	3	2	55

60 Questions	3 Correct	2 Incorrect	55 Unattempted
------------------------	---------------------	-----------------------	--------------------------

Show Answers

All	▼
-----	---

QUESTION 1 CORRECT

A company wants to store data that is not frequently accessed. What is the best and cost efficient solution that should be considered?

☐ A. Amazon Storage Gateway

- ☒ B. Amazon Glacier ✓
- ☐ C. Amazon EBS
- ☐ D. Amazon S3

Explanation :

Answer: B – Since the data is not required to be accessed frequently, the data can be stored on Amazon glacier for cheaper storage. Remember that the recovery time for getting data from Glacier is from 3-5 hours.

You can look at the FAQ section of aws glacier - <https://aws.amazon.com/glacier/faqs/>
(<https://aws.amazon.com/glacier/faqs/>)

Amazon Glacier is an extremely low-cost storage service that provides secure, durable, and flexible storage for data backup and archival. With Amazon Glacier, customers can reliably store their data for as little as \$0.007 per gigabyte per month. Amazon Glacier enables customers to offload the administrative burdens of operating and scaling storage to AWS, so that they don't have to worry about capacity planning, hardware provisioning, data replication, hardware failure detection and repair, or time-consuming hardware migrations.

Ask our Experts



QUESTION 2 CORRECT

You have an EC2 instance that is transferring data from S3 in the same region. The project sponsor is worried about the cost of the infrastructure. What can you do to convince him that you have a cost effective solution.

- ☐ A. You are going to be hosting only 4 instances, so you are minimizing on cost.
- ☒ B. There is no cost for transferring data from EC2 to S3 if they are in the same region. ✓
- ☐ C. AWS provides a discount if you transfer data from EC2 to S3 if they are in the same region.
- ☐ D. Both EC2 and S3 are in the same availability zone, so you can save via consolidated billing.

Explanation :

Answer – B – Please note that there is no cost when data is transferred from EC2 to S3 if they are in the same region. This is very important for an AWS Solution Architect to know.

Ask our Experts



QUESTION 3 INCORRECT

Which services allow the customer to retain full administrative privileges of the underlying EC2 instances?

- ☐ A. Amazon Relational Database Service
- ☒ B. Amazon Elastic Map Reduce ✓
- ☐ C. Amazon ElastiCache ✕
- ☐ D. Amazon DynamoDB

Explanation :

Answer – B

In Amazon EMR , you have the ability to work with the underlying instances wherein the EMR service allows you to associate the EC2 Key pair with the launched instances. This is also given in the AWS documentation.

Create an Amazon EC2 Key Pair and PEM File

Amazon EMR uses an Amazon EC2 key pair to ensure that you alone have access to the instances that you launch. The PEM file associated with this key pair is required to `ssh` directly to the master node of the cluster.

For more information on the access to EMR nodes , please visit the below URL

- <http://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-plan-access-ssh.html>
(<http://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-plan-access-ssh.html>)

Ask our Experts



QUESTION 4 CORRECT

Bucket names must be unique across all S3.

- ☒ A. True ✓
- ☐ B. False

Explanation :

Answer – A

Bucket names must be unique across all regions. Let's say you have created a bucket names devtoolslogging in the Singapore region.

Create a Bucket - Select a Bucket Name and Region Cancel

A bucket is a container for objects stored in Amazon S3. When creating a bucket, you can choose a Region to optimize for latency, minimize costs, or address regulatory requirements. For more information regarding bucket naming conventions, please visit the [Amazon S3 documentation](#).

Bucket Name:

Region:

Set Up Logging > Create Cancel

Now if you want to create a bucket of the same name in the Oregon region, you will get an error that the bucket already exists.

Create a Bucket - Select a Bucket Name and RegionCancel

That bucket already exists, you own it!

A bucket is a container for objects stored in Amazon S3. When creating a bucket, you can choose a Region to optimize for latency, minimize costs, or address regulatory requirements. For more information regarding bucket naming conventions, please visit the [Amazon S3 documentation](#).

Bucket Name:

devtoolslogging

Region:

Oregon

Set Up Logging >

Create

Cancel

Ask our Experts

QUESTION 5 INCORRECT

A customer has enabled website hosting on a bucket named "devtoolslogging" in the Singapore region. What website URL is assigned to your bucket?

- ☐ A. devtoolslogging.s3-website-ap-southeast-1.amazonaws.com ✓
- ☐ B. s3-website.devtoolslogging.amazonaws.com
- ☐ C. s3-website.devtoolslogging.website-ap-southeast-1.amazonaws.com
- ☐ D. devtoolslogging.ap-southeast-1.amazonaws.com ✗

Explanation :

Answer – A – You have the chance to enable static web site hosting for S3 buckets. This can be done via the properties option for the bucket. The end point of the bucket for static hosting will also be configured.

► Permissions

▼ Static Website Hosting

You can [host your static website](#) entirely on Amazon S3. Once you enable your bucket for static website hosting, all your content is accessible to web browsers via the Amazon S3 website endpoint for your bucket.

Endpoint: devtoolslogging.s3-website-ap-southeast-1.amazonaws.com

Each bucket serves a website namespace (e.g. "www.example.com"). Requests for your host name (e.g. "example.com" or "www.example.com") can be routed to the contents in your bucket. You can also redirect requests to another host name (e.g. redirect "example.com" to "www.example.com"). See our [walkthrough](#) for how to set up an Amazon S3 static website with your host name.

Ask our Experts



QUESTION 6 UNATTEMPTED

As a solutions architect, it is your job to design for high availability and fault tolerance. Company-A is utilizing Amazon S3 to store large amounts of file data. What steps would you take to ensure that if an availability zone was lost due to a natural disaster your files would still be in place and accessible

- ☐ A. Copy the S3 bucket to an EBS optimized backed EC2 instance
- ☐ B. Amazon S3 is highly available and fault tolerant by design and requires no additional configuration ✓
- ☐ C. Enable AWS Storage Gateway using gateway-stored setup
- ☐ D. None of the above

Explanation :

Answer – B – AWS S3 is already highly available and fault tolerant. This is very clearly mentioned in its FAQ's - <https://aws.amazon.com/s3/faqs/> (<https://aws.amazon.com/s3/faqs/>)

Q: How durable is Amazon S3?

Amazon S3 Standard and Standard - IA are designed to provide 99.99999999% durability of objects over a given year. This durability level corresponds to an average annual expected loss of 0.000000001% of objects. For example, if you store 10,000 objects with Amazon S3, you can on average expect to incur a loss of a single object once every 10,000,000 years. In addition, Amazon S3 is designed to sustain the concurrent loss of data in two facilities.

As with any environments, the best practice is to have a backup and to put in place safeguards against malicious or accidental users errors. For S3 data, that best practice includes secure access permissions, Cross-Region Replication, versioning and a functioning, regularly tested backup.

Q: How is Amazon S3 designed to achieve 99.99999999% durability?

Amazon S3 Standard and Standard - IA redundantly stores your objects on multiple devices across multiple facilities in an Amazon S3 Region. The service is designed to sustain concurrent device failures by quickly detecting and repairing any lost redundancy. When processing a request to store data, the service will redundantly store your object across multiple facilities before returning SUCCESS. Amazon S3 also regularly verifies the integrity of your data using checksums.

Ask our Experts



QUESTION 7 UNATTEMPTED

Which of the following are characteristics of Amazon VPC subnets? Choose 2 answers from the options given below

- ☐ A. Each subnet spans at least 2 Availability Zones to provide a high-availability environment.
- ☐ B. Each subnet maps to a single Availability Zone. ✓
- ☐ C. CIDR block mask of /25 is the smallest range supported.
- ☐ D. By default, all subnets can route between each other, whether they are private or public. ✓
- ☐ E. Instances in a private subnet can communicate with the Internet only if they have an Elastic IP.

Explanation :

Answer – B, D. Please see the below for further justification.

A subnet can only map to one availability zone. So from options A and B, B is correct.

Create Subnet

Use the CIDR format to specify your subnet's IP address block (e.g., 10.0.0.0/24). Note that block sizes must be between a /16 netmask and /28 netmask. Also, note that a subnet can be the same size as your VPC.

Name tag

VPC

Availability Zone

CIDR block

[Cancel](#) [Yes, Create](#)

When you create a CIDR block, the least allowable is /28, so option C is wrong.

Create Subnet

Use the CIDR format to specify your subnet's IP address block (e.g., 10.0.0.0/24). Note that block sizes must be between a /16 netmask and /28 netmask. Also, note that a subnet can be the same size as your VPC.

Name tag

VPC

Availability Zone

CIDR block

[Cancel](#) [Yes, Create](#)

Option E is wrong because EC2 instances in a private subnet will not be able to route anything on the internet even if they have an elastic IP.

Ask our Experts



QUESTION 8 UNATTEMPTED

What are the different options available when creating a VPC using the VPC

wizard? Please choose all options that apply.

- ☐ A. VPC with a Primary and Secondary subnet
- ☐ B. VPC with Public and Private Subnets ✓
- ☐ C. VPC with Public and Private Subnets and Hardware VPN Access ✓
- ☐ D. VPC with default settings

Explanation :

Answer – B and C. When you launch the VPC wizard, you will get the below options in the VPC wizard.

Step 1: Select a VPC Configuration

VPC with a Single Public Subnet

VPC with Public and Private Subnets

VPC with Public and Private Subnets and Hardware VPN Access

VPC with a Private Subnet Only and Hardware VPN Access

Your instances run in a private, isolated section of the AWS cloud with direct access to the Internet. Network access control lists and security groups can be used to provide strict control over inbound and outbound network traffic to your instances.

Creates:

A /16 network with a /24 subnet. Public subnet instances use Elastic IPs or Public IPs to access the Internet.

Select

Internet, S3, DynamoDB, SNS, SQS, etc.

Public Subnet

Amazon Virtual Private Cloud

Ask our Experts



QUESTION 9 UNATTEMPTED

In AWS, which security aspects are the customer's responsibility? Choose 4 answers

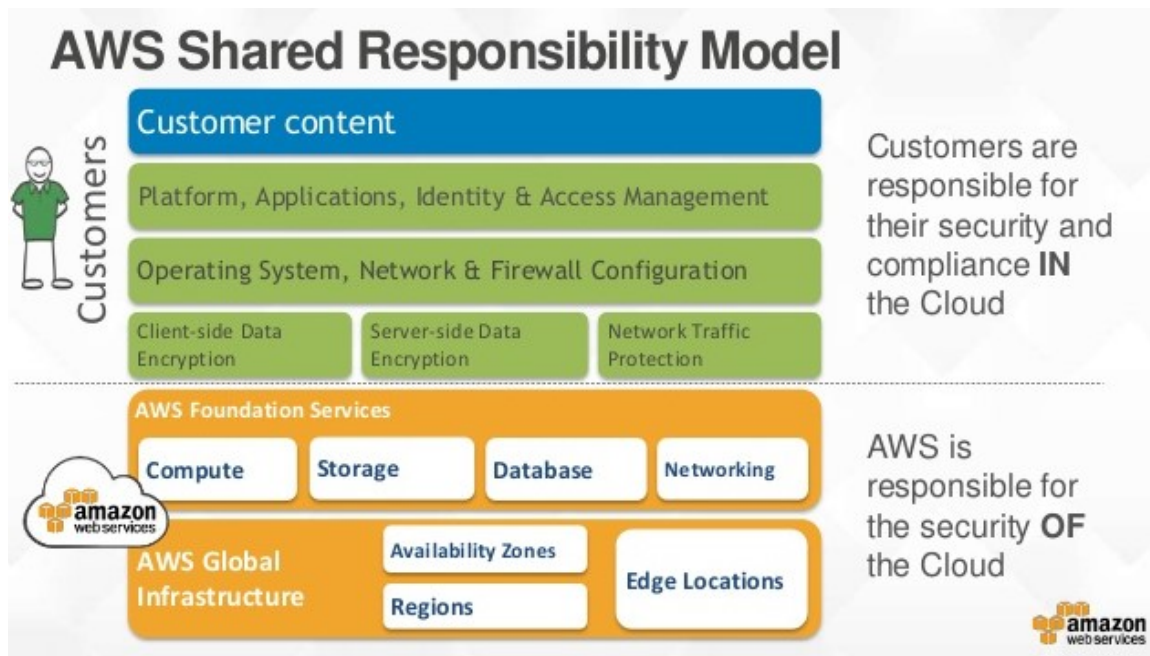
- ☐ A. Security Group and ACL (Access Control List) settings ✓
- ☐ B. Decommissioning storage devices
- ☐ C. Patch management on the EC2 instance's operating system ✓

- ☐ D. Life-cycle management of IAM credentials ✓
- ☐ E. Controlling physical access to compute resources
- ☐ F. Encryption of EBS (Elastic Block Storage) volumes ✓

Explanation :

Answer – A, C, D and F.

Please find the shared responsibility model as shared by AWS



Ask our Experts



QUESTION 10 UNATTEMPTED

When an EC2 EBS-backed (EBS root) instance is stopped, what happens to the data on any ephemeral store volumes?

- ☐ A. Data is automatically saved in an EBS volume.
- ☐ B. Data is unavailable until the instance is restarted.
- ☐ C. Data will be deleted and will no longer be accessible. ✓

- ☐ D. Data is automatically saved as an EBS snapshot.

Explanation :

Answer – C – ephemeral is temporary storage that is always deleted when an instance is restarted in aws.

When you stop or terminate an instance, every block of storage in the instance store is reset. Therefore, your data cannot be accessed through the instance store of another instance.

Find more details in AWS documentation here :

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/InstanceStorage.html?shortFooter=true#instance-store-lifetime>

(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/InstanceStorage.html?shortFooter=true#instance-store-lifetime>)

Data on the EBS volume is LOST only if the Root Volume is EBS backed and the Delete On Termination flag is checked (Checked by default)

Ask our Experts



QUESTION 11 UNATTEMPTED

SQS provides a timeout which is a period of time during which Amazon SQS prevents other consuming components from receiving and processing. What is this time period called?

- ☐ A. Component Timeout
- ☐ B. Visibility Timeout ✓
- ☐ C. Processing Timeout
- ☐ D. Receiving Timeout

Explanation :

Answer - B

Please refer to the AWS SQS FAQ section - <https://aws.amazon.com/sqs/faqs/>
(<https://aws.amazon.com/sqs/faqs/>)

Q: What is a visibility timeout?

The visibility timeout is a period of time during which Amazon SQS prevents other consuming components from receiving and processing a message. For more information, see [Visibility Timeout](#).

Ask our Experts



QUESTION 12 UNATTEMPTED

You are currently hosting an infrastructure and most of the EC2 instances are near 90 – 100% utilized. What is the type of EC2 instances you would utilize to ensure costs are minimized? Assume that the EC2 instance will be running continuously throughout the year.

- ☒ **A. Reserved instances** ✓
- ☐ **B. On-demand instances**
- ☐ **C. Spot instances**
- ☐ **D. Regular instances**

Explanation :

Answer: A – When you have instances that will be used continuously and throughout the year, the best option is to buy reserved instances. By buying reserved instances, you are actually allocated an instance for the entire year or the duration you specify with a reduced cost.

To understand more on reserved instances, please visit the link:

- <https://aws.amazon.com/ec2/pricing/reserved-instances/>
(<https://aws.amazon.com/ec2/pricing/reserved-instances/>)
- <https://blog.cloudability.com/maximizing-cost-savings-aws-reserved-instances/>
(<https://blog.cloudability.com/maximizing-cost-savings-aws-reserved-instances/>)
- <https://awsinsider.net/articles/2017/03/21/controlling-aws-costs.aspx>
(<https://awsinsider.net/articles/2017/03/21/controlling-aws-costs.aspx>)

Ask our Experts



QUESTION 13 UNATTEMPTED

What is the ability provided by AWS to enable fast, easy, and secure transfers of files over long distances between your client and your Amazon S3 bucket.

- ☐ A. File Transfer
- ☐ B. HTTP Transfer
- ☐ C. Transfer Acceleration ✓
- ☐ D. S3 Acceleration

Explanation :

Answer - C

Please refer to the AWS S3 FAQ section - <https://aws.amazon.com/s3/faqs/>
(<https://aws.amazon.com/s3/faqs/>)

Q. What is Transfer Acceleration?

Amazon S3 Transfer Acceleration enables fast, easy, and secure transfers of files over long distances between your client and your Amazon S3 bucket. Transfer Acceleration leverages Amazon CloudFront's globally distributed AWS Edge Locations. As data arrives at an AWS Edge Location, data is routed to your Amazon S3 bucket over an optimized network path.

Ask our Experts



QUESTION 14 UNATTEMPTED

What is one key difference between an Amazon EBS-backed and an instance-store backed instance?

- ☐ A. Amazon EBS-backed instances can be stopped and restarted. ✓
- ☐ B. Instance-store backed instances can be stopped and restarted.

- ☐ C. Auto scaling requires using Amazon EBS-backed instances.
- ☐ D. Virtual Private Cloud requires EBS backed instances.

Explanation :

Answer A

Amazon EBS-backed instances can be stopped and restarted. So we can say Instance-store backed instances cannot be restarted.

- Please see the url for the key differences between EBS and instance store volumes - <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html> (<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>)

Ask our Experts



QUESTION 15 UNATTEMPTED

You have an application hosted in AWS. The logs from the application are sent to Cloudwatch. The application has recently been encountering some errors. A patch needs to be developed for the error to be rectified. For the moment you need to automate the restart of the server whenever the error occurs. How can you achieve this?

- ☐ A. Check the Cloudwatch logs for the error keywords , create an alarm and then restart the server ✓
- ☐ B. Create a cloudwatch metric which looks at the CPU utilization and then restarts the server
- ☐ C. Create a cloudwatch metric which looks at the Memory utilization and then restarts the server
- ☐ D. Check the Cloudwatch logs for the error keywords, then send a notification to SQS to restart the server

Explanation :

Answer – A

The AWS Documentation mentions the following on Cloudwatch Logs

You can use CloudWatch Logs to monitor applications and systems using log data. For example, CloudWatch Logs can track the number of errors that occur in your application logs and send you a notification whenever the rate of errors exceeds a threshold you specify. CloudWatch Logs uses your log data for monitoring; so, no code changes are required. For example, you can monitor application logs for specific literal terms (such as "NullPointerException") or count the number of occurrences of a literal term at a particular position in log data (such as "404" status codes in an Apache access log). When the term you are searching for is found, CloudWatch Logs reports the data to a CloudWatch metric that you specify. Log data is encrypted while in transit and while it is at rest.

For more information on Cloudwatch logs , please visit the below URL:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/WhatIsCloudWatchLogs.html>
(<http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/WhatIsCloudWatchLogs.html>)

Q. Can I point my zone apex (example.com versus www.example.com) at my Elastic Load Balancer?

Yes. Amazon Route 53 offers a special type of record called an 'Alias' record that lets you map your zone apex (*example.com*) DNS name to your ELB DNS name (i.e. *elb1234.elb.amazonaws.com*). IP addresses associated with Amazon Elastic Load Balancers can change at any time due to scaling up, scaling down, or software updates. Route 53 responds to each request for an Alias record with one or more IP addresses for the load balancer. Queries to Alias records that are mapped to ELB load balancers are free. These queries are listed as "Intra-AWS-DNS-Queries" on the Amazon Route 53 usage report.

Q. Can I point my zone apex (example.com versus www.example.com) at my website hosted on Amazon S3?

Yes. Amazon Route 53 offers a special type of record called an 'Alias' record that lets you map your zone apex (*example.com*) DNS name to your Amazon S3 website bucket (i.e. *example.com.s3-website-us-west-2.amazonaws.com*). IP addresses associated with Amazon S3 website endpoints can change at any time due to scaling up, scaling down, or software updates. Route 53 responds to each request for an Alias record with one IP address for the bucket. Route 53 doesn't charge for queries to Alias records that are mapped to an S3 bucket that is configured as a website. These queries are listed as "Intra-AWS-DNS-Queries" on the Amazon Route 53 usage report.

Ask our Experts



QUESTION 16

UNATTEMPTED

A company wants to utilize aws storage. For them low storage cost is paramount, the data is rarely retrieved, and data retrieval times of several hours are acceptable for them. What is the best storage option to use?

- ☒ A. Glacier ✓
- ☐ B. Reduced Redundancy Storage
- ☐ C. EBS backed storage connected to EC2

☐ D. Cloud Front

Explanation :

Answer – A.

With the above requirements, the best option is to opt for Amazon Glacier. Please refer to the Glacier FAQ's

<https://aws.amazon.com/glacier/faqs/> (<https://aws.amazon.com/glacier/faqs/>)

Q: How should I choose between Amazon Glacier and Amazon Simple Storage Service (Amazon S3)?

Amazon S3 is a durable, secure, simple, and fast storage service designed to make web-scale computing easier for developers. Use Amazon S3 if you need low latency or frequent access to your data. Use Amazon Glacier if low storage cost is paramount, your data is rarely retrieved, and data retrieval times of several hours are acceptable.

Ask our Experts



QUESTION 17

UNATTEMPTED

A client application requires operating system privileges on a relational database server. What is an appropriate configuration for a highly available database architecture?

- ☐ A. Standalone Amazon EC2 instance
- ☐ B. Amazon RDS in a Multi-AZ configuration
- ☐ C. Amazon EC2 instances in a replication configuration utilizing a Single Availability Zone
- ☐ D. Amazon EC2 instances in a replication configuration utilizing two different Availability Zones ✓

Explanation :

Answer – D

You can not access OS of RDS Databases, as RDS is fully managed service by AWS. In case a customer wants to have access to OS for their Database for more granular control or other compliance reason, then they can install their Database engine in EC2 instance. In choice D, DB needs to be installed in EC2 for OS access with replication to support failover.

Please follow below link for reference, which shows steps to install and configure Oracle in EC2 instance

<https://oracle-base.com/articles/vm/aws-ec2-installation-of-oracle>

(<https://oracle-base.com/articles/vm/aws-ec2-installation-of-oracle>)

Since the client wants privilege on the RDS, option B is not valid. Since there is a requirement for highly availability, you cannot have just one AZ and one EC2 instance. Hence D is the right answer.

Please refer below link showing an architecture example to enable Oracle database high availability on EC2 server.

<http://docs.aws.amazon.com/quickstart/latest/oracle-database/architecture.html>

(<http://docs.aws.amazon.com/quickstart/latest/oracle-database/architecture.html>)

For more information, please read the below link:

http://docs.aws.amazon.com/dms/latest/userguide/CHAP_Introduction.ReplicationInstance.html

(http://docs.aws.amazon.com/dms/latest/userguide/CHAP_Introduction.ReplicationInstance.html)

Ask our Experts



QUESTION 18

UNATTEMPTED

Per the AWS Acceptable Use Policy, penetration testing of EC2 instances:

- ☐ A. May be performed by AWS, and will be performed by AWS upon customer request.
- ☐ B. May be performed by AWS, and is periodically performed by AWS.
- ☐ C. Are expressly prohibited under all circumstances.
- ☐ D. May be performed by the customer on their own instances with prior authorization from AWS. ✓
- ☐ E. May be performed by the customer on their own instances, only if performed from EC2 instances.

Explanation :

Answer – D. You need to take prior authorization from AWS before doing a penetration test on EC2

Instances. Please refer to the below url for more details.
<https://aws.amazon.com/security/penetration-testing/>
(<https://aws.amazon.com/security/penetration-testing/>)

Ask our Experts



QUESTION 19 UNATTEMPTED

How can you secure data at rest on an EBS volume?

- ☐ A. Attach the volume to an instance using EC2's SSL interface.
- ☐ B. Write the data randomly instead of sequentially.
- ☐ C. Encrypt the volume using the S3 server-side encryption service.
- ☐ D. Create an IAM policy that restricts read and write access to the volume.
- ☒ E. Use an encrypted file system on top of the EBS volume. ✓

Explanation :

Answer – E. In order to secure data at rest on an EBS volume, you either have to encrypt the volume when it is being created or encrypt the data after the volume is created.

Ask our Experts



QUESTION 20 UNATTEMPTED

Which aws service is used to monitor all API calls to AWS

- ☐ A. Amazon SES
- ☒ B. Amazon Cloudtrail ✓
- ☐ C. Amazon CloudFront

☐ D. Amazon S3

Explanation :

Answer – B

Please refer to the product description for AWS Cloudtrail at the URL -
<https://aws.amazon.com/cloudtrail/> (<https://aws.amazon.com/cloudtrail/>)

Ask our Experts



QUESTION 21 UNATTEMPTED

If you want to launch Amazon Elastic Compute Cloud (EC2) instances and assign each instance a private IP address in a predetermined range, you should: (choose one of the correct answer below)

- ☐ A. Launch the instance from a private Amazon Machine Image (AMI).
- ☐ B. Assign a group of sequential Elastic IP address to the instances.
- ☐ C. Launch the instances in the Amazon Virtual Private Cloud (VPC). ✓
- ☐ D. Launch the instances in a Placement Group.
- ☐ E. Use standard EC2 instances since each instance gets a private Domain Name Service (DNS) already.

Explanation :

Answer – C.

This is the default reason for a VPC to host your own subnet and have EC2 instances have a private IP when it is launched in a VPC. Below is an example of an EC2 instance having a Private IP.

Instance ID	i-794961f7
Instance state	running
Instance type	t2.micro
Private DNS	ip-10-0-0-68.ap-southeast-1.compute.internal
Private IPs	10.0.0.68
Secondary private IPs	
VPC ID	vpc-39ddd55c
Subnet ID	subnet-858373f3
Network interfaces	eth0
Source/dest. check	False

For more information on private IP addresses, please refer the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-instance-addressing.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-instance-addressing.html>)

Ask our Experts



QUESTION 22 UNATTEMPTED

A company needs to deploy virtual desktops to its customers in a virtual private cloud, leveraging existing security controls. Which set of AWS services and features will meet the company's requirements?

- ☐ A. Virtual Private Network connection, AWS Directory Services, and ClassicLink
- ☐ B. Virtual Private Network connection, AWS Directory Services, and Amazon Workspaces ✓
- ☐ C. AWS Directory Service, Amazon Workspaces, and AWS Identity and Access Management
- ☐ D. Amazon Elastic Compute Cloud, and AWS Identity and Access Management

Explanation :

Answer – B

Option B is the correct answer because AWS Directory Services are used to authenticate to an existing on-premises AD through VPN and AWS WorkSpaces service is used for Virtual desktops.

Option A is incorrect because a ClassicLink, within the same region, allows us to link an EC2-Classic instance to a VPC in our account.

Option C is incorrect because AWS Dictionary service needs a VPN connection to interact with an On-premise AD directory.

Option D is incorrect because we need WorkSpaces for virtual desktops.

Ask our Experts



QUESTION 23 UNATTEMPTED

Which of the following statements are true about Amazon Reduced Redundancy Storage (RRS)? Choose the correct 3 answers from the below options.

- ☐ A. RRS has the ability to provide eleven nines availability.
- ☐ B. RRS has the ability to provide 99.99% availability. ✓
- ☐ C. RRS has the ability to provide 99.99% durability. ✓
- ☐ D. If there is a requirement to store data that is easily reproducible or durably stored elsewhere, then RRS is the ideal option. ✓

Explanation :

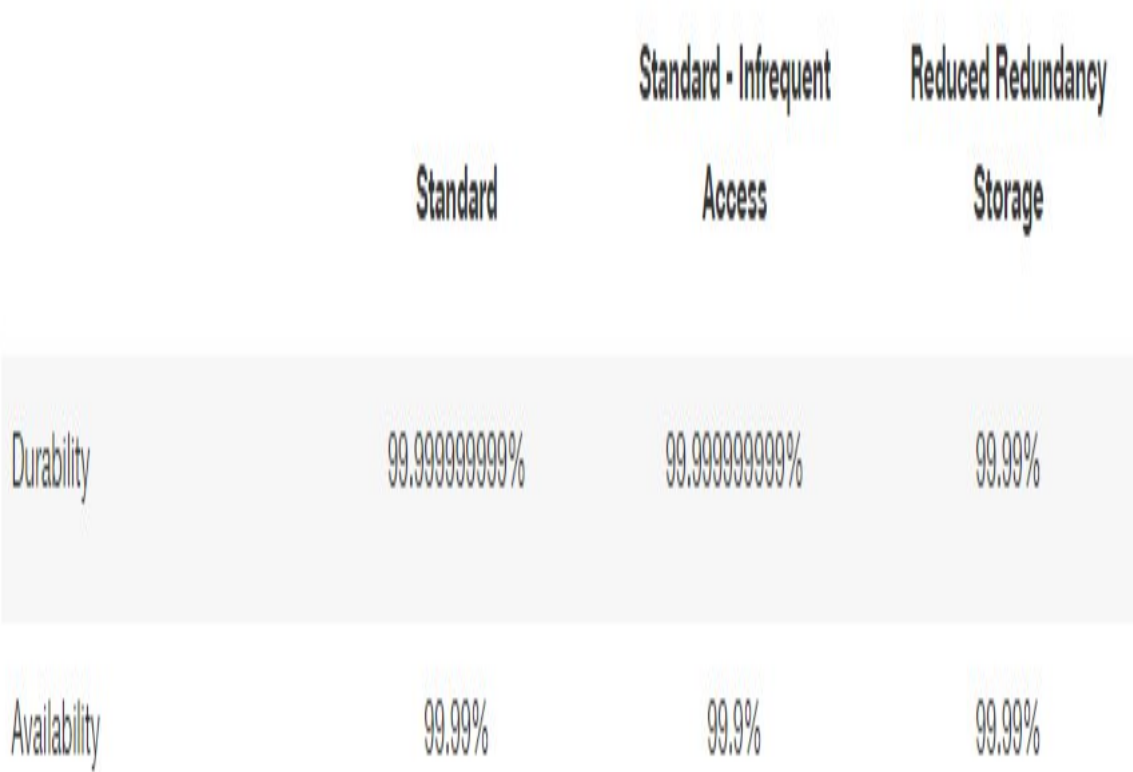
Answer: B,C and D. The Durability and availability are given in the aws site for RRS.

- Reduced Redundancy Storage (RRS) is an Amazon S3 storage option that enables customers to store noncritical, reproducible data at lower levels of redundancy than Amazon S3's standard storage.
- S3 is the most reliable and durable storage service from Amazon. Where as if you have data that is non-critical and can be easily reproducible if lost, then that can be stored in RRS to reduce the cost of your storage.
- The RRS option stores objects on multiple devices across multiple facilities, providing 400 times the durability of a typical disk drive, but does not replicate objects as many times as standard Amazon S3 storage.

You can read more about RRS in the below link:

- <https://aws.amazon.com/s3/reduced-redundancy/> (<https://aws.amazon.com/s3/reduced-redundancy/>)

Here is the snapshot of durability and availability of RRS compared to other standard storage options:



	Standard	Standard - Infrequent Access	Reduced Redundancy Storage
Durability	99.999999999%	99.999999999%	99.99%
Availability	99.99%	99.9%	99.99%
Concurrent facility fault tolerance	2	2	1
SSL support	Yes	Yes	Yes
First byte latency	Milliseconds	Milliseconds	Milliseconds
Lifecycle Management Policies	Yes	Yes	Yes

Ask our Experts



QUESTION 24 UNATTEMPTED

After creating a new IAM user which of the following must be done before they can successfully make API calls?

- ☐ A. Add a password to the user.
- ☐ B. Enable Multi-Factor Authentication for the user.
- ☐ C. Assign a Password Policy to the user.
- ☐ D. Create a set of Access Keys for the user. ✓

Explanation :

Answer – D

In IAM , when you create a user , you need to download the Access Key ID and Secret access key so that the user can access aws.

User	Access key ID	Secret access key
▼ ✓ demouser	AKIAJ6KYKT7WJUHHSHTQ	***** Show
<div>✓ Created user demouser</div> <div>✓ Created access key for user demouser</div>		

Ask our Experts



QUESTION 25 UNATTEMPTED

What is the AWS service provided which provides a fully managed NoSQL database service that provides fast and predictable performance with seamless scalability.

- ☐ A. AWSRDS
- ☐ B. DynamoDB ✓
- ☐ C. Oracle RDS
- ☐ D. Elastic Map Reduce

Explanation :

Answer: - B

DynamoDB is a fully managed NoSQL offering provided by AWS. It is now available in most regions for users to consume. The link provides the full details on the product

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Introduction.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Introduction.html>)

Ask our Experts



QUESTION 26 UNATTEMPTED

A company's application is intending to use Auto Scaling and has the requirement to store user state information.

Which of the following AWS services provides a shared data store with durability and low latency?

- ☐ A. AWS ElastiCache Memcached
- ☐ B. Amazon Simple Storage Service
- ☐ C. Amazon EC2 instance storage
- ☒ D. Amazon DynamoDB ✓

Explanation :

Answer – D

Amazon Dynamo DB is used for storing small amounts of data such as user state information. And this service offer's durability and low latency. Below is a snapshot of when to use S3 and DynamoDB from the DynamoDB FAQ's - <https://aws.amazon.com/dynamodb/faqs/>
(<https://aws.amazon.com/dynamodb/faqs/>)

Q: When should I use Amazon DynamoDB vs Amazon S3?

Amazon DynamoDB stores structured data, indexed by primary key, and allows low latency read and write access to items ranging from 1 byte up to 400KB. Amazon S3 stores unstructured blobs and suited for storing large objects up to 5 TB. In order to optimize your costs across AWS services, large objects or infrequently accessed data sets should be stored in Amazon S3, while smaller data elements or file pointers (possibly to Amazon S3 objects) are best saved in Amazon DynamoDB.

Note:

- Memcached is shared and low latency, but its durability is poor.
- Instance storage is low latency, but it is neither shared nor durable.
- S3 is both shared and amazingly durable, but it has trouble in the latency department: S3 really shines for static data, but it only offers eventual consistency for overwrite puts and deletes. Waiting for seconds or possibly even hours for a session state update to synchronise for all the servers is not low latency.
- DynamoDB is perfect for this: it is a "shared data store with durability and low latency", and it offers strongly consistent reads."

Ask our Experts



QUESTION 27 UNATTEMPTED

Which of the following requires a custom CloudWatch metric to monitor?

- ☒ A. Memory Utilization of an EC2 instance ✓
- ☐ B. CPU Utilization of an EC2 instance
- ☐ C. Disk usage activity of an EC2 instance
- ☐ D. Data transfer of an EC2 instance

Explanation :

Answer – A


Memory Utilization is a metric not offered directly by Cloudwatch.

So when you view the Cloudwatch metrics for your EC2 instance, you can see CPU Utilization and Disk Read Operations metrics.


Below are your CloudWatch metrics for the selected resources (a maximum of 10). Click on a graph to see an expanded view. All times shown are

› [View all CloudWatch metrics](#)


CPU Utilization (Percent)



Disk Reads (Bytes)



Disk Read Operations (Operations)



You can also see Network statistics for Data transfer, but you will not be able to see Memory Utilization. This will be a custom Cloudwatch metric.

Network Out (Bytes)



Network Packets In (Count)



Network Packets Out (Count)



Ask our Experts



QUESTION 28

UNATTEMPTED

Which of the following notification endpoints or clients are supported by Amazon Simple Notification Service? Choose 2 answers from the options below.

☐ A. Email ✓

☐ B. CloudFront distribution

- ☐ C. File Transfer Protocol
- ☐ D. Short Message Service ✓
- ☐ E. Simple Network Management Protocol

Explanation :

Answer – A and D

When you create a subscription in SNS , below are the protocols available.

Create Subscription

Building a mobile app? Try [AWS Mobile Hub](#).

Topic ARN

Protocol

Endpoint

Subscription

Ask our Experts



QUESTION 29 UNATTEMPTED

You have a read intensive application hosted in AWS. The application is currently using the MySQL RDS feature in AWS. Which of the following can be used to reduce the read throughput on the MySQL database

- ☐ A. Enable the Multi-AZ on the MySQL RDS
- ☐ B. Use Cold Storage Volumes for the MySQL RDS
- ☒ C. Enable Read Replica's and offload the reads to the replica's ✓
- ☐ D. Use SQS to queue up the reads

Explanation :

Answer – C

The AWS documentation mentions the following on Read Replica's

Amazon RDS Read Replicas provide enhanced performance and durability for database (DB) instances. This replication feature makes it easy to elastically scale out beyond the capacity constraints of a single DB Instance for read-heavy database workloads. You can create one or more replicas of a given source DB Instance and serve high-volume application read traffic from multiple copies of your data, thereby increasing aggregate read throughput. Read replicas can also be promoted when needed to become standalone DB instances.

For more information on Read Replica's , please visit the below URL:

- <https://aws.amazon.com/rds/details/read-replicas/>
(<https://aws.amazon.com/rds/details/read-replicas/>)

Ask our Experts



QUESTION 30 UNATTEMPTED

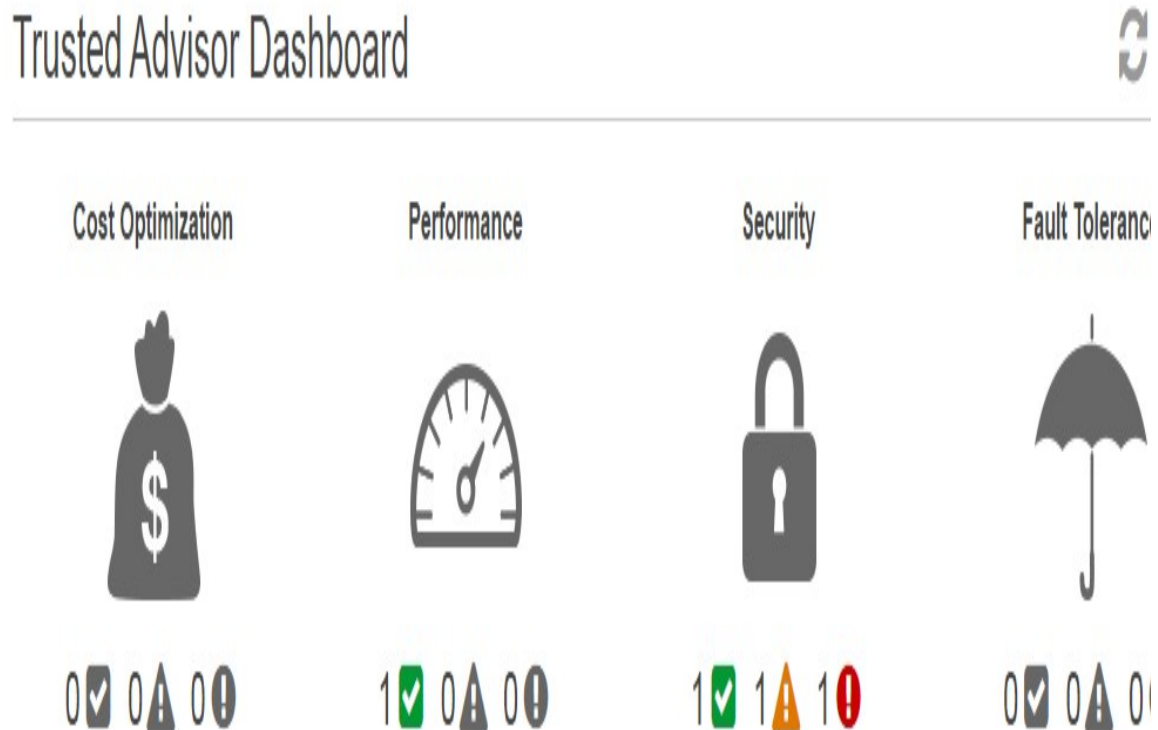
The Trusted Advisor service provides insight regarding which five categories of an AWS account?

- ☐ A. Security, fault tolerance, high availability, connectivity and service limits
- ☐ B. Security, access control, high availability, performance and service limits
- ☒ C. Performance, cost optimization, security, fault tolerance and service limits ✓
- ☐ D. Performance, cost optimization, access control, connectivity and service limits

Explanation :

Answer – C

Below is the screenshot of what services the Trusted Advisor Dashboard offers.



Recommended Actions

Please check the below AWS Docs for more details,

- <https://aws.amazon.com/premiumsupport/trustedadvisor/>
(<https://aws.amazon.com/premiumsupport/trustedadvisor/>)

Ask our Experts



QUESTION 31 UNATTEMPTED

A company is deploying a two-tier, highly available web application to AWS. Which service provides durable storage for static content while utilizing lower Overall CPU resources for the web tier?

- ☐ A. Amazon EBS volume
- ☒ B. Amazon S3 ✓
- ☐ C. Amazon EC2 instance store
- ☐ D. Amazon RDS instance

Explanation :

Answer – B

Amazon S3 is the default storage service that should be considered for companies. It provides durable storage for all static content.

Ask our Experts



QUESTION 32 UNATTEMPTED

When will you incur costs with an Elastic IP address (EIP)?

- ☐ A. When an EIP is allocated.
- ☐ B. When it is allocated and associated with a running instance.
- ☒ C. When it is allocated and associated with a stopped instance. ✓
- ☐ D. Costs are incurred regardless of whether the EIP is associated with a running instance.

Explanation :

Answer- C

The correct answer for this question is option "C". The option D is a little bit tricky and which will make us think that it might be correct even though it's not. Following AWS docs shows us when costs are not incurred.

An Elastic IP address doesn't incur charges as long as the following conditions are true:

- The Elastic IP address is associated with an Amazon EC2 instance.
- The instance associated with the Elastic IP address is running.
- The instance has only one Elastic IP address attached to it.

If you've stopped or terminated an EC2 instance with an associated Elastic IP address and you don't need that Elastic IP address any more, consider disassociating or releasing the Elastic IP address by following the instructions at Working with Elastic IP Addresses (<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/elastic-ip-addresses-eip.html#working-with-eips>).

Note: After an Elastic IP address is released, you can't provision that same Elastic IP address again, though you can provision a different Elastic IP address.

Reference link: <https://aws.amazon.com/premiumsupport/knowledge-center/elastic-ip-charges/>
(<https://aws.amazon.com/premiumsupport/knowledge-center/elastic-ip-charges/>)

AWS doesn't want you waste the static public IP's . You will be charged for elastic IP

1 - If EIP is created but not allocated to any instance.

2 - If EIP is attached to a stop instance.

Please find details below regarding Elastic IP Charges:

<https://aws.amazon.com/premiumsupport/knowledge-center/elastic-ip-charges/>
(<https://aws.amazon.com/premiumsupport/knowledge-center/elastic-ip-charges/>)



Ask our Experts



Which one of the following instances provides Non-Volatile Memory Express (NVMe) SSD-backed Instance storage?

- ☐ A. General Purpose T2
- ☒ B. Storage Optimised I3 ✓
- ☐ C. Memory Optimized M5
- ☐ D. Compute Optimized C5

Explanation :

Answer – B

I3 provides Non-Volatile Memory Express (NVMe) SSD-backed Instance storage optimized for low latency, very high random I/O performance, high sequential read throughput and provide high IOPS at a low cost.

For more information, Please check the following AWS Docs:

- <https://aws.amazon.com/ec2/instance-types/> (<https://aws.amazon.com/ec2/instance-types/>)

Ask our Experts





QUESTION 34 UNATTEMPTED

How many availability zones are mapped to a subnet?

- ☒ A. 1 ✓
- ☐ B. 2
- ☐ C. Depends on aws at the time of creating a subnet
- ☐ D. Depends on the number of instances you are going to host in the subnet.

Explanation :

Answer: A – Remember that when a subnet is created, it is always mapped to one availability zone. When you go to the VPC dashboard, and go to the Subnet section, you can click on Create Subnet

 Services ▾ Resource Groups ▾ 

VPC Dashboard

Filter by VPC: None ▾

Virtual Private Cloud

Your VPCs

Subnets

Route Tables

Internet Gateways

Create Subnet Subnet Actions ▾

X


<input type="checkbox"/>	Name ▴	Subnet ID ▾
<input type="checkbox"/>		subnet-b7ac40fe
<input type="checkbox"/>		subnet-e1665acb
<input type="checkbox"/>		subnet-e5d9e5bd
<input type="checkbox"/>		subnet-b1c7c88c

When you create the subnet, you can only attach one AZ to the subnet.


Create Subnet X

Use the CIDR format to specify your subnet's IP address block (e.g., 10.0.0.0/24). Note that block sizes must be between a /16 netmask and /28 netmask. Also, note that a subnet can be the same size as your VPC.


Name tag




VPC

vpc-6dcc550a (172.31.0.0/16) ▾ 

Availability Zone

us-east-1a ▾ 

CIDR block



Cancel

Yes, Create

Ask our Experts



A company is building a service in which Amazon EC2 worker instances process an uploaded audio file and generate a text file. You must store both of these files in the same durable storage until the text file is retrieved. You do not know what the storage capacity requirements are. Which storage option is both cost-efficient and scalable?

- ☐ A. Multiple Amazon EBS volume with snapshots
- ☐ B. A single Amazon Glacier vault
- ☒ C. A single Amazon S3 bucket ✓
- ☐ D. Multiple instance stores

Explanation :

Answer – C

For any sort of storage for file based system, it must be done in Amazon S3.

Ask our Experts



QUESTION 36 UNATTEMPTED

A custom script needs to be passed to a new Amazon Linux instances created in your Auto Scaling group. Which feature allows you to accomplish this?

- ☒ A. User data ✓
- ☐ B. EC2Config service
- ☐ C. IAM roles
- ☐ D. AWS Config

Explanation :

Answer – A

When you configure an instance during creation, you can add custom scripts to the User data section.

So in Step 3 of creating an instance, in the Advanced Details section, we can enter custom scripts in the User Data section. The below script installs Perl during the instance creation of the EC2 instance.

Step 3: Configure Instance Details

Additional charges apply.

Tenancy ⓘ

Shared - Run a shared hardware instance ▼

Additional charges will apply for dedicated tenancy.

▼ Network interfaces ⓘ

Device	Network Interface	Subnet	Primary IP	Secondary IP addresses
eth0	New network interface ▼	subnet-95ed8dd1 ▼	Auto-assign	Add IP
<div>Add Device</div>				

▼ Advanced Details

User data ⓘ

☒ As text ☐ As file ☐ Input is already base64 encoded

```
yum install perl-Switch perl-DateTime perl-Sys-Syslog perl-LWP-Protocol-https -y
```

Ask our Experts



QUESTION 37

UNATTEMPTED

A company is building software on AWS that require access to various AWS services. Which configuration should be used to ensure that AWS Credentials like Access Keys and Secret access keys are not compromised? (Choose Two Options)

☐ A. Enable Multi-Factor Authentication for your AWS root account. ✓

- ☐ B. Assign an IAM role to the Amazon EC2 instance. ✓
- ☐ C. Store the AWS Access Key ID/Secret Access Key combination in software comments.
- ☐ D. Assign an IAM user to the Amazon EC2 Instance.

Explanation :

Answer – A and B

It is the best practice to always create IAM roles which can be assigned to EC2 instances and enable MFA for the root account. This will help to not compromise the Access Key ID/Secret Access Key combination.

The correct answer is: Enable Multi-Factor Authentication for your AWS root account and Assign an IAM role to the Amazon EC2 instance.

Ask our Experts



QUESTION 38 UNATTEMPTED

A company has the requirement to store data using AWS storage services. The data is not frequently accessed. If data recovery time not an issue, Which of the below is the best and cost efficient solution to fulfil this requirement ?

- ☐ A. S3 Standard
- ☐ B. S3 Standard - IA (Infrequently Accessed)
- ☐ C. Glacier ✓
- ☐ D. Reduced Redundancy Storage?

Explanation :

Answer – C

The default time interval is one minute.

Note: Answer can also be B. S3 Standard - IA (Infrequently Accessed). However since other details are mentioned in question. we can say C. Glacier is most effective way of cost saving in this case.

Reference link: <https://aws.amazon.com/products/storage/>
(<https://aws.amazon.com/products/storage/>)

(<https://aws.amazon.com/cloudwatch/faqs/>)

Ask our Experts



QUESTION 39 UNATTEMPTED

Resources that are created in AWS are identified by a unique identifier which is known as what option given below

- ☐ A. Amazon Resource Number
- ☐ B. Amazon Resource Nametag
- ☐ C. Amazon Resource Name ✓
- ☐ D. Amazon Resource Namespace

Explanation :

Answer – C

Amazon Resource Names (ARNs) are used to uniquely identify AWS resources. For information on ARN's, refer to the link -

<http://docs.aws.amazon.com/general/latest/gr/aws-arns-and-namespaces.html>
(<http://docs.aws.amazon.com/general/latest/gr/aws-arns-and-namespaces.html>)

Ask our Experts



QUESTION 40 UNATTEMPTED

When you are using Route53 for a web site hosted in S3 , what are the one of the rules that must be adhered to? Choose the correct answer from the options below

- ☐ A. The S3 bucket name must be the same as the domain name ✓
- ☐ B. The record set cannot use an alias
- ☐ C. The record set must be of type "MX"
- ☐ D. The S3 bucket must be in the same region as the hosted zone

Explanation :

Answer – A

This is given in the aws documentation

Prerequisites

Before you get started, you need the following:

- An S3 bucket that is configured to host a static website. For more information, see [Configure a Bucket for Website Hosting](#) in the *Amazon Simple Storage Service Developer Guide*.

Important

The bucket must have the same name as your domain or subdomain. For example, if you want to use the subdomain `acme.example.com`, the name of the bucket must be `acme.example.com`.

For more information on using Route53 along with S3, please visit the link

<http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/RoutingToS3Bucket.html>

(<http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/RoutingToS3Bucket.html>)

Ask our Experts



QUESTION 41 UNATTEMPTED

What are some of the benefits of using the Cloudformation service? Choose 2 answers from the options given below

- ☐ A. Can automatically increase instance capacity
- ☐ B. A storage location for your applications code
- ☐ C. Infrastructure as Code ✓
- ☐ D. A great disaster recovery option ✓

Explanation :

Answer – C and D

The justification for Infrastructure as code is given in the aws documentation

Infrastructure as Code

A template can be used repeatedly to create identical copies of the same stack (or to use as a foundation to start a new stack). You can capture and control region-specific infrastructure variations such as Amazon EC2 AMIs, as well as Amazon EBS and Amazon RDS snapshot names. Templates are simple JSON or YAML formatted text files that can be placed under your normal source control mechanisms, stored in private or public locations such as Amazon S3, and exchanged via email. With AWS CloudFormation, you can "open the hood," to see exactly which AWS resources make up a stack. You retain full control and have the ability to modify any of the AWS resources created as part of a stack.

For the justification on disaster recovery, please visit the below link

<https://aws.amazon.com/blogs/aws/new-whitepaper-use-aws-for-disaster-recovery/>

(<https://aws.amazon.com/blogs/aws/new-whitepaper-use-aws-for-disaster-recovery/>)

For more information on Cloudformation, please visit the link

<https://aws.amazon.com/cloudformation/> (<https://aws.amazon.com/cloudformation/>)

Ask our Experts



QUESTION 42 UNATTEMPTED

AWS thrives on the concept of high availability. Which of the below follows the concept of high availability. Choose the correct answer from the options below

- ☐ A. Implementing security procedures
- ☐ B. Implementing multiple AWS services
- ☐ C. The ability of system to easily increase in size.
- ☐ D. A durable system that can operate for long periods of time without failure. ✓

Explanation :

Answer – D

High availability is a characteristic of a system, which aims to ensure an agreed level of operational performance, usually uptime, for a higher than normal period.

For more information on high availability, please refer to the following link

https://en.wikipedia.org/wiki/High_availability (https://en.wikipedia.org/wiki/High_availability)

Ask our Experts



QUESTION 43 UNATTEMPTED

What best describes the "Principal of Least Privilege"? Choose the correct answer from the options given below

- ☐ A. All users should have the same baseline permissions granted to them to use basic AWS services.
- ☐ B. Users should be granted permission to access only resources they need to do their assigned job. ✓
- ☐ C. Users should submit all access request in written so that there is a paper trail of who needs access to different AWS resources.
- ☐ D. Users should always have a little more access granted to them than they need, just in case they end up needing it in the future.

Explanation :

Answer – B

The principle means giving a user account only those privileges which are essential to perform its intended function. For example, a user account for the sole purpose of creating backups does not need to install software: hence, it has rights only to run backup and backup-related applications.

For more information on principle of least privilege, please refer to the following link

https://en.wikipedia.org/wiki/Principle_of_least_privilege
(https://en.wikipedia.org/wiki/Principle_of_least_privilege)

Ask our Experts



Which of the following best describes the purpose of an Elastic Load Balancer. Choose an answer from the options given below. Choose the correct answer from the options given below

- ☐ A. To scale more EC2 instances on demand
- ☐ B. To evenly distribute traffic among multiple EC2 instances located in single or different Availability Zones. ✓
- ☐ C. To distribute traffic to a second instance once the first instance capacity has reached it's limit.
- ☐ D. To evenly distribute traffic among multiple EC2 instances in the same Availability Zone.

Explanation :

Answer – B

Elastic Load Balancing automatically distributes incoming application traffic across multiple Amazon EC2 instances. It enables you to achieve fault tolerance in your applications, seamlessly providing the required amount of load balancing capacity needed to route application traffic.

And the ELB is used to distribute traffic between instances in Multiple AZ's.

For more information on Elastic Load Balancer, please refer to the following link

- <https://aws.amazon.com/elasticloadbalancing/>
(<https://aws.amazon.com/elasticloadbalancing/>)

Some more key points about ELB:

- Elastic Load Balancer (ELB) is used for routing traffic to various EC2 instances located across the multiple Availability Zones(AZs).
- ELB can detect the healthy and unhealthy EC2 instances. It will not route traffic to the unhealthy EC2 instances.
- If all the instances in the same AZ is not healthy, it will route the traffic to other AZ EC2 instances.
- Achieve higher levels of fault tolerance for your applications by using Elastic Load Balancing to automatically route traffic across multiple instances and multiple Availability Zones. Elastic Load Balancing ensures that only healthy Amazon EC2 instances receive traffic by detecting unhealthy instances and rerouting traffic across the remaining healthy instances. If all of your EC2 instances in one Availability Zone are unhealthy, and you have set up EC2 instances in multiple Availability Zones, Elastic Load Balancing will route traffic to your healthy EC2 instances in those other zones.

Elastic Load Balancing

Elastic Load Balancing automatically distributes incoming application traffic across multiple Amazon EC2 instances. It enables you to achieve fault tolerance in your applications, seamlessly providing the required amount of load balancing capacity needed to route application traffic.

Elastic Load Balancing offers two types of load balancers that both feature high availability, automatic scaling, and robust security. These include the [Classic Load Balancer](#) that routes traffic based on either application or network level information, and the [Application Load Balancer](#) that routes traffic based on advanced application level information that includes the content of the request. The

Ask our Experts



QUESTION 45 UNATTEMPTED

When AWS create a default VPC, what are the services you get by default in the VPC?

Select 2 options.

- ☐ A. An Elastic Load Balancer
- ☐ B. Default subnet in each Availability Zone ✓
- ☐ C. An Internet Gateway attached to the default VPC ✓

- ☐ D. A light weight rds such as SQL Server Express.

Explanation :

Answer: B and C.

Default VPC Components

When we create a default VPC, we do the following to set it up for you:

- Create a VPC with a size /16 IPv4 CIDR block (172.31.0.0/16). This provides up to 65,536 private IPv4 addresses.
- Create a size /20 default subnet in each Availability Zone. This provides up to 4,096 addresses per subnet, a few of which are reserved for our use.
- Create an **internet gateway** and connect it to your default VPC.
- Create a main route table for your default VPC with a rule that sends all IPv4 traffic destined for the internet to the internet gateway.
- Create a default security group and associate it with your default VPC.
- Create a default network access control list (ACL) and associate it with your default VPC.
- Associate the default DHCP options set for your AWS account with your default VPC.

For the list of default services given for a default VPC, follow the link

<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/default-vpc.html>

(<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/default-vpc.html>)

to get more information on what comes as part of a default VPC.

Ask our Experts



QUESTION 46 UNATTEMPTED

In an Autoscaling policy what best describes what the purpose of a scaling policy is. Choose an answer from the options below.

- ☒ A. A set of CloudWatch metric thresholds that dictate when to add or remove instances from the Auto Scaling group. ✓
- ☐ B. The IAM access policy granted to an Auto Scaling group.
- ☐ C. The percentage at which an ELB will send traffic to an instance before it sends traffic to a different instance.
- ☐ D. An SNS notification alert.

Explanation :

Answer – A

You can create a scaling policy that uses CloudWatch alarms to determine when your Auto Scaling group should scale out or scale in. Each CloudWatch alarm watches a single metric and sends messages to Auto Scaling when the metric breaches a threshold that you specify in your policy. You can use alarms to monitor any of the metrics that the services in AWS that you're using send to CloudWatch, or you can create and monitor your own custom metrics.

For more information on Scaling policies, please refer to the following link

http://docs.aws.amazon.com/autoscaling/latest/userguide/policy_creating.html

(http://docs.aws.amazon.com/autoscaling/latest/userguide/policy_creating.html)

Ask our Experts



QUESTION 47

UNATTEMPTED

A company has a solution hosted in AWS. This solution consists of a set of EC2 instances. They have been recently getting attacks as their IT security departments identified that attacks are from a set of IP addresses. Which of the following methods can be adopted to help in this situation.

- ☐ A. Place the EC2 instances into private subnets, and set up an NAT gateway so employees can access them.
- ☐ B. Remove the IGW from the VPC so that no outside traffic can reach the EC2 instances.
- ☐ C. Lock down of NACL for the set to IP address. ✓
- ☐ D. Place the EC2 instances into private subnets, and set up a bastion host so employees can access them.

Explanation :

Answer – C

The NACL's can be modified to be most secure by only denying the traffic from the set of IP addresses.

For more information on NACL, please refer to the following link

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLS.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLS.html)

Ask our Experts



QUESTION 48 UNATTEMPTED

You have an ELB distributing traffic a fleet of EC2 instances inside your VPC, evenly spread across two Availability Zones. However, you realize that only half of our instances are actually receiving traffic. What is the most likely cause of this problem? Choose the correct answer from the options given below

- ☐ A. The ELBs listener is not set to port 80.
- ☐ B. One or more security groups do not allow HTTP traffic.
- ☐ C. Cross-zone load balancing has not been enabled. ✓
- ☐ D. The health check ping port is set to port 80, but should be set to port 22.

Explanation :

Answer – C

For environments where clients cache DNS lookups, incoming requests might favor one of the Availability Zones. Using cross-zone load balancing, this imbalance in the request load is spread across all available instances in the region, reducing the impact of misbehaving clients.

By default, your Classic Load Balancer distributes incoming requests evenly across its enabled Availability Zones. For example, if you have ten instances in Availability Zone us-west-2a and two instances in us-west-2b, the requests are distributed evenly between the two Availability Zones. As a result, the two instances in us-west-2b serve the same amount of traffic as the ten instances in us-west-2a. To ensure that your load balancer distributes incoming requests evenly across all instances in its enabled Availability Zones, enable cross-zone load balancing.

For environments where clients cache DNS lookups, incoming requests might favor one of the Availability Zones. Using cross-zone load balancing, this imbalance in the request load is spread across all available instances in the region, reducing the impact of misbehaving clients.

Even with Cross-zone load balancing in place, AWS recommends you to maintain approximately

equivalent numbers of instances in each enabled Availability Zone for higher fault tolerance.

Since we are using an ELB in front of these ec2 instances to direct the traffic to all these instances we assume that these instances are all of the same type with the same application running on it and with the same Security groups assigned to it. So the option "One or more security groups do not allow HTTP traffic" won't be the right option since we are using the same security group for these instances.

However it seems that the word "evenly" is causing a bit of ambiguity in your mind. Consider that if few instances are failed in the health check, then ELB won't send traffic to those instances although it is not explicitly mentioned in the question. In those scenarios even though ec2 instances are evenly spread across AZ's the traffic won't be evenly distributed among instances if cross zone is not enabled.

More information is available at:

<https://docs.aws.amazon.com/elasticloadbalancing/latest/classic/enable-disable-crosszone-lb.html> (<https://docs.aws.amazon.com/elasticloadbalancing/latest/classic/enable-disable-crosszone-lb.html>)

Ask our Experts



QUESTION 49 UNATTEMPTED

You have an application currently running on five EC2 instances as part of an Auto Scaling group. For the past 30 minutes all five instances have been running at 100 CPU Utilization; however, the Auto Scaling group has not added any more instances to the group. What is the most likely cause?

Select 2 answers

- ☐ A. You already have 20 on-demand instances running. ✓
- ☐ B. The Auto Scaling group's MAX size is set at five. ✓
- ☐ C. The Auto Scaling group's scale down policy is too high.
- ☐ D. The Auto Scaling group's scale up policy has not yet been reached.

Explanation :

Answer – A and B

<number of instances> instance(s) are already running. Launching EC2 instance failed.

- **Cause:** The Auto Scaling group has reached the limit set by the `DesiredCapacity` parameter.
- **Solution:**
 - Update your Auto Scaling group by providing a new value for the `--desired-capacity` parameter using the `update-auto-scaling-group` command.
 - If you've reached your limit for number of EC2 instances, you can request an increase. For more information, see [AWS Service Limits](#).

This is provided in the aws documentation

Instance Limits	
Name	Current Limit
Running On-Demand EC2 instances ⓘ	20
Running On-Demand c4.2xlarge	The total number of running On-Demand instances that you can have in this region. Some instance types have different limits for this region that count against your total limit; these are listed below. Check the Current Limit column to find out how many instances per instance type you can run.
Running On-Demand c4.4xlarge	
Running On-Demand c4.8xlarge	
Running On-Demand c4.large	
Running On-Demand c4.xlarge	
Running On-Demand c5.18xlarge instances	0
Running On-Demand c5.2xlarge instances	5
Running On-Demand c5.4xlarge instances	1
Running On-Demand c5.9xlarge instances	1

For more information on troubleshooting Autoscaling, please refer to the following link

<http://docs.aws.amazon.com/autoscaling/latest/userguide/ts-as-capacity.html>
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/ts-as-capacity.html>)

Ask our Experts



If need to upload a file to S3 that is 600MB in size, which of the following is the best option to use? Choose an answer from the options below. Choose the correct answer from the options below

- ☐ A. Single operation upload
- ☐ B. Snowball
- ☐ C. AWS Import/Export
- ☒ D. Multi-part upload ✓

Explanation :

Answer – D

The Multipart upload API enables you to upload large objects in parts. You can use this API to upload new large objects or make a copy of an existing object (see Operations on Objects). Multipart uploading is a three-step process: You initiate the upload, you upload the object parts, and after you have uploaded all the parts, you complete the multipart upload. Upon receiving the complete multipart upload request, Amazon S3 constructs the object from the uploaded parts, and you can then access the object just as you would any other object in your bucket.

For more information on Multi-part file upload, please refer to the following link

<http://docs.aws.amazon.com/AmazonS3/latest/dev/mpuoverview.html>

(<http://docs.aws.amazon.com/AmazonS3/latest/dev/mpuoverview.html>)

Note:

This question is about "which of the following is the best option to use?".

Even you can upload up to 5GB in a single operation, it is recommended and the best way to choose multipart upload.

AWS says as mentioned below:

The multipart upload API is designed to improve the upload experience for larger objects. You can upload objects in parts. These object parts can be uploaded independently, in any order, and in parallel. You can use a multipart upload for objects from 5 MB to 5 TB in size.

We recommend that you use multipart uploading in the following ways:

- If you're uploading large objects over a stable high-bandwidth network, use multipart uploading to maximize the use of your available bandwidth by uploading object parts in parallel for multi-threaded performance.

- If you're uploading over a spotty network, use multipart uploading to increase resiliency to network errors by avoiding upload restarts. When using multipart uploading, you need to retry uploading only parts that are interrupted during the upload. You don't need to restart uploading your object from the beginning.

Ask our Experts



QUESTION 51 UNATTEMPTED

A company has the requirement to store files in S3. After a period of a month, these files can be archived. The archived files might be required after a period of 3-4 months. Which of the following suits the requirements

- ☐ A. Use EC2 instances with EBS volumes, one for normal storage and the other for archived storage
- ☐ B. Use S3 for normal file storage and use lifecycle policies for moving the files to glacier. ✓
- ☐ C. Use EC2 instances with EBS volumes and use lifecycle policies for moving the files to glacier.
- ☐ D. Use glacier for normal file storage and use lifecycle policies for moving the files to S3.

Explanation :

Answer – B

Lifecycle configuration enables you to specify the lifecycle management of objects in a bucket. The configuration is a set of one or more rules, where each rule defines an action for Amazon S3 to apply to a group of objects. These actions can be classified as follows:

- Transition actions – In which you define when objects transition to another storage class. For example, you may choose to transition objects to the STANDARD_IA (IA, for infrequent access)

storage class 30 days after creation, or archive objects to the GLACIER storage class one year after creation.

For more information on Lifecycle policies, please refer to the following link
<http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>)

Ask our Experts



QUESTION 52 UNATTEMPTED

Your team has an application hosted on Docker containers. You want to port that application in the easiest way possible onto AWS for your development community. Which of the following service can be used to fulfil this requirement

- ☐ A. AWS Elastic Load Balancer
- ☐ B. AWS SNS
- ☐ C. AWS SQS
- ☐ D. AWS Elastic Beanstalk ✓

Explanation :

Answer – D

The AWS documentation mentions the following

Elastic Beanstalk supports the deployment of web applications from Docker containers. With Docker containers, you can define your own runtime environment. You can choose your own platform, programming language, and any application dependencies (such as package managers or tools), that aren't supported by other platforms. Docker containers are self-contained and include all the configuration information and software your web application requires to run.

For more information on Elastic beanstalk and docker , please visit the below URL:

- http://docs.aws.amazon.com/elasticbeanstalk/latest/dg/create_deploy_docker.html
(http://docs.aws.amazon.com/elasticbeanstalk/latest/dg/create_deploy_docker.html)

Ask our Experts



QUESTION 53 UNATTEMPTED

Which of the following AWS System Manager for EC2 can be used to administer instances without the need to SSH or RDP into the instance.

- ☐ A. AWS Config
- ☐ B. AWS CodePipeline
- ☐ C. Run Command ✓
- ☐ D. EC2Config

Explanation :

Answer - C

You can use Run Command from the Amazon EC2 console to configure instances without having to login to each instance

For more information on the Run Command , please visit the below URL:

- <http://docs.aws.amazon.com/systems-manager/latest/userguide/rc-console.html>
(<http://docs.aws.amazon.com/systems-manager/latest/userguide/rc-console.html>)

Ask our Experts



QUESTION 54 UNATTEMPTED

If you wanted to extend your on-premise infrastructure with AWS, which of the below options would help. Choose 2 answers from the options given below

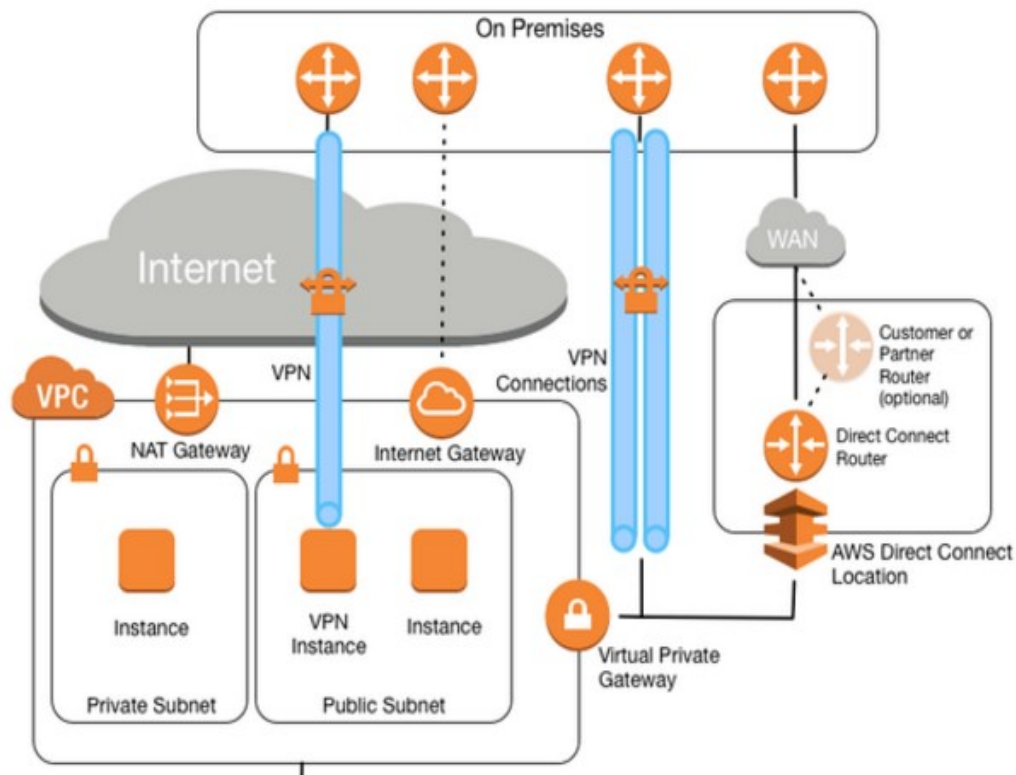
- ☐ A. Virtual Private Network ✓
- ☐ B. CloudFront Service
- ☐ C. Direct Connect ✓

☐ D. Primary Connection

Explanation :

Answer – A and C

You can either build a VPN or have a direct connect connection



For more information on VPC to on-premise networks, please refer to the following link
<https://aws.amazon.com/blogs/apn/amazon-vpc-for-on-premises-network-engineers-part-one/>
(<https://aws.amazon.com/blogs/apn/amazon-vpc-for-on-premises-network-engineers-part-one/>)

Ask our Experts



QUESTION 55

UNATTEMPTED

Why does stopping and starting an instance help in fixing a System Status Check error? Choose an answer from the options given below

- ☐ A. Stopping and starting an instance causes the instance to change the AMI.
- ☐ B. Stopping and starting an instance causes the instance to be provisioned on different AWS hardware. ✓
- ☐ C. Stopping and starting an instance reboots the operating system.
- ☐ D. None of the above

Explanation :

Answer – B

Refer below steps published by AWS support:

<https://aws.amazon.com/premiumsupport/knowledge-center/system-reachability-check/>
(<https://aws.amazon.com/premiumsupport/knowledge-center/system-reachability-check/>)

This is given in the aws documentation

When you stop a running instance, the following happens:

- The instance performs a normal shutdown and stops running; its status changes to `stopping` and then `stopped`.
- Any Amazon EBS volumes remain attached to the instance, and their data persists.
- Any data stored in the RAM of the host computer or the instance store volumes of the host computer is gone.
- In most cases, the instance is migrated to a new underlying host computer when it's started.

For more information on starting and stopping instances, please refer to the following link

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Stop_Start.html
(http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Stop_Start.html)

Note:

When you stop an instance, the resources provisioned for it is released. As soon as it is released and if there is no issue with the underlying resources it will be provisioned for some other users.

When you start your instance you will be provisioned with the resources that are currently available and in most cases it will not be the same compute resources. So stopping and starting an instance, in most cases, the instance will be provisioned in a new host with the appropriate compute and hardware resources associated with it.

Ask our Experts



In AWS Organization Consolidated Billing what are the 2 types of accounts?

- ☐ A. Main account and Subsidiary account
- ☒ B. Master account and Member account ✓
- ☐ C. Payer account and Linked account
- ☐ D. Domain account and Department account

Explanation :

Answer: B

Consolidated Billing Process

AWS Organizations provides consolidated billing so that you can track the combined costs of all the member accounts in your organization. The following steps provide an overview of the process for creating an organization and viewing your consolidated bill.

1. Open the AWS Organizations console or the AWS Billing and Cost Management console. If you open the AWS Billing and Cost Management console, select **Consolidated Billing** and then choose **Get started**. You will be redirected to the AWS Organizations console.
2. Choose **Create organization** on the AWS Organizations console.
3. Create an organization from the account that you want to be the master account of your new organization. For details, see [Creating an Organization](http://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_create.html) (http://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_create.html). The master account is responsible for paying the charges of all the member accounts.
4. (Optional) Create accounts that are automatically members of the organization. For details, see [Creating an AWS Account in Your Organization](http://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_accounts_create.html) (http://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_accounts_create.html).
5. (Optional) Invite existing accounts to join your organization. For details, see [Inviting an AWS Account to Join Your Organization](http://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_accounts_invites.html) (http://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_accounts_invites.html).
6. Each month AWS charges your master account for all the member accounts in a consolidated bill.

More information is available at:

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/consolidated-billing.html>
(<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/consolidated-billing.html>)

Ask our Experts



What is the term often used to describe a DR scenario in which a minimal version of an environment is always running in the cloud. Choose an answer from the options given below

- ☐ A. Backup and Restore
- ☒ B. Pilot Light ✓
- ☐ C. Warm standby
- ☐ D. Multi Site

Explanation :

Answer - B

This is given in a whitepaper published by AWS

Pilot Light for Quick Recovery into AWS

The term *pilot light* is often used to describe a DR scenario in which a minimal version of an environment is always running in the cloud. The idea of the pilot light is an analogy that comes from the gas heater. In a gas heater, a small flame that's always on can quickly ignite the entire furnace to heat up a house.

For more information on disaster recovery, please refer to the below link

https://media.amazonwebservices.com/AWS_Disaster_Recovery.pdf

(https://media.amazonwebservices.com/AWS_Disaster_Recovery.pdf)

Ask our Experts



Which of the following features ensures even distribution of traffic to Amazon EC2 instances in multiple Availability Zones registered with a load balancer?

- ☐ A. Elastic Load Balancing request routing
- ☐ B. An Amazon Route 53 weighted routing policy

- ☐ C. Elastic Load Balancing cross-zone load balancing ✓
- ☐ D. An Amazon Route 53 latency routing policy

Explanation :

Answer: C – To ensure that traffic is evenly distributed, you need to ensure the “Enable Cross-Zone Load balancing option” is chosen. This option comes up when you are creating a classic load balancer in Step 5 of Add EC2 instances.

Step 5: Add EC2 Instances

The table below lists all your running EC2 Instances. Check the boxes in the

VPC vpc-6dcc550a (172.31.0.0/16)

<input type="checkbox"/>	Instance	Name
--------------------------	----------	------

Availability Zone Distribution

- ☒ Enable Cross-Zone Load Balancing ⓘ
- ☒ Enable Connection Draining ⓘ 300 seconds

Ask our Experts



Currently you have a VPC with EC2 Security Group and several running EC2 instances. You change the Security Group rules to allow inbound traffic on a new port and protocol, and launch several new instances in the same Security Group. When will the Security Group changes be applied to the EC2 instances? Please choose the correct answer.

- ☒ A. Immediately to all instances in the security group. ✓
- ☐ B. Immediately to the new instances only.
- ☐ C. Immediately to the new instances, but old instances must be stopped and restarted before the new rules apply.
- ☐ D. To all instances, but it may take several minutes for old instances to see the changes.

Explanation :

Answer: A – By default whatever changes you make to security rules will be applied in all instances which are part of that security group.

When you add or remove rules, they are automatically applied to all instances associated with the security group.

For more information, please refer the below URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html)

Ask our Experts



QUESTION 60

UNATTEMPTED

What is the minimum size of an object that can be uploaded to Amazon S3?

- ☐ A. 1Megabyte
- ☒ B. 0Bytes ✓
- ☐ C. 1Byte

○ D. 5TB

Explanation :

Answer – B – The minimum size of an object in S3 can be 0 bytes. You can refer to the S3 FAQ's for more information on the allowable storage on S3.

<https://aws.amazon.com/s3/faqs/> (<https://aws.amazon.com/s3/faqs/>)

Q: How much data can I store?

The total volume of data and number of objects you can store are unlimited. Individual Amazon S3 objects can range in size from a minimum of 0 bytes to a maximum of 5 terabytes. The largest object that can be uploaded in a single PUT is 5 gigabytes. For objects larger than 100 megabytes, customers should consider using the [Multipart Upload](#) capability.

Ask our Experts



Finish Review (<https://www.whizlabs.com/learn/course/aws-csaa-practice-tests/quiz/13960>)

Certification

- ➔ Cloud Certification
(<https://www.whizlabs.com/cloud-certification-training-courses/>)
- ➔ Java Certification
(<https://www.whizlabs.com/oracle-java-certifications/>)
- ➔ PM Certification
(<https://www.whizlabs.com/project-management-certifications/>)
- ➔ Big Data Certification
(<https://www.whizlabs.com/big-data-certifications/>)

Company

- ➔ Support (<https://help.whizlabs.com/hc/en-us>)
- ➔ Discussions (<http://ask.whizlabs.com/>)
- ➔ Blog (<https://www.whizlabs.com/blog/>)

Mobile App





Follow us



(<https://www.facebook.com/whizlabs.software/>)



(<https://in.linkedin.com/company/whizlabs-software>)



(<https://twitter.com/whizlabs?lang=en>)



(<https://plus.google.com/+WhizlabsSoftware>)