



- [🏠 \(https://www.whizlabs.com/learn/\)](https://www.whizlabs.com/learn/) > [My Courses \(https://www.whizlabs.com/learn/my-courses\)](https://www.whizlabs.com/learn/my-courses)
- > [AWS Certified Solutions Architect Associate \(https://www.whizlabs.com/learn/course/aws-csaa-practice-tests#section-1\)](https://www.whizlabs.com/learn/course/aws-csaa-practice-tests#section-1)
- > [New Practice Test VII \(https://www.whizlabs.com/learn/course/aws-csaa-practice-tests/quiz/13955\)](https://www.whizlabs.com/learn/course/aws-csaa-practice-tests/quiz/13955) > **Report**

NEW PRACTICE TEST VII

Attempt 1

Marks Obtained 3 / 61

Your score is 4.92%

Completed on Monday , 28 January 2019 , 06:59 PM

Time Taken 00 H 00 M 27 S

Result Fail

Domains / Topics wise Quiz Performance Report

S.No.	Topic	Total Questions	Correct	Incorrect	Unattempted
1	Other	61	3	3	55

61 Questions	3 Correct	3 Incorrect	55 Unattempted
------------------------	---------------------	-----------------------	--------------------------

Show Answers

All	▼
-----	---

QUESTION 1 CORRECT

A company is trying to reduce their storage costs and want a more cost effective solution than Amazon S3. Secondly they claim that their data store is not frequently accessed. What is the best and cost efficient solution that should be considered?

- ☐ A. Amazon Storage Gateway
- ☒ B. Amazon Glacier ✓
- ☐ C. Amazon EBS
- ☐ D. Amazon S3

Explanation :

Answer: B – Since the data is not required to be accessed frequently, the data can be stored on Amazon glacier for cheaper storage. Remember that the recovery time for getting data from Glacier is from 3-5 hours.

All other options are not correct and expensive compared to Amazon Glacier service.

For more information on Glacier please visit the below URL:

- <https://aws.amazon.com/glacier/faqs/> (<https://aws.amazon.com/glacier/faqs/>)

Ask our Experts



QUESTION 2 INCORRECT

A company does not want to manage their databases. Which of the following services are fully managed databases provided by AWS? ^

- ☒ A. AWS RDS ✕
- ☐ B. DynamoDB ✓
- ☐ C. Oracle RDS
- ☐ D. Elastic Map Reduce

Explanation :

Answer: - B

DynamoDB is a fully managed NoSQL offering provided by AWS. It is now available in most regions for users to consume. **AWS RDS database is not fully managed database, it is partially managed. For RDS, we still need to specify the server capacity , security group etc.** This is the point most of them are confused, because they assume that RDS is the fully managed database. Even though the question doesn't ask about the type of database (NoSQL), the correct option is DynamoDB. For the fully managed option it is Aurora and DynamoDB. So, the correct option in this question is DynamoDB.

The link provides the full details on the product

1. <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Introduction.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Introduction.html>)
2. <https://aws.amazon.com/products/databases/>
(<https://aws.amazon.com/products/databases/>)

AWS Managed Services provides ongoing management of your AWS infrastructure so you can focus on your applications. By implementing best practices to maintain your infrastructure, AWS Managed Services helps to reduce your operational overhead and risk. AWS Managed Services automates common activities such as change requests, monitoring, patch management, security, and backup services, and provides full-lifecycle services to provision, run, and support your infrastructure. Our rigor and controls help to enforce your corporate and security infrastructure policies, and enable you to develop solutions and applications using your preferred development approach. AWS Managed Services improves agility, reduces cost, and unburdens you from infrastructure operations so you can direct resources toward differentiating your business.

For more information, please check the following link:

- <https://aws.amazon.com/products/databases/>
(<https://aws.amazon.com/products/databases/>)

Ask our Experts



Which of the following requires a custom CloudWatch metric to monitor?

- ☐ A. Memory Utilization of an EC2 instance ✓
- ☐ B. CPU Utilization of an EC2 instance
- ☐ C. Disk Reads activity of an EC2 instance
- ☒ D. Networks packets out of an EC2 instance ✕

Explanation :

Answer – A

Memory Utilization is a metric not offered directly by Cloudwatch.

So when you view the Cloudwatch metrics for your EC2 instance, you can see CPU Utilization and Disk Read Operations metrics.

You can also see Network statistics for Data transfer, but you will not be able to see Memory Utilization. This will be a custom Cloudwatch metric.

For more information on Cloudwatch, please refer the below URL:

<https://aws.amazon.com/cloudwatch/faqs/> (<https://aws.amazon.com/cloudwatch/faqs/>)

Ask our Experts



There is an urgent requirement to monitor few database metrics for a database hosted on AWS and send notifications. Which AWS services can accomplish this requirements?

Choose 2 answers from the options given below.

- ☐ A. Amazon Simple Email Service
- ☒ B. Amazon CloudWatch ✓
- ☐ C. Amazon Simple Queue Service (SQS)
- ☐ D. Amazon Route 53
- ☒ E. Amazon Simple Notification Service (SNS) ✓

Explanation :

Answer – B and E.

Amazon Cloudwatch will be used to monitor the IOP's metrics from the RDS instance and Amazon Simple Notification Service will be used to send the notification if any alarm is triggered.

For more information on Cloudwatch and SNS, please visit the below URLs:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)
- <https://aws.amazon.com/sns/> (<https://aws.amazon.com/sns/>)

Ask our Experts



QUESTION 5 INCORRECT

Which of the following instance types are available as SSD backed storage?

Choose 2 answers from the options below

- ☐ A. General purpose T2
- ☐ B. General purpose M3 ✓
- ☒ C. Compute-optimized C4 ✗



☒ **D. Compute-optimized C3** ✓

Explanation :

Answer – B and D

The below screenshots show the details for M3 and C3 instance types.

For details for all instance types, please visit the URL:

- <https://aws.amazon.com/ec2/instance-types/> (<https://aws.amazon.com/ec2/instance-types/>)

Ask our Experts



QUESTION 6 CORRECT

There is a requirement to install Perl on a Linux instance when it is launched.
Which feature allows you to accomplish this requirement?

- ☒ **A. User Data** ✓
- ☐ B. EC2Config Service
- ☐ C. IAM Roles
- ☐ D. AWS Config



Explanation :

Answer – A

When you configure an instance during creation, you can add custom scripts to the User data section.

So in Step 3 of creating an instance, in the Advanced Details section, we can enter custom scripts in the User Data section. The below script installs Perl during the instance creation of the EC2 instance.

For more information on user metadata and user data , please visit the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-metadata.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-metadata.html>)

Ask our Experts



An IAM user has been created in AWS. But the user is not able to perform any actions. What is the reason for this?

- ☐ A. IAM users are created by default with partial permissions
- ☐ B. IAM users are created by default with full permissions
- ☐ C. IAM users are created by default with no permissions ✓

Explanation :

Answer – C

By default no permissions are given to the user when they are created. Below is a snapshot of a newly created user. You can see that by default no permissions are assigned to the user.

For more information on IAM users , please visit the below URL:

- <https://aws.amazon.com/iam/details/manage-users/>
(<https://aws.amazon.com/iam/details/manage-users/>)

Ask our Experts



What happens when an instance behind an ELB fails a health check?

- ☐ A. The instance gets terminated automatically by the ELB.
- ☐ B. The instance gets quarantined by the ELB for root cause analysis.
- ☐ C. The instance is replaced automatically by the ELB.
- ☐ D. The ELB stops sending traffic to the instance that failed its health check ✓

Explanation :

Answer – D

To discover the availability of your EC2 instances, a load balancer periodically sends pings, attempts connections, or sends requests to test the EC2 instances. These tests are called health checks. The status of the instances that are healthy at the time of the health check is InService. The status of any instances that are unhealthy at the time of the health check is OutOfService. The load balancer performs health checks on all registered instances, whether the instance is in a healthy state or an unhealthy state.

The load balancer routes requests only to the healthy instances. When the load balancer determines that an instance is unhealthy, it stops routing requests to that instance. The load balancer resumes routing requests to the instance when it has been restored to a healthy state. You can see the status of the instance in the Registered Instances section of the load balancer.

For more information on ELB health checks , please visit the below URL:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-healthchecks.html>
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-healthchecks.html>)

Ask our Experts



In S3, what is the feature that is available to automatically transfer or archive data to Glacier?

- ☐ A. Use an EC2 instance and schedule a job to transfer the stale data from their S3 location to Amazon Glacier.
- ☒ B. Use Life-Cycle Policies ✓
- ☐ C. Use AWS SQS
- ☐ D. There is no option, the users will have to download the data and then transfer the data to AWS manually.

Explanation :

Answer – B

With Amazon lifecycle policies you can create transition actions in which you define when objects transition to another Amazon S3 storage class. For example, you may choose to transition objects to the STANDARD_IA (IA, for infrequent access) storage class 30 days after creation, or archive objects to the GLACIER storage class one year after creation.

Follow the below steps to get this in place:

Step 1) Go the Lifecycle section of the S3 bucket and click on Add Rule

Step 2) Choose what you want to export

Step 3) Choose the Action to perform and then confirm on the Rule creation in the next screen.



For more information on Lifecycle management, click on the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>)

Ask our Experts



QUESTION 10 UNATTEMPTED

Someone has initiated the snapshot creation of an EBS volume. One of the application still needs to use the same EBS volume. Which of the following scenarios are possible when it comes to usage of an EBS volume while the snapshot is initiated and not completed?

- ☒ A. Can be used while the snapshot is in progress. ✓
- ☐ B. Cannot be detached or attached to an EC2 instance until the snapshot completes



- ☐ C. Can be used in read-only mode while the snapshot is in progress.
- ☐ D. Cannot be used until the snapshot completes.

Explanation :

Answer – A

Snapshots occur asynchronously; the point-in-time snapshot is created immediately, but the status of the snapshot is pending until the snapshot is complete (when all of the modified blocks have been transferred to Amazon S3), which can take several hours for large initial snapshots or subsequent snapshots where many blocks have changed. While it is completing, an in-progress snapshot is not affected by ongoing reads and writes to the volume.

You can easily create a snapshot from a volume while the instance is running and the volume is in use. You can do this from the EC2 dashboard.

For more information on EBS snapshots, please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>)





QUESTION 11 UNATTEMPTED

There is a requirement to ensure that an EC2 instance can only be accessed from an IP address of 72.34.51.100. The users should be able to SSH into the instance. Which option will meet the customer requirement?

- ☐ A. Security Group Inbound Rule: Protocol – TCP, Port Range – 22, Source 72.34.51.100/32 ✓
- ☐ B. Security Group Inbound Rule: Protocol – UDP, Port Range – 22, Source 72.34.51.100/32
- ☐ C. Network ACL Inbound Rule: Protocol – UDP, Port Range – 22, Source 72.34.51.100/32
- ☐ D. Network ACL Inbound Rule: Protocol – TCP, Port Range-22, Source 72.34.51.100/0

Explanation :

Answer – A

For SSH access, the protocol has to be TCP, so Option B and C are wrong.

For Bastion host, only the IP of the client should be put and not the entire network of 72.34.51.100/0 as given in option D. So this option is also wrong.

A bastion host is a special purpose computer on a network specifically designed and configured to withstand attacks. The computer generally hosts a single application, for example a proxy server, and all other services are removed or limited to reduce the threat to the computer.

In AWS, A bastion host is kept on a public subnet. Users log on to the bastion host via SSH or RDP and then use that session to manage other hosts in the private subnets.

This is a security practice adopted by many organization to secure the assets in their private subnets.



For more information on security groups, please refer the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-network-security.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-network-security.html>)



Ask our Experts



QUESTION 12 UNATTEMPTED

Which of the following statements are true about Amazon Reduced Redundancy Storage (RRS) when it comes to availability?

- ☐ A. RRS has the ability to provide eleven nines availability.
- ☒ B. RRS has the ability to provide 99.99% availability. ✓
- ☐ C. RRS has the ability to provide 99% availability.
- ☐ D. RRS has the ability to provide 100% durability.

Explanation :

Answer: B. The Durability and availability are given in the aws site for RRS.



For more information on RRS please visit the URL:

- <https://aws.amazon.com/s3/reduced-redundancy/> (<https://aws.amazon.com/s3/reduced-redundancy/>)

Ask our Experts



QUESTION 13 UNATTEMPTED

Which service from AWS allows one to work with existing Chef server configuration?

- ☒ A. AWS OpsWorks ✓
- ☐ B. AWS Elastic Beanstalk



- ☐ C. AWS CloudFormation
- ☐ D. AWS SNS

Explanation :

Answer – A

AWS OpsWorks is a configuration management service that helps you configure and operate applications of all shapes and sizes using Chef. You can define the application's architecture and the specification of each component including package installation, software configuration and resources such as storage. Start from templates for common technologies like application servers and databases or build your own to perform any task that can be scripted. AWS OpsWorks includes automation to scale your application based on time or load and dynamic configuration to orchestrate changes as your environment scales.

For more information on Opswork, please visit the link:

- <https://aws.amazon.com/opsworks/> (<https://aws.amazon.com/opsworks/>)
- <https://aws.amazon.com/opsworks/chefautomate/>
(<https://aws.amazon.com/opsworks/chefautomate/>)



Ask our Experts



QUESTION 14 UNATTEMPTED

Which of the below AWS service can be used to deploy infrastructure using stacks and templates?

- ☐ A. Amazon Simple Workflow Service
- ☐ B. AWS Elastic Beanstalk
- ☐ C. AWS CloudFormation ✓
- ☐ D. AWS OpsWorks

Explanation :

Answer – C

AWS CloudFormation gives developers and systems administrators an easy way to create and manage a collection of related AWS resources, provisioning and updating them in an orderly and predictable fashion.

You can use AWS Cloud Formation's sample templates or create your own templates to describe the AWS resources, and any associated dependencies or runtime parameters, required to run your application. You don't need to figure out the order for provisioning AWS services or the subtleties of making those dependencies work. CloudFormation takes care of this for you. After the AWS resources are deployed, you can modify and update them in a controlled and predictable way, in effect applying version control to your AWS infrastructure the same way you do with your software. You can also visualize your templates as diagrams and edit them using a drag-and-drop interface with the AWS CloudFormation Designer.



For more information on Cloudformation, please visit the link:

- <https://aws.amazon.com/cloudformation/> (https://aws.amazon.com/cloudformation/)

Ask our Experts



QUESTION 15 UNATTEMPTED

Your company currently uses templates to deploy servers in their on-premise infrastructure. They want to have the same template configurations applied when deploying EC2 Instances. Which of the following can be done to ensure that EC2 Instances can be deployed as per the template standards defined by the organization. ^

- ☐ A. Use the EC2 metadata feature to deploy those features at runtime.
- ☐ B. Use the AWSConfig service to deploy updates to the EC2 Instances before they are launched.
- ☐ C. Create pre-built AMI's with the desired configuration as the organization templates. ✓
- ☐ D. It is not possible to define templates for EC2 Instances. You need to deploy the changes manually

Explanation :

Answer – C

The AWS Documentation mentions the following

An Amazon Machine Image (AMI) provides the information required to launch an instance, which is a virtual server in the cloud. You specify an AMI when you launch an instance, and you can launch as many instances from the AMI as you need. You can also launch instances from as many different AMIs as you need.

An AMI includes the following:

- A template for the root volume for the instance (for example, an operating system, an application server, and applications)
- Launch permissions that control which AWS accounts can use the AMI to launch instances
- A block device mapping that specifies the volumes to attach to the instance when it's launched

For more information on AMI's, please visit the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/AMIs.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/AMIs.html>)

Ask our Experts



QUESTION 16 UNATTEMPTED

What can be used for EC2 instances in a private subnet to connect to the internet? Choose an answer from the options below.

- ☐ A. WAF



- ☐ B. Direct Connect
- ☐ C. NAT Gateway ✓
- ☐ D. VPN

Explanation :

Answer – C

You can use a Network Address Translation (NAT) gateway to enable instances in a private subnet to connect to the Internet or other AWS services, but prevent the Internet from initiating a connection with those instances.

The below diagram from aws showcases how the NAT instance is used

For more information on NAT Gateways, please visit the URL:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-nat-gateway.html>
(<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-nat-gateway.html>)

Ask our Experts



QUESTION 17 UNATTEMPTED

Which AWS service allows businesses and web application developers an easy and cost effective way to distribute content with low latency and high data transfer speeds?

- ☐ A. Amazon SES
- ☐ B. Amazon Cloudtrail
- ☐ C. Amazon CloudFront ✓
- ☐ D. Amazon S3

Explanation :

Answer – C



Amazon CloudFront is a web service that gives businesses and web application developers an easy and cost effective way to distribute content with low latency and high data transfer speeds. Like other AWS services, Amazon CloudFront is a self-service, pay-per-use offering, requiring no long term commitments or minimum fees. With CloudFront, your files are delivered to end-users using a global network of edge locations.

For more information on CloudFront, please visit the link:

- <https://aws.amazon.com/cloudfront/> (<https://aws.amazon.com/cloudfront/>)

Ask our Experts



QUESTION 18 UNATTEMPTED

You try to connect to a newly created Amazon EC2 instance via SSH using PuTTY and get one of the following error messages

Error: Server refused our key (or)

Error: No supported authentication methods available

What steps should you take to identify the source of the behavior?

Choose 2 answers

- ☐ A. You should also verify that your private key (.pem) file has been correctly converted to the format recognized by PuTTY (.ppk). ✓
- ☐ B. Verify that your IAM user policy has permission to launch Amazon EC2 instances.
- ☐ C. Verify that you are connecting with the appropriate user name for your AMI. ✓
- ☐ D. Verify that the Amazon EC2 Instance was launched with the proper IAM role.

Explanation :

Answer - A and C

This is clearly given in the AWS documentation:

For more information on the connection errors to EC2 instances, please visit the link:



- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/TroubleshootingInstancesConnecting.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/TroubleshootingInstancesConnecting.html>)

Ask our Experts



QUESTION 19 UNATTEMPTED

Which feature in AWS is commonly used and best solution to store session data for web based applications?

- ☐ A. MySQL Installed on two Amazon EC2 Instances in a single Availability Zone
- ☐ B. Amazon RDS for MySQL with Multi-AZ
- ☐ C. Amazon ElastiCache ✓
- ☐ D. Amazon DynamoDB

Explanation :

Answer – C

Amazon ElastiCache is a web service that makes it easy to deploy, operate, and scale an in-memory data store or cache in the cloud. The service improves the performance of web applications by allowing you to retrieve information from fast, managed, in-memory data stores, instead of relying entirely on slower disk-based databases.



Elastic Cache is a better option when compared to DynamoDB. The main consideration would be the performance.

AWS Docs provides following details:

In order to address scalability and to provide a shared data storage for sessions that can be accessible from any individual web server, you can abstract the HTTP sessions from the web servers themselves. A common solution to for this is to leverage an In-Memory Key/Value store (<https://aws.amazon.com/elasticache/>) such as Redis (<https://aws.amazon.com/elasticache/what-is-redis/>) and Memcached.

While Key/Value data stores are known to be extremely fast and provide sub-millisecond latency, the added network latency and added cost are the drawbacks. An added benefit of leveraging Key/Value stores is that they can also be utilized to cache any data, not just HTTP sessions, which can help boost the overall performance of your applications.

For more information on Elastic cache, please visit the link:

- <https://aws.amazon.com/elasticache/> (<https://aws.amazon.com/elasticache/>)
- <https://aws.amazon.com/caching/session-management/> (<https://aws.amazon.com/caching/session-management/>)

Ask our Experts



QUESTION 20 UNATTEMPTED

Your application is having a very high traffic, so you have enabled autoscaling in multi availability zone to suffice the needs of your application but you observe that one of the availability zone is not receiving any traffic. What can be wrong here?

- ☐ A. Autoscaling only works for single availability zone
- ☐ B. Autoscaling can be enabled for multi AZ only in north Virginia region
- ☒ C. Availability zone is not added to Elastic load balancer ✓
- ☐ D. Instances need to manually added to availability zone



Explanation :

Answer – C

When you add an Availability Zone to your load balancer, Elastic Load Balancing creates a load balancer node in the Availability Zone. Load balancer nodes accept traffic from clients and forward requests to the healthy registered instances in one or more Availability Zones.

For more information on adding AZ's to ELB, please refer to the below URL:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/enable-disable-az.html>
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/enable-disable-az.html>)

Ask our Experts



QUESTION 21 UNATTEMPTED

Your company currently has an application hosted in their on-premise infrastructure. There is a mandate from management to move the application to the AWS Cloud. AS an architect you want to be cautious for the deployment of the application onto AWS. You have suggested to divert a percentage of the traffic from the users to the new application in AWS during the launch. Once it is confirmed that the cloud based application works with no issues , a full diversion to the new site can be implemented. Which of the following mechanisms can be used to ensure this scenario can be implemented.

- ☐ A. Use the Classic Elastic Load balancer to divert and proportion the traffic between the on-premise and AWS hosted application.
- ☐ B. Use the Application Elastic Load balancer to divert and proportion the traffic between the on-premise and AWS hosted application.
- ☐ C. Use Route53 with failover routing policy to divert and proportion the traffic between the on-premise and AWS hosted application.
- ☐ D. Use Route53 with Weighted routing policy to divert and proportion the traffic between the on-premise and AWS hosted application. ✓



Explanation :

Answer - D

The Weighted Routing policy is the best option here. You can ensure that the CNAME for your domain gets a lower proportion for the application hosted in AWS initially. Later on the percentage can be increased based on the application performance

The AWS documentation mentions the following on Route 53 Weighted Routing policy
Weighted routing lets you associate multiple resources with a single domain name (example.com) or subdomain name (acme.example.com) and choose how much traffic is routed to each resource. This can be useful for a variety of purposes, including load balancing and testing new versions of software.

For more information on Weighted Routing policy, please refer to the below URL:

- <http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html>
(<http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html>)

Ask our Experts



QUESTION 22 UNATTEMPTED

What step from the below options can be carried out to ensure that after an EBS volume is deleted , a similar volume with the same data can be created at a later stage.

- ☐ A. Create a copy of the EBS volume (not a snapshot)
- ☒ B. Store a snapshot of the volume ✓
- ☐ C. Download the content to an EC2 instance
- ☐ D. Back up the data in to a physical disk

Explanation :

Answer – B



Snapshots occur asynchronously; the point-in-time snapshot is created immediately, but the status of the snapshot is pending until the snapshot is complete (when all of the modified blocks have been transferred to Amazon S3), which can take several hours for large initial snapshots or subsequent snapshots where many blocks have changed. While it is completing, an in-progress snapshot is not affected by ongoing reads and writes to the volume.

You can easily create a snapshot from a volume while the instance is running and the volume is in use. You can do this from the EC2 dashboard.


For more information on EBS snapshots, please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>)

Ask our Experts



QUESTION 23 UNATTEMPTED

Which of the AWS Services following can be used to build an application based on 

a serverless architecture. Choose 3 answers from the options given below

- ☐ A. AWS API Gateway ✓
- ☐ B. AWS Lambda ✓
- ☐ C. AWS DynamoDB ✓
- ☐ D. AWS EC2

Explanation :

Answer – A,B and C

This is given in the AWS documentation

For more information on serverless platform, please refer to the below URL:

- <https://aws.amazon.com/serverless/> (<https://aws.amazon.com/serverless/>)

Ask our Experts



QUESTION 24 UNATTEMPTED

In the Amazon CloudWatch, which metric should I be checking to ensure that your DB Instance has enough free Memory?

- ☐ A. FreeStorage
- ☒ B. Freeable Memory ✓
- ☐ C. FreeStorageVolume
- ☐ D. FreeDBStorageSpace

Explanation :

Answer – B

When you go to the Monitoring tab for your AWS RDS instance, you will be able to see the Cloudwatch metrics.



For more information on Amazon Cloudwatch, please visit the below URL:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)

Ask our Experts



QUESTION 25 UNATTEMPTED

You have an Autoscaling Group which is launching a set of t2.small instances. You now need to replace those instances with a larger instance type. How would you go about making this change in an ideal manner?

- ☐ A. Change the Instance type in the current launch configuration to the new instance type.
- ☐ B. Create another Autoscaling Group and attach the new instance type.
- ☐ C. Create a new launch configuration with the new instance type and update your Autoscaling Group. ✓
- ☐ D. Change the Instance type of the Underlying EC2 instance directly.

Explanation :

Answer – C



The AWS Documentation mentions

A launch configuration is a template that an Auto Scaling group uses to launch EC2 instances. When you create a launch configuration, you specify information for the instances such as the ID of the Amazon Machine Image (AMI), the instance type, a key pair, one or more security groups, and a block device mapping. If you've launched an EC2 instance before, you specified the same information in order to launch the instance.

When you create an Auto Scaling group, you must specify a launch configuration. You can specify your launch configuration with multiple Auto Scaling groups. However, you can only specify one launch configuration for an Auto Scaling group at a time, and you can't modify a launch configuration after you've created it. Therefore, if you want to change the launch configuration for your Auto Scaling group, you must create a launch configuration and then update your Auto Scaling group with the new launch configuration.

For more information on launch configurations please see the below link:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/LaunchConfiguration.html> (<http://docs.aws.amazon.com/autoscaling/latest/userguide/LaunchConfiguration.html>)

Ask our Experts



QUESTION 26 UNATTEMPTED

In what events would cause Amazon RDS to initiate a failover to the standby replica?

Choose 3 answers from the options given below

- ☐ A. Loss of availability in primary Availability Zone ✓
- ☐ B. Loss of network connectivity to primary ✓
- ☐ C. Storage failure on secondary
- ☐ D. Compute unit failure on primary ✓

Explanation :

Answer - A, B and D



Amazon RDS detects and automatically recovers from the most common failure scenarios for Multi-AZ deployments so that you can resume database operations as quickly as possible without administrative intervention. Amazon RDS automatically performs a failover in the event of any of the following:

- Loss of availability in primary Availability Zone
- Loss of network connectivity to primary
- Compute unit failure on primary
- Storage failure on primary

Note: When operations such as DB Instance scaling or system upgrades like OS patching are initiated for Multi-AZ deployments, for enhanced availability, they are applied first on the standby prior to an automatic failover. As a result, your availability impact is limited only to the time required for automatic failover to complete. Note that Amazon RDS Multi-AZ deployments do not failover automatically in response to database operations such as long running queries, deadlocks or database corruption errors.

For more information on read replicas, please visit the below URL:

- <https://aws.amazon.com/rds/details/read-replicas/>
(<https://aws.amazon.com/rds/details/read-replicas/>)

Ask our Experts



QUESTION 27 UNATTEMPTED

Which of the following tools is available to send log data from EC2 Instances.

- ☒ A. CloudWatch Logs Agent ✓
- ☐ B. CloudWatch Agent
- ☐ C. Logs Stream

Explanation :

Answer - A

The AWS Documentation mentions the following

The CloudWatch Logs agent provides an automated way to send log data to CloudWatch Logs from Amazon EC2 instances. The agent is comprised of the following components:



- A plug-in to the AWS CLI that pushes log data to CloudWatch Logs.
- A script (daemon) that initiates the process to push data to CloudWatch Logs.
- A cron job that ensures that the daemon is always running.

For more information on Cloudwatch logs Agent, please see the below link:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/AgentReference.html>
(<http://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/AgentReference.html>)

Ask our Experts



QUESTION 28 UNATTEMPTED

You have a business-critical two tier web app currently deployed in 2 availability zones in a single region, using Elastic Load Balancing (ELB) and Auto-Scaling. The app depends on synchronous replication at the database layer. The application needs to remain fully available even if one application AZ goes off-line and AutoScaling cannot launch new instances in the remaining AZ. How can the current architecture be enhanced to ensure this requirement?

- ☐ A. Deploy in 2 regions using Weighted Round Robin with AutoScaling minimums set of 50% peak load per Region.
- ☐ B. Deploy in 3 AZ with Autoscaling minimum set to handle 33 percent peak load per zone.
- ☐ C. Deploy in 3 AZ with Autoscaling minimum set to handle 50 percent peak load per zone. ✓
- ☐ D. Deploy in 2 regions using Weighted Round Robin with AutoScaling minimums set of 100% peak load per Region.

Explanation :

Answer – C

Since the requirement is that the application should never go down even if an AZ is not available, we need to maintain 100% availability.

Option A and D are incorrect because region deployment is not possible for ELB. ELB's can



manage traffic within a region and not between regions.

Option B is incorrect because even if one AZ goes down, we would be operating at only 66% and not the required 100%.

For more information on Autoscaling please visit the below URL:

- <https://aws.amazon.com/autoscaling/> (<https://aws.amazon.com/autoscaling/>)

Ask our Experts



QUESTION 29 UNATTEMPTED

A VPC public subnet is one that (choose one of the correct option below):

- ☐ A. Has at least one route in its associated routing table that uses an Internet gateway ✓
- ☐ B. Includes a route in its associated routing table via a Network Address Translation(NAT) instance.
- ☐ C. Has a Network Access Control List(NACL) permitting outbound traffic to 0.0.0.0/0
- ☐ D. Has the public Subnet option selected in its configuration

Explanation :

Answer - A

The below diagram shows a public subnet. The public subnet has a route table that uses the internet gateway



For more information on public subnets please visit the below URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html)

Ask our Experts



QUESTION 30 UNATTEMPTED

A customer nightly EMR job processes a single 2-TB data file stored on S3. The EMR job runs on 2 on-demand core nodes and 3 on-demand task nodes. Which of the following may help reduce the EMR job completion time?

Choose 2 answers from the options below

- ☐ A. Use 3 spot instances rather than 3 on-demand instances for the task nodes.
- ☐ B. Change the input split size in the MapReduce job configuration ✓
- ☐ C. Use a bootstrap action to present the S3 bucket as a local filesystem
- ☐ D. Launch the core nodes and the task nodes with a VPC



☐ E. Adjust the number of simultaneous mapper tasks ✓

Explanation :

Answer - B and E

As per the AWS documentation, if you have too few tasks , then you have nodes sitting idle. You can increase the number of simultaneous mapper tasks and reduce the size of the MapReduce job configuration

For more information on EMR tasks please visit the below URL:

- http://docs.aws.amazon.com/emr/latest/DeveloperGuide/TaskConfiguration_H1.0.3.html
(http://docs.aws.amazon.com/emr/latest/DeveloperGuide/TaskConfiguration_H1.0.3.html)

Ask our Experts



QUESTION 31 UNATTEMPTED

What action is required to establish an VPC VPN connection between an on-premise data center and an VPC virtual private gateway?

- ☐ A. Assign a static internet-routable IP Address to an Amazon VPC customer gateway ✓



- ☐ B. Modify the main route table to allow traffic to a network address translation instance.
- ☐ C. Use a dedicated network address translation instance in the public subnet
- ☐ D. Establish a dedicated networking connection using Direct Connect

Explanation :

Answer - A

When defining a VPN connection between the on-premise network and the VPC, you need to have a customer gateway defined. Since this is accessed over the internet, it needs to have a static internet-routable IP Address.

For more information on VPC VPN connections please visit the below URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Introduction.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Introduction.html)
- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpn-connections.html>
(<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpn-connections.html>)
-



Ask our Experts



QUESTION 32 UNATTEMPTED

What combination of the following options will protect S3 objects from both accidental deletion and accidental overwriting?

Choose 2 answers from the options below

- ☐ A. Enable S3 versioning on the bucket ✓
- ☐ B. Access S3 data using only signed URL's
- ☐ C. Disable S3 delete using an IAM bucket policy
- ☐ D. Enable S3 RRS
- ☐ E. Enable MFA protected access ✓

Explanation :

Answer - A and E

This is clearly given in the AWS documentation:

For more information on S3 please visit the below URL:

- <https://aws.amazon.com/s3/faqs/> (<https://aws.amazon.com/s3/faqs/>)

Ask our Experts



A startup company hired you to help them build a mobile application that will ultimately store billions of images and videos on S3. The company is lean on funding and wants to minimize operational costs however they have an aggressive marketing plan, and expect to double their current installation base every six months. Due to the nature of their business they are expecting a sudden and large increases in traffic to and from S3 and need to ensure that it can handle the performance needs of their application. What other information must you gather from this customer in order to determine whether S3 is the right option?

- ☐ A. You must know how many customers the company has today because this is critical in understanding what their customer base will be in 2 years.
- ☒ B. You must find out the total number of requests per second at peak usage. ✓
- ☐ C. You must know the size of the individual objects being written to S3, in order to properly design the key namespace.
- ☐ D. In order to build the key namespace correctly you must understand the total amount of storage needs for each S3 bucket.

Explanation :

Answer - B

When you define an S3 bucket the billing is done on the requests. If you go to the URL:

- <http://calculator.s3.amazonaws.com/index.html>
(<http://calculator.s3.amazonaws.com/index.html>)

which is the calculator for S3 costs, you can see that the cost is related to the total number of requests in addition to the storage. In order to find the estimated cost for the S3 storage, you should get a number of requests.

Ask our Experts



QUESTION 34

UNATTEMPTED

You have configuring a solution which uses EC2 Instances and an Elastic Load Balancer. Which of the following protocols can be used to ensure that traffic is secure from the client machine to the Elastic Load Balancer. Choose 2 answers from the options given below

- ☐ A. HTTP
- ☐ B. HTTPS ✓
- ☐ C. TCP
- ☐ D. SSL ✓

Explanation :

Answer – B and D

The HTTPS protocol uses the SSL protocol to establish secure connections over the HTTP layer.

You can also use the SSL protocol to establish secure connections over the TCP layer.

For more information on ELB Listener configuration please see the below link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-listener-config.html>
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-listener-config.html>)

Ask our Experts



QUESTION 35

UNATTEMPTED

After creating a new AWS account, you use the API to request 40 on-demand EC2 instances in a single AZ. After 20 successful requests, subsequent requests failed. What could be a reason for this issue, and how can you resolve it?

- ☐ A. You encountered a soft limit of 20 instances per region. Submit the limit increase form and retry the failed requests once approved. ✓



- ☐ B. AWS allows you to provision no more than 20 instances per AZ. Select a different AZ and retry the failed request.
- ☐ C. You need to use VPC in order to provision more than 20 instances in a single AZ. Simply terminate the resources already provisioned and re-launch them all in a VPC.
- ☐ D. You encountered an API throttling situation and should try the failed requests using an exponential decay retry algorithm.

Explanation :

Answer – A

There is a soft limits of 20 instances. Since this is across an instance family, option B is wrong

because it will not work even if you try another availability zone.

For more information on all service limits please visit the below URL:

- https://aws.amazon.com/ec2/faqs/#How_many_instances_can_I_run_in_Amazon_EC2
(https://aws.amazon.com/ec2/faqs/#How_many_instances_can_I_run_in_Amazon_EC2)
- http://docs.aws.amazon.com/general/latest/gr/aws_service_limits.html
(http://docs.aws.amazon.com/general/latest/gr/aws_service_limits.html)

Ask our Experts



QUESTION 36 UNATTEMPTED

You have been tasked with creating a VPC network topology for your company. The VPC network must support both internet-facing applications and internally-facing applications accessed only over VPN. Both Internet-facing and internally-facing applications must be able to leverage at least 3 AZs for high availability. At a minimum, how many subnets must you create within your VPC to accommodate these requirements?



- ☐ A. 2
- ☐ B. 3
- ☐ C. 4
- ☐ D. 6 ✓

Explanation :

Answer - D

Internet as well as intranet(private) applications must be able to make use of at least three Availability Zones for high availability.

So 3 subnets for internet and 3 subnets for private is 6 subnets in total.

For more information on VPC and subnets please visit the below URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html)

Ask our Experts



QUESTION 37 **UNATTEMPTED**

You receive a Linux spot instance at a bit price of \$0.05/hour. After 30 minutes the Spot price is changed to \$0.06/hour and the Spot instance is terminated by you. What is the total cost associated with this EC2 instance?

- ☐ A. \$0.025 ✓
- ☐ B. \$0.00
- ☐ C. \$0.05
- ☐ D. \$0.06

Explanation :

Answer – A



So based on this, the total cost incurred will be \$0.025. We have mentioned in our question that after 30 minutes the spot price is increased to \$0.06/hr and the user terminated the instance.

But if the Amazon is terminating the instance then it will be charged for 30 mins which will be \$0.00.

If Amazon is terminating the instance, since the price change happened in the first hour of the instance itself, there won't be any charge associated with it .

- <https://www.amazonaws.cn/en/ec2/faqs/#spot-instances>
(<https://www.amazonaws.cn/en/ec2/faqs/#spot-instances>)

From 2nd October 2017, **per second billing** has come into effect for some EC2 instances and EBS. AWS per-second billing will apply to Linux On-Demand, Reserved, and Spot EC2 instances. However, Per-second billing is **not** applicable to **Microsoft Windows** instances or to **all Linux distributions**, so some Linux AMIs may still have an hourly charge.

- <https://aws.amazon.com/blogs/aws/new-per-second-billing-for-ec2-instances-and-ebs-volumes/> (<https://aws.amazon.com/blogs/aws/new-per-second-billing-for-ec2-instances-and-ebs-volumes/>)

With per-second billing in effect, if AWS stops your instance, you will be billed for exactly what you have used. For example, if your instance is stopped by you after a half an hour of use, you only pay for the 30 mins instead of a full hour.

For more information on spot instance pricing please visit the below URL:

- <https://aws.amazon.com/ec2/spot/pricing/> (<https://aws.amazon.com/ec2/spot/pricing/>)

Note:

For spot pricing you will be billed only for your usage of the instance. This means that if you have used an instance for 35 minutes and 20 seconds and after that if you have terminated it, you will be billed only for 35 minutes and 20 seconds. Previously with spot pricing you had to pay for the full hour instead of paying as per the seconds used.

However there are certain exceptions.

You have fixed a spot price of say \$.50 for an instance and say after 15 minutes if the spot price has gone up to \$.80 and amazon has terminated your instance due to this price hike. In that case you need not have to pay any amount since the instance has run for less than an hour.

But assume that the price has increased after 2 hours of your use, say the price has increased after 130 minutes of use and amazon terminated your instance then you need to pay for 130 minutes of usage.

Hope that this is clear now.

You can check one more similar question asked in our forum here:

- <http://ask.whizlabs.com/t/aws-csaa-exam-topics-preparation-tips/668/57>
(<http://ask.whizlabs.com/t/aws-csaa-exam-topics-preparation-tips/668/57>)



Ask our Experts



QUESTION 38 UNATTEMPTED

Which of the following is a durable key-value store?

- ☒ A. Amazon Simple Storage Service ✓
- ☐ B. Amazon Simple Queue Service
- ☐ C. Amazon Simple Workflow Service
- ☐ D. Amazon Simple Notification Service

Explanation :

Answer - A

This is clearly given in the AWS documentations:

Reference 1:



Reference 2:

For more information on S3 please visit the below URLs:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingObjects.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingObjects.html>)
- <https://aws.amazon.com/s3/details> (<https://aws.amazon.com/s3/details>)

Ask our Experts



QUESTION 39 UNATTEMPTED

In reviewing the Auto-Scaling events for your application you notice that your application is scaling up and down multiple times in the same hour. What design choice could you make to optimize for costs while preserving elasticity?

Select 2 options.

- ☐ A. Modify the Auto Scaling policy to use scheduled scaling actions
- ☐ B. Modify the Auto Scaling Group cool down timers ✓
- ☐ C. Modify the Amazon Cloudwatch alarm period that triggers your AutoScaling scale down policy. ✓
- ☐ D. Modify the Auto Scaling group termination policy to terminate the newest instance first.

Explanation :

Answer - B and C



The Auto Scaling cooldown period is a configurable setting for your Auto Scaling group that helps to ensure that Auto Scaling doesn't launch or terminate additional instances before the previous scaling activity takes effect. After the Auto Scaling group dynamically scales using a simple scaling policy, Auto Scaling waits for the cooldown period to complete before resuming scaling activities. When you manually scale your Auto Scaling group, the default is not to wait for the cooldown period, but you can override the default and honor the cooldown period. Note that if an instance becomes unhealthy, Auto Scaling does not wait for the cooldown period to complete before replacing the unhealthy instance.

For more information on Autoscale cool down timers please visit the URL:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/Cooldown.html>
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/Cooldown.html>)

You can also modify the Cloudwatch triggers to ensure the thresholds are appropriate for the scale down policy.

For more information on Autoscaling user guide please visit the URL:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/as-scale-based-on-demand.html> (<http://docs.aws.amazon.com/autoscaling/latest/userguide/as-scale-based-on-demand.html>)

Ask our Experts



QUESTION 40 UNATTEMPTED

You have an application running in us-west-2 that requires 6 EC2 instances running at all times. With 3 AZ available in that region, which of the following deployments provides 100% fault tolerance if any single AZ in us-west-2 becomes unavailable?

Choose 2 answers from the options below:

- ☐ A. us-west-2a with 2 instances, us-west-2b with 2 instances, us-west-2c with 2 instances
- ☐ B. us-west-2a with 3 instances, us-west-2b with 3 instances, us-west-2c with 0 instances
- ☐ C. us-west-2a with 4 instances, us-west-2b with 2 instances, us-west-2c with 2 instances



- ☐ D. us-west-2a with 6 instances, us-west-2b with 6 instances, us-west-2c with 0 instances ✓
- ☐ E. us-west-2a with 3 instances, us-west-2b with 3 instances, us-west-2c with 3 instances ✓

Explanation :

Answer - D and E

If you read the question carefully, it asks you the scenario when only one AZ goes down at a time. The requirement is to make 6 instances always running even if any one of the AZ is goes down. The questions doesn't ask you if any 2 or 3 AZ goes down at a time.

Hence D and E ensures that always 6 instances are running if any one AZ goes down at a time. I hope this clears your doubts. Since we need 6 instances running at all times, only D and E fulfill this option.

- Option A is invalid, because if any one of Availability zones goes down, then we are left with only 4 running instances.
- Option B is invalid because if either us-west-2a or us-west-2b goes down then we are left with less than 6 instances.
- Option C is invalid if us-west-2a goes down then we are left with less than 6 instances

For more information on building fault tolerant applications in AWS, please refer to the below link

- http://media.amazonwebservices.com/AWS_Building_Fault_Tolerant_Applications.pdf
(http://media.amazonwebservices.com/AWS_Building_Fault_Tolerant_Applications.pdf)
(http://media.amazonwebservices.com/AWS_Building_Fault_Tolerant_Applications.pdf)

Ask our Experts



QUESTION 41 UNATTEMPTED

Which route must be added to your routing table in order to allow connections to the internet from your subnet?

- ☐ A. Destination:0.0.0.0/0-->Target:your internet gateway ✓
- ☐ B. Destination:192.168.1.255/0-->Target:your internet gateway



- ☐ C. Destination:0.0.0.0/33-->Target:your virtual private gateway
- ☐ D. Destination:0.0.0.0/0--> Target:0.0.0.0/24

Explanation :

Answer – A

The below diagram shows a public subnet. The public subnet has a route table that uses the internet gateway.

For more information on public subnets please visit the below URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html)

Ask our Experts



QUESTION 42 UNATTEMPTED

You are deploying an application on Amazon EC2 that must call AWS API's. What is the method for securely passing credentials to the application that you use?^

- ☐ A. Embed the API credentials into your JAR files.
- ☐ B. Use the AWS Identity and Access Management (IAM) roles for EC2 instances ✓
- ☐ C. Store API credentials as an object in S3.
- ☐ D. Pass API credentials to the instance using instance userdata.

Explanation :

Answer – B

An IAM role is similar to a user, in that it is an AWS identity with permission policies that determine what the identity can and cannot do in AWS. However, instead of being uniquely associated with one person, a role is intended to be assumable by anyone who needs it. Also, a role does not have any credentials (password or access keys) associated with it. Instead, if a user is assigned to a role, access keys are created dynamically and provided to the user.

For more information on IAM role please visit the below URL:

- http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles.html
(http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles.html)

Ask our Experts



QUESTION 43 UNATTEMPTED

You are developing a highly available web application using stateless web servers. Which services are suitable for storing session state data?

Choose 3 answers from the options below

- ☐ A. Amazon RDS ✓
- ☐ B. Amazon Elastic Cache ✓
- ☐ C. Amazon Cloudwatch
- ☐ D. Elastic Load Balancing (ELB)
- ☐ E. Amazon DynamoDB ✓



Explanation :

Answer - A,B and E

Amazon RDS

Amazon Relational Database Service (Amazon RDS) makes it easy to set up, operate, and scale a relational database in the cloud. It provides cost-efficient and resizable capacity while managing time-consuming database administration tasks, freeing you up to focus on your applications and business.

For more information on Amazon RDS please visit the below URL:

- <https://aws.amazon.com/rds/> (<https://aws.amazon.com/rds/>)

Amazon Elastic Cache

Amazon ElastiCache is a web service that makes it easy to deploy, operate, and scale an in-memory data store or cache in the cloud. The service improves the performance of web applications by allowing you to retrieve information from fast, managed, in-memory data stores, instead of relying entirely on slower disk-based databases.

For more information on Amazon Elastic Cache please visit the below URL:

- <https://aws.amazon.com/elasticache/> (<https://aws.amazon.com/elasticache/>)

Amazon DynamoDB

Amazon DynamoDB is a fast and flexible NoSQL database service for all applications that need consistent, single-digit millisecond latency at any scale. It is a fully managed cloud database and supports both document and key-value store models. Its flexible data model and reliable performance make it a great fit for mobile, web, gaming, ad tech, IoT, and many other applications.

For more information on Amazon DynamoDB please visit the below URL:

- <https://aws.amazon.com/dynamodb/> (<https://aws.amazon.com/dynamodb/>)

Ask our Experts



QUESTION 44

UNATTEMPTED

What are some of the metrics that are monitored by AWS Lambda? Choose 3 answers from the options given below.

☐ A. Invocations ✓

☐ B. Duration ✓



- ☐ C. Errors ✓
- ☐ D. Database Changes

Explanation :

Answer – A,B and C

AWS Lambda automatically monitors functions on your behalf, reporting metrics through Amazon CloudWatch. These metrics include Invocations, Duration, and Errors.



For more information on Lambda metrics please visit the below URL:

- <https://docs.aws.amazon.com/lambda/latest/dg/monitoring-functions-metrics.html>
(<https://docs.aws.amazon.com/lambda/latest/dg/monitoring-functions-metrics.html>)
- <http://docs.aws.amazon.com/lambda/latest/dg/monitoring-functions-access-metrics.html>
(<http://docs.aws.amazon.com/lambda/latest/dg/monitoring-functions-access-metrics.html>)

Ask our Experts



QUESTION 45 **UNATTEMPTED**

There is a new facility from AWS which allows for fast, easy, and secure transfers of files over long distances between your client and your Amazon S3 bucket. What is this service called?

- ☐ A. File Transfer
- ☐ B. HTTP Transfer
- ☐ C. S3 Transfer Acceleration ✓
- ☐ D. Kinesis Acceleration

Explanation :

Answer - C

To know more about S3 transfer acceleration, please visit the below URL:



- <http://docs.aws.amazon.com/AmazonS3/latest/dev/transfer-acceleration.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/transfer-acceleration.html>)

Ask our Experts



QUESTION 46 UNATTEMPTED

What are the languages currently supported by AWS Lambda?

Choose 3 answers from the options given below.

- ☐ A. Node.js ✓
- ☐ B. Angular.js
- ☐ C. Java ✓
- ☐ D. C# ✓

Explanation :

Answer – A,C and D

AWS Lambda runs your code on a high-availability compute infrastructure and performs all of the administration of the compute resources, including server and operating system maintenance, capacity provisioning and automatic scaling, code monitoring and logging. All you need to do is supply your code in one of the languages that AWS Lambda supports (**currently Node.js, Java, C# and Python**).

For more information on Lambda please visit the below URL:

- <http://docs.aws.amazon.com/lambda/latest/dg/welcome.html>
(<http://docs.aws.amazon.com/lambda/latest/dg/welcome.html>)



Ask our Experts



QUESTION 47 UNATTEMPTED

Your company has an application hosted in AWS which makes use of DynamoDB. There is a requirement from the IT security department to ensure that all source IP addresses which make calls to the DynamoDB tables are recorded. Which of the following services can be used to ensure this requirement is fulfilled.

- ☐ A. AWS Code Commit
- ☐ B. AWS Code Pipeline
- ☒ C. AWS CloudTrail ✓
- ☐ D. AWS Cloudwatch

Explanation :

Answer - C

The AWS Documentation mentions the following

DynamoDB is integrated with CloudTrail, a service that captures low-level API requests made by or on behalf of DynamoDB in your AWS account and delivers the log files to an Amazon S3 bucket that you specify. CloudTrail captures calls made from the DynamoDB console or from the DynamoDB low-level API. Using the information collected by CloudTrail, you can determine what request was made to DynamoDB, the source IP address from which the request was made, who made the request, when it was made, and so on.

For more information on DynamoDB and Cloudtrail, please refer to the below link:

- <http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/logging-using-cloudtrail.html>
(<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/logging-using-cloudtrail.html>)

Ask our Experts



Which of the following statements are false about Amazon Glacier. Choose one answer from the options given below.

- ☐ A. It supports archive operations of Upload, Download and Delete
- ☐ B. The archives are mutable ✓
- ☐ C. Uploading an archive is a synchronous operation
- ☐ D. Archives can be as large as 40TB

Explanation :

Answer – B

This is clearly given in the AWS documentation.

A single archive can be as large as 40 terabytes. You can store an unlimited number of archives and an unlimited amount of data in Amazon Glacier. Each archive is assigned a unique archive ID at the time of creation, and the content of the archive is immutable, meaning that after an archive is created it cannot be updated.



For more information on AWS Glacier please visit the below URL:

- <https://aws.amazon.com/glacier/details/> (<https://aws.amazon.com/glacier/details/>)

Ask our Experts



Your company currently has a web application hosted on a single EC2 Instance. The load on the application has increased over time and now the users are complaining of slow response time. Which of the following implementations can help alleviate this issue.

- ☐ A. Attach an additional EBS Volume to the EC2 Instance and direct the application to make the reads from this new volume.
- ☐ B. Attach an additional network interface with an Elastic IP so that requests can be made onto multiple IP's.
- ☐ C. Launch additional EC2 Instances in a web server farm type configuration and place them behind an Elastic Load Balancer. ✓
- ☐ D. Launch additional EC2 Instances in a web server farm type configuration and place them behind Route53.

Explanation :

Answer – C

The AWS mentions the following about the Elastic Load balancer that can be used to help in this issue

A load balancer distributes incoming application traffic across multiple EC2 instances in multiple Availability Zones. This increases the fault tolerance of your applications. Elastic Load Balancing detects unhealthy instances and routes traffic only to healthy instances.

Your load balancer serves as a single point of contact for clients. This increases the availability of your application. You can add and remove instances from your load balancer as your needs change, without disrupting the overall flow of requests to your application. Elastic Load Balancing scales your load balancer as traffic to your application changes over time. Elastic Load Balancing can scale to the vast majority of workloads automatically.

For more information on the Elastic Load Balancer, please refer to the below link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/introduction.html>
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/introduction.html>)

Note:

AWS says that "If you host a website on multiple Amazon EC2 instances, you can distribute traffic to your website across the instances by using an Elastic Load Balancing (ELB) load balancer. The ELB service automatically scales the load balancer as traffic to your website changes over time.

The load balancer also can monitor the health of its registered instances and route domain traffic only to healthy instances." ✓

Further, in the question, we have mentioned "slow response time", is not because of latency. It is because of the heavy traffic, it can be resolved if we add one more instance.

Ask our Experts



QUESTION 50 UNATTEMPTED

Which of the following are used to get data records from Amazon Kinesis?
Choose an answer from the options below

- ☒ A. Consumer ✓
- ☐ B. Stream
- ☐ C. Producer
- ☐ D. None of the above

Explanation :

Answer – A

A consumer gets data records from Amazon Kinesis streams. A consumer, known as an *Amazon Kinesis Streams application*, processes the data records from a stream.

For more information on AWS Kinesis consumers please visit the below URL:

- <http://docs.aws.amazon.com/streams/latest/dev/amazon-kinesis-consumers.html>
(<http://docs.aws.amazon.com/streams/latest/dev/amazon-kinesis-consumers.html>)

Ask our Experts



QUESTION 51 UNATTEMPTED

What is the maximum possible retention period for data in Kinesis Streams?



Choose an answer from the options below.

- ☐ A. 5 days
- ☒ B. 7 days ✓
- ☐ C. 10 days
- ☐ D. 24 hours

Explanation :

Answer – B

For more information on AWS Kinesis consumers please visit the below URL:

- <http://docs.aws.amazon.com/streams/latest/dev/amazon-kinesis-consumers.html>
(<http://docs.aws.amazon.com/streams/latest/dev/amazon-kinesis-consumers.html>)

Data records are accessible for a default of 24 hours from the time they are added to a stream. This time frame is called the retention period and is configurable in hourly increments from 24 to 168 hours (1 to 7 days).

Ask our Experts



QUESTION 52 UNATTEMPTED

Which of the following is false when you create an encrypted EBS volume?

- ☐ A. Data is encrypted at rest inside the volume
- ☐ B. Data is encrypted when it is moved from one instance to another in the same subnet. ✓
- ☐ C. Data is encrypted when data is moved between the volume and the instance
- ☐ D. All snapshots created from the volume are encrypted

Explanation :

Answer – B

The AWS mentions the following about EBS Encryption



Amazon EBS encryption offers you a simple encryption solution for your EBS volumes without the need for you to build, maintain, and secure your own key management infrastructure. When you create an encrypted EBS volume and attach it to a supported instance type, the following types of data are encrypted:

- Data at rest inside the volume
- All data moving between the volume and the instance
- All snapshots created from the volume

For more information on EBS Encryption, please refer to the below link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSEncryption.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSEncryption.html>)

Ask our Experts



QUESTION 53 UNATTEMPTED

In AWS what constitutes temporary security credentials?

Choose 3 answers from the options given below

- ☐ A. AWS Access Key ID ✓
- ☐ B. Secret Access Key ✓
- ☐ C. Security Token ✓
- ☐ D. SSL Keys

Explanation :

Answer –A,B and C

This is given in the AWS documentation:



For more information on IAM please visit the below URL:

- <https://aws.amazon.com/iam/faqs/> (<https://aws.amazon.com/iam/faqs/>)

Ask our Experts



QUESTION 54 **UNATTEMPTED**

Your company has a set of resources hosted in AWS. Your IT Supervisor is concerned with the costs being incurred with the current set of AWS resources and wants to monitor the cost usage. Which of the following mechanisms can be used to monitor the costs of the AWS resources and also look at the possibility of cost optimization. Choose 3 answers from the options given below

- ☐ **A. Use the Cost Explorer to see the costs of AWS resources** ✓
- ☐ **B. Create budgets in billing section so that budgets are set beforehand** ✓
- ☐ **C. Send all logs to Cloudwatch logs and inspect the logs for billing details**
- ☐ **D. Consider using the Trusted Advisor** ✓

Explanation :

Answer - A,B and D

The AWS Documentation mentions the following



1) For a quick, high-level analysis use Cost Explorer, which is a free tool that you can use to view graphs of your AWS spend data. It includes a variety of filters and preconfigured views, as well as forecasting capabilities. Cost Explorer displays data from the last 13 months, the current month, and the forecasted costs for the next three months, and it updates this data daily.

2) Consider using budgets if you have a defined spending plan for a project or service and you want to track how close your usage and costs are to exceeding your budgeted amount. Budgets use data from Cost Explorer to provide you with a quick way to see your usage-to-date and current estimated charges from AWS. You can also set up notifications that warn you if you exceed or are about to exceed your budgeted amount.

3) Visit the AWS Trusted Advisor console regularly. Trusted Advisor works like a customized cloud expert, analyzing your AWS environment and providing best practice recommendations to help you save money, improve system performance and reliability, and close security gaps.

For more information on cost optimization, please visit the below URL:

- <https://aws.amazon.com/answers/account-management/cost-optimization-monitor/>
(<https://aws.amazon.com/answers/account-management/cost-optimization-monitor/>)

Ask our Experts



QUESTION 55 UNATTEMPTED

Who are federated users when it comes to AWS? Choose an answer from the options given below

- ☐ A. These are IAM users in aws
- ☐ B. These are IAM groups in aws
- ☐ C. These are Federated users (external identities) are users you manage outside of AWS in your corporate directory ✓
- ☐ D. None of the above

Explanation :

Answer - C

This is given in the AWS documentation:



For more information on IAM please visit the below URL:

- <https://aws.amazon.com/iam/faqs/> (<https://aws.amazon.com/iam/faqs/>)

Ask our Experts



QUESTION 56 UNATTEMPTED

As a solution architect, you have been asked to decide on whether to use Amazon EBS-backed or instance-store backed instance. What is one key difference between an Amazon EBS-backed and an instance-store backed instance that you need to keep in mind.

- ☐ A. Amazon EBS-backed instances can be stopped and restarted without lose the data. ✓
- ☐ B. Instance-store backed instances can be stopped and restarted without lose the data.
- ☐ C. Auto scaling requires using Amazon EBS-backed instances.
- ☐ D. Virtual Private Cloud (VPC) requires EBS backed instances.

Explanation :

Answer A - Amazon EBS-backed instances can be stopped and restarted.



Please visit the below URL for the key differences between EBS and instance store volumes:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ComponentsAMIs.html>)
- Amazon EBS-backed AMI can be placed in stopped state where instance is not running, but the root volume is persisted in Amazon EBS.
- Amazon Instance store -backed AMI cannot be in stopped state; instances are running or terminated

Ask our Experts



QUESTION 57

UNATTEMPTED

Which of the following are not supported in the classic load balancer service provided by AWS? Choose an answer from the options given below.

- ☐ A. Health Checks
- ☐ B. Cloudwatch Metrics
- ☐ C. Path Based Routing ✓
- ☐ D. Access Logs

Explanation :

Answer – C



This is clearly given in the AWS documentation:

For more information on ELB please visit the below URL:

- <https://aws.amazon.com/elasticloadbalancing/classicloadbalancer/faqs/>
(<https://aws.amazon.com/elasticloadbalancing/classicloadbalancer/faqs/>)

Ask our Experts



QUESTION 58 **UNATTEMPTED**

Your company has an on-premise Active Directory setup in place. The company has extended their footprint on AWS , but still want to have the ability to use their on-premise Active Directory for authentication. Which of the following AWS services can be used to ensure that AWS resources such as AWS Workspaces can continue to use the existing credentials stored in the on-premise Active Directory.



- ☐ A. Use the Active Directory service on AWS
- ☐ B. Use the AWS Simple AD service
- ☐ C. Use the Active Directory connector service on AWS ✓
- ☐ D. Use the ClassicLink feature on AWS

Explanation :

Answer – C

The AWS Documentation mentions the following

AD Connector is a directory gateway with which you can redirect directory requests to your on-premises Microsoft Active Directory without caching any information in the cloud. AD Connector comes in two sizes, small and large. A small AD Connector is designed for smaller organizations of up to 500 users. A large AD Connector can support larger organizations of up to 5,000 users.

For more information on the AD connector, please refer to the below URL:

- http://docs.aws.amazon.com/directoryservice/latest/admin-guide/directory_ad_connector.html
(http://docs.aws.amazon.com/directoryservice/latest/admin-guide/directory_ad_connector.html)

Ask our Experts



QUESTION 59 UNATTEMPTED

Which DNS record types does Amazon Route 53 support?

Select 3 options.

- ☐ A. A(address record) ✓
- ☐ B. AAAA(IPv6 address record) ✓
- ☐ C. TXT (txt record) ✓
- ☐ D. Host Information records (HINFO)

Explanation :



Answer – A, B and C.

For more information on Route53, please visit the below URL:

- <https://aws.amazon.com/route53/faqs/> (<https://aws.amazon.com/route53/faqs/>)

Ask our Experts



QUESTION 60 UNATTEMPTED

A user has been created in IAM but the user is still not able to make API calls. After creating a new IAM user which of the following must be done before they can successfully make API calls?

- ☐ A. Add a password to the user.
- ☐ B. Enable Multi-Factor Authentication for the user.
- ☐ C. Assign a Password Policy to the user.
- ☐ D. Create a set of Access Keys for the user. ✓



Explanation :

Answer – D

In IAM , when you create a user , you need to download the Access Key ID and Secret access key so that the user can access aws.

For more information on IAM please visit the following URL:

- <http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>
(<http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>)

Ask our Experts



QUESTION 61 UNATTEMPTED

An application needs to be setup on AWS. It consists of several components. Two primary components are required to run for 3 hours every day. The other components are required every day for more than 6-8 hours and **other servers also running in others instances**. Which of the following would you use to ensure COSTS are minimized for the underlying EC2 Instances?

- ☐ A. Reserved instances for the primary components and for the remaining components ,On-Demand Instances.
- ☐ B. Spot instances for the primary components and for the remaining components ,On-Demand Instances.
- ☐ C. On-Demand instances for the primary components and for the remaining components ,Spot Instances.
- ☐ D. On-Demand instances for the primary components and for the remaining components ,Reserved Instances. ✓

Explanation :

Answer – D

The AWS Documentation mentions the following



Reserved Instances provide you with a significant discount compared to On-Demand Instance pricing. Reserved Instances are not physical instances, but rather a billing discount applied to the use of On-Demand Instances in your account. These On-Demand Instances must match certain attributes in order to benefit from the billing discount.

On-Demand Instances – Pay, by the second, for the instances that you launch.

Option A is incorrect since the primary component just runs for 3 hours , you really don't need a reserved instance for it

Options B and C are incorrect since we don't know the type of workload to decide whether Spot Instances are required

For more information on On-Demand and Reserved Instances, please refer to the below URL

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-reserved-instances.html>

(<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-reserved-instances.html>)

Ask our Experts



Finish Review (<https://www.whizlabs.com/learn/course/aws-csaa-practice-tests/quiz/13955>)

Certification

- ➔ Cloud Certification
(<https://www.whizlabs.com/cloud-certification-training-courses/>)
- ➔ Java Certification
(<https://www.whizlabs.com/oracle-java-certifications/>)
- ➔ PM Certification
(<https://www.whizlabs.com/project-management-certifications/>)
- ➔ Big Data Certification
(<https://www.whizlabs.com/big-data-certifications/>)

Company

- ➔ Support (<https://help.whizlabs.com/hc/en-us>)
- ➔ Discussions (<http://ask.whizlabs.com/>)
- ➔ Blog (<https://www.whizlabs.com/blog/>)



Mobile App



Follow us



(<https://www.facebook.com/whizlabs.software/>)



(<https://in.linkedin.com/company/whizlabs-software>)



(<https://twitter.com/whizlabs?lang=en>)



(<https://plus.google.com/+WhizlabsSoftware>)

© Copyright 2018. Whizlabs Software Pvt. Ltd. All Rights Reserved.

