



[🏠 \(https://www.whizlabs.com/learn\)](https://www.whizlabs.com/learn) > [My Courses \(https://www.whizlabs.com/learn/my-courses\)](https://www.whizlabs.com/learn/my-courses)
> [AWS Certified SysOps Administrator Associate \(https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests#section-1\)](https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests#section-1)
> [Diagnostic Test \(https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests/quiz/12957\)](https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests/quiz/12957) > **Report**

DIAGNOSTIC TEST

Attempt	1	Completed on	Tuesday , 29 January 2019 , 02:27 PM
Marks Obtained	1 / 60	Time Taken	00 H 00 M 05 S
Your score is	1.67%	Result	Fail

Domains / Topics wise Quiz Performance Report

S.No.	Topic	Total Questions	Correct	Incorrect	Unattempted
1	Other	60	1	0	59

60 Questions	1 Correct	0 Incorrect	59 Unattempted	Show Answers	All	▼
------------------------	---------------------	-----------------------	--------------------------	--------------	-----	---

QUESTION 1 CORRECT

There are currently multiple applications hosted in a VPC. During monitoring it has been noticed that multiple port scans are coming in from a specific IP Address block. The internal security team has requested that all offending IP Addresses be denied for the next 24 hours. Which of the following is the best method to quickly and temporarily deny access from the specified IP Address's.

- ☐ A. Create an AD policy to modify the Windows Firewall settings on all hosts in the VPC to deny access from the IP Address block.
- ☒ B. Modify the Network ACLs associated with all public subnets in the VPC to deny access from the IP Address block. ✓
- ☐ C. Add a rule to all of the VPC Security Groups to deny access from the IP Address block.

- D. Modify the Windows Firewall settings on all AMI's that your organization uses in that VPC to deny access from the IP address block.

Explanation :

Answer – B

A network access control list (ACL) is an optional layer of security for your VPC that acts as a firewall for controlling traffic in and out of one or more subnets.

Option A and D are wrong because this is a tedious task and it only works for Windows systems. You need something that will work for Linux systems as well.

Option C is only adequate for EC2 instances, but you need rules that will apply to the whole subnet. Otherwise the task of having this done for all servers becomes a tedious task.

To change the Network ACL's follow the below steps

Step 1) Go to your VPC dashboard. You would probably already have a network ACL once the VPC is defined. Click on the Network ACL, go to Inbound Rules and click on Edit

The screenshot shows the AWS VPC console interface. On the left is a navigation menu with categories like 'Virtual Private Cloud', 'Your VPCs', 'Subnets', 'Route Tables', 'Internet Gateways', 'DHCP Options Sets', 'Elastic IPs', 'Endpoints', 'NAT Gateways', 'Peering Connections', 'Security', 'Network ACLs', and 'Security Groups'. The main area displays the 'VPC Dashboard' for 'vpc-6dcc550a'. It lists a single Network ACL, 'acl-05195662', which is associated with 2 subnets and is the default ACL. Below this, the 'acl-05195662' details page is shown, with the 'Inbound Rules' tab selected. A note states: 'Allows inbound traffic. Because network ACLs are stateless, you must create inbound and outbound rules.' An 'Edit' button is visible. Below the edit button is a table of inbound rules:

Rule #	Type	Protocol	Port Range	Source	Allow / Deny
100	ALL Traffic	ALL	ALL	0.0.0.0/0	ALLOW
*	ALL Traffic	ALL	ALL	0.0.0.0/0	DENY

Step 2) In the next screen, create a rule that will deny access to the offending IP Address. In the below snapshot, we are assuming that the IP Address are in the range of 7.7.7.7/32. Then ensure that the action is the Deny action. By putting a rule number of 95, this ensure that this rule is executed before the other rules in the NACL list.

VPC Dashboard

Filter by VPC: None

Virtual Private Cloud

Your VPCs

Subnets

Route Tables

Internet Gateways

DHCP Options Sets

Elastic IPs

Endpoints

NAT Gateways

Peering Connections

Security

Network ACLs

Security Groups

Create Network ACL Delete

Search Network ACLs and IT X

Name	Network ACL ID	Associated With	Default	VPC
	acl-05195662	2 Subnets	Yes	vpc-6dcc550a

acl-05195662

Summary Inbound Rules Outbound Rules Subnet Associations Tags

Allows inbound traffic. Because network ACLs are stateless, you must create inbound and outbound rules.

Cancel Save

Rule #	Type	Protocol	Port Range	Source	Allow / Deny	Remove
100	ALL Traffic	ALL	ALL	0.0.0.0/0	ALLOW	
95	ALL Traffic	ALL	ALL	7.7.7.7/32	DENY	

Add another rule

For more information on Network ACL's, please visit the URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLs.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLs.html)

Ask our Experts



QUESTION 2 UNATTEMPTED

You have two Elastic Compute Cloud (EC2) instances inside a Virtual Private Cloud (VPC) in the same Availability Zone (AZ) but in different subnets. One instance is running a database and the other instance an application that will interface with the database. You want to confirm that they can talk to each other for your application to work properly. Which two things do we need to confirm in the VPC settings so that these EC2 instances can communicate inside the VPC? Choose 2 answers.

- ☐ A. A network ACL that allows communication between the two subnets. ✓
- ☐ B. Both instances are the same instance class and using the same Key-pair.
- ☐ C. That the default route is set to a NAT instance or internet Gateway (IGW) for them to communicate.
- ☐ D. Security groups are set to allow the application host to talk to the database on the right port/protocol. ✓

Explanation :

Answer - A and D

When you design a web server and database server, the security groups must be defined so that the web server can talk to the database server. An example image from the AWS documentation is given below

Also when communicating between subnets you need to have the NACL's defined

Option B is wrong since the EC2 instances need not be of the same class or same key pair to communicate to each other.

Option C is wrong since there the NAT and Internet gateway is used for the subnet to communicate to the internet.

For more information on VPC and Subnets, please visit the below URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html)

Ask our Experts



QUESTION 3 UNATTEMPTED

You are managing a legacy application Inside VPC with hard coded IP addresses in its configuration. Which two mechanisms will allow the application to failover to new instances without the need for reconfiguration? Choose 2 answers

- ☐ A. Create an ELB to reroute traffic to a failover instance
- ☐ B. Create a secondary ENI that can be moved to a failover instance ✓
- ☐ C. Use Route53 health checks to fail traffic over to a failover instance
- ☐ D. Assign a secondary private IP address to the primary ENI that can be moved to a failover instance ✓

Explanation :

Answer - B and D

Here you can choose either the option of creating a Secondary Network interface which can moved to the failover instance or have a secondary IP Address which can be moved to the failover instance. For both cases, you can do this at the time of defining the EC2 instance as shown below.

▼ Network interfaces ⓘ

Device	Network Interface	Subnet	Primary IP	Secondary IP addresses
eth0	New network interface ▼	subnet-effe4ec2 ▼	Auto-assign	Auto-assign Add IP
eth1	New network interface ▼	subnet-effe4ec2 ▼	Auto-assign	Add IP

For more information on Network Interfaces, please visit the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-eni.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-eni.html>)

Ask our Experts



QUESTION 4 UNATTEMPTED

You are designing a system that has a Bastion host. This component needs to be highly available without human intervention. Which of the following approaches would you select?

- ☐ A. Run the bastion on two instances one in each AZ
- ☐ B. Run the bastion on an active Instance in one AZ and have an AMI ready to boot up in the event of failure
- ☐ C. Configure the bastion instance in an Auto Scaling group Specify the Auto Scaling group to include multiple AZs but have a min-size of 1 and max-size of 1 ✓
- ☐ D. Configure an ELB in front of the bastion instance

Explanation :

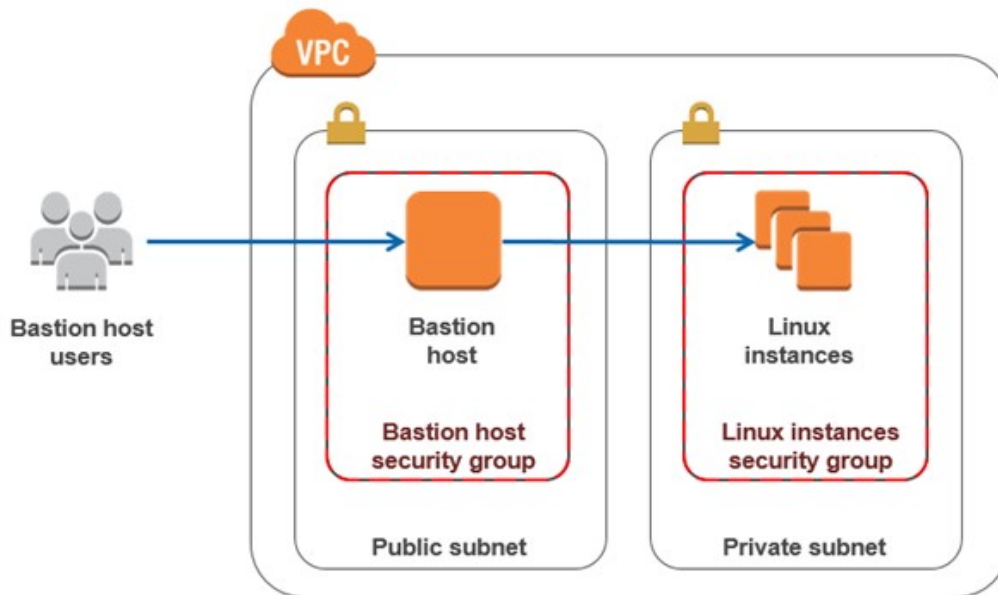
Answer – C

You can have a bastion host running in multiple AZ, but the recommendation is to have one running in each AZ. Hence you need to make sure that the Autoscaling group is set to a max-size of one.

A bastion host is a special purpose computer on a network specifically designed and configured to withstand attacks. The computer generally hosts a single application, for example a proxy server, and all other services are removed or limited to reduce the threat to the computer.

In AWS, A bastion host is kept on a public subnet. Users log on to the bastion host via SSH or RDP and then use that session to manage other hosts in the private subnets.

This is a security practice adopted by many organization to secure the assets in their private subnets.



For more information on AutoScaling, please visit the below URL:

- <https://aws.amazon.com/autoscaling/> (<https://aws.amazon.com/autoscaling/>)

Ask our Experts



QUESTION 5 UNATTEMPTED

It is possible to have a rollback window for objects in S3. If yes, then which of the below methods can help achieve this

- ☐ A. Data Encryption in S3.
- ☐ B. Using the lifecycle policy with Versioning. ✓
- ☐ C. Using S3 static site.
- ☐ D. This is not possible.

Explanation :

Answer – B

As per the AWS documentation, you can use the Lifecycle versioning policy in S3 to achieve the rollback window.

You can define the Lifecycle policy in S3 as shown below.

▼ Lifecycle

You can manage the lifecycle of objects by using [Lifecycle rules](#). Lifecycle rules enable you to automatically transition objects to the Standard - Infrequent Access Storage Class, and/or archive objects to the Glacier Storage Class, and/or remove objects after a specified time period. Rules are applied to all the objects that share the specified prefix.

Versioning is not currently enabled on this bucket.

You can use Lifecycle rules to manage all versions of your objects. This includes both the Current version and Previous versions.

 Add rule

Save

Cancel

For more information on S3, please visit the below URL:

- <https://aws.amazon.com/s3/faqs/> (<https://aws.amazon.com/s3/faqs/>)

Ask our Experts



QUESTION 6 UNATTEMPTED

Which feature in S3 allows one to analyze the identify access patterns whilst using the storage in S3?

Choose one answer from the options given below.

- ☒ A. S3 Analytics ✓
- ☐ B. S3 lifecycle policy
- ☐ C. S3 IA
- ☐ D. This is not possible

Explanation :

Answer - A

As per the AWS documentation, you can use the S3 storage analytics to see storage patterns.

With storage class analysis, you can analyze storage access patterns and transition the right data to the right storage class. This new S3 Analytics feature automatically identifies infrequent access patterns to help you transition storage to Standard-IA. You can configure a storage class analysis policy to monitor an entire bucket, a prefix, or object tag. Once an infrequent access pattern is observed, you can easily create a new lifecycle age policy based on the results. Storage class analysis also provides daily visualizations of your storage usage on the AWS Management Console that you can export to a S3 bucket to analyze using business intelligence tools of your choice such as Amazon QuickSight.

Option B is wrong because this is used for managing storage transfer to glacier.

Option C is wrong because this is used for Infrequent Access storage.

Option D is wrong because this is possible in S3.

For more information on S3, please visit the below URL:

- <https://aws.amazon.com/s3/faqs/> (<https://aws.amazon.com/s3/faqs/>)

Ask our Experts



QUESTION 7 UNATTEMPTED

You have been asked to automate many routine systems administrator backup and recovery activities. Your current plan is to leverage AWS-managed solutions as much as possible and automate the rest with the AWS CLI and scripts. Which task would be best accomplished with a script?

- ☒ A. Creating daily EBS snapshots with a monthly rotation of snapshots ✓
- ☐ B. Creating daily RDS snapshots with a monthly rotation of snapshots
- ☐ C. Automatically detect and stop unused or underutilized EC2 instances
- ☐ D. Automatically add Auto Scaled EC2 instances to an Amazon Elastic Load Balancer

Explanation :

Answer – A

EBS snapshots are the ones that can be managed via the CLI.

You can easily create a snapshot from a volume while the instance is running and the volume is in use.

You can do this from the EC2 dashboard.

Option B is wrong because this is managed by AWS via RDS.

Option C and D is wrong because this is managed by Autoscaling.

An example of doing it via the CLI is given below

```
AWS ec2 create-snapshot --volume-id vol-1234567890abcdef0 --description  
"This is my root volume snapshot."
```

For more information on the CLI , please visit the URL:

- <http://docs.aws.amazon.com/cli/latest/reference/ec2/create-snapshot.html>
(<http://docs.aws.amazon.com/cli/latest/reference/ec2/create-snapshot.html>)

For more information on EBS Snapshots, please visit the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSSnapshots.html>)



QUESTION 8 UNATTEMPTED

Amazon S3 is storage for the Internet. It's a simple storage service that offers software developers a highly-scalable, reliable, and low-latency data storage infrastructure at very low costs. From the below options which are true with regards to AWS S3. Choose 2 answers from the options given below.

- ☐ A. Objects are directly accessible via a URL ✓
- ☐ B. S3 should be used to host a relational database
- ☐ C. S3 allows you to store objects of virtually unlimited size
- ☐ D. S3 allows you to store virtually unlimited amounts of data ✓
- ☐ E. S3 offers Provisioned IOPS

Explanation :

Answer - A and D

Option A is correct since in S3, all objects are directly accessible via a URL. An example is shown below

Object: Doc1.docx

```
Bucket: devtoolslogging
Name: Doc1.docx
Link: https://s3-ap-southeast-1.amazonaws.com/devtoolslogging/Doc1.docx
Size: 106073
Last Modified: Fri Nov 25 08:24:43 GMT-800 2016
Owner: s362692
ETag: 1956a762be7f1239373eb9b89197bef4
Expiry Date: None
Expiration Rule: N/A
```

Option D is correct because as per the AWS documentation, S3 can store virtually unlimited amounts of data. Option C is incorrect since there is a limit on the object size.

Option B is incorrect because S3 cannot be used to host a relational database.

Option E is incorrect because EBS provides IOPS and not S3.

For more information on S3, please visit the below URL:

- <https://aws.amazon.com/s3/faqs/> (<https://aws.amazon.com/s3/faqs/>)



You receive a frantic call from a new DBA who accidentally dropped a table containing all your customers. Which Amazon RDS feature will allow you to reliably restore your database to within 5 minutes of when the mistake was made?

- ☐ A. Multi-AZ RDS
- ☐ B. RDS snapshots
- ☐ C. RDS read replicas
- ☒ D. RDS automated backup ✓

Explanation :

Answer - D

The Question is referring to an AWS RDS feature which will allow us to restore our db to a specified time which is just 5 minutes prior to the deletion of the table.

As per AWS documentation,

The Amazon RDS automated backup feature automatically creates a storage volume snapshot of your DB instance, backing up the entire DB instance and not just individual databases. This backup occurs during a daily user-configurable 30 minute period known as the *backup window*. Automated backups are kept for a configurable number of days (called the *backup retention period*). You can restore your DB instance to any specific time during this retention period, creating a new DB instance.

To determine the latest restorable time for a DB instance, use the AWS CLI `describe-db-instances` (<http://docs.aws.amazon.com/cli/latest/reference/rds/describe-db-instances.html>) command and look at the value returned in the **LatestRestorableTime** field for the DB instance. The latest restorable time for a DB instance is typically within 5 minutes of the current time.

For more information regarding this please use the link given below:

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_PIT.html

(https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_PIT.html)

Automated backups automatically back up your DB instance during a specific, user-definable backup window. Amazon RDS keeps these backups for a limited period that you can specify. You can later recover your database to any point in time during this backup retention period. And all of these backups get stored to S3 by default.

The backup retention period is present as a setting when you setup the RDS in AWS.

Backup

Please note that automated backups are currently supported for InnoDB storage engine only. If you are using MyISAM, refer to detail [here](#).

Backup Retention Period days

Backup Window

Option A and C are wrong because If the table is deleted the Multi-AZ and Read Replica will also have the same dropped tables because the synchronization would have already happened.

Option B is wrong because snapshots is a manual process.

For more information on Automated backups, please visit URL:

- http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_WorkingWithAutomatedBackups.html
(http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_WorkingWithAutomatedBackups.html)

Ask our Experts



QUESTION 10 UNATTEMPTED

A media company produces new video files on-premises every day with a total size of around 100GB after compression All files have a size of 1 -2 GB and need to be uploaded to Amazon S3 every night in a fixed time window between 3am and 5am Current upload takes almost 3 hours, although less than half of the available bandwidth is used. What step(s) would ensure that the file uploads are able to complete in the allotted time window?

- ☐ A. Increase your network bandwidth to provide faster throughput to S3
- ☒ B. Upload the files in parallel to S3 ✓
- ☐ C. Pack all files into a single archive, upload it to S3, and then extract the files in AWS
- ☐ D. Use AWS Import/Export to transfer the video files

Explanation :

Answer – B

When uploading large videos it's always better to make use of AWS multi part file upload.

So if you are using the Multi Upload option for S3, then you can resume on failure. Below are the advantage of Multi Part upload

- Improved throughput—you can upload parts in parallel to improve throughput.
- Quick recovery from any network issues—smaller part size minimizes the impact of restarting a failed upload due to a network error.
- Pause and resume object uploads—you can upload object parts over time. Once you initiate a multipart upload there is no expiry; you must explicitly complete or abort the multipart upload.
- Begin an upload before you know the final object size—you can upload an object as you are creating it.

For more information on Multi-part file upload for S3, please visit the URL:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/qfacts.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/qfacts.html>)

QUESTION 11 UNATTEMPTED

You are running a web-application on AWS consisting of the following components an Elastic Load Balancer (ELB) an Auto-Scaling Group of EC2 instances running Linux/PHP/Apache, and Relational DataBase Service (RDS) MySQL. Which security measures fall into AWS's responsibility?

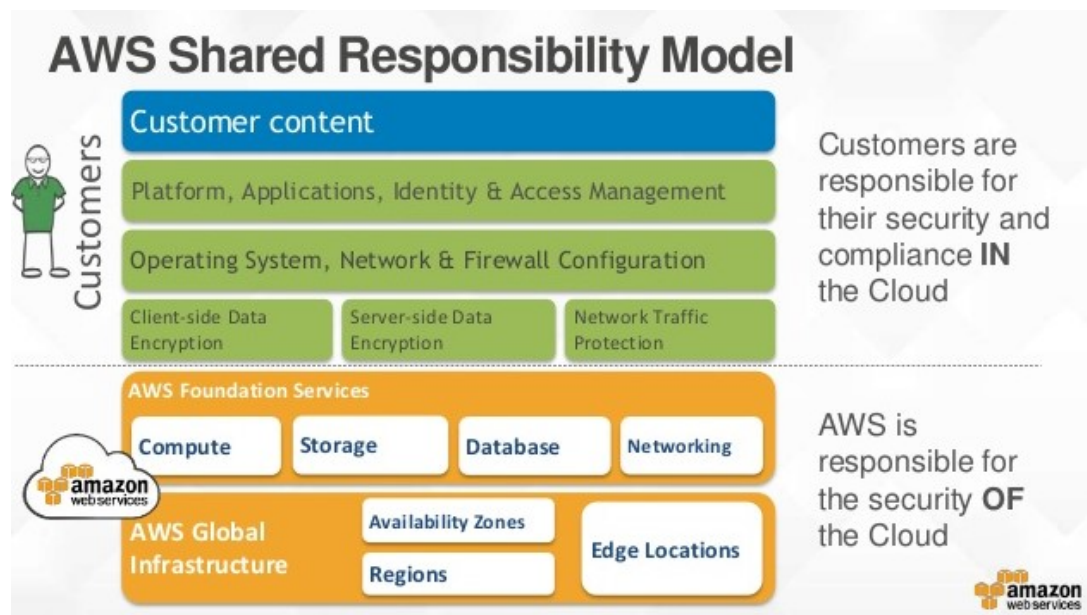
- ☐ A. Protect the EC2 instances against unsolicited access by enforcing the principle of least-privilege access
- ☒ B. Protect against IP spoofing or packet sniffing ✓
- ☐ C. Assure all communication between EC2 instances and ELB is encrypted
- ☐ D. Install latest security patches on ELB, RDS and EC2 instances

Explanation :

Answer – B

As per the shared responsibility shown below, the users are required to control the EC2 security via security groups and network access control layers.

Also it is the user's responsibility model, AWS takes care of the physical components and the infrastructure to provide Virtualization.



For more information on AWS shared responsibility model, please visit the link:

- <https://aws.amazon.com/blogs/security/tag/shared-responsibility-model/>
(<https://aws.amazon.com/blogs/security/tag/shared-responsibility-model/>)

Ask our Experts



QUESTION 12 UNATTEMPTED

When do you get billed for EC2 instances? Please choose one answer from the options given below.

- ☒ A. Running state ✓
- ☐ B. Terminated state
- ☐ C. Stopped state
- ☐ D. All of the above

Explanation :

Answer – A

Remember that you get charged for EC2 instances only when the instances are in a running state. This is also specified as per the documentation in AWS as shown in the snapshot below. All other options are invalid as per the documentation given below.

Instance-hours are billed for any time your instances are in a "running" state. If you no longer wish to be charged for your instance, you must "stop" or "terminate" the instance to avoid being billed for additional instance-hours. Billing starts when an instance transitions into the running state.

For more information on EC2 instances, please visit the URL:

- <https://aws.amazon.com/ec2/faqs/> (<https://aws.amazon.com/ec2/faqs/>)

Ask our Experts



QUESTION 13 UNATTEMPTED

Your team is excited about the use of AWS because now they have access to "programmable Infrastructure". You have been asked to manage your AWS infrastructure in a manner similar to the way you might manage application code. You want to be able to deploy exact copies of different versions of your infrastructure,

stage changes into different environments, revert back to previous versions, and identify what versions are running at any particular time (development test QA . production). Which approach addresses this requirement?

- ☐ A. Use cost allocation reports and AWS Opsworks to deploy and manage your infrastructure.
- ☐ B. Use AWS CloudWatch metrics and alerts along with resource tagging to deploy and manage your infrastructure.
- ☐ C. Use AWS Beanstalk and a version control system like GIT to deploy and manage your infrastructure.
- ☐ D. Use AWS CloudFormation and a version control system like GIT to deploy and manage your infrastructure. ✓

Explanation :

Answer – D

AWS CloudFormation gives developers and systems administrators an easy way to create and manage a collection of related AWS resources, provisioning and updating them in an orderly and predictable fashion.

You can use AWS Cloud Formation's sample templates

(<https://aws.amazon.com/cloudformation/aws-cloudformation-templates/>) or create your own templates to describe the AWS resources, and any associated dependencies or runtime parameters, required to run your application. You don't need to figure out the order for provisioning AWS services or the subtleties of making those dependencies work. CloudFormation takes care of this for you. After the AWS resources are deployed, you can modify and update them in a controlled and predictable way, in effect applying version control to your AWS infrastructure the same way you do with your software. You can also visualize your templates as diagrams and edit them using a drag-and-drop interface with the AWS CloudFormation Designer

(<https://aws.amazon.com/cloudformation/details/#designer>).

Option A is incorrect because Cost Allocation Reports is not helpful for the purpose of the question.

Option B is incorrect because Cloudwatch is used for monitoring.

Option C is incorrect because It abstracts away the underlying EC2 instances, Elastic Load Balancers, auto scaling groups, etc which is another word we don't have control over the infrastructure..

For more information on Cloudformation, please visit the link:

- <https://aws.amazon.com/cloudformation/> (<https://aws.amazon.com/cloudformation/>)

Ask our Experts



QUESTION 14

UNATTEMPTED

If you want to launch Amazon Elastic Compute Cloud (EC2) Instances and assign

each Instance a Predetermined private IP address you should:

- ☐ A. Assign a group or sequential Elastic IP address to the instances
- ☐ B. Launch the instances in a Placement Group
- ☒ C. Launch the instances in the Amazon virtual Private Cloud (VPC). ✓
- ☐ D. Use standard EC2 instances since each instance gets a private Domain Name Service (DNS) already
- ☐ E. Launch the Instance from a private Amazon Machine image.

Explanation :

Answer – C

Amazon Virtual Private Cloud (Amazon VPC) lets you provision a logically isolated section of the Amazon Web Services (AWS) cloud where you can launch AWS resources in a virtual network that you define. You have complete control over your virtual networking environment, including selection of your own IP address range, creation of subnets, and configuration of route tables and network gateways. You can use both IPv4 and IPv6 in your VPC for secure and easy access to resources and applications.

Option A is wrong because assigning Elastic IP's will not ensure that instances get a private IP Address.

Option B is wrong because placement groups does not guarantee that instances get a private IP Address.

Option D is wrong because DNS is a different settings and would not guarantee that instances get a private IP Address.

Option E is wrong because AMI's don't guarantee that instances get a private IP Address.

For more information on VPC, please visit the link:

- <https://aws.amazon.com/vpc/> (<https://aws.amazon.com/vpc/>)

Ask our Experts



QUESTION 15 UNATTEMPTED

What would happen to an RDS (Relational Database Service) multi-Availability Zone deployment of the primary DB instance fails?

- ☐ A. The IP of the primary DB instance is switched to the standby DB instance
- ☐ B. The RDS (Relational Database Service) DB instance reboots
- ☐ C. A new DB instance is created in the standby availability zone
- ☒ D. The canonical name record (CNAME) is changed from primary to standby ✓

Explanation :

Answer – D

Amazon RDS Multi-AZ deployments provide enhanced availability and durability for Database (DB) Instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB Instance, Amazon RDS automatically creates a primary DB Instance and synchronously replicates the data to a standby instance in a different Availability Zone (AZ). Each AZ runs on its own physically distinct, independent infrastructure, and is engineered to be highly reliable. In case of an infrastructure failure (for example, instance hardware failure, storage failure, or network disruption), Amazon RDS performs an automatic failover to the standby, so that you can resume database operations as soon as the failover is complete.

And as per the AWS documentation, the cname is changed to the standby DB when the primary one fails.

Q: What happens during Multi-AZ failover and how long does it take?

Failover is automatically handled by Amazon RDS so that you can resume database operations as quickly as possible without administrative intervention. When failing over, Amazon RDS simply flips the canonical name record (CNAME) for your DB Instance to point at the standby, which is in turn promoted to become the new primary. We encourage you to follow best practices and implement database connection retry at the application layer.

For more information on Multi-AZ RDS, please visit the link:

- <https://aws.amazon.com/rds/details/multi-az/> (<https://aws.amazon.com/rds/details/multi-az/>)

Ask our Experts



QUESTION 16 UNATTEMPTED

An organization has created 5 IAM users. The organization wants to give them the same login ID but different passwords. How can the organization achieve this?

- ☐ A. The organization should create a separate login ID but give the IAM users the same alias so that each one can login with their alias
- ☐ B. The organization should create each user in a separate region so that they have their own URL to login
- ☐ C. It is not possible to have the same login ID for multiple IAM users of the same account ✓
- ☐ D. The organization should create various groups and add each user with the same login ID to different groups. The user can login with their own group ID

Explanation :

Answer – C

AWS Identity and Access Management (IAM) enables you to securely control access to AWS services and resources for your users. Using IAM, you can create and manage AWS users and groups, and use permissions to allow and deny their access to AWS resources.

IAM is a feature of your AWS account offered at no additional charge. You will be charged only for use of other AWS services by your users.

It is not possible in any way to have the same ID and multiple passwords for different IAM users.

For more information on IAM, please visit the link:

- <http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>
(<http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>)

Ask our Experts



QUESTION 17 UNATTEMPTED

A user is planning to evaluate AWS for their internal use. The user does not want to incur any charge on his account during the evaluation. Which of the below mentioned AWS services would incur a charge if used?

- ☐ A. AWS S3 with 1 GB of storage
- ☐ B. AWS micro instance running 24 hours daily
- ☐ C. AWS ELB running 24 hours a day
- ☒ D. AWS Provisioned IOPS volume of 10 GB size ✓

Explanation :

Answer – D

As per the AWS documentation, Option A is wrong because 5GB is provided as part of S3.

As per the AWS documentation, Option B is wrong because you have 750 hours of EC2 Compute.

COMPUTE

750 HOURS

per month

Amazon EC2

Resizable compute capacity in the Cloud

[Learn more about Amazon EC2 »](#)

As per the AWS documentation, Option C is wrong because you have 750 hours of ELB.

For more information on the AWS Free tier, please visit the link:

- <https://aws.amazon.com/free/> (<https://aws.amazon.com/free/>)

Ask our Experts



QUESTION 18 UNATTEMPTED

A user has developed an application which is required to send the data to a NoSQL database. The user wants to decouple the data sending such that the application keeps processing and sending data but does not wait for an acknowledgement of DB. Which of the below mentioned applications helps in this scenario?

- ☐ A. AWS Simple Notification Service
- ☐ B. AWS Simple Workflow
- ☒ C. AWS Simple Queue Service ✓
- ☐ D. AWS Simple Query Service

Explanation :

Answer – C

SQS is the basic decoupling service provided by AWS.

Amazon Simple Queue Service (SQS) is a fast, reliable, scalable, fully managed message queuing service. Amazon SQS makes it simple and cost-effective to decouple the components of a cloud application. You can use Amazon SQS to transmit any volume of data, without losing messages or

requiring other services to be always available. Amazon SQS includes standard queues with high throughput and at-least-once processing, and FIFO queues that provide FIFO (first-in, first-out) delivery and exactly-once processing.

For more information on the SQS, please visit the link:

- <https://aws.amazon.com/sqs/> (<https://aws.amazon.com/sqs/>)

Ask our Experts



QUESTION 19 UNATTEMPTED

An organization has created 50 IAM users. The organization has introduced a new policy which will change the access of an IAM user. How can the organization implement this effectively so that there is no need to apply the policy at the individual user level?

- ☐ A. Use the IAM groups and add users as per their role to different groups and apply policy to group ✓
- ☐ B. The user can create a policy and apply it to multiple users in a single go with the AWS CLI
- ☐ C. Add each user to the IAM role as per their organization role to achieve effective policy setup
- ☐ D. Use the IAM role and implement access at the role level

Explanation :

Answer – A

For applying access across multiple users, you need to have IAM groups. This is the best practice from AWS for user management.

Groups let you specify permissions for multiple users, which can make it easier to manage the permissions for those users. For example, you could have a group called Administrators and give that group the types of permissions that administrators typically need. Any user in that group automatically has the permissions that are assigned to the group.

For more information on IAM, please visit the link:

- <http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>
(<http://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>)

Ask our Experts



A user is planning to use AWS Cloud formation for his automatic deployment requirements. Which of the below mentioned components are required as a part of the template?

- ☐ A. Parameters
- ☐ B. Outputs
- ☐ C. Template version
- ☐ D. Resources ✓

Explanation :

Answer – D

The resources section is required by the CloudFormation template. The other components are optional. An example cloudformation template is shown below. This template creates an EC2 instance based on the Image ID - ami-d6f32ab5

```
{
  "Resources": {
    "MyEC2Instance": {
      "Type": "AWS::EC2::Instance",
      "Properties": { "ImageId": "ami-d6f32ab5" }
    }
  }
}
```

For more information on Cloud Formation, please visit the link:

- <https://aws.amazon.com/cloudformation/> (<https://aws.amazon.com/cloudformation/>)

Ask our Experts



A user has created a web application with Auto Scaling. The user is regularly monitoring the application and he observed that the traffic is highest on Thursday and Friday between 8 AM to 6 PM. What is the best solution to handle scaling in this case?

- ☐ A. Add a new instance manually by 8 AM Thursday and terminate the same by 6 PM Friday
- ☒ B. Schedule Autoscaling to scale up by 8 AM and scale down after 6 PM on both Thursday and Friday
- ☐ C. Schedule a policy which may scale up every day at 8 AM and scales down by 6 PM
- ☐ D. Configure a batch process to add a instance by 8 AM and remove it by Friday 6 PM

Explanation :

Answer – B

To configure your Auto Scaling group to scale based on a schedule, you create a scheduled action, which tells Auto Scaling to perform a scaling action at specified times. To create a scheduled scaling action, you specify the start time when you want the scaling action to take effect, and the new minimum, maximum, and desired sizes for the scaling action. At the specified time, Auto Scaling updates the group with the values for minimum, maximum, and desired size specified by the scaling action.

For more information on AutoScaling, please visit the link:

- <https://aws.amazon.com/autoscaling/> (<https://aws.amazon.com/autoscaling/>)

Ask our Experts



QUESTION 22 UNATTEMPTED

A user has setup a CloudWatch alarm on an EC2 action when the CPU utilization is above 75%. The alarm sends a notification to SNS on the alarm state. If the user wants to simulate the alarm action how can he achieve this?

- ☐ A. Run activities on the CPU such that its utilization reaches above 75%
- ☐ B. From the AWS console change the state to 'Alarm'
- ☒ C. The user can set the alarm state to 'Alarm' using CLI
- ☐ D. Run the SNS action manually

Explanation :

Answer – C

The easiest way to set the ALARM in Cloudwatch is to trigger the alarm itself and that can be done via the CLI. To change the state of the alarm via the CLI with the set-alarm-state function. Below is an example of the CLI command which sets the state of the alarm.

```
AWS cloudwatch set-alarm-state --alarm-name "Testalarm" --state-value ALARM --state-reason "Demo purposes"
```

For more information on AWS Cloudwatch, please visit the link:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)
- <http://docs.aws.amazon.com/cli/latest/reference/cloudwatch/set-alarm-state.html>
(<http://docs.aws.amazon.com/cli/latest/reference/cloudwatch/set-alarm-state.html>)

Ask our Experts



QUESTION 23 UNATTEMPTED

When preparing for a compliance assessment of your system built inside of AWS. What are the 3 best practises for you to prepare for an audit?

Choose 3 answers from the options given below.

- ☐ A. Gather evidence of your IT Operational controls. ✓
- ☐ B. Request and obtain third party audited AWS compliance reports and certifications. ✓
- ☐ C. Request and obtain a compliance and security tour of an AWS data center for a pre-assessment security review.
- ☐ D. Request and obtain approval from AWS to perform relevant network scans and in-depth penetration testing of the user system's instances and endpoints. ✓
- ☐ E. Schedule meetings with AWS third-party auditors to provide evidence of AWS compliance that maps to your control objectives.

Explanation :

Answer – A,B and D

The first major requirement is for an organization to evaluate all the controls they have put in place for their AWS environment. So like who has access to what elements in AWS, how is data secured at rest etc.

One can also request AWS to perform network tests and penetration tests to ensure their environment is secure.

And finally there are third party's available for carrying out relevant audits.

Option C is incorrect because AWS does not allow data center tour.

Option E is incorrect because customers can request the reports and certifications produced by our third-party auditors or also can request more information about AWS Compliance.

Who performs the independent third-party audit of AWS for the SOC Reports? ^

Ernst & Young LLP performs the AWS SOC 1, SOC 2 and SOC 3 audits.

For more information on compliance and security, please visit the below URL's

- <https://aws.amazon.com/compliance/soc-faqs/> (<https://aws.amazon.com/compliance/soc-faqs/>)
- http://d0.awsstatic.com/whitepapers/compliance/AWS_Auditing_Security_Checklist.pdf (http://d0.awsstatic.com/whitepapers/compliance/AWS_Auditing_Security_Checklist.pdf)

Ask our Experts



QUESTION 24 UNATTEMPTED

A user has setup a billing alarm using CloudWatch for \$200. The usage of AWS exceeded \$200 after some days. The user wants to increase the limit from \$200 to \$400? What should the user do?

- ☐ A. Create a new alarm of \$400 and link it with the first alarm
- ☐ B. It is not possible to modify the alarm once it has crossed the usage limit
- ☒ C. Update the alarm to set the limit at \$400 instead of \$200 ✓
- ☐ D. Create a new alarm for the additional \$200 amount

Explanation :

Answer – C

Let's assume that an alarm has been created as shown below for any amounts exceeding 200 USD.

Create Alarm Modify Copy Delete		
Filter: All alarms AWS/Billing X		
State	Name	Threshold
<input checked="" type="checkbox"/> INSUFFICIENT_DATA	BillingAlarm	EstimatedCharges > 200 for 6 hours

To increase the limit, all you have to do is to click on the Modify option and you can change the value of the alarm in the next screen.

Modify Alarm

- [1. Select Metric](#) [2. Define Alarm](#)

Billing Alarm

You can create a billing alarm to receive e-mail alerts when your AWS charges exceed a threshold you choose. Simply:

1. Enter a spending threshold
2. Provide an email address
3. Check your inbox for a confirmation email and click the link provided

When my total AWS charges for the month

exceed: \$ USD

send a notification to: [New list](#)

For more information on AWS Cloudwatch, please visit the link:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)

Ask our Experts



QUESTION 25 UNATTEMPTED

A user is trying to save some cost on the AWS services. Which of the below mentioned options will not help him save cost?

- ☐ A. Delete the unutilized EBS volumes once the instance is terminated
- ☒ B. Delete the AutoScaling launch configuration after the instances are terminated
- ☐ C. Release the elastic IP if not required once the instance is terminated
- ☐ D. Delete the AWS ELB after the instances are terminated

Explanation :

Answer – B

Option A is wrong because EBS volumes does have a costing aspect and hence deleting the volumes will save on cost.

Option C is wrong because Elastic IP will consume cost if not removed.

Option D is wrong because ELB also incur costs.

Only Autoscaling groups are free of cost. It's only the underlying resources which you are charged for.

For more information on AWS Pricing, please visit the link:

- <https://aws.amazon.com/pricing/services/> (<https://aws.amazon.com/pricing/services/>)

Ask our Experts



QUESTION 26 UNATTEMPTED

A user is trying to aggregate all the CloudWatch metric data of the last 1 week. Which of the below mentioned statistics is not available for the user as a part of data aggregation?

- ☐ A. Aggregate ✓
- ☐ B. Sum
- ☐ C. Data Samples
- ☐ D. Average

Explanation :

Answer – A

If you go to Cloudwatch and go to any metric and see the statistic column you will see the below dimensions for statistics. And the only one which is not there from the above list is Aggregate.



For more information on AWS Cloudwatch, please visit the link:

- <https://aws.amazon.com/cloudwatch/> (https://aws.amazon.com/cloudwatch/)

Ask our Experts



QUESTION 27 UNATTEMPTED

An organization is planning to use AWS for their production roll out. The organization wants to implement automation for deployment such that it will automatically create a LAMP stack, download the latest PHP installable from S3 and setup the ELB. Which of the below mentioned AWS services meets the requirement for making an orderly deployment of the software?

- ☒ A. AWS Elastic Beanstalk ✓
- ☐ B. AWS Cloudfront
- ☐ C. AWS Cloudformation
- ☐ D. AWS DevOps

Explanation :

Answer – A

The Elastic Beanstalk is an easy-to-use service for deploying and scaling web applications and services.

We can simply upload code and Elastic Beanstalk automatically handles the deployment, from capacity provisioning, load balancing, auto-scaling to application health monitoring. Meanwhile we can retain full control over the AWS resources used in the application and can access the underlying resources at any time.

Hence, A is the CORRECT answer.

Although Cloud formation can also perform the same requirement of the scenario it would be easier to deploy the software using ElasticBeanstalk environment rather than Cloud formation templates.

For more information on launching a LAMP stack with Elastic Beanstalk:

- <https://aws.amazon.com/getting-started/projects/launch-lamp-web-app/>
(<https://aws.amazon.com/getting-started/projects/launch-lamp-web-app/>)

Note:

I do understand concern about Cloudformation which can also be answer.

We can modify the question as follows.

"An organization is planning to use AWS for their production roll out. The organization wants to implement automation for deployment such that it will automatically create a LAMP tack, download the latest PHP installable from S3 and set up the ELB. **The initial ease of deployment is also important.** Which of the below mentioned AWS services meets the requirement for making an orderly deployment of the software?"
?

However in the real exam, the question was asked like that. So its better if we keep it the same even though it makes hard for students to answer the question.

(<https://aws.amazon.com/cloudformation/>)

QUESTION 28

UNATTEMPTED

A user has created a subnet with VPC and launched an EC2 instance in that subnet with only default settings. Which of the below mentioned options is ready to use on the EC2 instance as soon as it is launched?

- ☐ A. Elastic IP
- ☒ B. Private IP ✓
- ☐ C. Public IP
- ☐ D. Internet gateway

Explanation :

Answer – B

When you create a subnet with the default settings, only the Private IP gets populated for EC2 instances. For Public IP, this is not possible because the Auto-assign Public IP will be 'no' by default. Also the Elastic IP and Internet gateway have to manually configured.

Subnet ID: subnet-dfd2a5f2 Default	Availability Zone: us-east-1a
CIDR: 192.168.0.0/28	Route table: rtb-c405f1bd
State: available	Network ACL: acl-7dd5ae1b
VPC: vpc-3e6dde58 192.168.0.0/24	Default subnet: no
Available IPs: 11	Auto-assign Public IP: no

For more information on VPC, please visit the link:

- <https://aws.amazon.com/vpc/> (<https://aws.amazon.com/vpc/>)

Ask our Experts



QUESTION 29

UNATTEMPTED

An organization is setting up programmatic billing access for their AWS account. Which of the below mentioned services is not required or enabled when the organization wants to use programmatic access?

- ☐ A. Programmatic access
- ☐ B. AWS bucket to hold the billing report

- ☒ C. AWS billing alerts ✓
- ☐ D. Monthly Billing report

Explanation :

Answer – C

Since the question is looking for service which is "**not**" required or enabled when the organization wants to use programmatic access , the correct answer is Option C. AWS Billing Alerts.

AWS gives an option to provide programmatic access to billing. Programmatic Billing Access takes the existing Amazon S3 APIs. So, the user will be able to build applications that reference the billing data from a CSV file which is stored in an Amazon S3 bucket.

In order to enable programmatic access, the user has to first enable the monthly billing report. Then he needs to provide an AWS bucket name in which the billing CSV will be uploaded. He must also enable the Programmatic access option.

For more information on AWS billing, please visit the link:

- <https://aws.amazon.com/documentation/account-billing/>
(<https://aws.amazon.com/documentation/account-billing/>)

Ask our Experts



QUESTION 30 UNATTEMPTED

A user is planning to use AWS Cloudformation.

Which functionality does not help him to correctly understand Cloudformation?

- ☒ A. Cloudformation follows the DevOps model for the creation of Dev & Test ✓
- ☐ B. AWS Cloudfromation does not charge the user for its service but only charges for the AWS resources created with it
- ☐ C. Cloudformation works with a wide variety of AWS services, such as EC2, EBS, VPC, IAM, S3, RDS, ELB, etc
- ☐ D. CloudFormation provides a set of application bootstrapping scripts which enables the user to install Software

Explanation :

Answer – A

So please refer to the below snapshots from the AWS documentation on what is true on Cloudformation

Q: How much does AWS CloudFormation cost?

There is no additional charge for AWS CloudFormation. You only pay for the AWS resources that are created (e.g., Amazon EC2 instances, Elastic Load Balancing load balancers etc.)

Here it mentions that Cloudformation does not have any explicit charge so Option B is correct.

Q: How is AWS CloudFormation different from AWS Elastic Beanstalk?

These services are designed to complement each other. [AWS Elastic Beanstalk](#) provides an environment to easily deploy and run applications in the cloud. It is integrated with developer tools and provides a one-stop experience for you to manage the lifecycle of your applications. AWS CloudFormation is a convenient provisioning mechanism for a broad range of [AWS resources](#). It supports the infrastructure needs of many different types of applications such as existing enterprise applications, legacy applications, applications built using a variety of AWS resources and container-based solutions (including those built using AWS Elastic Beanstalk).

Here it mentions that Cloudformation works with many AWS resources so Option C is correct.

Q: Can I use AWS CloudFormation with Chef?

Yes, AWS CloudFormation can be used to bootstrap both the Chef Server and Chef Client software on your EC2 instances. For more details and a how-to see [Integrating AWS CloudFormation with Chef](#).

Here it mentions that Cloudformation does have bootstrapping capabilities so Option D is correct.
For more information on Cloudformation please visit the link:

- <https://aws.amazon.com/cloudformation/> (<https://aws.amazon.com/cloudformation/>)

Note:

Please note "not" is the question - "A user is planning to use AWS Cloudformation. Which of the below mentioned functionalities does **not** help him to correctly understand Cloudformation?"

So Option B. AWS Cloudformation does not charge the user for its service but only charges for the AWS resources created with it -- Its True and It helps in understanding CloudFormation

Where as Option A. Cloudformation follows the DevOps model for the creation of Dev & Test - is false and does not help in understanding CloudFormation.

So Option A is correct.

Ask our Experts



QUESTION 31 UNATTEMPTED

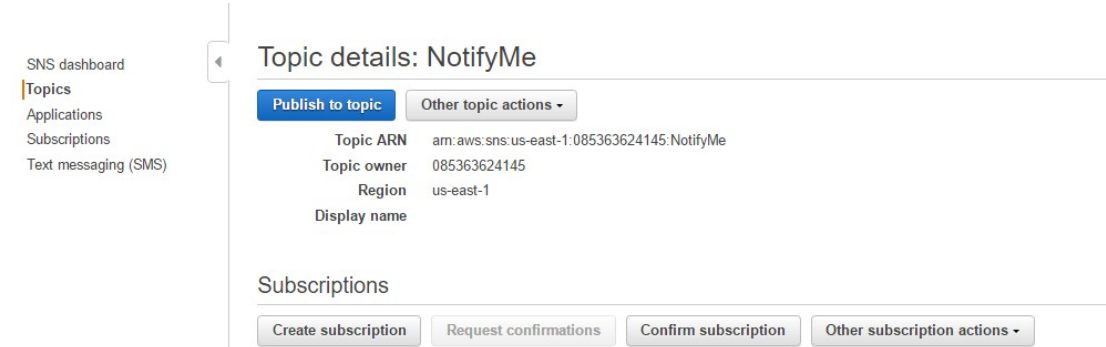
A user is trying to understand AWS SNS. To which of the below mentioned end points is SNS unable to send a notification?

- ☐ A. Email JSON
- ☐ B. HTTP
- ☐ C. AWSSQS
- ☒ D. AWSSES ✓

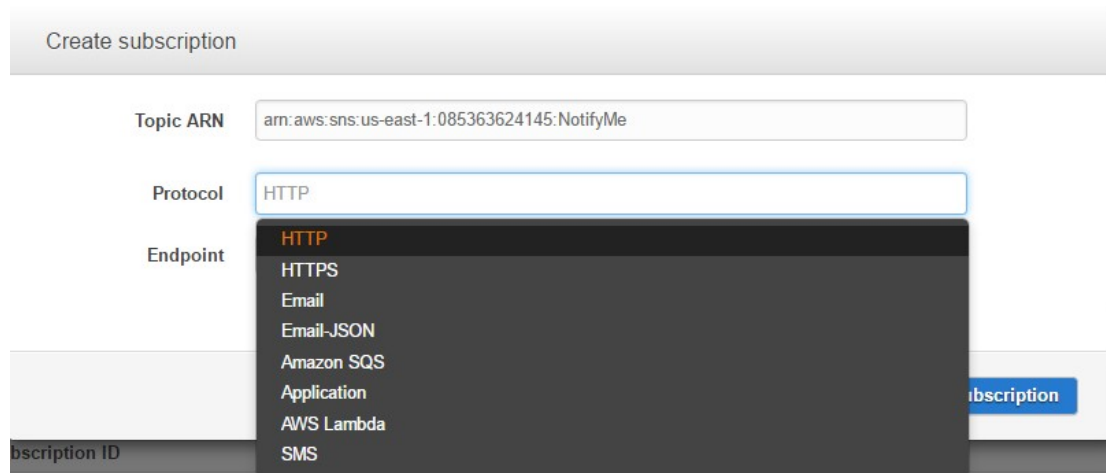
Explanation :

Answer – D

Let's see a sample Topic created in the SNS dashboard as shown below



When you create a subscription, you can see the different Protocols available



From the snapshot it is evident that option D is invalid.

For more information on SNS please visit the link:

- <https://aws.amazon.com/sns/> (https://aws.amazon.com/sns/)

Ask our Experts



QUESTION 32 UNATTEMPTED

You are building an online store on AWS that uses SQS to process your customer orders. Your backend system needs those messages in the same sequence the customer orders have been put in. How can you achieve that?

- ☐ A. It is not possible to do this with SQS
- ☒ B. You can use sequencing information on each message ✓

- ☐ C. You can do this with SQS but you also need to use SWF
- ☐ D. Messages will arrive in the same order by default

Explanation :

Answer – B

If you look at the AWS documentation, it is very clear that SQS messages does not guarantee the order of messages. So in order to do this, you need to add the sequencing information in each message itself.

Option D is wrong as per the below article.

Yes. FIFO (first-in-first-out) queues preserve the exact order in which messages are sent and received. If you use a FIFO queue, you don't have to place sequencing information in your messages. For more information, see [FIFO Queue Logic](#) in the *Amazon SQS Developer Guide*.

Standard queues provide a loose-FIFO capability that attempts to preserve the order of messages. However, because standard queues are designed to be massively scalable using a highly distributed architecture, receiving messages in the exact order they are sent is not guaranteed.

For more information on SQS please visit the link:

- <https://aws.amazon.com/sqs/faqs> (<https://aws.amazon.com/sqs/faqs>)

Note:

SQS FIFO queues preserves the order of messages while SQS does not guarantee the order or sequencing of messages. This is what has been elaborated on the explanation part of the solution too.

However the question is not mentioning about SQS FIFO queues . So inorder to preserve the sequence on SQS queues, we need to add additional sequencing information along with the messages.

Ask our Experts



QUESTION 33 UNATTEMPTED

An organisation wants to move their databases to Cloud. They are planning to use AWS Database Migration Service and is looking for a secure encrypted database storage option. Which of the following option will be suitable for their requirement?

- ☐ A. AWS MFA with EBS
- ☐ B. AWS EBS encryption
- ☐ C. Multi-tier encryption with Redshift ✓

☐ **D. AWS S3 server side storage**

Explanation :

Answer – C

In Amazon Redshift, you can enable database encryption for your clusters to help protect data at rest. When you enable encryption for a cluster, the data blocks and system metadata are encrypted for the cluster and its snapshots.

<https://aws.amazon.com/dms/faqs/> (<https://aws.amazon.com/dms/faqs/>)

Ask our Experts



QUESTION 34 UNATTEMPTED


You have been asked to leverage Amazon VPC ,EC2 and SQS to implement an application that submits and receives millions of messages per second to a message queue. You want to ensure your application has sufficient bandwidth between your EC2 instances and SQS. Which option will provide the most scalable solution for communicating between the application and SQS?

- ☐ **A. Ensure the application instances are properly configured with an Elastic Load Balancer**
- ☐ **B. Ensure the application instances are launched in private subnets with the EBS-optimized option enabled**
- ☐ **C. Ensure the application instances are launched in public subnets with the associate-public-IP-address=true option enabled**
- ☐ **D. Launch application instances in private subnets with an Auto Scaling group and Auto Scaling triggers configured to watch the SQS queue size ✓**

Explanation :

Answer – D

When you have an SQS configured with EC2 instances, the documented option is to scale up EC2 instances in an AutoScaling group based on length of the message queue.



Option A is wrong because the load balancer will not scale up EC2 instances. This will be done by the AutoScaling group.

Option B and C is wrong because subnets do not dictate the scalability of the solution.

For more information on SQS, please visit the below URL:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 35 UNATTEMPTED

A user wants to disable connection draining on an existing ELB. Which of the below mentioned statements helps the user disable connection draining on the ELB?

- ☐ A. The user can only disable connection draining from CLI
- ☐ B. It is not possible to disable the connection draining feature once enabled
- ☐ C. The user can disable the connection draining feature from EC2 -> ELB console or from CLI ✓
- ☐ D. The user needs to stop all instances before disabling connection draining

Explanation :

Answer – C

When you have an ELB , you can go to the Console, go to the Instances tab and edit the Connection draining time for instances.

You can also modify from the CLI via the below command. The below command will set the timeout to 100 seconds for the loadbalancer named my-loadbalancer.

```
AWS elb modify-load-balancer-attributes --load-balancer-name my-loadbalancer --load-balancer-attributes "{\"ConnectionDraining\":{\"Enabled\":true,\"Timeout\":100}}"
```

For more information on Amazon ELB Connection draining please visit the link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/config-conn-drain.html>
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/config-conn-drain.html>)

Ask our Experts



QUESTION 36 UNATTEMPTED

A user has created an S3 bucket which is not publicly accessible. The bucket is having thirty objects which are also private. If the user wants to make the objects public, how can he configure this with minimal efforts?

- ☐ A. The user should select all objects from the console and apply a single policy to mark them public
- ☐ B. The user can write a program which programmatically makes all objects public using S3 SDK
- ☐ C. Set the AWS bucket policy which marks all objects as public ✓
- ☐ D. Make the bucket ACL as public so it will also mark all objects as public

Explanation :

Answer – C

You can set AWS bucket policy to make everything public. You can do via the following steps

Step 1) Go to your S3 bucket and click on the Permissions section

Permissions

► Permissions

► Static Website Hosting

► Logging

► Events

► Versioning

► Lifecycle

► Cross-Region Replication

Step 2) Go to AWS Policy Generator

Bucket Policy Editor

Cancel

Policy for Bucket : "devtoolslogging"

Add a new policy or edit an existing bucket policy in the text area below. [Learn more.](#)

AWS Policy Generator | [Sample Bucket Policies](#)

Save

Delete

Close

Step 3) Add the S3 bucket policy to allow any Principal to access your bucket and choose the Actions accordingly.

Select Type of Policy

S3 Bucket Policy

Step 2: Add Statement(s)

A statement is the formal description of a single permission. See [a description of elements](#) that you can use in statements.

Effect

☒ Allow
 ☐ Deny

Principal

*

Use a comma to separate multiple values.

AWS Service

Amazon S3

All Services (*)

Use multiple statements to add permissions for more than one service.

Actions

-- Select Actions --

All Actions (*)

Amazon Resource Name (ARN)

arn:aws:s3:::bucketnm/*

ARN should follow the following format: arn:aws:s3:::<bucket_name>/<key_name>.

Use a comma to separate multiple values.

Add Conditions (Optional)

Add Statement

For more information on S3 Security please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/DataDurability.html>
 (http://docs.aws.amazon.com/AmazonS3/latest/dev/DataDurability.html)

Ask our Experts



QUESTION 37

UNATTEMPTED

A system admin is maintaining an application on AWS. The application is installed on EC2 and user has configured ELB and Auto Scaling. Considering future load increase, the user is planning to launch new servers proactively so that they get registered with ELB. How can the user add these instances with Auto Scaling?

- ☒ A. Increase the desired capacity of the Auto Scaling group ✓
- ☐ B. Increase the maximum limit of the Auto Scaling group
- ☐ C. Launch an instance manually and register it with ELB on the fly
- ☐ D. Decrease the minimum limit of the Auto Scaling group

Explanation :

Answer – A

To increase instances proactively you need to increase the desired limit.

Option B is wrong, because it indicates the maximum number of instances that we can launch in that group, however it will not get launched until the trigger occurs.

Option C is wrong because this is not a proactive measure by adding instances to an ELB.

Option D is wrong because the minimum limit will just tell what minimum scale the instances can terminate to, and does not look to launch new instances.

For more information on AutoScaling, please visit the link:

- <https://aws.amazon.com/autoscaling/> (<https://aws.amazon.com/autoscaling/>)

Ask our Experts



QUESTION 38

UNATTEMPTED

An organization, which has the AWS account ID as 999988887777, has created 50 IAM users. All the users are added to the same group demo. If the organization has enabled that each IAM user can login with the AWS console, which AWS login URL will the IAM users use?

- ☒ A. <https://999988887777.signin.aws.amazon.com/console/> ✓
- ☐ B. <https://signin.aws.amazon.com/demo/>
- ☐ C. <https://demo.signin.aws.amazon.com/999988887777/console/>
- ☐ D. <https://999988887777.aws.amazon.com/demo/>

Explanation :

Answer – A

When you go to the IAM dashboard, you can see the sign-in link which can be used. The sign in line is always prefixed by the account number. And the last keyword is console. By that aspect only option A is right.

For more information on IAM, please visit the link:

- <https://aws.amazon.com/iam/faqs/> (<https://aws.amazon.com/iam/faqs/>)

Ask our Experts



QUESTION 39 UNATTEMPTED

A user has setup connection draining with ELB to allow in-flight requests to continue while the instance is being deregistered through Auto Scaling. If the user has not specified the draining time, how long will ELB allow inflight requests traffic to continue?

- ☐ A. 600 seconds
- ☐ B. 3600 seconds
- ☒ C. 300 seconds ✓
- ☐ D. 0 seconds

Explanation :

Answer – C

When you have an ELB, you can go to the Console, go to the Instances tab and edit the Connection draining time for instances. BY default the Connection draining time limit is set to 300.

The screenshot shows the AWS Management Console interface for a load balancer named 'Demo'. At the top, there are buttons for 'Create Load Balancer' and 'Actions'. Below is a table with columns: Name, DNS name, State, VPC ID, Availability Zones, and Type. The table contains one entry: 'Demo' with DNS name 'internal-demo-824544633...', VPC ID 'vpc-3e6dde58', Availability Zones 'us-east-1a', and Type 'classic'. Below the table, there are tabs for 'Description', 'Instances', 'Health Check', 'Listeners', 'Monitoring', and 'Tags'. The 'Instances' tab is active, showing 'Connection Draining: Enabled, 300 seconds (Edit)'. Below this is an 'Edit Instances' button and a table with columns: Instance ID, Name, Availability Zone, Status, and Actions. The table contains one entry: Instance ID 'i-00acbc260371247b7', Availability Zone 'us-east-1a', Status 'OutOfService', and Action 'Remove from Load Balancer'.

Instance ID	Name	Availability Zone	Status	Actions
i-00acbc260371247b7		us-east-1a	OutOfService ⓘ	Remove from Load Balancer

For more information on ELB, please visit the link:

- <https://aws.amazon.com/elasticloadbalancing/>
(<https://aws.amazon.com/elasticloadbalancing/>)

Ask our Experts



QUESTION 40 UNATTEMPTED

A root AWS account owner is trying to understand various options to set the permission to AWS S3. Which of the below mentioned options is not the right option to grant permission for S3?

- ☐ A. User Access Policy
- ☐ B. S3 Object Access Policy ✓
- ☐ C. S3 Bucket Access Policy
- ☐ D. S3 ACL

Explanation :

Answer – B

In S3 when you go to the permissions section you can add the Grantee which is at the ACL level and then you can also add bucket permissions.

And then you can also create IAM policies at the user level to manage access to S3.

Hence Option A,C and D are valid. And hence Option B is the right answer to the question.

For more information on S3 Security, please visit the link:

- <https://docs.aws.amazon.com/AmazonS3/latest/dev/s3-access-control.html>
(<https://docs.aws.amazon.com/AmazonS3/latest/dev/s3-access-control.html>)

Ask our Experts



QUESTION 41 UNATTEMPTED

A system admin has created a shopping cart application and hosted it on EC2. The EC2 instances are running behind ELB. The admin wants to ensure that the end user request will always go to the EC2 instance where the user session has been created. How can the admin configure this?

- ☐ A. Enable ELB cross zone load balancing
- ☐ B. Enable ELB cookie setup
- ☐ C. Enable ELB sticky session ✓
- ☐ D. Enable ELB connection draining

Explanation :

Answer – C

To ensure that each end user request goes to the same EC2 instance as the session created, you need to enable stickiness at the ELB level

To enable stickiness go to the ELB and in the port configuration section you can enable the stickiness.

Create Load Balancer **Actions** ▾

Filter: X

<input type="checkbox"/>	Name	DNS name	State
<input type="checkbox"/>	Demo	internal-Demo-824544633...	

Scheme: internal

Availability Zones: subnet-dfd2a5f2 - us-east-1a

Port Configuration

Port Configuration: 80 (HTTP) forwarding to 80 (HTTP)
Stickiness: Disabled

Edit stickiness

For more information on ELB, please visit the link:

- <https://aws.amazon.com/elasticloadbalancing/>
(<https://aws.amazon.com/elasticloadbalancing/>)

Ask our Experts



An organization is using AWS since a few months. The finance team wants to visualize the pattern of AWS spending. Which of the below AWS tool will help for this requirement?

- ☐ A. AWS Cost Manager
- ☒ B. AWS Cost Explorer ✓
- ☐ C. AWS CloudWatch
- ☐ D. AWS Consolidated Billing

Explanation :

Answer – B

To get more details on your spending you can go to the Cost Explorer. Go to the Billing section and click on Cost Explorer.

Dashboard
Bills
Cost Explorer
Budgets
Reports
Cost Allocation Tags
Payment Methods
Payment History
Consolidated Billing
Preferences
Credits
Tax Settings
DevPay

Cost Explorer



[Launch Cost Explorer](#)

Graph, visualize, and analyze your spend. Filter what you see by specifying date ranges, services, tags, or a combination. [Learn more](#)

Preconfigured Views

Get started quickly with a view that is already set up to answer a common spend inquiry. Customize the view from there.

[Monthly Spend by Service View](#)

Monthly spend over last three months, grouped by AWS Services

[Monthly Spend by Linked Account View](#)

Monthly spend over last three months, grouped by Linked Accounts. Only available for Paying accounts in Consolidated Billing

[Daily Spend View](#)

Daily spend over last 60 days

And when you launch the Cost Explorer you can see more details on the usage of AWS services.

For more information on AWS billing, please visit the link:

- <https://aws.amazon.com/documentation/account-billing/>
(<https://aws.amazon.com/documentation/account-billing/>)

Ask our Experts



A user has launched an ELB which has 5 instances registered with it. The user deletes the ELB by mistake. What will happen to the instances?

- ☐ A. ELB will ask the user whether to delete the instances or not
- ☐ B. Instances will be terminated
- ☐ C. ELB cannot be deleted if it has running instances registered with it
- ☐ D. Instances will keep running ✓

Explanation :

Answer – D

From the AWS documentation on ELB deletion it clearly mentions that EC2 instances will continue to run even if the ELB is deleted.

Note that deleting a load balancer does not affect its registered targets. For example, your EC2 instances continue to run and are still registered to their target groups. To delete your target groups, see [Delete Your Target Group](#).

For more information on ELB deletion, please visit the link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-delete.html> (<http://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-delete.html>)

Ask our Experts



QUESTION 44 UNATTEMPTED

A user is planning to setup notifications on the RDS DB for a snapshot. Which of the below mentioned event categories is not supported by RDS for this snapshot source type?

- ☐ A. Backup ✓
- ☐ B. Creation
- ☐ C. Deletion
- ☐ D. Restoration

Explanation :

Answer – A

When you go to the Event Subscriptions section for an RDS and choose the source as snapshot, you can see the below options. Backup is not available as an option and hence A is the right option.

Instances
Clusters
Reserved Purchases
Snapshots
Security Groups
Parameter Groups
External Licenses
Option Groups
Subnet Groups
Events
Event Subscriptions
Notifications 1

Send notifications to - Select ARN - create topic

Source Type Snapshots

Enabled Yes No

Event Categories

Select All
Select specific

creation
deletion
notification
restoration

Snapshots

Select All
Select specific

arn:aws:rds:us-east-1:001573694396:snapshot:samah12121
arn:aws:rds:us-east-1:020014417079:snapshot:kony-devicedb-v12
arn:aws:rds:us-east-1:020014417079:snapshot:kony-devicedb-v13
arn:aws:rds:us-east-1:020783423197:snapshot:net1-final-snapshot
arn:aws:rds:us-east-1:026619168887:snapshot:cpe481-gp4-db-snapshot
arn:aws:rds:us-east-1:049165365843:snapshot:test-server-snap-

For more information on RDS snapshots, please visit the link:

- http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_CreateSnapshot.html
(http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_CreateSnapshot.html)

Ask our Experts

QUESTION 45 UNATTEMPTED

You have an Auto Scaling group associated with an Elastic Load Balancer (ELB). You have noticed that instances launched via the Auto Scaling group are being marked unhealthy due to an ELB health check, but these unhealthy instances are not being terminated. What do you need to do to ensure instances marked unhealthy by the ELB will be terminated and replaced?

- ☐ A. Change the thresholds set on the Auto Scaling group health check
- ☐ B. Add an Elastic Load Balancing health check to your Auto Scaling group ✓
- ☐ C. Increase the value for the Health check interval set on the Elastic Load Balancer
- ☐ D. Change the health check set on the Elastic Load Balancer to use TCP rather than HTTP checks

Explanation :

Answer – B

To discover the availability of your EC2 instances, a load balancer periodically sends pings, attempts connections, or sends requests to test the EC2 instances. These tests are called health checks. The status of the instances that are healthy at the time of the health check is InService. The status of any instances that are unhealthy at the time of the health check is OutOfService

The load balancer checks the health of the registered instances using either the default health check configuration provided by Elastic Load Balancing or a health check configuration that you configure. When configuring the Autoscaling group, you can choose either the option of EC2 or ELB health checks. Since Ec2 instances are being marked as unhealthy by ELB but not being terminated by Autoscaling it means that the check from the Autoscaling side is wrongly configured.

Create Auto Scaling Group

[Create new VPC](#)

Subnet [Create new subnet](#)
Each instance in this Auto Scaling group will be assigned a public IP address.

▼ Advanced Details

Load Balancing ☒ Receive traffic from one or more load balancers [Learn about Elastic Load Balancing](#)

Classic Load Balancers

Target Groups

Health Check Type ☐ ELB ☒ EC2

Health Check Grace Period seconds

Monitoring Amazon EC2 Detailed Monitoring metrics, which are provided at 1 minute frequency, are not enabled for the launch configuration Demo. Instances launched from it will use Basic Monitoring metrics, provided at 5 minute frequency. [Learn more](#)

Instance Protection

For more information on ELB, please visit the below URL:

- <https://aws.amazon.com/elasticloadbalancing/>
(<https://aws.amazon.com/elasticloadbalancing/>)

Ask our Experts



QUESTION 46

UNATTEMPTED

A customer is using AWS for Dev and Test. The customer wants to setup the Dev environment with Cloudformation. Which of the below mentioned steps are not required while using Cloudformation?

- ☐ A. Create a stack
- ☐ B. Configure a service ✓
- ☐ C. Create and upload the template

☐ D. Provide the parameters configured as part of the template

Explanation :

Answer – B

When you use Cloudformation you can create a stack, provide parameters when creating a template.

When creating a stack you can also upload a template as shown in the below screen.

But you cannot configure a service so Option B is the right answer.

For more information on Cloudformation please visit the link:

- <https://aws.amazon.com/cloudformation/> (<https://aws.amazon.com/cloudformation/>)

Note: Resources element is the only element which is mandatory for the creation of a Cloudformation template. However in this case we are creating a cloudformation template for a development environment. So we can assume that there might be various services like ec2, s3, elb, dynamo db etc used in this environment. So this template will most likely to have optional elements like Parameters, Mappings etc for it to function. So in this case scenario presented here in the question Parameters elements are required.

Below snapshot gives an overview of stack creation of a sample Wordpress template:

Review of stack creation

So the option which is not required in this case scenario is option B.

Ask our Experts



QUESTION 47 UNATTEMPTED

A user is accessing RDS from an application. The user has enabled the Multi AZ feature with the MS SQL RDS DB. During a planned outage how will AWS ensure that a switch from DB to a standby replica will not affect access to the application?

- ☐ A. RDS will have an internal IP which will redirect all requests to the new DB
- ☒ B. RDS uses DNS to switch over to stand by replica for seamless transition ✓
- ☐ C. The switch over changes Hardware so RDS does not need to worry about access
- ☐ D. RDS will have both the DBs running independently and the user has to manually switch over

Explanation :

Answer – B

Amazon RDS Multi-AZ deployments provide enhanced availability and durability for Database (DB) Instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB Instance, Amazon RDS automatically creates a primary DB Instance and synchronously

replicates the data to a standby instance in a different Availability Zone (AZ). Each AZ runs on its own physically distinct, independent infrastructure, and is engineered to be highly reliable. In case of an infrastructure failure (for example, instance hardware failure, storage failure, or network disruption), Amazon RDS performs an automatic failover to the standby, so that you can resume database operations as soon as the failover is complete.

And as per the AWS documentation, the cname is changed to the standby DB when the primary one fails.

Based on this, RDS Multi-AZ will use DNS to create the CNAME and hence B is the right option.

For more information on RDS Multi-AZ please visit the link:

- <http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZ.html>
(<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZ.html>)

Ask our Experts



QUESTION 48 UNATTEMPTED

An organization is generating digital policy files which are required by the admins for verification. Once the files are verified they may not be required in the future unless there is some compliance issue. If the organization wants to save them in a cost effective way, which is the best possible solution?

- ☐ A. AWS RRS
- ☐ B. AWS S3
- ☐ C. AWS RDS
- ☒ D. AWS Glacier ✓

Explanation :

Answer – D

If you look at the documentation for AWS Glacier it clearly mentions that glacier can be used for offline administrative storage and is low cost.

Amazon Glacier is an extremely low-cost storage service that provides secure, durable, and flexible storage for data backup and archival. With Amazon Glacier, customers can reliably store their data for as little as \$0.004 per gigabyte per month. Amazon Glacier enables customers to offload the administrative burdens of operating and scaling storage to AWS, so that they don't have to worry about capacity planning, hardware provisioning, data replication, hardware failure detection and repair, or time-consuming hardware migrations.

Option A and B are wrong because they are used for real time storage.

Option C is wrong because this is for database service.

For more information on Glacier please visit the link:

- <https://aws.amazon.com/glacier/details/> (<https://aws.amazon.com/glacier/details/>)

Ask our Experts



QUESTION 49 UNATTEMPTED

An organization has configured the custom metric upload with CloudWatch. The organization has given permission to its employees to upload data using CLI as well SDK. How can the user track the calls made to CloudWatch?

- ☐ A. The user can enable logging with CloudWatch which logs all the activities
- ☒ B. Use CloudTrail to monitor the API calls ✓
- ☐ C. Create an IAM user and allow each user to log the data using the S3 bucket
- ☐ D. Enable detailed monitoring with CloudWatch

Explanation :

Answer – B

AWS Cloudtrail is the defacto service provided by AWS for monitoring all API calls to AWS and is used for logging and monitoring purposes for compliance purposes. Amazon cloudtrail detects every call made to AWS and creates a log which can then be further used for analysis.

For more information on Cloudtrail please visit the link:

- <https://aws.amazon.com/cloudtrail/> (<https://aws.amazon.com/cloudtrail/>)

Ask our Experts



QUESTION 50 UNATTEMPTED

A user has created a queue named “demoqueue” with SQS. There are four messages published to queue which are not received by the consumer yet. If the user tries to delete the queue, what will happen?

- ☐ A. A user can never delete a queue manually. AWS deletes it after 30 days of inactivity on queue
- ☒ B. It will delete the queue ✓

- ☐ C. It will initiate the delete but wait for four days before deleting until all messages are deleted automatically.
- ☐ D. It will ask user to delete the messages first

Explanation :

Answer – B

This is a straightforward answer is that the queue will be deleted. The below screenshot shows the option of deleting a queue called Demo in SQS

It will just give a prompt message but post this the queue will be deleted.

For more information on SQS please visit the link:

- <https://aws.amazon.com/sqs/> (<https://aws.amazon.com/sqs/>)

Ask our Experts



QUESTION 51 UNATTEMPTED

A user has launched a large EBS backed EC2 instance in the US-East-1a region. The user wants to achieve Disaster Recovery (DR. for that instance by creating another small instance in Europe. How can the user achieve DR?

- ☐ A. Copy the running instance using the “Instance Copy” command to the EU region
- ☐ B. Create an AMI of the instance and copy the AMI to the EU region. Then launch the instance from the EU AMI ✓
- ☐ C. Copy the instance from the US East region to the EU region
- ☐ D. Use the “Launch more like this” option to copy the instance from one region to another

Explanation :

Answer – B

If you need an AMI across multiple regions, then you have to copy the AMI across regions. Note that by default AMI's that you have created will not be available across all regions.

To copy AMI's , follow the below steps

Step 1) The first step is to create an AMI from your running instance by choosing on Image->Create Image.

Step 2) Once the Image has been created, go to the AMI section in the EC2 dashboard and click on the

Copy AMI option.

Step 3) In the next screen , you can specify where to copy the AMI to.

For the entire details to copy AMI's , please visit the link:

.

(<https://www.whizlabs.com/lms/%20http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.htmr>

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>

(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>)

Ask our Experts



QUESTION 52 UNATTEMPTED

A user has created numerous EBS snapshots. What is the general limit for each AWS account for the maximum number of EBS snapshots that can be created by default?

- ☒ A. 100000 ✓
- ☐ B. 5000
- ☐ C. 100
- ☐ D. 1000

Explanation :

Answer – A

As per the AWS documentation, by default, the EBS Snapshots limit is 100000.

Amazon Elastic Block Store (Amazon EBS) Limits

Resource	Default Limit
Number of EBS snapshots	100,000
Concurrent snapshots allowed for a single volume	5 for io1, gp2, magnetic; 1 for st1, sc1
Concurrent snapshot copy requests to a single destination region	5
Total volume storage of General Purpose SSD (gp2) volumes	300 TiB
Total volume storage of Provisioned IOPS SSD (io1) volumes	300 TiB
Total volume storage of Throughput Optimized HDD (st1)	300 TiB
Total volume storage of Cold HDD (sc1)	300 TiB
Total volume storage of Magnetic volumes (standard)	300 TiB
Total provisioned IOPS	300,000

For more information on AWS Service limits please visit the link:

- http://docs.aws.amazon.com/general/latest/gr/aws_service_limits.html
(http://docs.aws.amazon.com/general/latest/gr/aws_service_limits.html)

Ask our Experts



QUESTION 53 UNATTEMPTED

A user has setup Auto Scaling with ELB on the EC2 instances. The user wants to configure that whenever the CPU utilization is below 10%, Auto Scaling should remove one instance. How can the user configure this?

- ☐ A. The user can get an email using SNS when the CPU utilization is less than 10%. The user can use the desired capacity of Auto Scaling to remove the instance
- ☐ B. Use CloudWatch to monitor the data and Auto Scaling to remove the instances using scheduled actions
- ☐ C. Configure CloudWatch to send a notification to Auto Scaling Launch configuration when the CPU utilization is less than 10% and configure the Auto Scaling policy to remove the instance

- ☐ **D. Configure CloudWatch to send a notification to the Auto Scaling group when the CPU Utilization is less than 10% and configure the Auto Scaling policy to remove the instance** ✓

Explanation :

Answer – D

This is the basic feature of Autoscaling.

Auto Scaling helps you maintain application availability and allows you to scale your Amazon EC2 (<https://aws.amazon.com/ec2/>) capacity up or down automatically according to conditions you define. You can use Auto Scaling to help ensure that you are running your desired number of Amazon EC2 instances. Auto Scaling can also automatically increase the number of Amazon EC2 instances during demand spikes to maintain performance and decrease capacity during lulls to reduce costs. For more information on AutoScaling, please visit the link:

- <https://aws.amazon.com/autoscaling/> (<https://aws.amazon.com/autoscaling/>)

Ask our Experts



QUESTION 54 UNATTEMPTED

A user is trying to configure the CloudWatch billing alarm. Which of the below mentioned steps should be performed by the user for the first time alarm creation in the AWS Account Management section?

- ☐ **A. Enable Receiving Billing Reports**
- ☐ **B. Enable Receiving Billing Alerts** ✓
- ☐ **C. Enable AWS billing utility**
- ☐ **D. Enable CloudWatch Billing Threshold**

Explanation :

Answer B

The pre-requisite is to ensure to go to Preferences in AWS and ensure that "Receive Billing Alerts" is enabled. Only then will you be able to define Clodwatch alarms on billing.

For more information on Cloudwatch, please visit the link:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)



Ask our Experts

QUESTION 55 UNATTEMPTED

A user is checking the CloudWatch metrics from the AWS console. The user notices that the CloudWatch data is coming in UTC. The user wants to convert the data to a local time zone. How can the user perform this?

- ☐ A. In the CloudWatch dashboard the user should set the local timezone so that CloudWatch shows the data only in the local time zone
- ☐ B. In the CloudWatch console select the local timezone under the Time Range tab to view the data as per the local timezone ✓
- ☐ C. The CloudWatch data is always in UTC; the user has to manually convert the data
- ☐ D. The user should have send the local timezone while uploading the data so that CloudWatch will show the data only in the local timezone

Explanation :

Answer – B

In cloudwatch, when you go to any metric. You can click on the Custom option and change the timezone from UTC to Local Timezone , hence Option B is the right answer.

For more information on Cloudwatch, please visit the link:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)

Ask our Experts



QUESTION 56 UNATTEMPTED

You are creating an Auto Scaling group whose Instances need to insert a custom metric into CloudWatch. Which method would be the best way to authenticate your CloudWatch PUT request?

- ☐ A. Create an IAM role with the Put MetricData permission and modify the Auto Scaling launch configuration to launch instances in that role ✓

- ☐ B. Create an IAM user with the PutMetricData permission and modify the Auto Scaling launch configuration to inject the userscredentials into the instance User Data
- ☐ C. Modify the appropriate Cloud Watch metric policies to allow the Put MetricData permission to instances from the Auto Scaling group
- ☐ D. Create an IAM user with the PutMetricData permission and put the credentials in a private repository and have applications on the server pull the credentials as needed

Explanation :

Answer – A

When providing permission to any AWS service from an EC2 instance, you need to always use IAM roles. So here you would create an IAM Role with permissions to add Cloudwatch metrics. To specify a role, follow the below steps

Step 1) Go to your security dashboard and click on Roles. Click on Create New Role.

Step 2) Provide a name to the role and select the Role Type as Amazon EC2.

Step 3) Attach the relevant policy

Step 4) Once the role is created, when launching an EC2 instance, you can mention the role the EC2 instance needs to be launched with.

Since other options have no mention of IAM roles, all of them are invalid.

For more information on IAM Roles, please visit the below URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/iam-roles-for-amazon-ec2.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/iam-roles-for-amazon-ec2.html>)

Ask our Experts



QUESTION 57 UNATTEMPTED

A user has configured Elastic Load Balancing by enabling a Secure Socket Layer - SSL. negotiation Configuration known as a Security Policy. Which of the below mentioned options is not part of this secure policy while negotiating the SSL connection between the user and the client?

- ☐ A. SSL Protocols
- ☐ B. Client Order Preference ✓

- ☐ C. SSL Ciphers
- ☐ D. Server Order Preference

Explanation :

Answer – B

If you see the AWS documentation for all possible SSL options in the below link you will see that SSL Protocols, SSL Ciphers and Server Order Preference are all part of the pre-defined policies. Only Client Order Preference is not present.

For more information on Secure ELB, please visit the link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-security-policy-table.html>
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-security-policy-table.html>)

Ask our Experts



QUESTION 58 UNATTEMPTED

When assessing an organization's use of AWS API access credentials which of the following three credentials should be evaluated? Choose 3 answers from the options given below

- ☐ A. Key pairs ✓
- ☐ B. Console passwords
- ☐ C. Access keys ✓
- ☐ D. Signing certificates ✓
- ☐ E. Security Group memberships

Explanation :

Answer - A, C and D

You use different types of security credentials depending on how you interact with AWS. For example, you use a user name and password to sign in to the AWS Management Console. You use access keys to make programmatic calls to AWS API actions.

Key pairs consist of a public key and a private key. You use the private key to create a digital signature, and then AWS uses the corresponding public key to validate the signature. You can create Amazon EC2 key pairs from the Amazon EC2 console, CLI, or API.

Access keys consist of an access key ID and a secret access key. You use access keys to sign programmatic requests that you make to AWS if you use the AWS SDKs, REST, or Query APIs.

One can use the IAM API to upload a certificate, via the UploadServerCertificate request.
Option B and E are wrong because this is done via the Console and not via API's.
For more information on Security, please visit the below URL:

- <http://docs.aws.amazon.com/general/latest/gr/aws-sec-cred-types.html>
(<http://docs.aws.amazon.com/general/latest/gr/aws-sec-cred-types.html>)

Ask our Experts



QUESTION 59 UNATTEMPTED

What is true about IAM groups? Please choose 3 answers from the options given below

- ☐ A. We can add users to or remove them from a group. ✓
- ☐ B. A user can belong to multiple groups. ✓
- ☐ C. Groups can belong to other groups.
- ☐ D. Groups can be granted permissions using access control policies ✓

Explanation :

Answer - A, B and D

As per the AWS documentation, the following is true about AWS groups

For more information on IAM, please visit the below URL:

- <https://aws.amazon.com/iam/faqs/> (<https://aws.amazon.com/iam/faqs/>)

Ask our Experts



QUESTION 60 UNATTEMPTED

Which service allows one to issue temporary credentials in AWS? Choose one answer from the options below.

- ☐ A. AWS SQS
- ☒ B. AWS STS ✓
- ☐ C. AWS SES

- ☐ D. None of the above. You need to use a third party software to achieve this.

Explanation :

Answer – B

You can use the AWS Security Token Service (AWS STS) to create and provide trusted users with temporary security credentials that can control access to your AWS resources. Temporary security credentials work almost identically to the long-term access key credentials that your IAM users can use.

Option A is wrong because this is the queuing service provided by AWS.

Option C is wrong because this is the emailing service provided by AWS.

Option D is wrong because there is a service which exists from AWS.

For more information on STS, please visit the below URL:

- http://docs.aws.amazon.com/IAM/latest/UserGuide/id_credentials_temp.html
(http://docs.aws.amazon.com/IAM/latest/UserGuide/id_credentials_temp.html)

Ask our Experts



Finish Review (<https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests/quiz/12957>)

Certification

- Cloud Certification
(<https://www.whizlabs.com/cloud-certification-training-courses/>)
- Java Certification
(<https://www.whizlabs.com/oracle-java-certifications/>)
- PM Certification
(<https://www.whizlabs.com/project-management-certifications/>)
- Big Data Certification
(<https://www.whizlabs.com/big-data-certifications/>)

Company

- Support
(<https://help.whizlabs.com/hc/en-us>)
- Discussions (<http://ask.whizlabs.com/>)
- Blog (<https://www.whizlabs.com/blog/>)

Mobile App



Android Coming Soon



Follow us



(<https://www.facebook.com/whizlabs.software/>)



(<https://in.linkedin.com/company/whizlabs-software>)



(<https://twitter.com/whizlabs?lang=en>)



(<https://plus.google.com/+WhizlabsSoftware>)