

[🏠 \(https://www.whizlabs.com/learn\)](https://www.whizlabs.com/learn) > [My Courses \(https://www.whizlabs.com/learn/my-courses\)](https://www.whizlabs.com/learn/my-courses)> [AWS Certified SysOps Administrator Associate \(https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests#section-1\)](https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests#section-1)> [Practice Test I \(https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests/quiz/12743\)](https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests/quiz/12743) > **Report**

PRACTICE TEST I

Attempt 1

Marks Obtained 1 / 60

Your score is 1.67%

Completed on Tuesday , 29 January 2019 , 02:09 PM

Time Taken 00 H 00 M 14 S

Result Fail

Domains / Topics wise Quiz Performance Report

S.No.	Topic	Total Questions	Correct	Incorrect	Unattempted
1	Other	60	1	0	59

60 Questions	1 Correct	0 Incorrect	59 Unattempted
------------------------	---------------------	-----------------------	--------------------------

Show Answers

All	▼
-----	---

QUESTION 1 CORRECT

A user has configured CloudWatch monitoring on an EBS backed EC2 instance. Which of the below mentioned metrics will always show a 0 value?



- ☒ A. DiskReadBytes ✓
- ☐ B. NetworkIn
- ☐ C. NetworkOut
- ☐ D. CPU Utilization

Explanation :

Answer – A

DiskReadBytes is the number of bytes read from an instance store, and this number is zero when there are no instance store disks.

- <https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/ec2-metricscollected.html>
(<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/ec2-metricscollected.html>)

The below snapshot shows the Cloudwatch metric's for simple EC2 Instance. By default there would always be some NetworkIn , NetworkOut , CPU Utilization. But the DiskRead(Bytes) would remain 0.



For more information on cloudwatch metrics , please visit the URL:

- http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CW_Support_For_AWS.html
(http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CW_Support_For_AWS.html)

Ask our Experts



QUESTION 2 UNATTEMPTED

A user is trying to setup a recurring Auto Scaling process. The user has setup one process to scale up every day at 8 am and scale down at 7 PM. The user is trying to setup another recurring process which scales up on the 1st of every month at 8 AM and scales down the same day at 7 PM. What will Auto Scaling do in this scenario?

- ☐ A. Auto Scaling will execute both processes but will add just one instance on the 1st
- ☐ B. Auto Scaling will add two instances on the 1st of the month
- ☐ C. Auto Scaling will schedule both the processes but execute only one process randomly
- ☐ D. Auto Scaling will throw an error since there is a conflict in the schedule of two separate Auto Scaling Processes ✓

Explanation :

Answer – D



Assume that you have an alarm which is scheduled to start scaling at 8 AM and scale down at 19:00.

Now if we try to create a second scheduling action with the same start time you will get an error as shown below. Since there is already a scheduling process, this will result in an error.

For more information on Scheduled scaling group, please visit the URL:

- http://docs.aws.amazon.com/autoscaling/latest/userguide/schedule_time.html
(http://docs.aws.amazon.com/autoscaling/latest/userguide/schedule_time.html)

Ask our Experts



QUESTION 3 UNATTEMPTED

Your client is planning to setup infrastructure on AWS for the Christmas sales. The client is planning to use Auto Scaling based on the schedule for proactive scaling. What specific suggestion would you offer to the client?

- ☐ A. It is not possible to schedule the auto-scaling group in advance.
- ☐ B. The scaling only can be setup 7 days in advance.
- ☐ C. In AWS Management Console you can schedule autoscaling group to any future date. But if you use CLI you can schedule autoscaling group only up to 30 days in advance. ✓
- ☐ D. It is not advisable to use scheduled scaling.

Explanation :

Answer – C

As per AWS, When we use CLI for setting up scheduling for your Auto Scaling group it will allow you to schedule it one month in advance. Using console you can schedule it for a future date with out any such limitations.

StartTime -> (timestamp)



The date and time that the action is scheduled to begin. This date and time can be up to one month in the future.

Please refer the following link for more information.

- <https://docs.aws.amazon.com/cli/latest/reference/autoscaling/describe-scheduled-actions.html> (<https://docs.aws.amazon.com/cli/latest/reference/autoscaling/describe-scheduled-actions.html>)

Auto Scaling based on a schedule allows the user to scale the application in response to predictable load changes. The user can specify any date in the future to scale up or down during that period. As per Auto Scaling the user can schedule an action for up to a month in the future. Thus, it is recommended to wait until end of November before scheduling for Christmas. For more information on Auto scaling group, please visit the URL:

- <https://aws.amazon.com/autoscaling/> (<https://aws.amazon.com/autoscaling/>)

Ask our Experts



QUESTION 4 UNATTEMPTED

A user is trying to understand the ACL and policy for an S3 bucket. Which of the below mentioned policy permissions is equivalent to the WRITE ACL on a bucket?

- ☐ A. s3:GetObjectAcl
- ☐ B. s3:GetObjectVersion
- ☐ C. s3:ListBucketVersions
- ☒ D. s3:DeleteObject ✓

Explanation :

Answer – D

As per the AWS documentation, the WRITE permission corresponds to either the PutObject or



DeleteObject. Hence D is the only right option.

For more information on S3 ACL, please visit the URL:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/acl-overview.html>
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/acl-overview.html>)

Ask our Experts



QUESTION 5 UNATTEMPTED

A user has created an ELB with Auto Scaling. Which of the below mentioned offerings from ELB helps the user to stop sending new requests traffic from the load balancer to the EC2 instance when the instance is being deregistered while continuing in-flight requests?

- ☐ A. ELB sticky session
- ☐ B. ELB deregistration check
- ☒ C. ELB connection draining ✓
- ☐ D. ELB auto registration Off

Explanation :

Answer – C

To ensure that a Classic Load Balancer stops sending requests to instances that are de-registering or unhealthy, while keeping the existing connections open, use connection draining. This enables the load balancer to complete in-flight requests made to instances that are de-registering or unhealthy.

You will see the connection draining option in the Instances Tab of your load balancer.

For more information on ELB connection draining, please visit the URL:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/config-conn-drain.html>
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/config-conn-drain.html>)



Ask our Experts



QUESTION 6 UNATTEMPTED

A user has launched an EC2 instance from an instance store backed AMI. The infrastructure team wants to create an AMI from the running instance. Which of the below mentioned steps will not be performed while creating the AMI?

- ☐ A. Define the AMI launch permissions ✓
- ☐ B. Upload the bundled volume
- ☐ C. Register the AMI
- ☐ D. Bundle the volume

Explanation :

Answer – A

Below is the snapshot from the AWS documentation on the Instance store lifecycle of an AMI.

From here it is evident that there is no stage of 'Define the AMI launch permissions'

For more information on Instance store EC2 volumes, please visit the URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/creating-an-ami-instance-store.html> (<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/creating-an-ami-instance-store.html>)

Ask our Experts



QUESTION 7 UNATTEMPTED



You are managing the AWS account of a big organization. The organization has more than 1000+ employees and they want to provide access to the various services to most of the employees. Which of the below mentioned options is the best possible solution in this case?

- ☐ A. The user should create a separate IAM user for each employee and provide access to them as per the policy
- ☐ B. The user should create an IAM role and attach STS with the role. The user should attach that role to the EC2 instance and setup AWS authentication on that server
- ☐ C. The user should create IAM groups as per the organization's departments and add each user to the group for better access control
- ☐ D. Attach an IAM role with the organization's authentication service to authorize each user for various AWS services ✓

Explanation :

Answer – D

The best practise for IAM is to create roles which has specific access to an AWS service and then give the user permission to the AWS service via the role.

If its a combination of various services, there will be large number of groups and might not be maintainable.

Role for service and ability for user to take up the role and ability to change roles would be much cleaner.

Below is an example to create a role which has access to the ELB service.

Step 1) Create a role which has the required ELB access

Step 2) You need to provide permissions to the underlying EC2 instances in the Elastic Load Balancer

For the best practises on IAM policies, please visit the link:

- <http://docs.aws.amazon.com/IAM/latest/UserGuide/best-practices.html>
(<http://docs.aws.amazon.com/IAM/latest/UserGuide/best-practices.html>)



Ask our Experts



QUESTION 8 UNATTEMPTED

A user is launching an instance, he is prompted to add 'Tags' to the instance. Which of the following option is not true about the 'tags'?

- ☐ A. Each tag will have a key and value
- ☐ B. The user can apply tags to the EBS Volume
- ☒ C. The maximum value of the tag key length is 64 unicode characters ✓
- ☐ D. AWS tags are used to find the cost distribution of various resources

Explanation :

Answer – C

Tags enable you to categorize your AWS resources in different ways, for example, by purpose, owner, or environment. As per the AWS documentation below is the restrictions on tags. From the below text it can be seen that the maximum key length is 128 and not 64.

- Maximum number of tags per resource—50
- Maximum key length—128 Unicode characters in UTF-8
- Maximum value length—256 Unicode characters in UTF-8
- Tag keys and values are case sensitive.
- Do not use the AWS: prefix in your tag names or values because it is reserved for AWS use. You can't edit or delete tag names or values with this prefix. Tags with this prefix do not count against your tags per resource limit.
- If your tagging schema will be used across multiple services and resources, remember that other services may have restrictions on allowed characters. Generally allowed characters are: letters, spaces, and numbers representable in UTF-8, plus the following special characters: + - = . _ : / @.

For more information on tagging resources, please visit the link:

- http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Using_Tags.html
(http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Using_Tags.html)



Ask our Experts



QUESTION 9 UNATTEMPTED

A user has created a VPC with CIDR 10.0.0.0/24. The user has created a public subnet with CIDR 10.0.0.0/25 and a private subnet with CIDR 10.0.0.128/25. The user has launched one instance each in the private and public subnets. Which of the below mentioned options cannot be the correct IP address (private IP. assigned to an instance in the public or private subnet)?

- ☒ A. 10.0.0.255 ✓
- ☐ B. 10.0.0.132
- ☐ C. 10.0.0.122
- ☐ D. 10.0.0.55

Explanation :

Answer – A

As per the AWS documentation there is a reservation of IP addresses. Hence option A is right because this IP address will be reserved by AWS.

The first four IP addresses and the last IP address in each subnet CIDR block are not available for you to use, and cannot be assigned to an instance. For example, in a subnet with CIDR block 10.0.0.0/24, the following five IP addresses are reserved:

- 10.0.0.0: Network address.
- 10.0.0.1: Reserved by AWS for the VPC router.
- 10.0.0.2: Reserved by AWS. The IP address of the DNS server is always the base of the VPC network range plus two; however, we also reserve the base of each subnet range plus two. 10.0.0.3: Reserved by AWS for future use.
- 10.0.0.255: Network broadcast address. We do not support broadcast in a VPC, therefore we reserve this address.
- For more information on IP Reservation, please visit the link:
- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html)



Ask our Experts



QUESTION 10 UNATTEMPTED

A user has launched an EBS backed EC2 instance. The user has rebooted the instance. Which of the below mentioned statements is not true with respect to the reboot action?

- ☐ A. The private and public address remains the same
- ☐ B. The Elastic IP remains associated with the instance
- ☐ C. The volume is preserved
- ☐ D. The instance runs on a new host computer ✓

Explanation :

Answer – D

When you reboot your instance, as per the AWS documentation the following actions occur
An instance reboot is equivalent to an operating system reboot. In most cases, it takes only a few minutes to reboot your instance. When you reboot an instance, it remains on the same physical host, so your instance keeps its public DNS name (IPv4), private IPv4 address, IPv6 address (if applicable), and any data on its instance store volumes.

From the above explanation it is pretty straightforward to understand that Option D is the right option.

For more information on rebooting instance, please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-reboot.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-reboot.html>)

Ask our Experts



A user is using a small MySQL RDS DB. The user is experiencing high latency due to the Multi AZ feature. Which of the below mentioned options may not help the user in this situation?

- ☐ A. Schedule the automated back up in non-working hours
- ☐ B. Use a large or higher size instance
- ☐ C. Use Provisioned IOPS
- ☐ D. Take a snapshot from standby Replica ✓

Explanation :

Answer – D

AWS Docs says :

Q: How do DB Snapshots and automated backups work with my Multi-AZ deployment?

You interact with automated backup and DB Snapshot functionality in the same way whether you are running a standard deployment in a Single-AZ or Multi-AZ deployment. If you are running a Multi-AZ deployment, automated backups and DB Snapshots are simply taken from the standby to avoid I/O suspension on the primary. ***Please note that you may experience increased I/O latency (typically lasting a few minutes) during backups for both Single-AZ and Multi-AZ deployments.***

<https://aws.amazon.com/rds/faqs/#multi-az-deployments>

(<https://aws.amazon.com/rds/faqs/#multi-az-deployments>)

In a Multi-AZ deployment, Amazon RDS automatically provisions and maintains a synchronous standby replica in a different Availability Zone. The primary DB instance is synchronously replicated across Availability Zones to a standby replica to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups

As per AWS the below are the best practices for multiAZ

For production workloads, we recommend you use Provisioned IOPS and DB instance classes (m1.large and larger) that are optimized for Provisioned IOPS for fast, consistent performance. Hence option B and C are valid.

I/O activity is not suspended on your primary during backup for Multi-AZ deployments for the MySQL, MariaDB, Oracle, and PostgreSQL engines, because the backup is taken from the standby. However, note that you may still experience elevated latencies for a few minutes during backups for Multi-AZ deployments. This can be a problem for some applications. Hence it is better to schedule backups during out of office hours.

For more information on MultiAZ RDS, please visit the link:



- <http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZ.html>
(<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZ.html>)

Ask our Experts



QUESTION 12 UNATTEMPTED

A user has created a queue named “myqueue” in US-East region with AWS SQS. The user’s AWS account ID is 123456123456. If the user wants to perform some action on this queue, which of the below Queue URL should he use?

- ☐ A. <https://sqs.us-east-1.amazonaws.com/123456123456/myqueue> ✓
- ☐ B. <http://sqs.amazonaws.com/123456123456/myqueue>
- ☐ C. <http://sqs.123456123456.us-east-1.amazonaws.com/myqueue>
- ☐ D. <http://123456123456.sqs.us-east-1.amazonaws.com/myqueue>

Explanation :

Answer – A

Below is an example of a queue and the URL endpoint. The SQS endpoint normally has the format of URL:

- <https://sqs.RegionName.amazonaws.com/AWS-account-number/QueueName>

By this reference the only right answer is A.

For more information on SQS , please visit the URL:

- <https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-general-identifiers.html>
(<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-general-identifiers.html>)
- <https://aws.amazon.com/sqs/> (<https://aws.amazon.com/sqs/>)

Ask our Experts



**QUESTION 13****UNATTEMPTED**

A user is displaying the CPU utilization, and Network in and Network out CloudWatch metrics data of a single instance on the same graph. The graph uses one Y-axis for CPU utilization and Network in and another Y-axis for Network out. Since Network in is too high, the CPU utilization data is not visible clearly on graph to the user. How can the data be viewed better on the same graph?

- ☐ A. It is not possible to show multiple metrics with the different units on the same graph
- ☐ B. Add a third Y-axis with the console to show all the data in proportion
- ☐ C. On the Graphed metrics tab of the Metric window of the Cloudwatch select the Y axis that plots CPU Utilization data. ✓
- ☐ D. Change the units of CPU utilization so it can be shown in proportion with Network

Explanation :

Answer – C

The steps below describes how to switch from the Left Y Axis to Right Y axis.

You can switch between two different Y axes for your graph. ie. a Left Y Axis and a right Y axis for your graph.

To modify the Y axis on a graph

1. Open the CloudWatch console at <https://console.aws.amazon.com/cloudwatch/> (<https://console.aws.amazon.com/cloudwatch/>).
2. In the navigation pane, choose **Metrics**.
3. Select a metric namespace (for example, EC2) and then a metric dimension (for example, Per-Instance Metrics).
4. The **All metrics** tab displays all metrics for that dimension in that namespace. To graph a metric, select the check box next to the metric.
5. On the **Graph options** tab, specify the **Min** and **Max** values for **Left Y Axis**. The value of **Min** cannot be greater than the value of **Max**



6. To create a second Y axis, specify the **Min** and **Max** values for **Right Y Axis**.
7. To switch between the two Y axes, choose the **Graphed metrics** tab. For **Y Axis**, choose **Left Y Axis** or **Right Y Axis**.

Ask our Experts



QUESTION 14 UNATTEMPTED

A user is planning to use AWS services for his web application. If the user is trying to set up his own billing management system for AWS, how can he configure it?



- ☐ A. Set up programmatic billing access. Download and parse the bill as per the requirement ✓
- ☐ B. It is not possible for the user to create his own billing management service with AWS
- ☐ C. Enable the AWS CloudWatch alarm which will provide APIs to download the alarm data
- ☐ D. Use AWS billing APIs to download the usage report of each service from the AWS billing console

Explanation :

Answer – A

Billing reports provide information about your usage of AWS resources and estimated costs for that usage. You can have AWS generate billing reports that break down your estimated costs in different ways:

- By the hour, day, or month
- By each account in your organization
- By product or product resource
- By tags that you define yourself

For more information on programmatic billing, please visit the link:

- <http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/billing-reports.html>
(<http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/billing-reports.html>)

Ask our Experts



QUESTION 15 UNATTEMPTED

A user has created a public subnet with VPC and launched an EC2 instance within it. The user is trying to delete the subnet. What will happen in this scenario?



- ☐ A. It will delete the subnet and make the EC2 instance as a part of the default subnet
- ☒ B. It will not allow the user to delete the subnet until the instances are terminated ✓
- ☐ C. It will delete the subnet as well as terminate the instances
- ☐ D. The subnet can never be deleted independently, but the user has to delete the VPC first

Explanation :

Answer – B

In AWS, when you try to delete a subnet which has instances it will not allow to delete it. The below error message will be shown when you try to delete a subnet with instances.

For more information on VPC and subnets please visit the link:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html)

Ask our Experts



QUESTION 16 **UNATTEMPTED**

A user has setup an EBS backed instance and attached 2 EBS volumes to it. The user has setup a CloudWatch alarm on each volume for the disk data. The user has stopped the EC2 instance and detached the EBS volumes. What will be the status of the alarms on the EBS volume?

- ☐ A. OK
- ☒ B. Insufficient Data ✓
- ☐ C. Alarm
- ☐ D. The EBS cannot be detached until all the alarms are removed



Explanation :

Answer – B

An alarm has three possible states:

- OK—The metric is within the defined threshold
- ALARM—The metric is outside of the defined threshold
- INSUFFICIENT_DATA—The alarm has just started, the metric is not available, or not enough data is available for the metric to determine the alarm state

Since the EC2 instance was stopped and EBS volumes are no longer available, the alarms will not have enough information so the state of the alarm will be Insufficient Data.

For more information on Cloudwatch alarms, please visit the link:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/AlarmThatSendsEmail.html> (http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/AlarmThatSendsEmail.html)

Ask our Experts



QUESTION 17 UNATTEMPTED

A user has launched an EC2 instance from an instance store backed AMI. The infrastructure team wants to create an AMI from the running instance. Which of the below mentioned credentials is not required while creating the AMI?

- ☐ A. AWS account ID
- ☐ B. X.509 certificate and private key
- ☐ C. AWS login ID to login to the console ✓
- ☐ D. Access key and secret access key

Explanation :

Answer – C

In the AWS documentation it clearly mentions what needs to be in place when creating the AMI's. As per the below text, the option C is not required.



Note: As the process is done using CLI the login ID is not needed.

For more information on creating instance store AMI's, please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/creating-an-ami-instance-store.html> (<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/creating-an-ami-instance-store.html>)

Ask our Experts



QUESTION 18 UNATTEMPTED

A user has configured a SSL listener at ELB as well as on the back-end instances. Which of the below mentioned statement is correct with regard to ELB and SSL integration?

- ☐ A. It is not possible to have the SSL listener both at ELB and back-end instances
- ☐ B. ELB will modify headers to add requestor details
- ☐ C. ELB will intercept the request to add the cookie details if sticky session is enabled
- ☒ D. ELB can not modify the headers ✓

Explanation :

Answer – D

As per the AWS documentation, please find the below excerpts

Option A is invalid because of the below documentation.

Option B is invalid because of the below documentation.

Option C is invalid because of the below documentation.

For more information on ELB, please visit the link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-listener-config.html> (<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-listener-config.html>) ^

Note:

When you use TCP (layer 4) for both front-end and back-end connections, your load balancer forwards the request to the back-end instances without modifying the headers. After your load balancer receives the request, it attempts to open a TCP connection to the back-end instance on the port specified in the listener configuration.

Hence, headers will not be modified until you use Proxy Protocol.

You can enable Proxy Protocol, which adds a header with the connection information of the client, such as the source IP address, destination IP address, and port numbers. The header is then sent to the back-end instance as a part of the request.

Ask our Experts

**QUESTION 19 UNATTEMPTED**

A user has created a Cloudformation stack. The stack creates AWS services, such as EC2 instances, ELB, AutoScaling, and RDS. While creating the stack it created EC2, ELB and AutoScaling but failed to create RDS. What will Cloudformation do in this scenario?

- ☐ A. Cloudformation can never throw an error after launching a few services since it verifies all the steps before launching
- ☐ B. It will warn the user about the error and ask the user to manually create RDS
- ☐ C. Rollback all the changes and terminate all the created services ✓
- ☐ D. It will wait for the user's input about the error and correct the mistake after the input

Explanation :

Answer – C

As per the AWS documentation it clearly states that the cloudformation stack will be rolled



back.

For more information on Cloudformation, please visit the link:

- <https://aws.amazon.com/cloudformation/faqs/>
(<https://aws.amazon.com/cloudformation/faqs/>)

Ask our Experts



QUESTION 20 UNATTEMPTED

A user has created a VPC with public and private subnets using the VPC wizard with NAT instance. The user has not launched any instance manually and is trying to delete the VPC. What will happen in this scenario?

- ☐ A. It will not allow to delete the VPC as it has subnets with route tables
- ☐ B. It will not allow to delete the VPC since it has a running route instance
- ☐ C. It will terminate the VPC along with all the instances launched by the wizard
- ☐ D. It will not allow to delete the VPC since it has a running NAT instance ✓

Explanation :

Answer – D

Since the VPC will contain a NAT instance because of the private/public subnet combination , when you try to delete the VPC you will get the below error message.

For more information on VPC and subnets please visit the link:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html)

Ask our Experts



An organization is measuring the latency of an application every minute and storing data inside a file in the JSON format. The organization wants to send all latency data to AWS CloudWatch. How can the organization achieve this?

- ☐ A. The user has to parse the file before uploading data to CloudWatch
- ☐ B. It is not possible to upload the custom data to CloudWatch
- ☐ C. The user can supply the file as an input to the CloudWatch command ✓
- ☐ D. The user can use the CloudWatch Import command to import data from the file to CloudWatch

Explanation :

Answer – C

You can publish your own metrics to CloudWatch using the AWS CLI or an API.

2 examples of how metrics can be published is given below

Reference link:

- <https://docs.aws.amazon.com/cli/latest/reference/cloudwatch/put-metric-data.html>
(<https://docs.aws.amazon.com/cli/latest/reference/cloudwatch/put-metric-data.html>)

```
AWS cloudwatch put-metric-data --metric-name PageViewCount --namespace MyService --value 2 --timestamp 2016-10-14T12:00:00.000Z
```

```
AWS cloudwatch get-metric-statistics--namespace MyService --metric-name PageViewCount \
```

```
--statistics "Sum" "Maximum" "Minimum" "Average" "SampleCount" \
```

```
--start-time 2016-10-20T12:00:00.000Z --end-time 2016-10-20T12:05:00.000Z --period 60
```

For more information on publishing metrics please visit the link:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/publishingMetrics.html>
(<http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/publishingMetrics.html>)

Ask our Experts



QUESTION 22

UNATTEMPTED

A user has launched an EBS backed instance with EC2-Classical. The user stops and starts the instance. Which of the below mentioned statements is not true with respect to the stop/start action?

- ☐ A. The instance gets new private and public IP addresses
- ☐ B. The volume is preserved
- ☐ C. The Elastic IP remains associated with the instance ✓
- ☐ D. The instance may run on a new host computer

Explanation :

Answer – C

As per the AWS documentation, below is the difference stated when working with EC2 classic. Hence C is the right option.

For more information on the differences please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-vpc.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-vpc.html>)

Ask our Experts



QUESTION 23

UNATTEMPTED

A system admin is trying to understand the Auto Scaling activities. Which of the below mentioned processes is not performed by Auto Scaling?

- ☐ A. Reboot Instance ✓
- ☐ B. Schedule Actions



- ☐ C. Replace Unhealthy
- ☐ D. Availability Zone Balancing

Explanation :

Answer – A

The below snapshot shows the options of an Autoscaling group. In the Suspended processes you can see the different processes for Autoscaling. Here you can see that there are processes for Schedule Actions, Replace unhealthy and AZ Rebalance. But no Reboot Instance, Hence Option A is the right answer.

For more information on Autoscaling , please visit the URL:

- <https://aws.amazon.com/autoscaling/> (<https://aws.amazon.com/autoscaling/>)

Ask our Experts



QUESTION 24 UNATTEMPTED

A user has launched an RDS postgresSQL DB with AWS. The user did not specify the maintenance window during creation. The user has configured RDS to update the DB instance type from micro to large. If the user wants to have it during the maintenance window, what will AWS do?

- ☐ A. AWS will not allow to update the DB until the maintenance window is configured
- ☐ B. AWS will select the default maintenance window if the user has not provided it ✓
- ☐ C. AWS will ask the user to specify the maintenance window during the update
- ☐ D. It is not possible to change the DB size from micro to large with RDS

Explanation :

Answer – B



As per the AWS documentation it clearly mentions that if you don't specify a maintenance window the default one will be selected.

For more information on aws rds please visit the link:

- <https://aws.amazon.com/rds/faqs/> (<https://aws.amazon.com/rds/faqs/>)

Ask our Experts



QUESTION 25 UNATTEMPTED

A user has launched multiple EC2 instances for the purpose of development and testing in the same region. The user wants to find the separate cost for the production and development instances. How can the user find the cost distribution?

- ☐ A. The user should download the activity report of the EC2 services as it has the instance ID wise data
- ☐ B. It is not possible to get the AWS cost usage data of single region instances separately
- ☐ C. The user should use Cost Distribution Metadata and AWS detailed billing
- ☐ D. The user should use Cost Allocation Tags and AWS billing reports ✓

Explanation :

Answer – D

A tag is a label that you or AWS assigns to an AWS resource. Each tag consists of a *key* and a *value*. A key can have more than one value. You can use tags to organize your resources, and cost allocation tags to track your AWS costs on a detailed level. After you activate cost allocation tags, AWS uses the cost allocation tags to organize your resource costs on your cost allocation report, to make it easier for you to categorize and track your AWS costs

For more information on cost allocation tags please visit the link:

- <http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/cost-alloc-tags.html> (<http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/cost-alloc-tags.html>)



Ask our Experts



QUESTION 26 UNATTEMPTED

A user has created a VPC with the public subnet. The user has created a security group for that VPC. Which of the below mentioned statements is true when a security group is created?

- ☐ A. It can connect to the AWS services, such as S3 and RDS by default
- ☐ B. It allows all the inbound traffic by default
- ☐ C. It allows all the outbound traffic by default ✓
- ☐ D. It will by default allow traffic to the internet gateway

Explanation :

Answer – C

When you create a security group , by default all Outbound rules will be allowed.

For more information refer to AWS Documentation @ URL:

- <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-network-security.html#default-security-group>
(<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-network-security.html#default-security-group>)

Allows all inbound traffic from other instances associated with the default security group (the security group specifies itself as a source security group in its inbound rules)

Allows all outbound traffic from the instance.

The correct Answer is C: It will have all the outbound traffic by default.

Ask our Experts



QUESTION 27

UNATTEMPTED

A user has setup an Auto Scaling group. The group has failed to launch a single instance for more than 24 hours. What will happen to Auto Scaling in this condition?

- ☐ A. Auto Scaling will keep trying to launch the instance for 72 hours
- ☐ B. Auto Scaling will suspend the scaling process ✓
- ☐ C. Auto Scaling will start an instance in a separate region
- ☐ D. The Auto Scaling group will be terminated automatically

Explanation :

Answer – B

As per the AWS documentation it clearly mentions that the scaling process will be suspended if no instances launch in 24 hours.

For more information on autoscaling suspension please visit the link:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/as-suspend-resume-processes.html> (<http://docs.aws.amazon.com/autoscaling/latest/userguide/as-suspend-resume-processes.html>)

Ask our Experts



QUESTION 28

UNATTEMPTED

A user is planning to set up the Multi AZ feature of RDS. Which of the below mentioned conditions won't take advantage of the Multi AZ feature?

- ☐ A. Availability zone outage
- ☐ B. A manual failover of the DB instance using Reboot with failover option
- ☐ C. Region outage ✓



☐ D. When the user changes the DB instance's server type

Explanation :

Answer – C

As per the AWS documentation all the failover conditions are given and Region Outage will not make use of multi AZ.

For more information on multiAZ RDS please visit the link:

- <https://aws.amazon.com/rds/details/multi-az/>
(<https://aws.amazon.com/rds/details/multi-az/>)

Ask our Experts



QUESTION 29 UNATTEMPTED

An organization has configured Auto Scaling with ELB. One of the instance health check returns the status as Impaired to Auto Scaling. What will Auto Scaling do in this scenario?

- ☐ A. Perform a health check until cool down before declaring that the instance has failed
- ☐ B. Terminate the instance and launch a new instance ✓
- ☐ C. Notify the user using SNS for the failed state
- ☐ D. Notify ELB to stop sending traffic to the impaired instance

Explanation :

Answer – B

Below is the diagram for Autoscaling. It clearly shows that if an instance has failed a health check it will be terminated.

For more information on Autoscaling lifecycle please visit the link:



- <http://docs.aws.amazon.com/autoscaling/latest/userguide/AutoScalingGroupLifecycle.html> (<http://docs.aws.amazon.com/autoscaling/latest/userguide/AutoScalingGroupLifecycle.html>)

For more information on Auto Scaling Maintain instance level, please visit the link:

- <https://docs.aws.amazon.com/autoscaling/latest/userguide/as-maintain-instance-levels.html> (<https://docs.aws.amazon.com/autoscaling/latest/userguide/as-maintain-instance-levels.html>)

By default, the Auto Scaling group determines the health state of each instance by periodically checking the results of EC2 instance status checks. If the instance status is any state other than running or if the system status is impaired, Auto Scaling considers the instance to be unhealthy and launches a replacement.

Ask our Experts



QUESTION 30 UNATTEMPTED

A user is using Cloudformation to launch an EC2 instance and then configure an application after the instance is launched. The user wants the stack creation of ELB and AutoScaling to wait until the EC2 instance is launched and configured properly. How can the user configure this?

- ☐ A. It is not possible that the stack creation will wait until one service is created and launched
- ☐ B. The user can use the HoldCondition resource to wait for the creation of the other dependent resources
- ☐ C. The user can use the DependentCondition resource to hold the creation of the other dependent resources
- ☐ D. The user can use the WaitCondition resource to hold the creation of the other dependent resources ✓

Explanation :

Answer – D

You can use a wait condition for situations like the following:



- To coordinate stack resource creation with configuration actions that are external to the stack creation
- To track the status of a configuration process

Note: There is no condition in Cloudformation called "**Dependent Condition**".

For more information on Cloudformation Wait condition please visit the link:

- <http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/aws-properties-waitcondition.html>
(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/aws-properties-waitcondition.html>)

Ask our Experts



QUESTION 31 UNATTEMPTED

An organization has configured two single availability zones. The Auto Scaling groups are configured in separate zones. The user wants to merge the groups such that one group spans across multiple zones. How can the user configure this?

- ☐ A. Run the command `as-join-auto-scaling-group` to join the two groups
- ☐ B. Run the command `update-auto-scaling-group` to configure one group to span across zones and delete the other group ✓
- ☐ C. Run the command `as-copy-auto-scaling-group` to join the two groups
- ☐ D. Run the command `as-merge-auto-scaling-group` to merge the groups

Explanation :

Answer – B

The CLI command `update-auto-scaling-group` can be used to combine 2 groups. An example of the command is given below.

`AWS autoscaling update-auto-scaling-group --auto-scaling-group-name my-auto-scaling-group --health-check-type ELB --health-check-grace-period 60`

For more information on the CLI command please visit the link:



- <http://docs.aws.amazon.com/cli/latest/reference/autoscaling/update-auto-scaling-group.html> (<http://docs.aws.amazon.com/cli/latest/reference/autoscaling/update-auto-scaling-group.html>)

Ask our Experts



QUESTION 32 UNATTEMPTED

An AWS account wants to be part of the consolidated billing of his organization's payer account. How can the owner of that account achieve this?

- ☐ A. The payer account has to request AWS support to link the other accounts with his account
- ☐ B. The owner of the linked account should add the payer account to his master account list from the billing console
- ☐ C. The payer account will send a request to the linked account to be a part of consolidated billing ✓
- ☐ D. The owner of the linked account requests the payer account to add his account to consolidated billing

Explanation :

Answer – C

You can use the Consolidated Billing feature to consolidate payment for multiple Amazon Web Services (AWS) accounts or multiple Amazon International Services Pvt. Ltd (AISPL) accounts within your organization by designating one of them to be the payer account. With Consolidated Billing, you can see a combined view of AWS charges incurred by all accounts, as well as get a cost report for each individual account associated with your payer account. In the billing section you need to Send a request for the Linked accounts.

For more information on consolidated billing please visit the link:

- <http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/consolidated-billing.html> (<http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/consolidated-billing.html>)



Ask our Experts



QUESTION 33 UNATTEMPTED

A sysadmin has created the below mentioned policy on an S3 bucket named testdemo321. What does this policy define?

```
"Statement": [{  
  "Sid": "Stmt1388811069831",  
  "Effect": "Allow",  
  "Principal": { "AWS": "*" },  
  "Action": [ "s3:GetObjectAcl", "s3:ListBucket" ],  
  "Resource": [ "arn:aws:s3:::testdemo321" ]  
}]
```

- ☐ A. It will make the demo bucket as well as all its objects as public
- ☐ B. It will allow everyone to view the ACL of the bucket
- ☐ C. It will give an error as no object is defined as part of the policy while the action defines the rule about the object ✓
- ☐ D. It will make the demo bucket as public

Explanation :

Answer – C

As per the aws documentation , the List bucket and GetobjectACL will give read access to users.

For more information, please visit the link:

- <https://docs.aws.amazon.com/AmazonS3/latest/dev/s3-arn-format.html>
(<https://docs.aws.amazon.com/AmazonS3/latest/dev/s3-arn-format.html>)

JSON policy documents are made up of elements.



- Version - The Version policy element is used within a policy and defines the version of the policy language.
- ID - The Id element specifies an optional identifier for the policy
- Statement - The Statement element is the main element for a policy. This element is required. It can include multiple elements
- Sid - You can assign a Sid value to each statement in a statement array. It is an optional identifier.
- Action - The Action element describes the specific action or actions that will be allowed or denied. Statements must include either an Action or NotAction element. Each AWS service has its own set of actions that describe tasks that you can perform with that service.
- Effect - The Effect element is required and specifies whether the statement results in an allow or an explicit deny. Valid values for Effect are Allow and Deny.
- Resource - The Resource element specifies the object or objects that the statement covers. Statements must include either a Resource or a NotResource element. You specify a resource using an ARN.
- Principal- Use the Principal element to specify the user (IAM user, federated user, or assumed-role user), AWS account, AWS service, or other principal entity that is allowed or denied access to a resource.

Note:

If you see the example provided in the question, "Resource" just has only the bucket name. No object was specified.

It should be like mentioned below:

arn:aws:s3::bucket_name/key_name (or)

arn:aws:s3::bucket_name/*

("*" means all the objects inside the bucket)"

<https://docs.aws.amazon.com/AmazonS3/latest/dev/s3-arn-format.html>

(<https://docs.aws.amazon.com/AmazonS3/latest/dev/s3-arn-format.html>)

Ask our Experts



An organization has setup consolidated billing with 3 different AWS accounts. Which of the below mentioned advantages will organization receive in terms of the AWS pricing?

- ☐ A. The consolidated billing does not bring any cost advantage for the organization
- ☐ B. All AWS accounts will be charged for S3 storage by combining the total storage of each account ✓
- ☐ C. The EC2 instances of each account will receive a total of 750*3 micro instance hours free
- ☐ D. The free usage tier for all the 3 accounts will be 3 years and not a single year

Explanation :

Answer – B

As per the AWS documentation for billing purposes, AWS treats all the accounts on the consolidated bill as if they were one account. Some services, such as Amazon EC2 and Amazon S3, have volume pricing tiers across certain usage dimensions that give you lower prices when you use the service more. With consolidated billing, AWS combines the usage from all accounts to determine which volume pricing tiers to apply, giving you a lower overall price whenever possible.

Hence based on this fact, Option A ,C and D are the wrong options.

For more information on Consolidated billing , please visit the URL:

- <http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/consolidated-billing.html> (<http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/consolidated-billing.html>)

Ask our Experts



A user is trying to connect to a running EC2 instance using SSH. However, the user gets a Permission denied (publickey) along with the Host key not found in [directory] error message? Which of the below mentioned options is a possible reason for rejection?

- ☒ A. The user has provided the wrong user name for the OS login ✓
- ☐ B. The instance CPU is heavily loaded
- ☐ C. The security group is not configured properly
- ☐ D. The access key to connect to the instance is wrong

Explanation :

Answer – A

As per the AWS documentation , you need to ensure you are entering the right user name if you get the error “Host key not found”

For more information on EC2 trouble shooting please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-troubleshoot.html> (<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-troubleshoot.html>)

Ask our Experts



QUESTION 36

UNATTEMPTED

A user has hosted an application on EC2 instances. The EC2 instances are configured with ELB and Auto Scaling. The application server session time out is 2 hours. The user wants to configure connection draining to ensure that all in-flight requests are supported by ELB even though the instance is being deregistered. What time out period should the user specify for connection draining?



- ☐ A. 5 minutes
- ☒ B. 1 hour ✓
- ☐ C. 30 minutes
- ☐ D. 2 hours

Explanation :

Answer – B

As per the AWS documentation the maximum timeout value can be just 3600 secs=1 hour.

For more information on connection draining please visit the link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/config-conn-drain.html>
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/config-conn-drain.html>)

Ask our Experts



QUESTION 37 UNATTEMPTED

A user is using the AWS EC2. The user wants to make sure that when there is an issue with this EC2 server such as an instance status failed then he would like to have it recovered with out much effort and inform the user about the changes that have happened. Which of the following AWS service helps to achieve this automation?

- ☐ A. AWS CloudWatch + Cloudformation
- ☐ B. AWS CloudWatch + AWS AutoScaling + AWS ELB
- ☐ C. AWS CloudWatch + AWS VPC
- ☒ D. AWS CloudWatch + AWS SNS ✓

Explanation :



Answer – D

When you create an ALARM, you have the option to add a notification as well. So the SNS service can be used to trigger a service which can then be used to trigger an instance in the private cloud.

For a Status check failed (System) error Cloudwatch EC2 Action can perform the recover action for the instance.

As per AWS,

You can create an Amazon CloudWatch alarm that monitors an Amazon EC2 instance and automatically recovers the instance if it becomes impaired due to an underlying hardware failure or a problem that requires AWS involvement to repair.

For more information please refer:

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-recover.html>
(<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-recover.html>)

For more information on Cloudwatch , please visit the URL:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)

Ask our Experts



A sys admin has enabled logging on ELB. Which of the below mentioned fields will not be a part of the log file name?

- ☐ A. Load Balancer Id
- ☒ B. EC2 instance IP ✓
- ☐ C. S3 bucket name
- ☐ D. Random string

Explanation :

Answer – B

As per the AWS documentation below is format of the ELB access log file. In this format the only thing which is not there is the EC2 instance IP.

bucket[/prefix]/AWSLogs/AWS-account-id/elasticloadbalancing/region/yyyy/mm/dd/AWS-account-id_elasticloadbalancing_region_load-balancer-id_end-time_ip-address_random-string.log.gz

For more information on ELB access logs , please visit the URL:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-access-logs.html>
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-access-logs.html>)

Note:

The following is an example log file name.

s3://my-bucket/prefix/AWSLogs/123456789012/elasticloadbalancing/us-east-2/2016/05/01/123456789012_elasticloadbalancing_us-east-2_my-loadbalancer_20140215T2340Z_172.160.001.192_20sg8hgm.log.gz

which comprises of

bucket - my-bucket

The name of the S3 bucket.

prefix

The prefix (logical hierarchy) in the bucket. If you don't specify a prefix, the logs are placed at the root level of the bucket.

aws-account-id

The AWS account ID of the owner.

region



The region for your load balancer and S3 bucket.

yyyy/mm/dd

The date that the log was delivered.

load-balancer-id

The resource ID of the load balancer. If the resource ID contains any forward slashes (/), they are replaced with periods (.).

end-time

The date and time that the logging interval ended. For example, an end time of 20140215T2340Z contains entries for requests made between 23:35 and 23:40.

ip-address

The IP address of the load balancer node that handled the request. For an internal load balancer, this is a private IP address.

random-string

A system-generated random string.

Ask our Experts



QUESTION 39 UNATTEMPTED

A user has created a queue named "AWSmodule" with SQS. One of the consumers of queue is down for 3 days and then becomes available. Will that component receive message from queue?

- ☐ A. Yes, since SQS by default stores message for 4 days ✓
- ☐ B. No, since SQS by default stores message for 1 day only
- ☐ C. No, since SQS sends message to consumers who are available that time
- ☐ D. Yes, since SQS will not delete message until it is delivered to all consumers

Explanation :

Answer – A

As per the AWS documentation it clearly mentions that the default setting for message retention is 4 days hence option A is correct.



For more information on SQS , please visit the URL:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Note:

The question's focus is on the default retention period. However as you have mentioned if there are more than 1 consumer it can poll out the message from the queue to process it , but option C states that SQS will send the message to another consumer which is not correct.

For more information please refer:

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-dg.pdf>

(<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-dg.pdf>)

Ask our Experts



QUESTION 40

UNATTEMPTED

An organization has setup multiple IAM users. The organization wants that each IAM user accesses the IAM console only within the organization and not from outside. How can it achieve this?

- ☐ A. Create an IAM policy with the security group and use that security group for AWS console login
- ☐ B. Create an IAM policy with a condition which denies access when the IP address range is not from the organization ✓
- ☐ C. Configure the EC2 instance security group which allows traffic only from the organization's IP range
- ☐ D. Create an IAM policy with VPC and allow a secure gateway between the organization and AWS Console



Explanation :

Answer – B

You can actually use a Deny condition which will not allow the person to log in from outside. The below example shows the Deny condition to ensure that any address specified in the source address is not allowed to access the resources in AWS.

```
{
  "Version": "2012-10-17",
  "Statement": {
    "Effect": "Deny",
    "Action": "*",
    "Resource": "*",
    "Condition": { "NotIpAddress": { "AWS:SourceIp": [
      "192.0.2.0/24",
      "203.0.113.0/24"
    ] } }
  }
}
```

For more information on IAM policy conditions , please visit the URL:

- http://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_examples.html#iam-policy-example-ec2-two-conditions
(http://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_examples.html#iam-policy-example-ec2-two-conditions)

Ask our Experts



QUESTION 41

UNATTEMPTED

An organization has created one IAM user and applied the below mentioned policy to the user. What entitlements do the IAM users avail with this policy?

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
```



```
"Effect": "Allow",
"Action": "ec2:Describe*",
"Resource": "*"
},
{
"Effect": "Allow"
"Action": [
"cloudwatch:ListMetrics",
"cloudwatch:GetMetricStatistics",
"cloudwatch:Describe*"
],
"Resource": "*"
},
{
"Effect": "Allow",
"Action": "autoscaling:Describe*",
"Resource": "*"
}
}]}
```

- ☐ A. The policy will allow the user to perform all read only activities on the EC2 services
- ☐ B. The policy will allow the user to list all the EC2 resources except EBS
- ☐ C. The policy will allow the user to perform all read and write activities on the EC2 services
- ☐ D. The policy will allow the user to perform all read only activities on the EC2 services except load Balancing ✓



Explanation :

Answer – D

The above policy will allow the user to view EC2 instances , look at Autoscaling and Cloudwatch but not allow the user access to Load Balancing. For access to load balancing you need to have the following statements as well.

```
"Effect": "Allow",  
"Action": "elasticloadbalancing:Describe*",  
"Resource": "*"
```

For more information on EC2 Read Only Access, please visit the link:

- https://docs.aws.amazon.com/directoryservice/latest/admin-guide/role_ec2_read_only.html
(https://docs.aws.amazon.com/directoryservice/latest/admin-guide/role_ec2_read_only.html)
- <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/iam-roles-for-amazon-ec2.html> (<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/iam-roles-for-amazon-ec2.html>)

Ask our Experts



QUESTION 42

UNATTEMPTED

A user has enabled session stickiness with ELB. The user does not want ELB to manage the cookie; instead he wants the application to manage the cookie. What will happen when the server instance, which is bound to a cookie, crashes?

- ☐ A. The response will have a cookie but stickiness will be deleted
- ☐ B. The session will not be sticky until a new cookie is inserted ✓
- ☐ C. ELB will throw an error due to cookie unavailability
- ☐ D. The session will be sticky and ELB will route requests to another server as ELB keeps replicating the Cookie



Explanation :

Answer – B

As per the AWS documentation if the instances crashes the session will no longer be sticky until a new cookie is inserted.

For more information on ELB sticky sessions, please visit the URL:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-sticky-sessions.html> (<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-sticky-sessions.html>)

Ask our Experts



QUESTION 43

UNATTEMPTED

A user is observing the EC2 CPU utilization metric on CloudWatch. The user has observed some interesting patterns while filtering over the 1 week period for a particular hour. The user wants to zoom that data point to a more granular period. How can the user do that easily with CloudWatch?

- ☐ A. The user can zoom a particular period by selecting that period with the mouse and then releasing the mouse ✓
- ☐ B. The user can zoom a particular period by double clicking on that period with the mouse
- ☐ C. The user can zoom a particular period by specifying the aggregation data for that period
- ☐ D. The user can zoom a particular period by specifying the period in the Time Range

Explanation :

Answer – A

So in cloudwatch for the metric, you can actually zoom a particular period by selecting that period with the mouse and then releasing the mouse.



For more information on Cloudwatch , please visit the URL:

- <https://aws.amazon.com/cloudwatch/> (https://aws.amazon.com/cloudwatch/)

Ask our Experts



QUESTION 44 UNATTEMPTED

A user has created an Auto Scaling group with default configurations from CLI. The user wants to setup the CloudWatch alarm on the EC2 instances, which are launched by the Auto Scaling group. The user has setup an alarm to monitor the CPU utilization every minute. Which of the below mentioned statements is true?

- ☐ A. It will fetch the data at every minute but the four data points [corresponding to 4 minutes] will not have value since the EC2 basic monitoring metrics are collected every five minutes
- ☐ B. It will fetch the data at every minute as detailed monitoring on EC2 will be enabled by the default launch configuration of Auto Scaling ✓
- ☐ C. The alarm creation will fail since the user has not enabled detailed monitoring on the EC2 instances
- ☐ D. The user has to first enable detailed monitoring on the EC2 instances to support alarm monitoring at every minute

Explanation :

Answer – B

As per AWS documentation when you enable Auto Scaling via CLI the detailed monitoring will be enabled.

For more information on Launch configurations , please visit the URL:



- <https://docs.aws.amazon.com/cli/latest/reference/autoscaling/create-launch-configuration.html> (<https://docs.aws.amazon.com/cli/latest/reference/autoscaling/create-launch-configuration.html>)
- <http://docs.aws.amazon.com/autoscaling/latest/userguide/LaunchConfiguration.html> (<http://docs.aws.amazon.com/autoscaling/latest/userguide/LaunchConfiguration.html>)

Ask our Experts



QUESTION 45 UNATTEMPTED

A user has launched two EBS backed EC2 instances in the US-East-1a availability zone of region. The user wants to change the availability zone of one of the instances. How can the user change it?

- ☐ A. Stop one of the instances and change the availability zone
- ☐ B. The zone can only be modified using the AWS CLI
- ☐ C. From the AWS EC2 console, select the Actions – > Change zones and specify new zone
- ☐ D. Create an AMI of the running instance and launch the instance in a separate AZ ✓

Explanation :

Answer – D

If you need an AMI across multiple regions or AZ's , then you have to copy the AMI across regions. Note that by default AMI's that you have created will not be available across all regions. So option A is automatically invalid.

Next you can share AMI's with other users, but they will not be available across regions. So option B and C is invalid. You have to copy the AMI across regions.

To copy AMI's , follow the below steps

Step 1) The first step is to create an AMI from your running instance by choosing on Image->Create Image.

Step 2) Once the Image has been created, go to the AMI section in the EC2 dashboard and click



on the Copy AMI option.

Step 3) In the next screen , you can specify where to copy the AMI to.

For the entire details to copy AMI's , please visit the link :

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>)

Option A,B and C is wrong because you cannot change the AZ of an instance once it is running.

Ask our Experts



QUESTION 46 UNATTEMPTED

A user has created a VPC with public and private subnets using the VPC wizard. Which of the below mentioned statements is not true in this scenario?

- ☐ A. The VPC will create a routing instance and attach it with a public subnet ✓
- ☐ B. The VPC will create two subnets
- ☐ C. The VPC will create one internet gateway and attach it to VPC
- ☐ D. The VPC will launch one NAT instance with an elastic IP

Explanation :

Answer – A

Below is the general diagram of what is created when you have a private and public subnet used when using the VPC wizard. So you will get the below options

- 1) 2 subnets – one private and one public
- 2) One NAT instance to route traffic from the public to private subnet
- 3) One internet gateway attached to the VPC.

For more information on VPC and subnets , please visit the URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario2.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario2.html)



Ask our Experts



QUESTION 47

UNATTEMPTED

A user has created a subnet in VPC and launched an EC2 instance within it. The user is not able to access the instance through the internet. Which of the below mentioned statements is true and best with respect to this scenario?

- ☐ A. The instance will always have a public DNS attached to the instance by default
- ☐ B. The user can directly attach an elastic IP to the instance
- ☐ C. The instance will never launch if the public IP is not assigned
- ☐ D. The user would need to create an internet gateway and then attach an elastic IP to the instance to connect from internet ✓

Explanation :

Answer – D

When you create a simple subnet and VPC , there is no internet gateway attached to the VPC. So to access it from the internet you need to ensure an Internet gateway is attached to the VPC and an elastic or public IP is assigned to the EC2 instance. For more information on VPC and subnets , please visit the URL:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html)
- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Internet_Gateway.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Internet_Gateway.html)



Ask our Experts



QUESTION 48

UNATTEMPTED

An organization has applied the below mentioned policy on an IAM group which has selected the IAM users. What entitlements do the IAM users avail with this policy?

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": "*",  
      "Resource": "*"   
    }  
  ]  
}
```

- ☐ A. The policy is not created correctly. It will throw an error for wrong resource name
- ☐ B. The policy is for the group. Thus, the IAM user cannot have any entitlement to this
- ☐ C. It allows full access to all AWS services for the IAM users who are a part of this group ✓
- ☐ D. If this policy is applied to the EC2 resource, the users of the group will have full access to the EC2 Resources



Explanation :

Answer – C

This is like having the complete administrator access. If you see the below snapshot you can see the policy content is the same as the question.

For more information on IAM policies , please visit the URL:

- http://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies.html
(http://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies.html)

Ask our Experts



QUESTION 49 UNATTEMPTED

A user has configured a Cloudwatch alarm for RDS in order to receive a notification whenever the CPU utilization of RDS is higher than 50%. But, the user would like to receive notification if the RDS instance is in unknown or unavailable state. How can this be achieved?

- ☐ A. Setup the notification when the CPU is more than 75% on RDS
- ☐ B. Setup the notification when the state is Insufficient Data ✓
- ☐ C. Setup the notification when the CPU utilization is less than 10%
- ☐ D. It is not possible to setup the alarm on RDS

Explanation :

Answer – B

An alarm has three possible states:

- OK—The metric is within the defined threshold
- ALARM—The metric is outside of the defined threshold
- INSUFFICIENT_DATA—The alarm has just started, the metric is not available, or not enough data is available for the metric to determine the alarm state



The best way to configure the alarm is to configure it for insufficient data.

For more information on Cloudwatch alarms, please visit the link:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/AlarmThatSendsEmail.html>
(<http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/AlarmThatSendsEmail.html>)

Ask our Experts



QUESTION 50 UNATTEMPTED

George has shared an EC2 AMI created in the US East region from his AWS account with Stefano. George copies the same AMI to the US West region. Can Stefano access the copied AMI of George's account from the US West region?

- ☐ A. No, copy AMI does not copy the permission ✓
- ☐ B. It is not possible to share the AMI with a specific account
- ☐ C. Yes, since copy AMI copies all private account sharing permissions
- ☐ D. Yes, since copy AMI copies all the permissions attached with the AMI

Explanation :

Answer – A

You can copy an Amazon Machine Image (AMI) within or across an AWS region using the AWS Management Console, the AWS command line tools or SDKs, or the Amazon EC2 API, all of which support the CopyImage action. You can copy both Amazon EBS-backed AMIs and instance store-backed AMIs. You can copy AMIs with encrypted snapshots and encrypted AMIs. AWS does not copy launch permissions, user-defined tags, or Amazon S3 bucket permissions from the source AMI to the new AMI. After the copy operation is complete, you can apply launch permissions, user-defined tags, and Amazon S3 bucket permissions to the new AMI.

For more information on copying AMI's, please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>)

Ask our Experts





QUESTION 51 UNATTEMPTED

A user has created a VPC with a subnet and a security group. The user has launched an instance in that subnet and attached a public IP. The user is still unable to connect to the instance. The internet gateway has also been created. What can be the reason for the error?

- ☐ A. The internet gateway is not configured with the route table ✓
- ☐ B. The private IP is not present
- ☐ C. The outbound traffic on the security group is disabled
- ☐ D. The internet gateway is not configured with the security group

Explanation :

Answer – A

The below diagram shows a simple public subnet in AWS. You have to ensure the internet gateway is configured in the route table.

For more information on VPC with a public subnet, please visit the link:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html
(http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Scenario1.html)

Ask our Experts



QUESTION 52 UNATTEMPTED

A user is trying to setup a security policy for ELB. The user wants ELB to meet the cipher supported by the client by configuring the server order preference in ELB security policy. Which of the below mentioned preconfigured policies supports this feature? ^

- ☒ A. ELBSecurityPolicy-2016-08 ✓
- ☐ B. ELBSecurity Policy-2011-08
- ☐ C. ELBDefault Negotiation Policy
- ☐ D. ELBSample- OpenSSLDefault Cipher Policy

Explanation :

Answer – A

Currently the default ELB Predefined Security Policy is ELBSecurityPolicy-2016-08

Elastic Load Balancing supports the *Server Order Preference* option for negotiating connections between a client and a load balancer. During the SSL connection negotiation process, the client and the load balancer present a list of ciphers and protocols that they each support, in order of preference. By default, the first cipher on the client's list that matches any one of the load balancer's ciphers is selected for the SSL connection. If the load balancer is configured to support Server Order Preference, then the load balancer selects the first cipher in its list that is in the client's list of ciphers. This ensures that the load balancer determines which cipher is used for SSL connection. If you do not enable Server Order Preference, the order of ciphers presented by the client is used to negotiate connections between the client and the load balancer.

For more information on ELB security policies, please visit the link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-security-policy-options.html> (<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-security-policy-options.html>)





QUESTION 53

UNATTEMPTED

A user has configured ELB with Auto Scaling. The user suspended the Auto Scaling AlarmNotification which notifies Auto Scaling for CloudWatch alarms process for a while. What will Auto Scaling do during this period?

- ☐ A. AWS will not receive the alarms from CloudWatch
- ☐ B. AWS will receive the alarms but will not execute the Auto Scaling policy ✓
- ☐ C. Auto Scaling will execute the policy but it will not launch the instances until the process is resumed
- ☐ D. It is not possible to suspend the AlarmNotification process

Explanation :

Answer – B

If you suspend AlarmNotification, Auto Scaling does not automatically execute policies that would be triggered by an alarm. If you suspend Launch or Terminate, Auto Scaling would not be able to execute scale out or scale in policies, respectively.

For more information on Suspend process autoscaling, please visit the link:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/as-suspend-resume-processes.html> (<http://docs.aws.amazon.com/autoscaling/latest/userguide/as-suspend-resume-processes.html>)

Ask our Experts



QUESTION 54

UNATTEMPTED



A user had aggregated the CloudWatch metric data on the AMI ID. The user observed some abnormal behaviour of the CPU utilization metric while viewing the last 2 weeks of data. The user wants to share that data with his manager. How can the user achieve this easily with the AWS console?

- ☐ A. The user can use the copy URL functionality of CloudWatch to share the exact details ✓
- ☐ B. The user can use the export data option from the CloudWatch console to export the current data point
- ☐ C. The user has to find the period and data and provide all the aggregation information to the manager
- ☐ D. The user can use the CloudWatch data copy functionality to copy the current data points

Explanation :

Answer – A

You can share an graph in Cloudwatch from the Share Graph URL option in cloudwatch , hence option A is the right option.

For more information on cloudwatch, please visit the link:

- <https://aws.amazon.com/cloudwatch/> (https://aws.amazon.com/cloudwatch/)

Ask our Experts



QUESTION 55 UNATTEMPTED

A user has setup a CloudWatch alarm on the EC2 instance for CPU utilization. The user has setup to receive a notification on email when the CPU utilization is higher than 60%. The user is running a virus scan on the same instance at a particular time. The user wants to avoid receiving an email at this time. What should the user do?

- ☐ A. Remove the alarm
- ☒ B. Disable the alarm for a while using CLI ✓
- ☐ C. Modify the CPU utilization by removing the email alert
- ☐ D. Disable the alarm for a while using the console

Explanation :

Answer – B

You can use the below CLI command to disable a cloudwatch alarm.

AWS cloudwatch disable-alarm-actions --alarm-names myalarm

Disables the actions for the specified alarms. When an alarm's actions are disabled, the alarm actions do not execute when the alarm state changes.

For more information on cloudwatch, please visit the link:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)

Ask our Experts



QUESTION 56 UNATTEMPTED

A user would like to be notified if the CPU utilization of his EC2 instances exceed 90%. Which of the following option would be able to address his concerns?

- ☐ A. AWS CloudWatch + AWS SES
- ☒ B. AWS CloudWatch + AWS SNS ✓
- ☐ C. None. It is not possible to notify the user with the AWS services.
- ☐ D. Use a dedicated third party software alongside Cloudwatch for notifying the user.

Explanation :

Answer – B



When you create an ALARM, you have the option to add a notification.

Option A is wrong because SES is just an emailing service which will not serve the purpose.
Option D is wrong because Cloudwatch can't be directly used to interact with a dedicated software to turn on the light.

For more information on Cloudwatch , please visit the URL:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)

Ask our Experts



QUESTION 57 UNATTEMPTED

A user has configured ELB with SSL using a security policy for secure negotiation between the client and load balancer. Which of the below mentioned SSL protocols is not supported by the security policy?

- ☒ A. TLS1.3 ✓
- ☐ B. TLS1.2
- ☐ C. TLS1.1
- ☐ D. TLS1

Explanation :

Answer – A

As per the AWS documentation below are the supported protocols , hence Option A is right.



For more information, check the following AWS document link:

- <https://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-ssl-security-policy.html#ssl-protocols>
(<https://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-ssl-security-policy.html#ssl-protocols>)

For more information on ELB security polices, please visit the link:



- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-security-policy-table.html> (<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/elb-security-policy-table.html>)

Ask our Experts



QUESTION 58

UNATTEMPTED

An organization has added 3 of his AWS accounts to consolidated billing. One of the AWS accounts has purchased a Reserved Instance (RI) of a small instance size in the US-East-1a zone. All other AWS accounts are running instances of a small size in the same zone. What will happen in this case for the RI pricing?

- ☐ A. Only the account that has purchased the RI will get the advantage of RI pricing
- ☐ B. One instance of a small size and running in the US-East-1a zone of each AWS account will get the benefit of RI pricing
- ☐ C. Any single instance from all the three accounts can get the benefit of AWS RI pricing if they are running in the same zone and are of the same size ✓
- ☐ D. If there are more than one instances of a small size running across multiple accounts in the same zone no one will get the benefit of RI

Explanation :

Answer – C

As per the AWS documentation for billing purposes, AWS treats all the accounts on the consolidated bill as if they were one account. Some services, such as Amazon EC2 and Amazon S3, have volume pricing tiers across certain usage dimensions that give you lower prices when you use the service more. With consolidated billing, AWS combines the usage from all accounts to determine which volume pricing tiers to apply, giving you a lower overall price whenever possible.

For more information on Consolidated billing, please visit the URL:



- <http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/consolidated-billing.html> (<http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/consolidated-billing.html>)

Ask our Experts



QUESTION 59

UNATTEMPTED

A user has setup an EBS backed instance and a CloudWatch alarm when the CPU utilization is more than 65%. The user has setup the alarm to watch it for 5 periods of 5 minutes each. The CPU utilization is 60% between 9 AM to 6 PM. The user has stopped the EC2 instance for 15 minutes between 11 AM to 11:15 AM. What will be the status of the alarm at 11:30 AM?

- ☐ A. Alarm
- ☒ B. OK ✓
- ☐ C. Insufficient Data
- ☐ D. Error

Explanation :

Answer – B

Since the alarm has been set for 5 periods of 5 minutes each, that means if there is no activity for 25 minutes then the state of the Alarm will change. Now since the CPU utilization was 60%, that means the alarm was not triggered and the state should be OK. And since the instance was just stopped for 15 minutes and did not exceed the threshold for 25 minutes that means the state of the alarm stays the same.

For more information on Cloudwatch alarms, please visit the URL:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/AlarmThatSendsEmail.html> (<http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/AlarmThatSendsEmail.html>)

Note:



AS per AWS,

Similar to how each alarm is always in one of three states, each specific data point reported to CloudWatch falls under one of three categories:

- Not breaching (within the threshold)
- Breaching (violating the threshold)
- Missing

You can specify how alarms handle missing data points. Choose whether to treat missing data points as:

- *missing* (The alarm looks back farther in time to find additional data points)
- *notBreaching* (Treated as a data point that is within the threshold)
- *breaching* (Treated as a data point that is breaching the threshold)
- *ignore* (The current alarm state is maintained)

In our case discussed here we have set the alarm for 5 consecutive periods of 5 minutes each. So the data was missing for a period of 15 minutes. So if the data is missing then the alarm looks back further to find previous data points to see the alarm status. Prior to the instance being stopped, instance's Alarm Status was **OK**. Hence it will maintain the same Alarm status. Hence the answer is **OK**.

Ask our Experts



A user has setup an RDS DB with Oracle. The user wants to get notifications when someone modifies the security group of that DB. How can the user configure that?

- ☐ A. It is not possible to get the notifications on a change in the security group
- ☐ B. Configure SNS to monitor security group changes
- ☒ C. Configure event notification on the DB security group ✓
- ☐ D. Configure the CloudWatch alarm on the DB for a change in the security group

Explanation :

Answer – C

When you go to the Event dashboard in RDS, you can actually create event subscriptions based on Security groups.

You would not choose Option B and D because there is already in-built for RDS.

For more information on Event Notification group, please visit the URL:

- http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_Events.html
(http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_Events.html)

Ask our Experts




Finish Review (<https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests/quiz/12743>)

Certification

- 🔗 Cloud Certification
(<https://www.whizlabs.com/cloud-certification-training-courses/>)

Company

- 🔗 Support (<https://help.whizlabs.com/hc/en-us>)
- 🔗 Discussions (<http://ask.whizlabs.com/>) 

- ➔ Java Certification
(<https://www.whizlabs.com/oracle-java-certifications/>)
- ➔ PM Certification
(<https://www.whizlabs.com/project-management-certifications/>)
- ➔ Big Data Certification
(<https://www.whizlabs.com/big-data-certifications/>)

Mobile App

 Android Coming Soon

 iOS Coming Soon

- ➔ Blog (<https://www.whizlabs.com/blog/>)

Follow us



(<https://www.facebook.com/whizlabs.software/>)



(<https://in.linkedin.com/company/whizlabs-software>)



(<https://twitter.com/whizlabs?lang=en>)



(<https://plus.google.com/+WhizlabsSoftware>)

