

[🏠 \(https://www.whizlabs.com/learn\)](https://www.whizlabs.com/learn) > [My Courses \(https://www.whizlabs.com/learn/my-courses\)](https://www.whizlabs.com/learn/my-courses)  
> [AWS Certified SysOps Administrator Associate \(https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests#section-1\)](https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests#section-1)  
> [Practice Test V \(https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests/quiz/13327\)](https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests/quiz/13327) > **Report**

## PRACTICE TEST V

<b>Attempt</b>	2	<b>Completed on</b>	Tuesday , 29 January 2019 , 02:22 PM
<b>Marks Obtained</b>	1 / 60	<b>Time Taken</b>	00 H 00 M 33 S
<b>Your score is</b>	1.67%	<b>Result</b>	Fail

### Domains / Topics wise Quiz Performance Report

S.No.	Topic	Total Questions	Correct	Incorrect	Unattempted
1	Other	60	1	1	58

60 Questions	1 Correct	1 Incorrect	58 Unattempted	Show Answers	All	▼
-----------------	--------------	----------------	-------------------	--------------	-----	---

### QUESTION 1 CORRECT

In Autoscaling what are some of the error messages that can occur. Choose 3 answers from the options below

- ☒ A. AutoScalingGroup <Auto Scaling group name> not found. ✓
- ☒ B. The requested configuration is currently not supported. ✓
- ☒ C. The requested Availability Zone is no longer supported. Please retry your request ..... ✓
- ☐ D. The requested Region is no longer supported. Please retry your request .....

#### Explanation :

Answer – A,B and C

The cause for error in Option A is The Auto Scaling group might have been deleted.

The cause for error in Option B is some options in your launch configuration might not be currently supported.

The cause for error in Option C is the requested Availability Zone is no longer supported. Please retry your request by not specifying an Availability Zone or choosing <list of available Availability Zones>. Launching EC2 instance failed.  
For more information on Autoscaling troubleshooting, please visit the below URL:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/ts-as-instancelaunchfailure.html>  
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/ts-as-instancelaunchfailure.html>)

Ask our Experts



## QUESTION 2      INCORRECT

What is the largest size of an object which can be uploaded by a PUT request in S3?  
Choose the correct answer from the options below

- ☒ A. 5 GB ✓
- ☐ B. 1 TB ✗
- ☐ C. 100 MB
- ☐ D. 1 GB

### Explanation :

Answer – A

This is clearly mentioned in the AWS documentation

The total volume of data and number of objects you can store are unlimited. Individual Amazon S3 objects can range in size from a minimum of 0 bytes to a maximum of 5 terabytes. The largest object that can be uploaded in a single PUT is 5 gigabytes. For objects larger than 100 megabytes, customers should consider using the [Multipart Upload](#) capability.

For more information on S3 please refer to the below URL:

- <https://aws.amazon.com/s3/faqs/> (<https://aws.amazon.com/s3/faqs/>)

Ask our Experts



## QUESTION 3      UNATTEMPTED

Which one of AWS RDS features listed below is supposed to allow 'a point in time restore' of your database?

Please select the best answer.

- ☐ A. Multi-AZ RDS
- ☐ B. AWS S3
- ☐ C. RDS read replicas
- ☒ D. RDS automated backup ✓

#### Explanation :

Answer – D

Automated backups automatically back up your DB instance during a specific, user-definable backup window. Amazon RDS keeps these backups for a limited period that you can specify. You can later recover your database to any point in time during this backup retention period. And all of these backups get stored to S3 by default.

The backup retention period is present as a setting when you setup the RDS in AWS.

#### Backup

Please note that automated backups are currently supported for InnoDB storage engine only. If you are using MyISAM, refer to detail [here](#).

Backup Retention Period  days

Backup Window

Option A and C are wrong because If the table is deleted the Multi-AZ and Read Replica will also have the same dropped tables because the synchronization would have already happened.

For more information on Automated backups, please visit URL:

- [http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER\\_WorkingWithAutomatedBackups.html](http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_WorkingWithAutomatedBackups.html)  
([http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER\\_WorkingWithAutomatedBackups.html](http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_WorkingWithAutomatedBackups.html))

Ask our Experts



QUESTION 4 UNATTEMPTED

If there is a requirement to upload a 6GB file to S3 what is the best option to use from the below?

- ☐ A. Increase your network bandwidth to provide faster throughput to S3
- ☒ B. Use Multipart upload feature of S3 and upload the parts in parallel ✓
- ☐ C. Pack all files into a single archive, upload it to S3, and then extract the files in AWS
- ☐ D. Use AWS Import/Export to transfer the video files

### Explanation :

Answer – B

When uploading large videos it's always better to make use of AWS multi part file upload.

So if you are using the Multi Upload option for S3, then you can resume on failure. Below are the advantage of Multi Part upload

- Improved throughput—you can upload parts in parallel to improve throughput.
- Quick recovery from any network issues—smaller part size minimizes the impact of restarting a failed upload due to a network error.
- Pause and resume object uploads—you can upload object parts over time. Once you initiate a multipart upload there is no expiry; you must explicitly complete or abort the multipart upload.
- Begin an upload before you know the final object size—you can upload an object as you are creating it.

For more information on Multi-part file upload for S3, please visit the URL:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/qfacts.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/qfacts.html>)

Ask our Experts



### QUESTION 5 UNATTEMPTED

Which of the below services is used as infrastructure as a code in AWS. Choose an answer from the options below

- ☐ A. SNS
- ☐ B. SES
- ☒ C. Cloudformation ✓
- ☐ D. None of the above

### Explanation :

Answer – C

AWS CloudFormation gives developers and systems administrators an easy way to create and manage a collection of related AWS resources, provisioning and updating them in an orderly and predictable fashion.

You can use AWS CloudFormation's sample templates

(<https://aws.amazon.com/cloudformation/aws-cloudformation-templates/>) or create your own templates to describe the AWS resources, and any associated dependencies or runtime parameters, required to run your application.

For more information on cloudformation, please visit the URL:

- <https://aws.amazon.com/cloudformation/> (<https://aws.amazon.com/cloudformation/>)



## QUESTION 6 UNATTEMPTED

A user is planning to evaluate AWS for their internal use. The user does not want to incur any charge on his account during the evaluation. Which of the below mentioned AWS services would incur a charge if used?

- ☐ A. 1000 hours of RDS usage ✓
- ☐ B. 50GB of Amazon Cloudfront storage
- ☐ C. 750 hours of ELB usage
- ☐ D. 30GB of EBS

**Explanation :**

Answer – A

The below documentation from AWS gives what are offered free from AWS. In this , only 750 hours of RDS storage is allowed.

<p><b>DATABASE</b></p> <p><b>750 HOURS</b> per month of database usage</p> <p><b>Amazon RDS</b></p> <p>Managed Relational Database Service for MySQL, PostgreSQL, MariaDB, Oracle BYOL, or SQL Server</p> <p><a href="#">Learn more about Amazon RDS »</a></p>	<p><b>INTERNET OF THINGS</b></p> <p><b>250K MSG.</b> (published or delivered) per month</p> <p><b>AWS IoT</b></p> <p>Connect devices to the cloud</p> <p><a href="#">Learn more about AWS IoT »</a></p>	<p><b>COMPUTE</b></p> <p><b>500MB</b> per month of storage</p> <p><b>Amazon EC2 Container Registry</b></p> <p>Store and retrieve Docker images</p> <p><a href="#">Learn more about Amazon EC2 Container Registry »</a></p>
<p><b>COMPUTE</b></p> <p><b>750 HOURS</b> per month shared between Classic and Application load balancers</p> <p><b>Elastic Load Balancing</b></p> <p>Automatic distribution of incoming application traffic across multiple Amazon EC2 instances</p> <p><a href="#">Learn more about Elastic Load Balancing »</a></p>	<p><b>STORAGE &amp; CONTENT DELIVERY</b></p> <p><b>50GB</b> of storage</p> <p><b>Amazon CloudFront</b></p> <p>Web service to distribute content to end users with low latency and high data transfer speeds</p> <p><a href="#">Learn more about Amazon CloudFront »</a></p>	<p><b>STORAGE &amp; CONTENT DELIVERY</b></p> <p><b>30GB</b> any combination of General Purpose (SSD) or Magnetic</p> <p><b>Amazon Elastic Block Storage</b></p> <p>Persistent, durable, low-latency block-level storage volumes for EC2 instances</p> <p><a href="#">Learn more about Amazon Elastic Block Storage »</a></p>

For more information on the AWS Free tier, please visit the link:

- <https://aws.amazon.com/free/> (<https://aws.amazon.com/free/>)

Ask our Experts



QUESTION 7 UNATTEMPTED

Which of the below services is a fully managed MySQL Solution from AWS. Choose an answer from the options below

- ☐ A. DynamoDB
- ☒ B. Aurora ✓
- ☐ C. SQLServer
- ☐ D. Oracle

**Explanation :**

Answer – B

Amazon Aurora is a relational database engine that combines the speed and reliability of high-end commercial databases with the simplicity and cost-effectiveness of open source databases. It delivers up to five times the throughput of standard MySQL running on the same hardware.

For more information on aurora, please visit the link:

- <https://aws.amazon.com/rds/aurora/details/> (https://aws.amazon.com/rds/aurora/details/)

Ask our Experts



QUESTION 8 UNATTEMPTED

Is it possible to increase the limit of a billing alarm in cloudwatch?

- ☒ A. Yes ✓
- ☐ B. No

**Explanation :**

Answer – A

Let's assume that an alarm has been created as shown below for any amounts exceeding 200 USD.

<a href="#">Create Alarm</a> <a href="#">Modify</a> <a href="#">Copy</a> <a href="#">Delete</a>		
Filter: <a href="#">All alarms</a> <input type="text" value="AWS/Billing"/>		
State	Name	Threshold
<input checked="" type="checkbox"/> INSUFFICIENT_DATA	BillingAlarm	EstimatedCharges > 200 for 6 hours

To increase the limit, all you have to do is to click on the Modify option and you can change the value of the alarm in the next screen.

## Modify Alarm

1. [Select Metric](#) 2. [Define Alarm](#)

### Billing Alarm

You can create a billing alarm to receive e-mail alerts when your AWS charges exceed a threshold you choose. Simply:

1. Enter a spending threshold
2. Provide an email address
3. Check your inbox for a confirmation email and click the link provided

When my total AWS charges for the month

exceed: \$  USD

send a notification to:  [New list](#)

For more information on AWS Cloudwatch, please visit the link:

- <https://aws.amazon.com/cloudwatch/> (<https://aws.amazon.com/cloudwatch/>)

Ask our Experts



## QUESTION 9 UNATTEMPTED

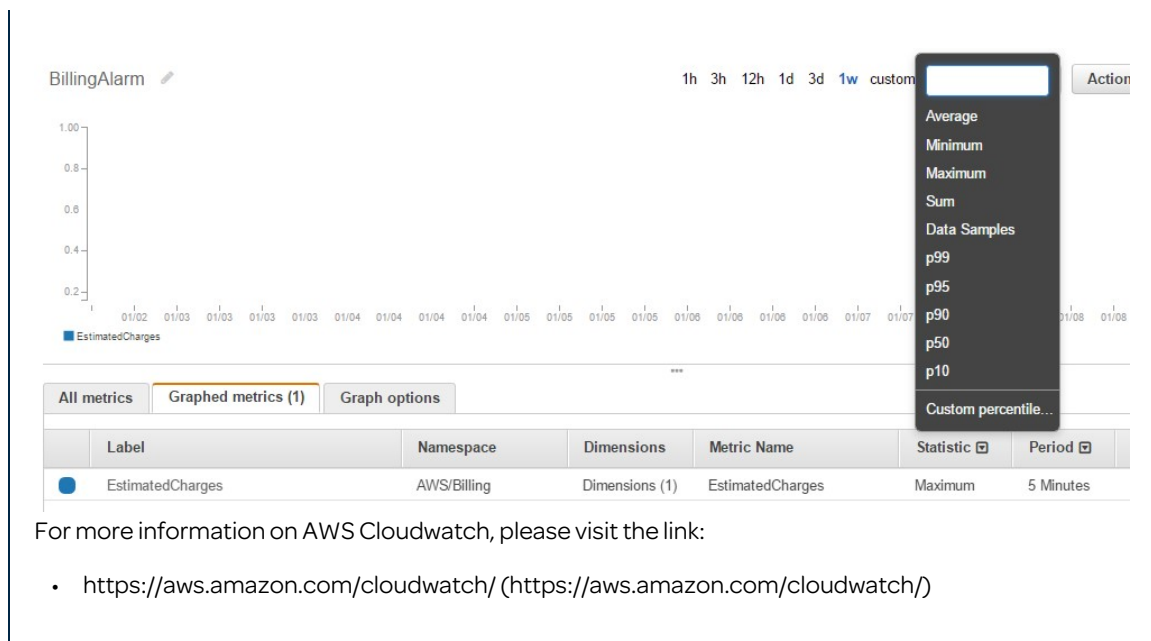
A user is trying to aggregate all the CloudWatch metric data of the last 1 week. Which of the below mentioned statistics is not available for the user as a part of data aggregation?

- ☐ A. Minimum
- ☐ B. Maximum
- ☒ C. Mean ✓
- ☐ D. Average

### Explanation :

Answer – C

If you go to Cloudwatch and go to any metric and see the statistic column you will see the below dimensions for statistics. And the only one which is not there from the above list is Mean.



Ask our Experts



## QUESTION 10 UNATTEMPTED

Which of the below is used to direct traffic to subnets

- ☐ A. SubnetID
- ☒ B. Route table ✓
- ☐ C. VPC
- ☐ D. Network ACL

### Explanation :

Answer - B

Subnet ID: subnet-dfd2a5f2 | Default  
 CIDR: 192.168.0.0/28  
 State: available  
 VPC: vpc-3e6dde58 | 192.168.0.0/24  
 Available IPs: 11

Availability Zone: us-east-1a  
 Route table: rtb-c405f1bd  
 Network ACL: acl-7dd5ae1b  
 Default subnet: no  
 Auto-assign Public IP: no

A route table contains a set of rules, called routes that are used to determine where network traffic is directed.

Each subnet in your VPC must be associated with a route table; the table controls the routing for the subnet. A subnet can only be associated with one route table at a time, but you can associate multiple subnets with the same route table.



For more information on Route tables, please visit the link:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Route\\_Tables.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Route_Tables.html)  
([http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Route\\_Tables.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Route_Tables.html))

**Note:**

AWS says about NACL that "A *network access control list (ACL)* is an optional layer of security for your VPC that acts as a firewall for controlling traffic in and out of one or more subnets."

However, "A *route table* contains a set of rules, called *routes*, that are used to determine where network traffic is directed."

Please check the below links to know more about it:

- [https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_ACLs.html](https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLs.html)  
([https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_ACLs.html](https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLs.html))
- [https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Route\\_Tables.html](https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Route_Tables.html)  
([https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Route\\_Tables.html](https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Route_Tables.html))

Ask our Experts



QUESTION 11 UNATTEMPTED

A user is trying to understand AWS SNS. To which of the below mentioned end points is SNS unable to send a notification?

- ☐ A. Application
- ☐ B. SQS
- ☐ C. AWS Lambda
- ☐ D. AWS SES ✓

**Explanation :**

Answer – D

Let's see a sample Topic created in the SNS dashboard as shown below.

SNS dashboard

Topics

Applications

Subscriptions

Text messaging (SMS)

Topic details: NotifyMe

Publish to topic

Other topic actions ▾

Topic ARN

am:aws:sns:us-east-1:085363624145:NotifyMe

Topic owner

085363624145

Region

us-east-1

Display name

Subscriptions

Create subscription

Request confirmations

Confirm subscription

Other subscription actions ▾

When you create a subscription, you can see the different Protocols available

Create subscription

Topic ARN

am:aws:sns:us-east-1:085363624145:NotifyMe

Protocol

HTTP

Endpoint

HTTP

HTTPS

Email

Email-JSON

Amazon SQS

Application

AWS Lambda

SMS

Subscription ID

Subscription

From the snapshot it is evident that option D is invalid.  
For more information on SNS please visit the link:

- <https://aws.amazon.com/sns/> (<https://aws.amazon.com/sns/>)

Ask our Experts

👍

👎

QUESTION 12      UNATTEMPTED

Can Autoscaling be used to launch Spot Instances.

☒ A. Yes ✓

☐ B. No

Explanation :

Answer – A

Spot instances are a cost-effective choice compared to On-Demand instances, if you can be flexible about when your applications run and if your applications can be interrupted. You can set up Auto Scaling to launch Spot instances instead of On-Demand instances.

For more information on Autoscaling, please visit the below URL:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/US-SpotInstances.html>  
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/US-SpotInstances.html>)

Ask our Experts



## QUESTION 13 UNATTEMPTED

When you set up a static web site in S3 what are the most important steps to carry out? Choose 3 correct answer from the options below

- ☐ A. Upload an index document to your S3 bucket ✓
- ☐ B. Select the "Make Public" permission for your bucket's objects ✓
- ☐ C. Create an Alias record in Route 53
- ☐ D. Enable static website hosting in your S3 bucket properties ✓

### Explanation :

Answer – A, B and D

When you go to S3, make sure that static website hosting is enabled. Also upload the index.html document , make sure it is Public and also ensure that it is put as the Index document.

#### ▼ Static Website Hosting

You can [host your static website](#) entirely on Amazon S3. Once you enable your bucket for static website hosting, all your content is accessible to web browsers via the Amazon S3 website endpoint for your bucket.

**Endpoint:** `unique-bucket-1021.s3-website-ap-southeast-1.amazonaws.com`

Each bucket serves a website namespace (e.g. "www.example.com"). Requests for your host name (e.g. "example.com" or "www.example.com") can be routed to the contents in your bucket. You can also redirect request to another host name (e.g. redirect "example.com" to "www.example.com"). See our [walkthrough](#) for how to set up an Amazon S3 static website with your host name.

☐ Do not enable website hosting

☒ Enable website hosting

**Index Document:**

**Error Document:**

► **Edit Redirection Rules:** You can set custom rules to automatically redirect web page requests for specific content.

For more information on all S3 web site hosting please refer to the below URL:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/WebsiteHosting.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/WebsiteHosting.html>)



## QUESTION 14 UNATTEMPTED

There is an S3 bucket which is private. There are also objects present in this bucket. Which of the below steps is the recommended one to make the objects in the bucket accessible to other users?

- ☐ A. The user should select all objects from the console and apply a single policy to mark them public
- ☐ B. The user can write a program which programmatically makes all objects public using S3 SDK
- ☒ C. Set the AWS bucket policy which marks all objects as public ✓
- ☐ D. Make the bucket ACL as public so it will also mark all objects as public

**Explanation :**

Answer – C

You can set AWS bucket policy to make everything public. You can do via the following steps

Step 1) Go to your S3 bucket and click on the Permissions section

PERMISSIONS

► Permissions

► Static Website Hosting

► Logging

► Events

► Versioning

► Lifecycle

► Cross-Region Replication

Step 2) Go to AWS Policy Generator

Bucket Policy Editor

Cancel

Policy for Bucket : "devtoolslogging"

Add a new policy or edit an existing bucket policy in the text area below. [Learn more.](#)

[AWS Policy Generator](#) | [Sample Bucket Policies](#)

SaveDeleteClose

Step 3) Add the S3 bucket policy to allow any Principal to access your bucket and choose the Actions accordingly.

Select Type of Policy S3 Bucket Policy

## Step 2: Add Statement(s)

A statement is the formal description of a single permission. See [a description of elements](#) that you can use in statement

Effect ☒ Allow ☐ Deny

Principal \*

Use a comma to separate multiple values.

AWS Service Amazon S3 ☐ All Services (\*\*)

Use multiple statements to add permissions for more than one service.

Actions -- Select Actions -- ☒ All Actions (\*\*)

Amazon Resource Name (ARN) arn:aws:s3:::bucketnm/\*

ARN should follow the following format: arn:aws:s3:::<bucket\_name>/<key\_name>.  
Use a comma to separate multiple values.

[Add Conditions \(Optional\)](#)

Add Statement

For more information on S3 Security please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/DataDurability.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/DataDurability.html>)

Ask our Experts



Which of the following is the maximum allowable time for connection draining in AWS elb?

- ☐ A. 600 seconds
- ☒ B. 3600 seconds ✓
- ☐ C. 300 seconds
- ☐ D. 0 seconds

### Explanation :

Answer – B

When you have an ELB, you can go to the Console, go to the Instances tab and edit the Connection draining time for instances. By default the Connection draining time limit is set to 300.

The screenshot shows the AWS Management Console for a load balancer named 'Demo'. The 'Instances' tab is active, displaying a table with one instance. The instance ID is 'i-00acbc260371247b7', located in 'us-east-1a', and its status is 'OutOfService'. The 'Connection Draining' status is 'Enabled, 300 seconds'.

Instance ID	Name	Availability Zone	Status	Actions
i-00acbc260371247b7		us-east-1a	OutOfService ⓘ	<a href="#">Remove from Load Balancer</a>

When you enable connection draining, you can specify a maximum time for the load balancer to keep connections alive before reporting the instance as de-registered. The maximum timeout value can be set between 1 and 3,600 seconds (the default is 300 seconds). When the maximum time limit is reached, the load balancer forcibly closes connections to the de-registering instance.

For more information on connection draining, please refer to the below link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/config-conn-drain.html>  
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/config-conn-drain.html>)

Ask our Experts



QUESTION 16 UNATTEMPTED

In S3 , when giving permissions via ACL what are the specific permissions that can be

given via the console? Choose 3 answers from the options below

- ☐ A. List Objects ✓
- ☐ B. Full Control
- ☐ C. Write Objects ✓
- ☐ D. Read Bucket Permissions ✓

#### Explanation :

Answer – A,C and D

S3 Console:

##### ▼ Permissions

You can control access to the bucket and its contents using access policies. [Learn more.](#)

Grantee:  ☐ List ☐ Upload/Delete ☐ View Permissions ☐ Edit Permissions

For more information on S3 Security, please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/DataDurability.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/DataDurability.html>)

Ask our Experts



#### QUESTION 17 UNATTEMPTED

When configuring an ELB what setting can be made to ensure that the user request always goes to the same EC2 instance?

- ☐ A. Enable ELB cross zone load balancing
- ☐ B. Enable ELB cookie setup
- ☐ C. Enable ELB sticky session ✓

☐ D. Enable ELB connection draining

**Explanation :**

Answer – C

To ensure that each end user request goes to the same EC2 instance as the session created, you need to enable stickiness at the ELB level

To enable stickiness go to the ELB and in the port configuration section you can enable the stickiness.

Create Load Balancer Actions ▾

Filter:  X

<input type="checkbox"/>	Name	DNS name	State
<input type="checkbox"/>	Demo	internal-Demo-824544633....	

Scheme: internal

Availability Zones: subnet-dfd2a5f2 - us-east-1a

**Port Configuration**

Port Configuration: 80 (HTTP) forwarding to 80 (HTTP)  
Stickiness: Disabled

Edit stickiness

For more information on ELB, please visit the link:

- <https://aws.amazon.com/elasticloadbalancing/>  
(<https://aws.amazon.com/elasticloadbalancing/>)

Ask our Experts



QUESTION 18 UNATTEMPTED

When EC2 instances are registered with an AWS ELB, they are registered in a group. What is that group called?

- ☐ A. Placement Group
- ☐ B. Primary Group
- ☐ C. Secondary Group



## ☐ D. Target Group ✓

### Explanation :

Answer - D

You register targets, such as EC2 instances, with a target group. To route requests to the targets in a target group, specify the target group in a rule for one of the listeners for your load balancer.

You define health check settings for your load balancer on a per target group basis. Each target group uses the default health check settings, unless you override them when you create the target group or modify them later on. After you specify a target group in a rule for a listener, the load balancer continually monitors the health of all targets registered with the target group that are in an Availability Zone enabled for the load balancer. The load balancer routes requests to the registered targets that are healthy.

For more information on Target groups please visit the link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-target-groups.html> (<http://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-target-groups.html>)

Ask our Experts



## QUESTION 19 UNATTEMPTED

Is the below cloudformation template a valid one?

```
{
  "Parameters": {
    "KeyName": {
      "Description": "The EC2 Key Pair to allow SSH access to the instance",
      "Type": "AWS::EC2::KeyPair::KeyName"
    }
  }
}
```

- ☐ A. Yes
- ☒ B. No ✓

### Explanation :

Answer - B

```
{
```

```
"Parameters": {  
  "KeyName": {  
    "Description": "The EC2 Key Pair to allow SSH access to the instance",  
    "Type": "AWS::EC2::KeyPair::KeyName"  
  }  
}
```

There has to be a resources section in the template, this is a mandatory section.

For more information on template basics please visit the link:

- <http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/gettingstarted.templatebasics.html>  
(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/gettingstarted.templatebasics.html>)

Ask our Experts



QUESTION 20

UNATTEMPTED

An SQS queue has been created with the default settings. There are 3 messages published and not been consumed since 3 days. Will the messages still be available for consumption?

- ☒ A. True ✓
- ☐ B. False

#### Explanation :

Answer – A

Yes, since the default retention period for messages in an SQS queue is 4 days.

##### **Q: How long can I keep my messages in Amazon SQS message queues?**

Longer message retention provides greater flexibility to allow for longer intervals between message production and consumption.

You can configure the Amazon SQS message retention period to a value from 1 minute to 14 days. The default is 4 days. Once the message retention limit is reached, your messages are automatically deleted.

For more information on SQS please visit the link:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



QUESTION 21

UNATTEMPTED

What is the ideal step to enable disaster recovery for EC2 instances?

- ☐ A. Copy the running instance using the “Instance Copy” command to the another region
- ☐ B. Create an AMI of the instance and copy the AMI to the another region. Then launch the instance from the AMI itself. ✓
- ☐ C. Copy the instance from the US East region to the EU region
- ☐ D. Use the “Launch more like this” option to copy the instance from one region to another

### Explanation :

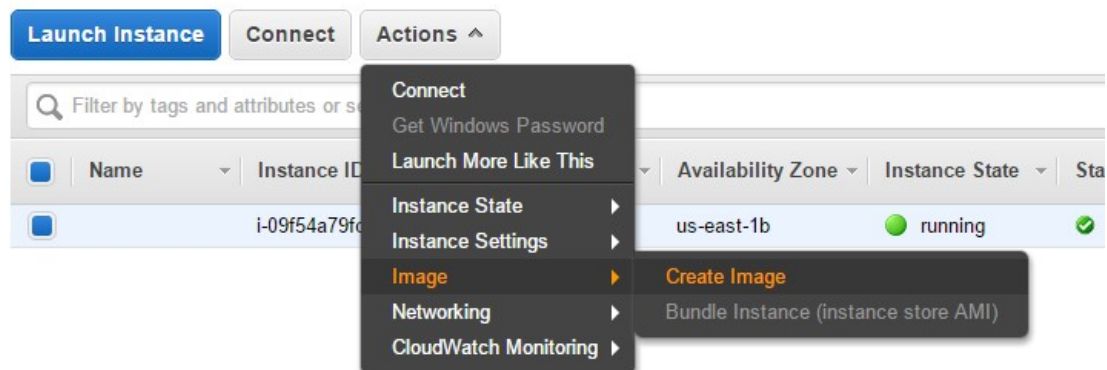
Answer – B

The first step is to ensure that you create an AMI of the current EC2 instance so that you have all of the required configuration settings.

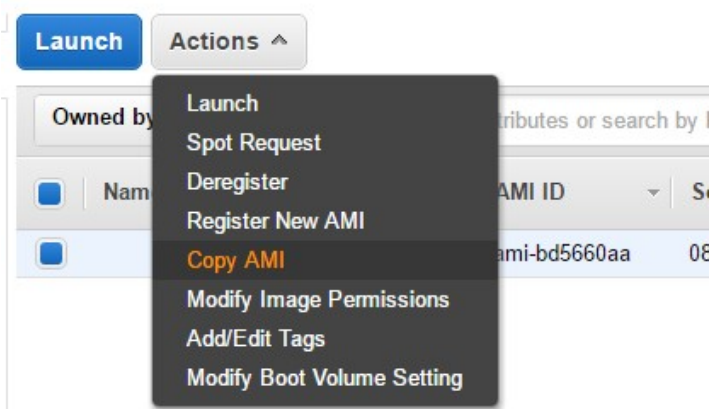
If you need an AMI across multiple regions, then you have to copy the AMI across regions. Note that by default AMI's that you have created will not be available across all regions.

To copy AMI's , follow the below steps

Step 1) The first step is to create an AMI from your running instance by choosing on Image->Create Image.



Step 2) Once the Image has been created, go to the AMI section in the EC2 dashboard and click on the Copy AMI option.



Step 3) In the next screen , you can specify where to copy the AMI to.

### Copy AMI

AMI ami-bd5660aa will be copied to a new AMI. Set the new AMI settings below.

**Destination region\*** Asia Pacific (Tokyo)

**Name**

**Description**

**Encryption** ☐ Encrypt target EBS snapshots ⓘ

Cancel Copy AMI

For the entire details to copy AMI's , please visit the link:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>  
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/CopyingAMIs.html>)

Ask our Experts



QUESTION 22 UNATTEMPTED

What is the general AWS limit for number of EBS snapshots?

- ☐ A. 10000 ✓
- ☐ B. 5000
- ☐ C. 100
- ☐ D. 1000

**Explanation :**

Answer – A

As per the AWS documentation it is clear that 10,000 is the limit.

**Amazon Elastic Block Store (Amazon EBS) Limits**

Resource	Default Limit
Number of EBS volumes	5,000
Number of EBS snapshots	10,000
Total volume storage of General Purpose SSD ( <i>gp2</i> ) volumes	20 TiB
Total volume storage of Provisioned IOPS SSD ( <i>io1</i> ) volumes	20 TiB
Total volume storage of Throughput Optimized HDD ( <i>st1</i> )	20 TiB
Total volume storage of Cold HDD ( <i>sc1</i> )	20 TiB
Total volume storage of Magnetic volumes	20 TiB
Total provisioned IOPS	40,000

For more information on AWS Service limits please visit the link:

- [http://docs.aws.amazon.com/general/latest/gr/aws\\_service\\_limits.html](http://docs.aws.amazon.com/general/latest/gr/aws_service_limits.html)  
([http://docs.aws.amazon.com/general/latest/gr/aws\\_service\\_limits.html](http://docs.aws.amazon.com/general/latest/gr/aws_service_limits.html))

Ask our Experts



QUESTION 23 UNATTEMPTED

Which of the below CLI commands can be used to merge autoscaling groups into single multi-zone group? Choose an answer from the options below

- ☐ A. update-auto-scaling-group ✓
- ☐ B. Create-auto-scaling-group
- ☐ C. Create-launch-configuration
- ☐ D. Describe-auto-scaling-groups

### Explanation :

Answer – A

An example is provided in the AWS documentation in which the first command used is update-auto-scaling-group

#### To merge separate single-zone groups into a single multi-zone group

1. Use the following `update-auto-scaling-group` command to add the `us-west-2c` Availability Zone to the supported Availability Zones for `my-group-a` and increase the maximum size of this group to allow for the instances from both single-zone groups:

```
aws autoscaling update-auto-scaling-group --auto-scaling-group-name my-group-a --avai
```

For more information on Autoscaling merging, please visit the below URL:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/merge-auto-scaling-groups.html>  
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/merge-auto-scaling-groups.html>)

Ask our Experts



### QUESTION 24 UNATTEMPTED

A system admin is trying to understand the Auto Scaling activities. Which of the below mentioned processes is performed by Auto Scaling? Select 3 Options

- ☐ A. Reboot Instance
- ☐ B. Schedule Actions ✓
- ☐ C. HealthCheck ✓
- ☐ D. AddToLoadBalancer ✓

### Explanation :

Answer – B,C and D

The below snapshot shows the options of an Autoscaling group. In the Suspended processes you can see the different processes for Autoscaling.

Max

1

Health Check Type

EC2

Health Check Grace Period

300

Termination Policies

Default x

Creation Time

Sun Jan 15 07:59:03 GMT-800 2017

Default Cooldown

300

Placement Group

Suspended Processes

Launch

Terminate

HealthCheck

ReplaceUnhealthy

AZRebalance

AlarmNotification

ScheduledActions

AddToLoadBalancer

Enabled Metrics

Instance Protection

Ask our Experts



## QUESTION 25 UNATTEMPTED

A user wants to ensure that whenever the CPU utilization of the AWS EC2 instance is above 90% he gets an email notification. Which of the below mentioned AWS services is helpful for this purpose?

- ☐ A. AWS CloudWatch + AWS SES
- ☐ B. AWS CloudWatch + AWS SQS
- ☒ C. AWS CloudWatch + AWS SNS ✓
- ☐ D. AWS CloudWatch + AWS SWF

### Explanation :

Answer – C

CloudWatch + SNS can be used to send email.

Refer AWS Documentation @ URL:

- [http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/US\\_SetupSNS.html](http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/US_SetupSNS.html)  
([http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/US\\_SetupSNS.html](http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/US_SetupSNS.html))

Ask our Experts



A user is trying to understand the ACL and policy for an S3 bucket. Which of the below mentioned policy permissions is equivalent to the READ permission on a bucket object?

- ☒ A. s3:GetObject ✓
- ☐ B. s3:PutObject
- ☐ C. s3:ReadObject
- ☐ D. s3:DeleteObject

#### Explanation :

Answer – A

As per the AWS documentation, the READ permission corresponds to the GetObject permission.

ACL Permission	Corresponding access policy permissions when the ACL permission is granted on a bucket	Corresponding access policy permissions when the ACL permission is granted on an object
READ	s3:ListBucket, s3:ListBucketVersions, and s3:ListBucketMultipartUploads	s3:GetObject, s3:GetObjectVersion, and s3:GetObjectTorrent
WRITE	s3:PutObject and s3:DeleteObject.  In addition, when the grantee is the bucket owner, granting WRITE permission in a bucket ACL allows the s3:DeleteObjectVersion action to be performed on any version in that bucket.	Not applicable

For more information on S3 ACL, please visit the URL:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/acl-overview.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/acl-overview.html>)

Ask our Experts



A user has launched an EC2 instance from an instance store backed AMI. The infrastructure team wants to create an AMI from the running instance. Which of the below mentioned steps is performed while creating the AMI? Choose 3 answers from the options below

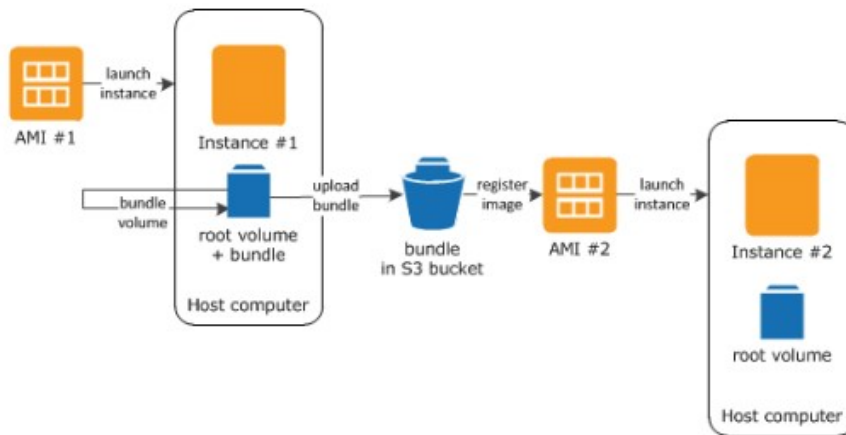


- ☐ A. Define the AMI launch permissions
- ☒ B. Upload the bundled volume ✓
- ☒ C. Register the AMI ✓
- ☒ D. Bundle the volume ✓

#### Explanation :

Answer – B, C and D

Below is the snapshot from the AWS documentation on the Instance store lifecycle of an AMI.



From here it is evident that there is no stage of 'Define the AMI launch permissions'

For more information on Instance store EC2 volumes, please visit the URL:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/creating-an-ami-instance-store.html>  
(<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/creating-an-ami-instance-store.html>)

Ask our Experts



#### QUESTION 28 UNATTEMPTED

In a subnet with CIDR block 10.0.0.0/24 , what are the IP Addresses reserved by AWS? Choose 3 answers from the options below

- ☒ A. 10.0.0.0 ✓
- ☒ B. 10.0.0.1 ✓
- ☒ C. 10.0.0.255 ✓
- ☐ D. 10.0.0.4

### Explanation :

Answer – A,B and C

The first four IP addresses and the last IP address in each subnet CIDR block are not available for you to use, and cannot be assigned to an instance. For example, in a subnet with CIDR block 10.0.0.0/24, the following five IP addresses are reserved:

- 10.0.0.0: Network address.
- 10.0.0.1: Reserved by AWS for the VPC router.
- 10.0.0.2: Reserved by AWS. The IP address of the DNS server is always the base of the VPC network range plus two; however, we also reserve the base of each subnet range plus two. 10.0.0.3: Reserved by AWS for future use.
- 10.0.0.255: Network broadcast address. We do not support broadcast in a VPC, therefore we reserve this address.

For more information on IP Reservation, please visit the link:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Subnets.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html)  
([http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Subnets.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html))

Ask our Experts



QUESTION 29

UNATTEMPTED

A company is experiencing high latency due to the Multi AZ feature for their MySQL database. Which of the below mentioned options can help alleviate the situation?

Select three options.

- ☐ A. Schedule the automated back up in non-working hours ✓
- ☐ B. Use a large or higher size instance ✓
- ☐ C. Use PIOPS ✓
- ☐ D. Take a snapshot from standby Replica

### Explanation :

Answer – A,B and C

In a Multi-AZ deployment, Amazon RDS automatically provisions and maintains a synchronous standby replica in a different Availability Zone. The primary DB instance is synchronously replicated across Availability Zones to a standby replica to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups

As per AWS the below are the best practices for multiAZ

For production workloads, we recommend you use Provisioned IOPS and DB instance classes (m1.large and larger) that are optimized for Provisioned IOPS for fast, consistent performance.

Also if backups are scheduled during working hours , then I/O can be suspended and increase the latency of the DB, hence it is better to schedule outside of office hours.

For more information on MultiAZ RDS, please visit the link:

- <http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZ.html>  
(<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZ.html>)

Ask our Experts



QUESTION 30

UNATTEMPTED


Can you delete a subnet which has instances in it?

- ☐ A. True
- ☒ B. False ✓

#### Explanation :

Answer – B

In AWS , when u try to delete a subnet which has instances it will not allow to delete it. The below error message will be shown when u try to delete a subnet with instances.

**Delete Subnet** 

The following subnets contain one or more instances or network interfaces. You cannot delete these subnets until those instances have been terminated, and the network interfaces have been deleted.

- subnet-dfd2a5f2 | Default

[Click here to view your instances.](#)  
[Click here to view your network interfaces.](#)

[Cancel](#) [Yes, Delete](#)

For more information on VPC and subnets please visit the link:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Subnets.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html)  
([http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Subnets.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html))

Ask our Experts



QUESTION 31

UNATTEMPTED

Which of the below alarm states corresponds to “metric is outside of the defined threshold”. Choose an answer from the options below

- ☐ A. OK
- ☒ B. ALARM ✓
- ☐ C. INSUFFICIENT\_DATA
- ☐ D. None of the above

**Explanation :**

Answer - B

An alarm has three possible states:

- OK—The metric is within the defined threshold
- ALARM—The metric is outside of the defined threshold
- INSUFFICIENT\_DATA—The alarm has just started, the metric is not available, or not enough data is available for the metric to determine the alarm state

For more information on Cloudwatch alarms, please visit the link:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/AlarmThatSendsEmail.html>  
(<http://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/AlarmThatSendsEmail.html>)

Ask our Experts



QUESTION 32 UNATTEMPTED

Which of the following error codes relates to “A malformed or canceled request from the client”. Choose an answer from the options below

- ☐ A. HTTPCode\_Backend\_2XX
- ☐ B. HTTPCode\_Backend\_3XX
- ☒ C. HTTPCode\_ELB\_4XX ✓
- ☐ D. HTTPCode\_ELB\_5XX

**Explanation :**

Answer – C

This is clearly given in the AWS documentation

## HTTPCode\_ELB\_4XX

**Cause:** A malformed or canceled request from the client.

### Solutions

- See [HTTP 400: BAD\\_REQUEST](#).
- See [HTTP 405: METHOD\\_NOT\\_ALLOWED](#).
- See [HTTP 408: Request Timeout](#).

For more information on troubleshooting ELB, please visit the link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/ts-elb-http-errors.html>  
(<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/ts-elb-http-errors.html>)

Ask our Experts



## QUESTION 33 UNATTEMPTED

Does the Basic HTTP load balancer support the X-Forwarded-For header?

- ☐ A. True ✓
- ☐ B. False

### Explanation :

Answer - A

This is clearly given in the AWS documentation

#### HTTP/HTTPS Load Balancer

Use Case	Front-End Protocol	Front-End Options	Back-End Protocol	Back-End Options	Notes
Basic HTTP load balancer	HTTP	NA	HTTP	NA	<ul style="list-style-type: none"><li>• Supports the <a href="#">X-Forwarded-For</a> header</li></ul>

For more information on ELB configuration, please visit the link:

- <http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/using-elb-listenerconfig-quickref.html> (<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/using-elb-listenerconfig-quickref.html>)

Ask our Experts



## QUESTION 34

UNATTEMPTED

You are trying to access AWS web services using your Facebook credentials from an application running on a mobile device . Which of the following will allow you to do so?

- ☐ A. AWS Credentials
- ☐ B. Active Directory
- ☐ C. LDAP
- ☒ D. Web Identity Federation ✓

**Explanation :**

Answer - D

This is mentioned in the AWS documentation with regards to using web identity federation

Imagine that you are creating a mobile app that accesses AWS resources, such as a game that runs on a mobile device and stores player and score information using Amazon S3 and DynamoDB.

When you write such an app, you'll make requests to AWS services that must be signed with an AWS access key. However, we **strongly** recommend that you do **not** embed or distribute long-term AWS credentials with apps that a user downloads to a device, even in an encrypted store. Instead, build your app so that it requests temporary AWS security credentials dynamically when needed using *web identity federation*. The supplied temporary credentials map to an AWS role that has only the permissions needed to perform the tasks required by the mobile app.

For more information on web identity federation, please visit the below URL:

- [http://docs.aws.amazon.com/IAM/latest/UserGuide/id\\_roles\\_providers\\_oidc.html](http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_providers_oidc.html)  
([http://docs.aws.amazon.com/IAM/latest/UserGuide/id\\_roles\\_providers\\_oidc.html](http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_providers_oidc.html))

Ask our Experts



## QUESTION 35

UNATTEMPTED

Which API call is used to describe the state of the specified instances with respect to the specified load balancer. Choose an answer from the options below

- ☒ A. DescribeInstanceHealth ✓
- ☐ B. DescribeLoadBalancerAttributes
- ☐ C. DescribeLoadBalancerPolicies

☐ D. None of the above

**Explanation :**

Answer - A

This is clearly given in the AWS documentation

**DescribeInstanceHealth**

Describes the state of the specified instances with respect to the specified load balancer. If no instances are specified, the call describes the state of all instances that are currently registered with the load balancer. If instances are specified, their state is returned even if they are no longer registered with the load balancer. The state of terminated instances is not returned.

For more information on the API call, please visit the link:

- [http://docs.aws.amazon.com/elasticloadbalancing/2012-06-01/APIReference/API\\_DescribeInstanceHealth.html](http://docs.aws.amazon.com/elasticloadbalancing/2012-06-01/APIReference/API_DescribeInstanceHealth.html)  
([http://docs.aws.amazon.com/elasticloadbalancing/2012-06-01/APIReference/API\\_DescribeInstanceHealth.html](http://docs.aws.amazon.com/elasticloadbalancing/2012-06-01/APIReference/API_DescribeInstanceHealth.html))

Ask our Experts



**QUESTION 36 UNATTEMPTED**

When you define a security group rule for EC2 instances, which of the below form part of the rule. Choose 3 answers form the options below.

- ☐ A. Protocol ✓
- ☐ B. Port Range ✓
- ☐ C. Destination ✓
- ☐ D. Rule ID

**Explanation :**

Answer – A,B and C

The below snapshot shows the rules in the AWS console. From the below you can see that Protocol, Port Range and Destination form part of the Rule.

Create Security GroupActions

Group ID : sg-590cf325Add filter

Name	Group ID	Group Name	VPC ID	Description
	sg-590cf325	Demo	vpc-3e6dde58	Demo

Security Group: sg-590cf325

DescriptionInboundOutboundTags

Edit

Type	Protocol	Port Range	Destination
All traffic	All	All	0.0.0.0/0

For more information on the VPC security rules, please visit the link:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_SecurityGroups.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html)  
([http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_SecurityGroups.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html))

Ask our Experts

QUESTION 37      UNATTEMPTED

Does your VPC come with a security group?

- ☒ A. Yes ✓
- ☐ B. No

Explanation :

Answer – A

The VPC comes with a default security group. This is mentioned in the AWS documentation



Your VPC automatically comes with a default security group. Each EC2 instance that you launch in your VPC is automatically associated with the default security group if you don't specify a different security group when you launch the instance.

The following table describes the default rules for a default security group.

Inbound			
Source	Protocol	Port Range	Comments
The security group ID (sg-xxxxxxx)	All	All	Allow inbound traffic from instances assigned to the same security group.
Outbound			
Destination	Protocol	Port Range	Comments
0.0.0.0/0	All	All	Allow all outbound IPv4 traffic.
::/0	All	All	Allow all outbound IPv6 traffic. This rule is added by default if you create a VPC with an IPv6 CIDR block or if you associate an IPv6 CIDR block with your existing VPC.

For more information on the default security rules, please visit the link:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_SecurityGroups.html#DefaultSecurityGroup](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html#DefaultSecurityGroup) ([http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_SecurityGroups.html#DefaultSecurityGroup](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_SecurityGroups.html#DefaultSecurityGroup))

Ask our Experts



#### QUESTION 38 UNATTEMPTED

What is the default limit for the number of VPC's per region?

- ☒ A. 5 ✓
- ☐ B. 10
- ☐ C. 100
- ☐ D. There is no limit

#### Explanation :

Answer - A

This is clearly mentioned in the AWS documentation.

## VPC and Subnets

Resource	Default limit	Comments
VPCs per region	5	If you need to increase this limit, <a href="#">submit a request</a> . The limit for Internet gateways per region is directly correlated to this one. Increasing this limit will increase the limit on Internet gateways per region by the same amount.
Subnets per VPC	200	If you need to increase this limit, <a href="#">submit a request</a> .

For more information on the AWS limits for VPC please visit the link:

- [http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Appendix\\_Limits.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Appendix_Limits.html) ([http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Appendix\\_Limits.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Appendix_Limits.html))

Ask our Experts



### QUESTION 39 UNATTEMPTED

What is the range for CIDR blocks for subnets for IPv4?

- ☐ A. From /12 to /28
- ☐ B. From /14 to /28
- ☒ C. From /16 to /28 ✓
- ☐ D. From /10 to /28

#### Explanation :

Answer - C

This is clearly mentioned in the AWS documentation.

#### IPv4 and IPv6 Characteristics and Restrictions

IPv4	IPv6
The format is 32-bit, 4 groups of 4 numerical digits.	The format is 128-bit, 8 groups of 4 hexadecimal digits.
Default and required for all VPCs; cannot be removed.	Opt-in only.
The VPC CIDR block size can be from /16 to /28.	The VPC CIDR block size is fixed at /56.
The subnet CIDR block size can be from /16 to /28.	The subnet CIDR block size is fixed at /64.

For more information on IP addressing please visit the link:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-ip-addressing.html>  
(<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-ip-addressing.html>)

Ask our Experts



QUESTION 40 UNATTEMPTED

What is the feature in S3 that helps replicates data across AWS regions. Choose an answer from the options below

- ☐ A. Cross site replication
- ☒ B. Cross region replication ✓
- ☐ C. Secondary region replication
- ☐ D. Primary region replication

**Explanation :**

Answer - B

CRR is an Amazon S3 feature that automatically replicates data across AWS regions. With CRR, every object uploaded to an S3 bucket is automatically replicated to a destination bucket in a different AWS region that you choose. You can use CRR to provide lower-latency data access in different geographic regions. CRR can also help if you have a compliance requirement to store copies of data hundreds of miles apart.

For more information CRR please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/crr.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/crr.html>)

Ask our Experts



QUESTION 41 UNATTEMPTED

To store data in S3, what is the first thing you need to do? Choose an answer from the options below

- ☐ A. Create an EC2 instance
- ☒ B. Create a bucket ✓
- ☐ C. Mark the bucket as public

### Explanation :

Answer - B

This is clearly mentioned in the AWS documentation.

Amazon S3 is cloud storage for the Internet. To upload your data (photos, videos, documents etc.), you first create a bucket in one of the AWS Regions. You can then upload any number of objects to the bucket.

In terms of implementation, buckets and objects are resources, and Amazon S3 provides APIs for you to manage them. For example, you can create a bucket and upload objects using the Amazon S3 API. You can also use the Amazon S3 console to perform these operations. The console internally uses the Amazon S3 APIs to send requests to Amazon S3.

For more information on S3 buckets please visit the link:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingBucket.html>  
(<http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingBucket.html>)

Ask our Experts



### QUESTION 42 UNATTEMPTED

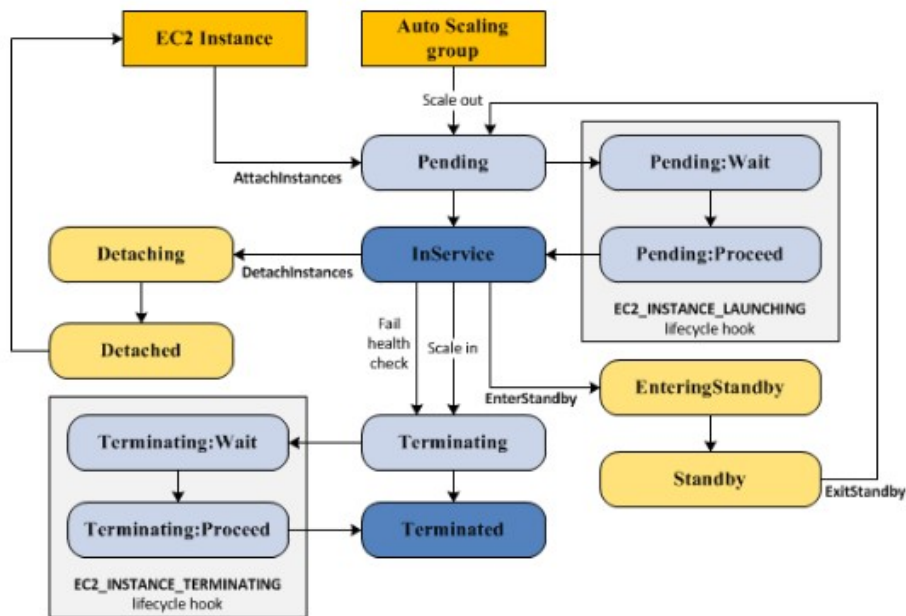
What is the state of any instance when it is fully configured and passes the Amazon EC2 health checks, it is attached to the Auto Scaling group

- ☐ A. InState
- ☒ B. InService ✓
- ☐ C. OutService
- ☐ D. OutState

### Explanation :

Answer – B

Below is the diagram for Autoscaling. When any instance is fully configured and passes the Amazon EC2 health checks, it is attached to the Auto Scaling group and enters the InService state.



For more information on Autoscaling lifecycle please visit the link:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/AutoScalingGroupLifecycle.html>  
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/AutoScalingGroupLifecycle.html>)

Ask our Experts



#### QUESTION 43 UNATTEMPTED

In Autoscaling what is the set of instructions that tells Auto Scaling how to respond to alarm messages.

- ☒ A. Policy ✓
- ☐ B. Document
- ☐ C. Cloudformation
- ☐ D. JSON

#### Explanation :

Answer - A

This is clearly given in the AWS documentation

An Auto Scaling group uses a combination of alarms and policies to determine when the conditions for scaling are met. An *alarm* is an object that watches over a single metric (for example, the average CPU utilization of the EC2 instances in your Auto Scaling group) over a specified time period. When the value of the metric breaches the threshold that you defined, for the number of time periods that you specified, the alarm performs one or more actions (such as sending messages to Auto Scaling). A *policy* is a set of instructions that tells Auto Scaling how to respond to alarm messages.

For more information on Dynamic scaling please visit the link:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/as-scale-based-on-demand.html>  
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/as-scale-based-on-demand.html>)

Ask our Experts



#### QUESTION 44 UNATTEMPTED

In Autoscaling what are the different types of checks carried out on EC2 Instances.  
Choose 3 answers from the options below

- ☐ A. Random checks
- ☐ B. Status checks ✓
- ☐ C. Health Checks ✓
- ☐ D. Custom checks ✓

#### Explanation :

Answer – B,C and D

This is clearly given in the AWS documentation

An Auto Scaling instance is either healthy or unhealthy. Auto Scaling determines the health status of an instance using one or more of the following:

- Status checks provided by Amazon EC2. For more information, see [Status Checks for Your Instances](#) in the *Amazon EC2 User Guide for Linux Instances*.
- Health checks provided by Elastic Load Balancing. For more information, see [Health Checks for Your Target Groups](#) in the *Application Load Balancer Guide* or [Configure Health Checks for Your Classic Load Balancer](#) in the *Classic Load Balancer Guide*.
- Custom health checks. For more information, see [Instance Health Status and Custom Health Checks](#).

For more information on health checks scaling please visit the link:

- <http://docs.aws.amazon.com/autoscaling/latest/userguide/healthcheck.html>  
(<http://docs.aws.amazon.com/autoscaling/latest/userguide/healthcheck.html>)

Ask our Experts



## QUESTION 45

UNATTEMPTED

Which of the following databases supports Multi-AZ deployments without any dependency on native Mirroring? Choose 3 answers from the options below

- ☐ A. MySQL ✓
- ☐ B. Oracle ✓
- ☐ C. MS SQL
- ☐ D. PostgreSQL ✓

**Explanation :**

Answer – A,B and D

As per the AWS documentation it is very clear that MultiAZ is supported for MySQL, MariaDb, Oracle and PostgreSQL. With Microsoft SQL server, you need to use the native mirroring to achieve High Availability.

Multi-AZ deployments for the [MySQL](#), [MariaDB](#), [Oracle](#), and [PostgreSQL](#) engines utilize synchronous physical replication to keep data on the standby up-to-date with the primary. Multi-AZ deployments for the [SQL Server](#) engine use synchronous logical replication to achieve the same result, employing SQL Server-native Mirroring technology. Both approaches safeguard your data in the event of a DB Instance failure or loss of an Availability Zone.

For more information on RDS MultiAZ, please visit the link:

- <https://aws.amazon.com/rds/details/multi-az/> (<https://aws.amazon.com/rds/details/multi-az/>)

Ask our Experts



## QUESTION 46

UNATTEMPTED

Which of the following functions are condition Intrinsic functions in Cloudformation? Choose 3 answers from the options below

- ☐ A. Fn::And ✓
- ☐ B. Fn::Xor
- ☐ C. Fn::If ✓
- ☐ D. Fn::Equals ✓

**Explanation :**

Answer – A,C and D

This is clearly given in the AWS documentation

### Condition Intrinsic Functions

You can use the following intrinsic functions to define conditions:

- Fn::And
- Fn::Equals
- Fn::If
- Fn::Not
- Fn::Or

For more information on Cloudformation functions please visit the link:

- <http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/conditions-section-structure.html> (<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/conditions-section-structure.html>)

Ask our Experts



#### QUESTION 47 UNATTEMPTED

Which of the below combination of services can be used in conjunction for continuous delivery to automatically build and test changes to your AWS CloudFormation templates before promoting them to production stacks.

- ☐ A. AWS Cloudformation and SNS
- ☒ B. AWS Cloudformation and AWS CodePipeline ✓
- ☐ C. AWS Cloudformation and EC2
- ☐ D. AWS Cloudformation and SES

Explanation :

Answer - B



## Continuous Delivery with AWS CodePipeline

Continuous delivery is a release practice in which code changes are automatically built, tested, and prepared for release to production. With AWS CloudFormation and AWS CodePipeline, you can use continuous delivery to automatically build and test changes to your AWS CloudFormation templates before promoting them to production stacks. This release process lets you rapidly and reliably make changes to your AWS infrastructure.

For more information on Continuous Delivery please visit the link:

- <http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/continuous-delivery-codepipeline.html>  
(<http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/continuous-delivery-codepipeline.html>)

Ask our Experts



### QUESTION 48 UNATTEMPTED

At the moment what are the operating systems supported for EC2 instances.  
Choose 3 answers from the options below?

- ☐ A. Mac OS
- ☐ B. Windows server ✓
- ☐ C. Red Hat Enterprise ✓
- ☐ D. SUSE Linux Enterprise Server ✓

#### Explanation :

Answer – B,C and D

This is given in the AWS documentation for the currently supported Operating systems

#### Q: What operating system environments are supported?

Amazon EC2 currently supports a variety of operating systems including: Amazon Linux, Ubuntu, Windows Server, Red Hat Enterprise Linux, SUSE Linux Enterprise Server, Fedora, Debian, CentOS, Gentoo Linux, Oracle Linux, and FreeBSD. We are looking for ways to expand it to other platforms.

For more information on EC2 please visit the link:

- <https://aws.amazon.com/ec2/faqs/> (<https://aws.amazon.com/ec2/faqs/>)

Ask our Experts



What is state of the EC2 instance that makes it billable for EC2 instance hours.

- ☒ A. Running ✓
- ☐ B. Stopped
- ☐ C. Terminated

**Explanation :**

Answer - A

This is given in the AWS documentation

**Q: What defines billable EC2 instance-hours?**

Instance-hours are billed for any time your instances are in a "running" state. If you no longer wish to be charged for your instance, you must "stop" or "terminate" the instance to avoid being billed for additional instance-hours. Billing starts when an instance transitions into the running state.

For more information on EC2 please visit the link:

- <https://aws.amazon.com/ec2/faqs/> (<https://aws.amazon.com/ec2/faqs/>)

Ask our Experts



If data has to be transferred between EC2 instances in different regions , would the data rate transfer be charged at Internet Data transfer rates?

- ☒ A. Yes ✓
- ☐ B. No

**Explanation :**

Answer – A

This is given in the AWS documentation

**Q. If I have two instances in different regions, how will I be charged for data transfer?**

Each instance is charged for its data in and data out at Internet Data Transfer rates. Therefore, if data is transferred between these two instances, it is charged at Internet Data Transfer Out for the first instance and at Internet Data Transfer In for the second instance.

For more information on EC2 please visit the link:

- <https://aws.amazon.com/ec2/faqs/> (<https://aws.amazon.com/ec2/faqs/>)

Ask our Experts



QUESTION 51 UNATTEMPTED

In order to use Enhanced Networking for EC2 do you need to pay extra?

- ☐ A. True
- ☐ B. False ✓

**Explanation :**

Answer – B

This is given in the AWS documentation

**Q: Do I need to pay an additional fee to use Enhanced Networking?**

No, there is no additional fee for Enhanced Networking. To take advantage of Enhanced Networking you need to launch the appropriate AMI on a supported instance type in a VPC.

For more information on EC2 please visit the link:

- <https://aws.amazon.com/ec2/faqs/> (<https://aws.amazon.com/ec2/faqs/>)

Ask our Experts



QUESTION 52 UNATTEMPTED

Which of the below instances are available as Spot Instances. Choose 3 answers from the options below

- ☐ A. Unix servers ✓
- ☐ B. Linux servers ✓
- ☐ C. Windows servers ✓
- ☐ D. Windows Server with SQL Server

**Explanation :**

Answer – A,B and C

This is given in the AWS documentation

**Q. Which operating systems are available as Spot instances?**

Linux/UNIX and Windows Server are available. Windows Server with SQL Server is not currently available.

For more information on EC2 please visit the link:

- <https://aws.amazon.com/ec2/faqs/> (<https://aws.amazon.com/ec2/faqs/>)

Ask our Experts



**QUESTION 53 UNATTEMPTED**

Can you connect your corporate data center using a Hardware VPN connection to a VPC?

- ☒ A. True ✓
- ☐ B. False

**Explanation :**

Answer – A

As per the documentation below are the options for connecting to a VPC

**Q. What are the connectivity options for my VPC?**

You may connect your VPC to:

- The Internet (via an Internet gateway)
- Your corporate data center using a Hardware VPN connection (via the virtual private gateway)
- Both the Internet and your corporate data center (utilizing both an Internet gateway and a virtual private gateway)
- Other AWS services (via Internet gateway, NAT, virtual private gateway, or VPC endpoints)
- Other VPCs (via VPC peering connections)

For more information on VPC please visit the link:

- <https://aws.amazon.com/vpc/faqs/> (<https://aws.amazon.com/vpc/faqs/>)

Ask our Experts



**QUESTION 54 UNATTEMPTED**

Can you change the size of a VPC once created?

☒ A. True ✓

☐ B. False

### Explanation :

Answer – A

As per the documentation this is possible

#### Q. Can I change a VPC's size?

Yes. You can expand your existing VPC by adding four (4) secondary IPv4 IP ranges (CIDRs) to your VPC. You can shrink your VPC by deleting the secondary CIDR blocks you have added to your VPC. You cannot however change the size of the IPv6 address range of your VPC.

For more information on VPC please visit the link:

- <https://aws.amazon.com/vpc/faqs/> (<https://aws.amazon.com/vpc/faqs/>)

The question is as straightforward as it is, it is asking whether are we able to change the size of the VPC after it is created. The answer is yes, this is a new feature which was announced recently.

AWS Docs says:

## Adding IPv4 CIDR Blocks to a VPC

You can associate secondary IPv4 CIDR blocks with your VPC. When you associate a CIDR block with your VPC, a route is automatically added to your VPC route tables to enable routing within the VPC (the destination is the CIDR block and the target is *local*).

In the following example, the VPC on the left has a single CIDR block (*10.0.0/16*) and two subnets. The VPC on the right represents the architecture of the same VPC after you've added a second CIDR block (*10.2.0.0/16*) and created a new subnet from the range of the second CIDR.

#### Q. Can I change a VPC's size?

No. To change the size of a VPC you must terminate your existing VPC and create a new one.

You may refer to the link below for the news and the documentations:

- <https://aws.amazon.com/about-aws/whats-new/2017/08/amazon-virtual-private-cloud-vpc-now-allows-customers-to-expand-their-existing-vpcs/> (<https://aws.amazon.com/about-aws/whats-new/2017/08/amazon-virtual-private-cloud-vpc-now-allows-customers-to-expand-their-existing-vpcs/>)
- [https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Subnets.html#vpc-resize](https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html#vpc-resize) ([https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_Subnets.html#vpc-resize](https://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html#vpc-resize))

Ask our Experts



QUESTION 55 UNATTEMPTED

Does VPC support multicast?

- ☐ A. True
- ☐ B. False ✓

**Explanation :**

Answer – B

As per the documentation this is not possible

**Q. Does Amazon VPC support multicast or broadcast?**

No.

For more information on VPC please visit the link:

- <https://aws.amazon.com/vpc/faqs/> (<https://aws.amazon.com/vpc/faqs/>)

Ask our Experts



QUESTION 56 UNATTEMPTED


Which of the below AWS services is normally used in a decoupling scenario and controlling of EC2 instances in an Autoscaling group

- ☐ A. SQS ✓
- ☐ B. SNS
- ☐ C. SES
- ☐ D. S3

**Explanation :**

Answer – A

When you have an SQS configured with EC2 instances, the documented option is to scale up EC2 instances in an AutoScaling group based on length of the message queue.



For more information on SQS, please visit the below URL:

- <https://aws.amazon.com/sqs/faqs/> (<https://aws.amazon.com/sqs/faqs/>)

Ask our Experts



#### QUESTION 57 UNATTEMPTED

Is it possible for a VPC to span multiple Availability zones?

- ☒ A. True ✓
- ☐ B. False

#### Explanation :

Answer – A

As per the documentation this is possible

#### Q. Can a VPC span multiple Availability Zones?

Yes.

For more information on VPC please visit the link:

- <https://aws.amazon.com/vpc/faqs/> (<https://aws.amazon.com/vpc/faqs/>)

Ask our Experts



Which of the below API calls is used to add data points to Cloudwatch. Choose an answer from the options below

- ☒ A. PutMetricData ✓
- ☐ B. putMetricAlarm
- ☐ C. SetAlarmState
- ☐ D. ListMetrics

**Explanation :**

Answer - A

This is clearly given in the AWS documentation

**PutMetricData**

Publishes metric data points to Amazon CloudWatch. Amazon CloudWatch associates the data points with the specified metric. If the specified metric does not exist, Amazon CloudWatch creates the metric. When Amazon CloudWatch creates a metric, it can take up to fifteen minutes for the metric to appear in calls to [ListMetrics](#).

For more information on PutMetricData, please visit the below URL:

- [http://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API\\_PutMetricData.html](http://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API_PutMetricData.html)  
([http://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API\\_PutMetricData.html](http://docs.aws.amazon.com/AmazonCloudWatch/latest/APIReference/API_PutMetricData.html))

Ask our Experts



Which of the below instances is normally used as a jump server to access EC2 instances in a private subnet in a VPC. Choose an answer from the options below

- ☒ A. Bastion Host ✓
- ☐ B. Primary Host
- ☐ C. Backup Host
- ☐ D. Secondary Host



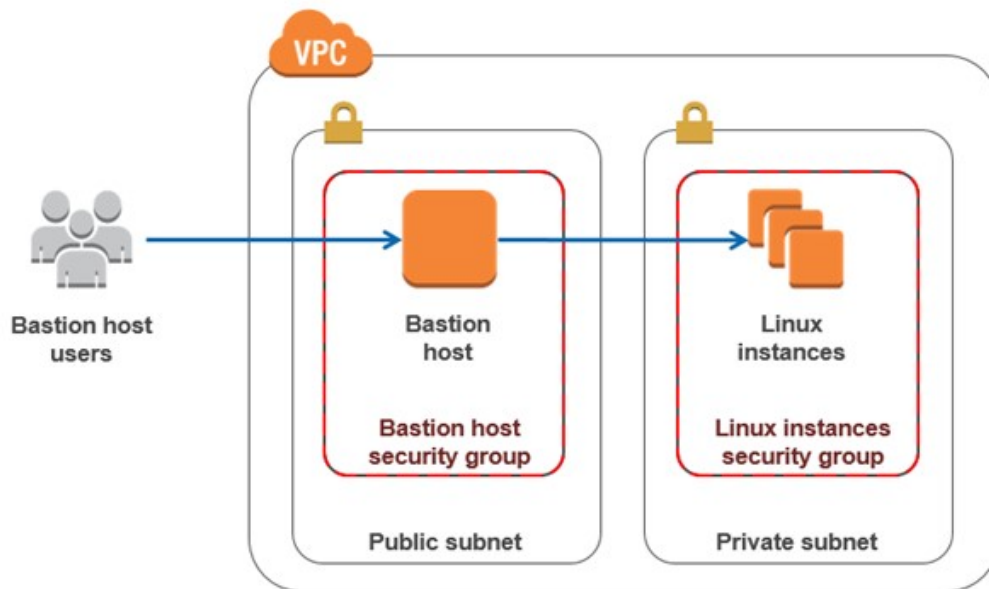
### Explanation :

Answer – A

A bastion host is a special purpose computer on a network specifically designed and configured to withstand attacks. The computer generally hosts a single application, for example a proxy server, and all other services are removed or limited to reduce the threat to the computer.

In AWS, A bastion host is kept on a public subnet. Users log on to the bastion host via SSH or RDP and then use that session to manage other hosts in the private subnets.

This is a security practice adopted by many organization to secure the assets in their private subnets.



Ask our Experts



QUESTION 60

UNATTEMPTED

How can you ensure maximum protection of preserved versions in S3? Choose one answer from the options below.

- ☐ A. Versioning
- ☐ B. Encryption
- ☐ C. MFA Delete ✓
- ☐ D. None of the above

### Explanation :

Answer – C

This is clearly given in the AWS documentation

**Q: How can I ensure maximum protection of my preserved versions?**

Versioning's MFA Delete capability, which uses [multi-factor authentication](#), can be used to provide an additional layer of security. By default, all requests to your Amazon S3 bucket require your AWS account credentials. If you enable Versioning with MFA Delete on your Amazon S3 bucket, two forms of authentication are required to permanently delete a version of an object: your AWS account credentials and a valid six-digit code and serial number from an authentication device in your physical possession. To learn more about enabling Versioning with MFA Delete, including how to purchase and activate an authentication device, please refer to the [Amazon S3 Technical Documentation](#).

For more information on S3 please visit the below link:

- <https://aws.amazon.com/s3/faqs/> (<https://aws.amazon.com/s3/faqs/>)

**Note:**

In the question itself we are emphasising the fact that versioning is already enabled and how we can ensure maximum protection on these preserved versions of the files. Thus by enabling MFA delete capability it ensures an additional layer of security to these previous versions of the files.

Ask our Experts



Finish Review (<https://www.whizlabs.com/learn/course/aws-csyopaa-practice-tests/quiz/13327>)

## Certification

- 🔗 Cloud Certification  
(<https://www.whizlabs.com/cloud-certification-training-courses/>)
- 🔗 Java Certification  
(<https://www.whizlabs.com/oracle-java-certifications/>)
- 🔗 PM Certification  
(<https://www.whizlabs.com/project-management-certifications/>)
- 🔗 Big Data Certification  
(<https://www.whizlabs.com/big-data-certifications/>)

## Company

- 🔗 Support  
(<https://help.whizlabs.com/hc/en-us>)
- 🔗 Discussions (<http://ask.whizlabs.com/>)
- 🔗 Blog (<https://www.whizlabs.com/blog/>)

## Mobile App

 Android Coming Soon

 iOS Coming Soon

## Follow us



(<https://www.facebook.com/whizlabs.software/>)



(<https://in.linkedin.com/company/whizlabs-software>)



(<https://twitter.com/whizlabs?lang=en>)



(<https://plus.google.com/+WhizlabsSoftware>)